

THE ART OF AI — BOOK 1

A Practical Guide to Better
Prompts, Better Context, and
Better Workflows


THE ART OF AI — BOOK 1

The Art of AI

How to Get Better Results From
Every AI Conversation

A Practical Guide to Better Prompts, Better Context,
and Better Workflows

THE ART OF AI · ELEANOR MERCER





THE ART OF AI — BOOK 1

The Art of AI

How to Get Better Results From Every
AI Conversation

A Practical Guide to Better Prompts, Better Context,
and Better Workflows

ELEANOR MERCER

FIRST EDITION • 2026

*For every person who typed something
into ChatGPT, got a vague wall of text,
and quietly wondered whether
everyone else was getting something
they weren't.*

COPYRIGHT

Copyright © 2026 The Art of AI

All rights reserved. No part of this publication may be reproduced, distributed, or transmitted in any form or by any means, including photocopying, recording, or other electronic or mechanical methods, without the prior written permission of the publisher, except in the case of brief quotations embodied in critical reviews and certain other non-commercial uses permitted by copyright law.

Published by: The Art of AI

Series: The Art of AI, Book 1

First edition: 2026

AI Disclosure: AI tools were used in the research, drafting, and editing of this book. All factual claims were verified against primary sources at the time of writing. Examples were reviewed for clarity and intended use. Editorial judgment, structure, and final decisions are the author's own.

Amazon KDP AI disclosure completed in the publishing dashboard per current requirements. KDP's AI content policy should be rechecked before any future edition.

CONTENTS

INTRODUCTION — You Are Not Bad at AI	7
PART I — A DIFFERENT WAY TO THINK ABOUT AI	10
Chapter 1 — Why AI Gives You Generic Answers (and How to Change That)	11
Chapter 2 — The AI Interaction Stack — What Actually Determines Output Quality	17
Chapter 3 — The 7 Questions — Your Portable AI Framework	25
Chapter 4 — How to Talk to an AI (It's Not Google)	31
PART II — BUILD BETTER INTERACTIONS	38
Chapter 5 — Give AI a Job, Not Just a Question	39
Chapter 6 — Context Is Everything	45
Chapter 7 — Instructions That Actually Work	51
Chapter 8 — Examples Change Everything	57
Chapter 9 — Control the Output	62
PART III — BUILD AN AI ENVIRONMENT	69
Chapter 10 — Stop Starting From Zero Every Time	70
Chapter 11 — Projects and Persistent Instructions	74
Chapter 12 — Files, Documents, and Source Material	79
Chapter 13 — Tools — What AI Can Do Beyond Text	83
Chapter 14 — The Memory Problem (and How to Work Around It)	90
Chapter 15 — Building Your Personal AI Workspace	94
PART IV — FROM PROMPTS TO WORKFLOWS	99
Chapter 16 — When One Prompt Isn't Enough	100

Chapter 17 — Breaking Tasks Into Steps	105
Chapter 18 — Critique, Edit, and Improve	109
Chapter 19 — Chains, Loops, and Handoffs	114
PART V — 25 REAL-WORLD WORKFLOWS	119
BUSINESS AND PROFESSIONAL	
W1 — Write and Improve a Professional Email	120
W2 — Prepare for a Difficult Conversation	122
W3 — Research a Topic for a Meeting	123
W4 — Summarise a Long Document	125
W5 — Draft a Project Proposal	127
W6 — Write a Job Posting	129
W7 — Prepare a Client Report	130
W8 — Create a Meeting Agenda and Follow-Up	132
MARKETING AND CONTENT	
W9 — Write a LinkedIn Post	134
W10 — Draft a Blog Post from Notes	135
W11 — Repurpose Content Across Formats	137
W12 — Write Website Copy	138
RESEARCH AND LEARNING	
W13 — Research a Topic You Know Nothing About	139
W14 — Understand a Complex Document	141
W15 — Compare Options Before a Decision	142
W16 — Learn a New Skill Faster	144
PERSONAL PRODUCTIVITY	
W17 — Plan a Project	146
W18 — Process a Backlog of Tasks	147
W19 — Draft a Difficult Personal Message	148
W20 — Prepare for a Job Interview	150
HOME AND LIFE	
W21 — Plan a Trip	152

W22 — Understand a Contract or Legal Document	153
W23 — Get a Second Opinion on a Health Question	154
WRITING AND EDITING	
W24 — Edit and Improve Your Own Writing	155
W25 — Turn Rough Notes into a Polished Document	157
PART VI — THE AI USER'S PLAYBOOK	160
Chapter 20 — How to Verify AI Output	161
Chapter 21 — What AI Gets Wrong (Consistently)	168
Chapter 22 — Calibrating Trust	176
Chapter 23 — Building Your AI Habits	180
Chapter 24 — Keeping Up Without Burning Out	186
APPENDICES	190
Appendix A — The 7 Questions — Quick Reference Card	190
Appendix B — The Before → Better → Best Worksheet	191
Appendix C — Verification Checklist	194
Appendix D — AI Use Cases by Profession	197
Appendix E — Product Notes — Platform-Specific Information	203
E1 — ChatGPT (OpenAI)	204
E2 — Claude (Anthropic)	206
E3 — Gemini (Google)	207
E4 — Microsoft Copilot	208
E5 — Cross-Platform Notes	209
E6 — Image Generation (All Platforms)	210
A Note on Sources	211
About the Author	213
About This Book	214

INTRODUCTION

You Are Not Bad at AI

There is a particular kind of frustration that comes from using a tool that is supposed to be extraordinary and finding it merely annoying.

You type a question into ChatGPT or Claude or Gemini. You wait. You read the response. It is, somehow, both very long and not quite what you wanted. It is technically correct in the way that a weather forecast saying "conditions may vary" is technically correct. It is not wrong. It is not useful. You close the tab and do the task yourself.

This happens to most people most of the time. Not because AI is bad — it isn't — but because the interaction was badly set up from the start.

Here is what no one tells you: the quality of what AI produces is determined less by the AI and more by what you give it. Not your technical knowledge. Not your ability to write special commands. What you give it: the information, the instructions, the structure, the context, the constraints, the documents, the clear job description.

This book is about that system.

WHAT YOU WILL ACTUALLY LEARN

By the end of Part I, you will understand why AI produces generic output and exactly what changes when it produces something good. That understanding alone is worth more than any list of prompts, because it means you can diagnose what went wrong in any interaction and fix it — rather than just trying different words until something works.

By the end of Part II, you will know how to build an AI interaction deliberately: giving AI the role, the context, the instructions, the examples, and the output format that get you the result you

need.

By the end of Part III, you will know how to build an AI environment — a persistent workspace where you stop repeating yourself, where your files and preferences carry forward, and where AI actually knows something about your situation.

By the end of Part IV, you will be able to design multi-step workflows: decomposing a complex task into stages, running them in sequence, and using AI's output from one step as the input to the next.

Part V gives you twenty-five tested workflows to use immediately or adapt. Part VI gives you the habits and verification practices that distinguish people who use AI reliably from people who use it hopefully.

THE FRAMEWORK

Running through all of it is a single seven-question framework you can carry in your head. Before any significant AI interaction, these seven questions identify what you need to provide. They work on every AI platform — ChatGPT, Claude, Gemini, Copilot, and every tool built on these models. They will still work when the current models are replaced by the next generation.

The framework is called The 7 Questions. You will encounter it in Chapter 3, and you will use it for the rest of the book.

A WORD ABOUT TECHNICAL KNOWLEDGE

None of this requires technical knowledge. You do not need to understand how large language models work. You do not need to know what a token is, or an API, or an embedding. When technical concepts appear in this book, they appear because they directly affect how you use AI, and they are always explained in ordinary language before the technical term is introduced.

What you need is the willingness to think about AI interactions differently — not as typing questions into a search engine, but as briefing a capable colleague who starts every conversation

knowing nothing about you, your work, or what you need.

Once you understand that, everything else follows.

Let's begin.

A Different Way to Think About AI

The four chapters in this part build the mental model that makes everything else in the book work. They are short by design. If you understand what this part is teaching, you will already be getting noticeably better results before you reach Part II.

CHAPTER 1

Why AI Gives You Generic Answers (and How to Change That)

THE PROBLEM EVERYONE HAS BUT NOBODY EXPLAINS

Ask ten people what frustrates them about AI and you will hear the same answers:

"It's too generic."

"It doesn't understand what I actually want."

"It's confident but wrong."

"The answer doesn't fit my situation."

"It gives me five paragraphs when I needed two sentences."

These are not problems with the AI. They are problems with the interaction.

This chapter explains why generic inputs produce generic outputs — and what specifically changes when they don't.

THE SLOT MACHINE MODEL (AND WHY IT'S WRONG)

Most people approach AI the way they might approach a vending machine: insert request, press button, hope for the right thing to fall out. When it doesn't, they try slightly different words. Sometimes this works. Usually it's a matter of luck.

This approach treats AI as a slot machine. It also produces slot-machine results: random, unpredictable, and mostly disappointing.

The people who get consistently good results from AI are not luckier. They are not using secret prompts. They understand something that slot-machine users don't: AI output is shaped by

input, and almost every dimension of that input is under your control.

WHY AI IS GENERIC BY DEFAULT

When you send AI a vague request — "Write me a summary of this report" — it has to make assumptions about everything it doesn't know:

- Who is going to read this summary?
- How long should it be?
- What are the most important points to include?
- What should be left out?
- What tone is appropriate?
- What does the reader already know?
- What is the purpose of the summary?

AI will fill in all of these gaps. It will fill them in with the most statistically average answers it can construct — the kind of summary that would be acceptable to the largest number of readers in the largest number of situations. This is why AI output so often feels as though it was written for everyone and therefore fits no one particularly well.

The good news: every assumption AI makes is one you could have specified yourself. The summary would look completely different if you told AI: *"This is for our CEO, who hasn't read the report and needs to decide whether to approve a £200,000 budget. She wants two paragraphs, direct language, and the headline number up front."*

Same document. Same AI. Dramatically different output. The only difference is what you provided.

WHY IT WORKS

AI tools generate text by predicting what words should follow what came before, based on patterns learned from enormous quantities of written material. This means the quality and specificity of what they produce is directly shaped by the quality and specificity of what they receive. Providing more relevant information doesn't confuse AI — it constrains the prediction space to something closer to what you actually need.

THE THREE MOST COMMON INPUT FAILURES

After looking at hundreds of AI interactions that produced poor results, three input problems appear again and again:

1. THE VAGUE JOB DESCRIPTION

"Help me with my presentation" tells AI you want a presentation. It doesn't tell it the topic, the audience, the desired length, your current draft, what help means, or what "better" would look like. AI has no choice but to respond generically to a generic request.

2. THE MISSING CONTEXT

"Reply to this email" gives AI a text to work with. It doesn't tell AI who sent it, what the relationship is, what a good reply should accomplish, whether you need to decline or accept, how formal the tone should be, or what your constraints are. AI will guess. Sometimes it guesses right. Usually the first reply it produces needs significant editing.

3. THE WRONG ASSUMED AUDIENCE

"Explain compound interest" gives AI a topic. But who is asking? A secondary school student? A financial professional? Someone deciding whether to take out a mortgage? AI will produce an

explanation pitched at a generic reader — which means it may be too basic, too technical, or simply beside the point for your actual need.

REAL-WORLD EXAMPLE

Sarah is a marketing manager who needs to send a follow-up email to a potential agency partner after an introductory call. She types:

"Write a follow-up email after a business meeting."

AI produces a template with placeholders like [Name] and [Company] and generic phrases like "It was great connecting with you." It's technically a follow-up email. It has nothing to do with Sarah's actual situation, the tone of the meeting, what was discussed, or what the next step should be.

She then provides: who the meeting was with, what was discussed, what her specific goal for the follow-up is, and how she wants to come across. The email AI produces next is one she could send without changes.

Same AI. Same model. The variable was what she gave it.

WHAT ACTUALLY CHANGES GOOD AI OUTPUT

Five inputs move the needle most consistently. They appear in greater detail throughout this book, but in brief:

- 1. A clear job:** What specifically do you want AI to do? (Not "help with" — do.)
- 2. Relevant context:** What does AI need to know about your situation, audience, and constraints?

- 3. Source material:** What documents, notes, or data should AI work from?
- 4. Output specifications:** What should the result look like — length, format, tone, structure?
- 5. Verification instructions:** How should AI flag uncertainty or potential errors?

This is not a complete list. The next chapter introduces the full model. But if you apply only these five, you will notice an immediate and significant improvement.

TRY THIS

Take a prompt you have used before that produced a disappointing result. Write it down.

Now ask yourself: What did I assume AI would know that it couldn't have known? What context did I leave out? What does "good" actually look like for this task?

Rewrite the prompt with those gaps filled in and run both versions. Notice what changed.

WHAT THIS CHAPTER ESTABLISHED

Generic inputs produce generic outputs. This is a feature, not a bug — it means every improvement you make to your input produces a corresponding improvement in output. You are not at the mercy of the AI. You are in control of the most important variable in every interaction.

The next chapter introduces a model for understanding exactly which variables are in play — and which ones matter most for any given task.

WATCH OUT

More detail is not always better. Adding irrelevant information doesn't improve output — it can dilute it. The skill is not writing longer prompts; it's identifying what AI actually needs to know to do this specific task well. You'll develop this judgment through the exercises in Part II.

CHAPTER 2

The AI Interaction Stack — What Actually Determines Output Quality

INTRODUCING A MAP

When a piece of work goes wrong, it helps to know where in the process the problem occurred. A cake that doesn't rise failed at a different point than a cake that doesn't taste right. The diagnosis matters.

AI interactions fail at predictable places. This chapter introduces a model — the AI Interaction Stack — that maps those places. It gives you a vocabulary for diagnosing what went wrong and a checklist for getting it right before you start.

THE AI INTERACTION STACK

The AI Interaction Stack has eleven layers. Each layer represents something that shapes what AI produces. When an interaction produces poor results, at least one of these layers is inadequate or missing.

You do not need all eleven layers for every task. A quick factual question might only need layers 1-4. A complex document workflow might need all eleven. The point is knowing which layers exist so you can deliberately choose which ones to activate.

Here they are, in order:

THE AI INTERACTION STACK

11	Iteration	What happens next?
10	Verification	How will you check it?
9	Output	What should the result look like?
8	Workflow	One step or several?
7	Tools	What capabilities are available?
6	Workspace	What context can persist?
5	Source material	What can AI work from?
4	Examples	What does good look like?
3	Instructions	What exactly should AI do?
2	Context	What does AI need to know?
1	Intent	What are you trying to achieve?

Read the stack from the bottom up. Start with the real goal, then add only the layers that this task needs. The upper layers help you check, improve, and continue the work; they do not replace the clarity of the layers below them.

LAYER 1 — INTENT

What are you actually trying to achieve? Not the surface task ("write a summary") but the underlying goal ("help my CEO decide quickly"). These are often different, and the difference matters.

A summary intended to help someone decide looks different from a summary intended to brief someone who already decided. Same task, different intent, different output.

Before you start: Can you state your actual goal in one sentence?

LAYER 2 — CONTEXT

What does AI need to know about your situation? Who is involved? What is the background? What constraints apply? What has already been tried?

Context is the layer most people consistently underinvest in. It is also the layer where providing more almost always pays off.

Before you start: What does AI not know that would change its answer?

LAYER 3 — INSTRUCTIONS

What exactly should AI do? This is distinct from intent. Intent is the goal; instructions are the specific action you want AI to take to achieve it.

"Summarise this document" is an instruction. "Summarise this document in three bullet points, prioritising financial implications, written for a non-specialist reader" is a better one.

Before you start: How specific are your instructions? Could a capable human follow them without asking a question?

LAYER 4 — EXAMPLES

What does good look like? Sometimes the most efficient way to specify a tone, structure, or style is to show AI an example of it — a previous email you wrote, a format you liked, a sample of the style you want.

Examples do more than description. Telling AI to be "warm but professional" is imprecise. Showing it a warm-but-professional email is not.

Before you start: Is there an example that would make the target output clearer?

LAYER 5 — SOURCE MATERIAL

What documents, data, notes, or materials should AI work from? AI can read files — contracts, reports, research papers, meeting notes, spreadsheets — and its output will be grounded in what

you provide rather than in generalisation.

When you ask AI to help with a specific document or dataset, provide it. Don't describe it — provide it.

Before you start: What source materials could improve the specificity and accuracy of the output?

LAYER 6 — WORKSPACE

Have you set up a persistent context — a project, a set of standing instructions, a shared background — so you don't have to repeat essential information every time?

For tasks you do repeatedly, a workspace eliminates repetition and builds consistency. Part III is dedicated to this layer.

Before you start: Is this a recurring task where a workspace would save significant setup time?

LAYER 7 — TOOLS

What capabilities does AI have access to in this context? Can it search the web? Run calculations? Read uploaded files? Generate images? Access calendar or email?

Knowing which tools are available — and which aren't — shapes what you can reasonably ask for and where you need to verify externally.

Before you start: What tools are available? Does this task benefit from any of them?

LAYER 8 — WORKFLOW

Is this a single-step task or a multi-step process? Complex work almost always benefits from decomposition: breaking the task into stages, running each stage separately, reviewing and adjusting between stages.

Single-prompt attempts at complex tasks produce mediocre results that are hard to improve. Staged workflows produce better outputs at each stage and make errors easier to catch.

Before you start: Is this a single-step task, or should it be broken into stages?

LAYER 9 — OUTPUT

What should the result look like? Format, length, structure, tone, audience, medium — all of these shape the output and all are specifiable.

Without output specifications, AI defaults to whatever format its training suggests is most common for this type of request. That default is often wrong for your specific situation.

Before you start: Have you specified format, length, tone, and structure?

LAYER 10 — VERIFICATION

How will you check the output? What specifically might be wrong? AI makes errors — factual, logical, and interpretive — with calm confidence. Your verification approach should match the stakes of the task and the risk profile of the content.

Telling AI to flag its own uncertainties is a useful first step. It doesn't replace your review, but it helps.

Before you start: How will you verify this output before you use or share it?

LAYER 11 — ITERATION

How will you use this output as the starting point for the next stage? AI interactions are rarely single exchanges — they are conversations where each response informs the next prompt.

Good output at layer 11 becomes input at layer 2 of the next pass.

Before you start: Is this a one-shot task, or will you need to iterate? What will the next step be?

WHY IT WORKS

The stack is not a checklist to complete in full before every interaction. It is a diagnostic framework. When output is disappointing, the stack tells you where to look. When you're designing a complex task, the stack tells you what to prepare. Most interactions need four or five layers well-activated; a few need all eleven.

THE STACK IN PRACTICE

Here is the same task at two different levels of stack activation:

LOW STACK (LAYERS 1, 3 ONLY):

"Summarise this meeting."

Result: A generic summary with equal weight given to all topics discussed. Possibly bulleted, possibly in paragraphs. Pitched at a generic reader. Takes five minutes to edit into something useful.

HIGH STACK (LAYERS 1-5, 9-10):

"I need a meeting summary for three people who weren't present: my manager, the client's account manager, and one of our developers. Each person needs to see different things. My manager wants the key decisions and next steps only. The client account manager wants what the client committed to. The developer wants the technical requirements we discussed. The attached transcript is the

source. Write three separate summaries — one for each person — no longer than half a page each. Flag anything that seemed unresolved or ambiguous."

Result: Three targeted summaries, immediately usable, with potential problems flagged.

The difference is not length of prompt. It is activation of the right layers for the task.

TRY THIS

Think of an AI task you need to do this week. Write down what layers of the stack you would normally activate — most people use layers 1, 3, and maybe 9 by default. Now consider which additional layers are relevant to this task. What would you add if you invested two more minutes in the setup?

WHICH LAYERS MATTER MOST?

The research on what makes AI interactions succeed points consistently to three layers as the highest-leverage investments for most people in most tasks:

Context (Layer 2) is the most consistently underinvested layer and the one with the highest return. Providing relevant context costs little and changes output significantly.

Instructions (Layer 3) are the most frequently underdeveloped layer. Most instructions are too vague. The skill of writing specific, actionable instructions is developed in Chapter 7.

Verification (Layer 10) is the most commonly skipped layer. Part VI covers it in depth. Skipping it is also the source of most AI-related professional embarrassment.

WATCH OUT

The stack is a mental model, not a procedure. It does not mean you should spend twenty minutes preparing every AI prompt. A quick factual question needs layers 1 and 3, briefly addressed, and that's correct. The stack is useful when you are planning a complex task or diagnosing a poor result — not as mandatory overhead for every interaction.

CHAPTER 3

The 7 Questions — Your Portable AI Framework

A FRAMEWORK YOU CAN CARRY IN YOUR HEAD

The AI Interaction Stack gives you a complete map. But maps aren't always what you need. Sometimes you need a quick checklist — something you can run through in sixty seconds before an important interaction.

The 7 Questions is that checklist.

It is derived from the stack but simplified into seven questions that cover the most important dimensions of any significant AI interaction. It is short enough to memorise. It is structured enough to catch the most common input failures. And it is abstract enough to work on any AI platform, with any model, now or in the future.

THE 7 QUESTIONS

Before any significant AI interaction:

1. WHAT AM I TRYING TO ACHIEVE?

State the underlying goal, not just the surface task. "I want my CEO to approve the budget" is more useful than "I want a summary." The goal shapes every other decision.

2. WHAT DOES AI NEED TO KNOW?

What background information, constraints, relationships, or context is essential to getting a good output? What would a capable human assistant need to know before starting this task?

3. WHAT INFORMATION CAN I PROVIDE?

What documents, data, notes, examples, or materials exist that would make AI's output more specific and more accurate? What can you attach or paste in?

4. WHAT EXACTLY SHOULD AI DO?

What is the specific action? Not "help with" — do. Write, summarise, extract, compare, analyse, rewrite, draft, plan. Be precise about the verb and the scope.

5. WHAT SHOULD THE RESULT LOOK LIKE?

Format, length, tone, structure, audience, medium. If you do not specify, AI will choose. Specify.

6. HOW WILL I CHECK IT?

What might be wrong? What specifically needs verification before you use or share this output? Are there numbers, dates, names, claims, or recommendations that require checking?

7. WHAT SHOULD HAPPEN NEXT?

Is this one step in a larger process? What will you do with this output? Will it become input to another AI prompt, a document you share, or a decision you make?

WHY IT WORKS

Each question addresses a different failure mode. Question 1 catches the wrong goal. Question 2 catches the missing context. Question 3 catches unused resources. Question 4 catches the vague instruction. Question 5 catches the wrong format. Question 6 catches the unverified output. Question 7 catches the single-step tunnel vision.

You can have a perfect answer to six questions and fail on the seventh. The framework is complete by design.

USING THE 7 QUESTIONS: A WORKED EXAMPLE

Sarah needs to create a briefing document on a potential new supplier for her director.

Let's run through the questions:

Q1 — WHAT AM I TRYING TO ACHIEVE?

My director needs to decide whether to invite this supplier to a formal tender. I need to give her the information required to make that decision, not a complete company profile.

Q2 — WHAT DOES AI NEED TO KNOW?

The decision context (tender shortlist), the supplier's name (TechFlow Ltd), what the tender is for (CRM software implementation), our current system (Salesforce), our rough budget (£80,000), and what factors matter most to us (integration capability, support quality, UK presence).

Q3 — WHAT INFORMATION CAN I PROVIDE?

I have the supplier's website, a brochure they sent us, and notes from a call I took with their sales contact last week.

Q4 — WHAT EXACTLY SHOULD AI DO?

Create a supplier briefing document using the materials I provide. Organised under the headings: Company Overview, Relevant Experience, Key Capabilities, Potential Risks, Recommended Questions to Ask.

Q5 — WHAT SHOULD THE RESULT LOOK LIKE?

One page, maximum. Professional but plain language. My director doesn't have time for anything longer. Bullet points under each heading.

Q6 — HOW WILL I CHECK IT?

I'll verify any specific facts against the brochure and my call notes. I'll flag for my director any claims from the supplier that couldn't be independently verified.

Q7 — WHAT SHOULD HAPPEN NEXT?

This document will go to my director for her review. If she decides to proceed, I'll use AI to help draft the invitation to tender based on the same materials.

Now Sarah turns this into a prompt:

"I need a one-page supplier briefing document for my director, who is deciding whether to include TechFlow Ltd in a CRM software tender shortlist. The tender is for an implementation project with a rough budget of £80,000. Our current system is Salesforce.

*

The document should be organised under these five headings: Company Overview, Relevant Experience, Key Capabilities, Potential Risks, Recommended Questions to Ask.

*

Please work from the materials I am attaching: [brochure], [call notes]. Plain language, bullet points, maximum one page. Flag any claims in the supplier's materials that you cannot independently verify from what I've provided."

This is a complete interaction. It took about three minutes to prepare using the 7 Questions. It would have taken considerably longer to produce a good result any other way.

REAL-WORLD EXAMPLE

James is a recent graduate preparing for a job interview. He runs through the 7 Questions:

Q1 — I want to walk into this interview well-prepared, able to give specific answers about why I want this role.

Q2 — The company is Meridian Digital, a mid-size marketing agency. The role is Junior Account Manager. My background is a marketing degree and six months of internship experience. The interview is in two days.

Q3 — I have the job description, the company's website, their LinkedIn page, and a Glassdoor page with recent employee reviews.

Q4 — I want AI to help me prepare for likely interview questions, draft answers to the hardest ones, and identify things about the company I should know.

Q5 — I want a structured prep document: company background summary, ten likely questions with draft answers, and a section on questions I should ask them.

Q6 — I'll verify any company facts against their current website before the interview.

Q7 — I'll use the document to do a practice run the night before.

The resulting prompt takes forty seconds to write. The output he gets is interview prep he can actually use, not generic "tell me about a time you overcame a challenge" boilerplate.

THE 7 QUESTIONS AS A PORTABLE CARD

The 7 Questions are reproduced as a reference card in Appendix A. You can photograph it, print it, or simply remember it. The goal is that it becomes automatic — something you run through mentally before any AI interaction that matters.

You do not need to answer every question in full every time. For a simple, low-stakes task, three or four seconds on each question is sufficient. For a complex, high-stakes task — a document going to a senior stakeholder, a recommendation with significant consequences, a piece of work that will be published — the full seven questions deserve a full minute each.

PRO TIP

Once you have a well-developed answer to the 7 Questions for a recurring task — a weekly report, a client email format, a type of document you produce regularly — save it as a template. The 7 Questions become a workspace setup (Chapter 11). Run once; reuse indefinitely.

WATCH OUT

The 7 Questions are a preparation tool, not a rigid protocol. If you find yourself spending ten minutes answering the questions before a thirty-second task, you are misapplying the framework. Use judgment about scale. The questions exist to catch the gaps that consistently cause problems — not to add overhead to interactions that don't need it.

CHAPTER 4

How to Talk to an AI (It's Not Google)

THE SEARCH ENGINE HABIT

Most people's first instinct when using AI is to type a search query. Short. Keyword-heavy. Punchy.

"best email subject lines"

"how to negotiate salary"

"presentation tips"

This works for Google because Google is retrieving existing documents that match your keywords. The shorter and cleaner the query, the more precisely it matches what's out there.

AI tools are not search engines. They are not retrieving existing content. They are generating new content based on what you've given them. The search-engine reflex is the single most common reason people get worse results than they could.

WHAT AI ACTUALLY DOES

Without requiring a technical explanation: an AI language model generates text by working out what words are most likely to follow what came before, based on patterns learned from enormous quantities of text. It is not looking up answers. It is constructing them.

This has three practical implications:

1. MORE RELEVANT INPUT, MORE RELEVANT OUTPUT.

A longer, more specific prompt does not confuse AI. It constrains the range of possible responses toward what you actually need. Specificity is an advantage.

2. A FRESH CONVERSATION SHOULD NOT BE ASSUMED TO KNOW YOUR PREVIOUS WORK.

Unless the platform explicitly provides and uses a memory, project, history, or connected-source feature, a new conversation begins without dependable knowledge of who you are, what you've discussed before, or what your preferences are. This is why context must be provided explicitly — and why workspaces (Part III) matter so much for regular users.

3. AI GENERATES PLAUSIBLE TEXT, NOT NECESSARILY ACCURATE TEXT.

AI's training teaches it what good writing looks like in a given context. This means it can generate fluent, confident text on subjects it is wrong about. The confidence in the prose is not correlated with the accuracy of the content. This is the most important habit to build: treating AI output as a draft that requires verification, not a source that can be trusted.

WHY IT WORKS

Understanding the generation model — rather than thinking of AI as an oracle or a search engine — changes how you interact with it. It explains why providing examples works (they constrain the generated text toward the style you want). It explains why AI can be confidently wrong (plausibility and accuracy are different things). It explains why asking AI to explain its reasoning helps (it forces more systematic text generation). These are not quirks to work around; they are features of the model to use deliberately.

THE CONVERSATION, NOT THE QUERY

AI tools work best as conversations, not as single-query exchanges. The first response is rarely the last one.

This is not because AI needs multiple chances to get things right (though sometimes it does). It is because a conversation allows you to:

- Add context you forgot in the first prompt
- Adjust the output format based on what you see
- Request a deeper dive on one section
- Push back on something that seems wrong
- Ask AI to regenerate with a different constraint

Think of it as briefing a capable but uninformed colleague. The first briefing gets you a first draft. The subsequent conversation refines it. This is normal, efficient, and often faster than trying to write the perfect prompt in one go.

THE ANATOMY OF AN EFFECTIVE PROMPT

Not every prompt needs all of these elements. But a complete prompt for a significant task includes:

ROLE (optional but often useful):

Act as a [role] with expertise in [domain].

CONTEXT:

Background information AI needs to know.

Who is involved, what the situation is, what constraints apply.

TASK:

Exactly what you want AI to do.

Use a clear action verb. Be specific.

SOURCE MATERIAL (when relevant):

[Attached document / pasted text / data]

OUTPUT REQUIREMENTS:

Format, length, tone, audience, structure.

VERIFICATION REQUEST (for factual or high-stakes output):

Please flag any claims you are uncertain about,
any assumptions you've made, and any information

I should verify independently.

This is not a template to apply rigidly. It is a map of what a complete prompt includes. For quick tasks, you might cover all of this in two sentences. For complex tasks, each element might be a paragraph.

TRY THIS

Take a standard query you might type into a search engine and rewrite it as an AI prompt using the anatomy above. You don't need all six elements — pick the three most relevant.

For example:

Search query: "how to negotiate salary"

AI prompt: "I am about to have a salary negotiation conversation with my manager. I have been in this role for two years, my performance reviews have both been 'exceeds expectations', and I know from a conversation with a colleague that my salary is about 12% below the market rate for this role in London. I am asking for a 15% increase. Help me prepare for this conversation: what arguments are strongest, what objections might I face, and how should I handle the moment when they counter with a lower number? Tone: calm and factual, not aggressive."

Notice the difference in what AI has to work with.

ASKING AI TO THINK, NOT JUST WRITE

One of the most underused techniques is asking AI to reason through a problem before producing an answer.

When AI generates text quickly, it sometimes takes the most obvious route. Asking AI to think first — to break down the problem, identify considerations, or outline before drafting — tends to produce more thoughtful, more accurate results.

Phrases that invoke this:

- *"Before answering, think through the main considerations."*
- *"First give me an outline, then draft the full document."*

- *"Walk me through your reasoning."*
- *"What are the main risks or failure points here?"*

This works because it changes the pattern of text generation — the model is working through intermediate steps rather than jumping to a conclusion. You can think of it as asking someone to show their work rather than just hand you an answer.

PRO TIP

When you receive an output that is close to what you want but not quite right, do not start again from scratch. Respond to the output with a specific refinement instruction: "That's good but the tone is too formal — can you make it sound more conversational?" or "The summary is accurate but too long — cut it to three bullet points and keep only the most important point in each." Iteration on an existing draft is almost always faster than rewriting.

WATCH OUT

AI confidently produces wrong answers. The fluency of AI prose — the fact that it reads well, sounds authoritative, and doesn't hedge — can make it easy to accept without checking. Any output that contains specific facts, numbers, names, dates, quotes, citations, or recommendations deserves verification before you rely on it. Chapter 20 covers verification in detail. For now, adopt the rule: polished prose is not evidence of accuracy.

WHAT PART I HAS ESTABLISHED

You now have the foundation:

- Generic inputs produce generic outputs because AI fills gaps with averages
- The AI Interaction Stack identifies the eleven layers that determine output quality
- The 7 Questions give you a portable framework for activating the right layers
- A basic AI chat is a generation tool, not a retrieval tool; search and connected-source tools change what it can access — and conversation is the natural mode of interaction

Part II takes each element of a strong interaction and develops it into a skill.

Build Better Interactions

The five chapters in Part II develop the core skills that apply to every AI interaction you will ever have. Each chapter takes one element of the AI Interaction Stack and teaches you to use it deliberately. Together they account for the majority of the quality difference between people who get good results from AI and people who don't.

CHAPTER 5

Give AI a Job, Not Just a Question

THE QUESTION VERSUS THE JOB

There is a meaningful difference between asking AI a question and giving AI a job.

A question invites a response. A job assigns a task with a clear deliverable, a scope, and implicit quality criteria.

Question: "What do I need to know about redundancy letters?"

Job: "You are an HR communications specialist. Draft a redundancy letter for a member of our sales team who has been with the company for three years. The reason is restructuring, not performance. Tone: respectful and clear. The letter must include: confirmation of the decision, notice period (one month), details of the settlement package (to be inserted), available support resources, and the process for appeal. UK employment law applies."

The question gets you information about redundancy letters — a textbook answer pitched at an anonymous reader. The job gets you a document you can actually edit and use.

WHY IT WORKS

Role assignment and job framing do not change what AI "knows." They change the frame through which AI selects and structures what it produces. A prompt that says "you are an experienced financial analyst" causes AI to draw more heavily from the patterns associated with financial analysis — more technical vocabulary, more rigorous structure, more appropriate caution. A prompt that specifies a deliverable (a letter, a report, an agenda) constrains the format. Together, role and job framing channel AI's generation toward something more precisely useful.

THE BEFORE → BETTER → BEST: GIVING AI A JOB

BEFORE (the vague question)

"What should I include in a client proposal?"

What AI produces: A generic list of proposal sections that any business book might contain. Executive summary, scope, timeline, pricing, terms. Correct. Applicable to no one's specific situation.

Why it's weak: AI has no information about the client, the industry, the type of work, the relationship, or what distinguishes this proposal from a hundred others.

BETTER (the specific task)

"Write an outline for a proposal for a digital marketing retainer, targeting a mid-size UK retail company. The agency specialises in SEO and paid social. The prospect has expressed interest but is also talking to two other agencies. The proposal needs to address: our methodology, relevant case studies, pricing structure (monthly retainer, tiered options), and what makes us different."

What AI produces: A structured outline specific to this type of proposal with relevant sections and brief notes on what each should contain.

What improved: AI now has the industry, the type of work, the competitive context, and the key elements to address.

BEST (role + job + context + output specifications)

"You are a senior account director at a boutique UK digital marketing agency. You are writing a proposal for a monthly SEO and paid social retainer for Thornton's, a mid-size UK garden centre chain with 12 locations. The prospect is also considering two larger agencies. The decision maker is the Marketing Director, who has told us she values transparency and long-term partnerships over flashy pitches."

Write a proposal outline that: leads with their specific situation rather than our credentials; demonstrates understanding of the garden centre sector and its seasonal patterns; presents three pricing tiers (Starter / Growth / Full Service); and ends with a section on how we measure success and report to clients.

Structure: section headings with one-sentence notes on content. Not the full proposal — the outline I will then develop. Max two pages when written up."

What AI produces: A highly specific outline that could only work for this client, this agency, and this competitive context. The structure respects the decision maker's stated values. The pricing presentation is appropriately tiered. The language is right for the relationship.

WHAT CHANGED?

WHAT WAS ADDED	WHY IT MATTERED
Role: senior account director	Contextualised the voice and level of expertise
Specific client name and context	Anchored content in a real situation
Decision maker's stated values	Shaped the strategic angle of the proposal
Sector knowledge requirement	Prompted sector-relevant thinking
Clear deliverable	"Outline I will develop" prevented over-generation
Output constraints	Kept length appropriate

CHOOSING THE RIGHT ROLE

Role assignment is optional but often valuable. The rule of thumb: assign a role when the task has a domain, a voice, or a level of expertise that you want AI to reflect.

Useful role frames:

- Professional identity: *"You are an experienced employment solicitor"*
- Expertise level: *"You are a secondary school science teacher"*
- Perspective: *"You are a sceptical reader encountering this argument for the first time"*
- Function: *"You are a copy editor reviewing this draft"*

Role frames that don't help much: overly elaborate backstories, fictional identities with no relevance to the task, vague superlatives ("you are the world's greatest expert"). Keep roles grounded in what kind of expertise or perspective would actually be useful for this task.

ASSIGNING A SPECIFIC, ACTIONABLE JOB

The instruction "help me with X" is not a job. It is an open invitation that AI will answer with whatever seems most plausibly helpful — which is usually the most generic interpretation of the task.

A job has:

- A clear action verb: write, summarise, extract, compare, identify, structure, rewrite, analyse
- A clear deliverable: a document, a list, a table, a set of questions, a draft
- A clear scope: this document, these notes, this specific situation

Compare:

- *"Help me with my CV"* → AI will ask what kind of help or produce something generic
 - *"Rewrite the experience section of my CV to emphasise project management and stakeholder communication, using active verbs, targeting senior operations manager roles in the financial services sector"* → AI has a job
-

TRY THIS

Write down a task you regularly do that takes more than an hour. Now describe it as a job for AI using this structure:

[Role, if relevant]. [Action verb] [specific deliverable] using [source material]. The output should [format/length/tone specifications]. The intended reader is [audience]. Do not [explicit exclusion if needed].

Notice how much this forces you to clarify what you actually want before AI starts.

WATCH OUT

Assigning an expert role does not make AI's output more accurate. It may make it more fluently confident — which, as noted in Chapter 4, can be more convincing and equally wrong. Role framing improves format, voice, and structure. It does not improve factual accuracy. Verify subject-matter claims regardless of the role you assigned.

CHAPTER 6

Context Is Everything

THE CONTEXT GAP

Here is the practical starting point for most AI interactions: unless you provide relevant context or explicitly enable a persistence feature, AI should be treated as knowing nothing dependable about you, your work, your organisation, your history with this client, your professional constraints, or the specific situation you are dealing with.

You know all of these things. They seem obvious. So you don't mention them.

This is the context gap — the difference between what AI needs to know to give you a genuinely useful answer, and what you actually provide. It is the single most valuable gap to close.

WHAT COUNTS AS CONTEXT

Context is not just background information. It is any fact about your situation that would change how a knowledgeable person would approach the task. It includes:

Situational context — What is happening? What is the history? Why does this matter now?

Audience context — Who is going to read, use, or act on this output? What do they know? What do they need? What are their constraints?

Relationship context — What is the relationship between parties? Long-standing client? New contact? Internal colleague? Adversarial negotiation?

Constraint context — What must the output not include? What word count, format, or style requirements apply? What legal, regulatory, or policy constraints apply?

Resource context — What documents, data, or previous work is relevant? What has already been tried?

Goal context — What does success look like? What decision will this output inform? What action will it prompt?

REAL-WORLD EXAMPLE

David runs a small consulting firm. He is writing a proposal for a client renewal and asks AI to help draft the pricing section.

Without context: AI produces a generic template for presenting consulting fees — day rates, retainer structures, and payment terms that could apply to any consulting firm anywhere.

With context: David tells AI that this client has been with the firm for four years, has three times mentioned "feeling nickel-and-dimed" by itemised billing, is a family business where the owner and CFO are the same person, and that David wants to move to a monthly retainer to create predictable revenue on both sides. He also notes the client's previous annual spend was around £42,000.

The output AI produces with this context is a specifically targeted pricing section that addresses the client's stated frustration, proposes a retainer at a price point consistent with their history, and frames the retainer in terms of what the client gets (predictability, priority access, no surprises) rather than what David gets. It is a usable draft. The generic version was not.

The difference was entirely in the context David provided.

THE CONTEXT TEST

A useful diagnostic: before sending a prompt, ask yourself:

"If I had just hired a very good human assistant who was new to this job, what would they need to know before I asked them to do this?"

That question surfaces most of the context you are probably omitting. A good assistant would ask about the audience, the history, the constraints, the goal, and the format before starting work. Providing that context in your prompt removes the need for them to ask — and removes the iterations required when AI guesses wrong.

THE BEFORE → BETTER → BEST: CONTEXT

BEFORE

"Write an email to our supplier about a late delivery."

Output: A generic, somewhat formal complaint email about a late delivery with placeholder information. Polite but weightless.

BETTER

"Write an email to our supplier, BuildRight Ltd, about a delivery that is now 12 days late. The order was for 500 units of packaging material needed for a Christmas promotional campaign. The delay means we cannot start production on time and we risk missing the campaign window."

Output: A specific email that references the order type, the operational impact, and the time pressure. Usable with minor edits.

BEST

"Write an email to BuildRight Ltd, our packaging supplier of five years. Our contact is James Thornton in their sales team. The delivery of 500 units of promotional packaging (order number BRTX-2241) is now 12 days late. This delivery was critical for our Christmas campaign — production was due to start on Monday and we have an external print deadline of 18th November.

We do not want to damage the relationship, which has been excellent until this point. The tone should be: firm and clear about the business impact, but not accusatory. We want two things: a confirmed delivery date in writing and an explanation of what happened.

Do not threaten to cancel the order or use legal language at this stage."

Output: An email that reflects the history of the relationship, is specific about the business impact and timeline, makes two clear requests, and observes the explicit constraint about not escalating. Might be sent with no editing.

WHAT CHANGED?

CONTEXT ADDED	WHY IT MATTERED
Five-year relationship	Shaped the non-accusatory tone
Named contact	Made the email personal
Order number	Made the reference specific and professional
Production deadline	Explained why timing matters so much
Print deadline date	Created concrete urgency without exaggeration
Two specific requests	Focused the email on clear outcomes
Explicit constraint (no legal language)	Prevented escalation AI might otherwise include

PROVIDING CONSTRAINT CONTEXT

Constraints are particularly easy to omit because they feel obvious to you. "Obviously I can't mention the legal proceedings." "Obviously I shouldn't mention the budget number." "Obviously the tone needs to be formal."

These are obvious to you. They are invisible to AI. State every constraint that would make the output unusable if violated.

Useful constraint phrasing:

- *"Do not include [X]."*
- *"This cannot mention [X] for legal reasons."*
- *"The tone must not suggest [X] — this would damage the relationship."*
- *"UK spelling and law applies throughout."*
- *"This will be reviewed by [X] so should not contain [Y]."*

PRO TIP

For recurring tasks where the context is always the same, you don't need to retype it every time. Chapter 11 covers persistent instructions — a way of establishing your standard context once so it applies automatically to every conversation. The context you would type for every client email, every report, every team communication becomes a standing brief that AI carries forward.

WATCH OUT

Providing irrelevant context can dilute the output as much as providing too little. If you paste a twenty-page document when AI only needs three paragraphs, you risk getting a response that is unfocused across all twenty pages. Context should be relevant to this task. If in doubt, introduce the document and tell AI which section matters: *"I'm attaching our full company handbook, but for this task only the redundancy policy section (pages 14-17) is relevant."*

CHAPTER 7

Instructions That Actually Work

WHY MOST INSTRUCTIONS ARE TOO VAGUE

Instructions are where most users put most of their effort — and still get disappointing results. The reason is usually not that the instructions are wrong. It is that they are underspecified.

"Make it more professional" is an instruction. But "more professional" in what dimension? More formal vocabulary? Shorter sentences? No contractions? Different paragraph structure? Fewer qualifications? More precise language?

AI will choose. And it will choose the version of "more professional" that appears most frequently in its training data for this type of text — which may not be what you meant at all.

Effective instructions are specific enough that a capable human could follow them without asking a clarifying question.

THE SPECIFICITY TEST

Before sending an instruction, ask: *"Could two different people interpret this differently?"*

If yes, it is underspecified.

"Make it shorter" → How much shorter? Half? A third? A specific word count?

"Make it more engaging" → More engaging how? More conversational? Shorter paragraphs? An opening question? Personal examples?

"Make it suitable for our audience" → Who is the audience and what do they need?

"Improve the structure" → More headings? Different order? Remove sections? Add an introduction?

Each of these pairs shows an instruction that fails the specificity test, and the questions that reveal what was actually meant. Answer those questions in the instruction.

POSITIVE AND NEGATIVE INSTRUCTIONS

Instructions can specify what to do (positive) or what not to do (negative). Both are important. Most people use only positive instructions and then wonder why the output contains things they didn't want.

Negative instructions are particularly useful for:

- Tone and register: *"Do not use exclamation marks. Do not use filler phrases like 'I hope this finds you well.'"*
- Content constraints: *"Do not mention the merger — this is not public."*
- Style: *"Do not use bullet points — this should be flowing prose."*
- Claims: *"Do not make specific claims about delivery times — these vary by region."*

A well-instructed prompt often includes both: *"Write this in a direct, confident tone. Do not use passive voice. Do not hedge claims with phrases like 'it could be argued' or 'some might say.'"*

INSTRUCTION TYPES AND WHEN TO USE THEM

Different tasks benefit from different instruction types:

Format instructions — What structure should the output take? *"Use a numbered list." / "Present this as a table with columns for X, Y, Z." / "Use three short paragraphs, no headings." / "One page maximum, A4, with a heading and four bullet points."*

Tone instructions — How should it sound?

"Direct and professional." / "Warm but concise." / "Technical but accessible — assume the reader has a science degree but is not a specialist in this field." / "In the voice of the existing document I'm attaching."

Priority instructions — What matters most?

"Lead with the cost saving, not the process improvement." / "The most important point must be in the first sentence." / "The client's concern about timeline is more important than the features list — address it first."

Exclusion instructions — What to leave out?

"Do not include implementation details — this is for a non-technical audience." / "Exclude anything older than 2024 — the reader will already know the history." / "Do not include a summary — the reader will read the whole thing."

Role instructions — From whose perspective?

"Write this as if the reader has never heard of our company." / "The reader is a sceptic — address their likely objections before they ask."

THE BEFORE → BETTER → BEST: INSTRUCTIONS

BEFORE

"Rewrite this paragraph to make it better."

[Paragraph: "Our team has extensive experience in the delivery of complex technology projects. We have worked with a wide range of clients across multiple sectors and have a proven track record of delivering projects on time and on budget."]

Output: AI produces a slightly more polished version of the same paragraph. Same ideas, slightly varied vocabulary. The reader would not notice a meaningful difference.

BETTER

"Rewrite this paragraph to remove clichés and make it more specific and credible. Replace phrases like 'extensive experience' and 'proven track record' with concrete language."

Output: AI produces a version that avoids the worst clichés and becomes more specific in tone, though it invents specifics it doesn't have.

BEST

"Rewrite this paragraph for inclusion in a proposal to a housing association considering a digital transformation project. The audience is a director who has seen many similar proposals and is sceptical of vague claims."

Instructions:

- Remove all clichés: 'extensive experience', 'wide range', 'proven track record', 'multiple sectors' must all go.*
- Replace vague claims with specific patterns of evidence. I will provide specifics in a moment — for now, leave [BRACKETS] as placeholders for: number of years in the sector, number of completed projects, one specific client sector, and one specific result we achieved.*
- First sentence should establish relevance to housing sector specifically, not generic technology experience.*
- Maximum two sentences. Punchy, not promotional."*

Output: A tightly constructed two-sentence paragraph that is specific in structure, leaves appropriate placeholders, and is pitched correctly at the sceptical reader.

WHAT CHANGED?

INSTRUCTION ELEMENT	WHY IT MATTERED
Audience specified	Shaped the pitch level and scepticism antidote
Banned phrases listed explicitly	Prevented the most common failure
Placeholder pattern	Kept AI honest about missing specifics
Structural requirement (first sentence)	Prevented the generic opener
Length constraint (two sentences)	Forced concision
Tone instruction (not promotional)	Prevented the default enthusiastic pitch

TRY THIS

Take the last AI output you were dissatisfied with. Identify one thing that was wrong with it — too long, wrong tone, too generic, incorrect emphasis, wrong format.

Write a follow-up instruction that specifically addresses that one thing, using the specificity test: could two people interpret this instruction differently?

If yes, refine it until the answer is no.

WATCH OUT

Instructions that conflict produce confused output. If you say "be brief and comprehensive," "be formal but conversational," or "be specific but applicable to all situations," AI will attempt to honour all instructions simultaneously and usually succeed at none of them. When instructions conflict, decide which takes priority and say so: *"Be brief — prioritise this over comprehensiveness if necessary."*

CHAPTER 8

Examples Change Everything

SHOWING VERSUS TELLING

There are two ways to specify what you want from AI. You can describe it, or you can demonstrate it.

Describing works for many things. But when tone, format, or style is what you care about, demonstrating is almost always more effective.

This is because language about language is inherently imprecise. "Warm but professional" means something different to you than it does to AI. "Clear and engaging" is interpreted through thousands of examples in AI's training data, most of which are not exactly what you meant. The gap between your description and AI's interpretation is where most style failures live.

An example collapses that gap. Instead of asking AI to match your description of a style, you are showing it the style and asking it to match that.

TYPES OF EXAMPLES

Examples can take several forms:

TONE EXAMPLES

A piece of your own writing, or a piece of writing in the voice you want. "Write in the same tone as this email I wrote last year" or "match the style of this example I'm pasting in."

FORMAT EXAMPLES

A previous document in the format you need. "Use the same structure as this report section." "Produce output that looks like this."

QUALITY STANDARD EXAMPLES

A high-quality output you want to match or improve on. "This is an example of the level of specificity I want."

NEGATIVE EXAMPLES

Something you don't want. "Do not write like this" is sometimes as effective as "write like this."

WHY IT WORKS

Providing examples is what AI researchers call "few-shot prompting" — giving AI a small number of examples from which to infer the desired pattern. The technical reason it works is that examples constrain the generation space toward outputs with similar characteristics: similar vocabulary choices, similar sentence length, similar structural decisions. Your description of "professional" is abstract; an example of it is concrete.

THE BEFORE → BETTER → BEST: EXAMPLES

The scenario: Margaret wants a weekly AI-generated summary of news relevant to her industry (sustainable packaging) to send to her team. She wants it to sound as if she wrote it, not AI.

BEFORE

"Write a summary of this week's sustainability news for my team."

Output: A formal, comprehensive news summary. Heavy paragraphs. AI-standard prose. Margaret would have to rewrite it entirely before it sounds like her.

BETTER

"Write a summary of this week's sustainability packaging news for my team. Tone: conversational, direct, not overly formal."

Output: Less formal, but still not quite right. "Conversational" is interpreted differently by different people, and AI doesn't know what Margaret sounds like.

BEST

"I write a weekly news summary for my team of 8. Here is an example of one from three months ago that captured the tone exactly right:

[EXAMPLE SUMMARY]

'Three things from this week:

- 1. The EU is tightening its position on recycled content claims — this matters for our German clients more than anywhere else. Worth watching.*
- 2. Paperboard prices stabilised in October after six months of increases. Our procurement team will already know this, but the sector analysis I've linked to is good background.*
- 3. Maersk announced a new bio-methanol shipping route. Too early to know if this is real progress or PR. I'll share more when I know more.'*

Write this week's summary in the same format and voice. Source material I'm attaching covers: [pasted or uploaded articles]. Keep it to three items unless something genuinely

warrants a fourth."

Output: A summary that reads like it could have been written by Margaret — same three-item structure, same direct tone, same characteristic qualification ("too early to know"). She edits two words and sends it.

WHAT CHANGED?

WHAT WAS ADDED	WHY IT MATTERED
Example of actual past output	Gave AI a concrete target instead of abstract description
Format from example	Three numbered items, short, no preamble
Voice characteristics from example	Opinionated, direct, honest about uncertainty
Source material provided	Grounded in real content rather than general knowledge
Constraint on length	Prevented overgeneration

USING YOUR OWN WRITING AS AN EXAMPLE

One of the most practical applications of example-giving is using your own previous writing to teach AI your voice. This is particularly useful for:

- Regular communications: weekly updates, team emails, client newsletters
- Documents with a house style: reports, proposals in a format you've established
- Personal correspondence: emails where you want AI to draft and you want it to sound like you

The method: provide two or three examples of your writing in the relevant genre, and ask AI to match the style. Longer examples give AI more to work from. If you have a document that perfectly

captures your voice or style, use that as the primary reference.

PRO TIP

If you save a "style reference" document in your AI workspace (Chapter 11), you don't need to paste examples into every prompt. You reference the document once in your standing instructions and it's available for every subsequent interaction.

WHEN NOT TO USE EXAMPLES

Examples constrain AI toward the example. This is usually what you want — but not always.

If you want AI to be creative, novel, or to produce something better than your existing examples, providing examples can limit the output. In those cases, describe the qualities you want and explicitly say: *"Feel free to go beyond this starting point."*

Similarly: if the examples you have are not actually that good, AI may faithfully replicate their weaknesses. Use examples you would genuinely like to match.

WATCH OUT

Examples are most effective for style and format. They are less effective for accuracy. Providing an example of a data table you want does not teach AI to fill it with correct data. You are specifying format, not content. Verify content as normal regardless of how good the format looks.

CHAPTER 9

Control the Output

THE OUTPUT PROBLEM

You have provided excellent context, a clear job description, specific instructions, and a strong example. AI produces a response.

It is too long.

It is the wrong format.

The most important point is buried in paragraph four.

It uses a slightly different structure than you specified.

This happens because output defaults are baked into AI's training. Without explicit specifications, AI produces the type, length, and format of output most commonly associated with your type of request. That default is often close but rarely exactly right.

Controlling output is the final step in building an effective interaction — specifying not just what to produce, but what the result should look like when it arrives.

THE FIVE OUTPUT DIMENSIONS

Five dimensions of output are almost always worth specifying for significant tasks:

1. LENGTH

The most commonly underspecified dimension. How long? Not "short" — how many words, sentences, bullet points, or paragraphs? "A one-page briefing," "a three-bullet summary," "a maximum of 150 words," "five sentences" — these are unambiguous. "Brief" is not.

2. FORMAT

How should the information be structured? Prose paragraphs, numbered list, bullet points, table, structured document with headings, slide deck content, script, dialogue? Different formats serve different purposes. The format should match how the reader will use the output.

3. STRUCTURE

If the output has sections or parts, what are they and in what order? Specifying structure prevents AI from inventing a structure that may not serve your purpose. "The document should have four sections: Background, Key Findings, Recommendation, Next Steps" is better than allowing AI to choose.

4. TONE

How should it sound? The vocabulary level, the degree of formality, the use of first person or third person, the amount of hedging or qualification, the presence or absence of technical language — all are specifiable. Use examples when description alone is insufficient (Chapter 8).

5. AUDIENCE

Who will read this? What do they know? What do they need? The same information presented to a technical expert and a first-time reader should be substantially different. Specifying audience tells AI how to calibrate complexity, assumed knowledge, and depth of explanation.

WHY IT WORKS

AI has no inherent preference for one format over another. It defaults to the format most common in its training data for a given type of request. Explicit output specifications don't constrain AI in ways that affect quality — they redirect its generation toward the format you need. Specifying output is not about limiting AI; it is about using what AI has already generated to its best effect for your purpose.

THE BEFORE → BETTER → BEST: OUTPUT CONTROL

The scenario: James needs a summary of a forty-page government report on graduate employment for a blog post he is writing.

BEFORE

"Summarise this report on graduate employment."

Output: Six paragraphs, dense, comprehensive, presenting findings in the order they appear in the report. Difficult to read. Not formatted for a blog audience. James would have to restructure it entirely.

BETTER

"Summarise the key findings of this report for a general audience interested in graduate careers. Keep it short and clear."

Output: Shorter, more readable, but still not structured for a blog post. "Short and clear" was under-specified.

BEST

"Summarise the key findings of this graduate employment report for a blog post aimed at final-year students. They are interested in practical takeaways about their job prospects, not policy analysis.

Format: Three sections with short headings:

- 1. What the data says (two to three bullet points with the most significant statistics)*
- 2. What this means for you (two to three practical implications in plain language)*
- 3. What to do about it (one or two concrete suggestions)*

Maximum 300 words total. Active voice. No jargon. Read level: a confident 20-year-old who reads the Guardian.

Do not include methodology, policy recommendations, or anything that doesn't directly affect a job-seeking student."

Output: A structured, audience-appropriate blog summary. The sections are clear, the language is right for the audience, and the length is correct. James reads it once, adds a paragraph of his own at the end, and posts it.

WHAT CHANGED?

OUTPUT SPECIFICATION	WHY IT MATTERED
Three named sections	Imposed a usable structure
Bullet count per section	Prevented overgeneration
Maximum 300 words	Made length unambiguous
"Active voice"	Improved readability
Reading level specified	Calibrated vocabulary and complexity
Explicit exclusion list	Prevented irrelevant policy content

SPECIFYING FORMAT: A PRACTICAL REFERENCE

When you need:

THIS	SAY
A concise summary	"Three to five bullet points, one sentence each"
A structured document	List the sections explicitly with brief descriptions
A table	"Present this as a table with columns for [X], [Y], [Z]"
Prose, not bullet points	"Write in flowing paragraphs — no bullet points"
An email	"Subject line + body, professional tone, under 150 words"
An outline	"Numbered outline with second-level sub-points"
A list	"Numbered list, [n] items, no more"
A one-pager	"One A4 page of content: heading, three sections"

CONTROLLING OUTPUT MID-CONVERSATION

Output control is not only an upfront specification. It is also a live tool in conversation.

When the output is close but not quite right, you can redirect it with a specific instruction:

- *"Good structure but too long — cut it by half."*
- *"Remove the executive summary section — I don't need it."*
- *"This is in bullet points — rewrite as prose paragraphs."*
- *"The third point is the most important — lead with it."*
- *"Change the tone — this reads too formally for the audience."*

Each of these is a specific, actionable correction. AI will act on them precisely. The instruction "make it better" is not a specific correction. Be specific about the dimension you want changed.

TRY THIS

For your next significant AI task, write out your output specification separately before you write the main prompt. Answer these five questions:

- 1.** How long?
- 2.** What format?
- 3.** What structure?
- 4.** What tone?
- 5.** Who is the audience?

Then build those answers into the prompt. Notice whether the first output is closer to what you need before you have to ask for revisions.

WATCH OUT

Excessive output specification can overconstrain the response. If you specify every sentence, paragraph, and word choice, you are writing the document yourself and asking AI to type it. Specify the frame — length, format, structure, tone, audience — and leave AI to fill it in appropriately. The goal is a usable first draft, not a perfect document that required more work to specify than it would have taken to write.

WHAT PART II HAS ESTABLISHED

You can now build an effective AI interaction from the ground up:

- Give AI a clear role and a specific job, not a vague question
- Provide the context that closes the gap between what AI assumes and what you actually need
- Write instructions specific enough that a capable person could follow them without asking clarifying questions
- Show AI examples when style, format, or tone matter
- Specify the output in terms of length, format, structure, tone, and audience

Part III moves from individual interactions to the environment in which they happen — and the significant gains available from building that environment deliberately.

Build an AI Environment

Most people use AI as if it were a calculator: pick it up, use it once, put it down. Every new calculation starts from zero. Part III shows you how to change that — building a persistent environment where AI carries your context forward, knows your situation, and becomes genuinely more useful the more you invest in setting it up.

CHAPTER 10

Stop Starting From Zero Every Time

THE REPETITION PROBLEM

If you use AI more than occasionally, you have encountered this situation: you are about to start a new conversation and you find yourself typing, again, the things AI should already know.

"I work in marketing at a mid-size UK financial services company. I write for a non-specialist audience. I prefer a formal-but-accessible tone. UK spelling throughout. Never use exclamation marks..."

This is not how a useful working relationship operates. A good colleague doesn't need to be reminded of your name and job every time you speak to them.

The reason this happens is that a fresh AI conversation should not be assumed to have dependable knowledge of previous conversations. Some platforms now offer memory, projects, searchable history, or connected sources, but availability and behaviour vary. Without one of those features being explicitly enabled and used, the context you built in Tuesday's session is not reliably available on Wednesday.

This is one of AI's most significant practical limitations — and one that is largely solvable.

THREE STRATEGIES FOR NOT STARTING FROM ZERO

STRATEGY 1 — SYSTEM INSTRUCTIONS (PERSISTENT INSTRUCTIONS)

Most AI platforms now offer a way to set persistent instructions that apply to every conversation. Sometimes called "custom instructions," "system prompts," or "memory settings," these are a set of facts and preferences you provide once and that AI carries forward automatically.

This is the most powerful strategy for regular users. Chapter 11 covers it in depth.

STRATEGY 2 — PROJECTS

ChatGPT, Claude, and several other platforms offer a "Projects" feature that groups conversations, persists context, and maintains uploaded documents within a dedicated workspace. For recurring work — client accounts, ongoing projects, regular document types — projects eliminate most setup repetition.

STRATEGY 3 — CONTEXT DOCUMENTS

When platform features are unavailable or insufficient, a context document you paste in at the start of relevant conversations serves the same purpose. A single well-crafted paragraph covering who you are, what you're doing, and what your preferences are can be copied and pasted in seconds.

A WORKED FIRST SETUP

David runs a small consultancy and keeps starting client proposals from a blank chat. His first persistent-instructions draft is short:

I run a six-person UK consultancy. I use AI mainly for proposals, client communications, meeting preparation, and research briefings. Write in plain, professional UK English. Assume the reader is busy and commercially aware. Do not invent client facts, figures, case studies, or commitments. Flag anything that needs checking.

This does not try to describe every preference David has. It captures the context he repeats most often and the errors that would make an output unusable. After a week, he notices that AI still makes proposals too long, so he adds a default: lead with the decision the reader needs to make and keep first drafts under 800 words unless he asks for more. The instructions improve through use rather than through an attempt to predict every future task.

WHY IT WORKS

Persistent context doesn't improve AI's underlying capability. It changes what AI has to work with. An AI that knows your role, your audience, your style preferences, and your constraints produces better first drafts and needs fewer iterations. The investment is front-loaded — you write the context once — and the return compounds with every subsequent conversation.

WHAT BELONGING TO THE SAME WORKSPACE GIVES YOU

When you work within a project or persistent environment:

- Your uploaded documents remain available without re-uploading
- Your style and tone preferences apply without re-specifying
- Your role and context are understood without re-explaining
- Outputs from previous conversations can inform new ones
- Multiple AI conversations about the same work form a coherent body of output

This is a qualitative shift in how AI integrates into professional work. It moves AI from an occasional tool you use for one-off tasks to something closer to a working partner for ongoing projects.

REAL-WORLD EXAMPLE

Sarah manages a marketing account for three different clients. She sets up a separate Project for each:

Project: Hartley Financial — contains the brand guidelines document, a style reference email, a note on the client's preferences ("informal tone in client-facing copy; formal in regulatory documents"), the previous quarter's campaign materials, and a standing instruction that the client's compliance team must review any copy before publication.

Project: Meridian Retail — contains the target audience profile, previous copy that worked well, a note on what the Marketing Director has rejected in the past, and the seasonal campaign calendar.

When Sarah opens the Hartley Financial project to write new email copy, AI already knows the brand, the style, the constraints, and the history. The first draft is usable. Without the project, she would spend the first five minutes of every session re-briefing context she's already established.

WATCH OUT

Persistent context can become stale. If your role changes, a client's preferences shift, or a project evolves, the standing context needs to be updated. Review your projects and persistent instructions every few months. Outdated context is better than no context, but not by much — AI will confidently apply the wrong assumptions.

CHAPTER 11

Projects and Persistent Instructions

SETTING UP PERSISTENT INSTRUCTIONS

Persistent instructions are the highest-leverage single investment most regular AI users can make. Writing them well takes thirty minutes. The return on those thirty minutes compounds with every subsequent AI conversation.

WHAT TO INCLUDE IN PERSISTENT INSTRUCTIONS

The most valuable things to specify in persistent instructions:

1. YOUR PROFESSIONAL CONTEXT

Who you are, what you do, what organisation or sector you work in. This calibrates AI's assumptions about what background knowledge is relevant and what terminology is appropriate.

Example: "I am a senior solicitor at a UK mid-market law firm. My clients are typically owner-managed businesses. I communicate with clients, not other lawyers."

2. YOUR COMMUNICATION PREFERENCES

Tone, vocabulary level, things you never want to see in output. This is where you codify the things you would otherwise correct in every first draft.

Example: "Write in plain English. No jargon unless I request it. No excessive hedging. No bullet lists unless I specify them. UK English throughout. Professional but accessible."

3. YOUR STANDARD AUDIENCE

Who reads what you produce? This prevents AI from pitching responses at the wrong level.

Example: "Assume I am writing for a non-specialist reader unless I specify otherwise. Never assume technical knowledge of legal process or financial instruments."

4. KEY CONSTRAINTS THAT ALWAYS APPLY

Anything that, if AI ignored it, would make the output unusable. Legal constraints, brand constraints, regulatory constraints, relationship sensitivities.

Example: "I work in financial services. Do not make specific investment recommendations. Always include appropriate caveats where content could be interpreted as financial advice."

5. THINGS YOU WANT AI TO ALWAYS DO

Behaviours you want as default, not just when you remember to ask.

Example: "At the end of any factual output, flag any claims I should verify independently. If you are uncertain about any fact, say so explicitly rather than stating it with false confidence."

A TEMPLATE FOR PERSISTENT INSTRUCTIONS

Here is a starting template. Adapt it to your situation:

About me:

[Your role, sector, organisation type, years of experience if relevant]

What I typically need AI for:

[Your most common use cases – writing, analysis, research, editing, planning]

My communication preferences:

[Tone, vocabulary level, formatting preferences, things you never want]

My standard audience:

[Who reads your output – clients, colleagues, the public, specialists, generalists]

Key constraints that always apply:

[Legal, regulatory, brand, relationship, or accuracy constraints]

What I always want AI to do:

[Default behaviours you want in every interaction – verification flags, format defaults, etc.]

TRY THIS

Set up your persistent instructions today. Use the template above. Start with the sections that feel most obvious — your professional context and communication preferences — and add the others as you notice what you keep re-specifying in individual conversations.

After two weeks, review the instructions. What are you still correcting in every conversation? That's a gap in your persistent instructions. Add it.

SETTING UP A PROJECT

Projects are more powerful than persistent instructions because they can also hold documents, ongoing outputs, and accumulated context. Here is how to approach setting one up:

STEP 1 — DEFINE THE PROJECT'S PURPOSE

What work is this project for? One client? One recurring task type? One ongoing work project? The answer determines what belongs in it.

STEP 2 — WRITE THE PROJECT INSTRUCTIONS

These are like persistent instructions but scoped to this project. Include the context specific to this work that doesn't belong in your general instructions.

STEP 3 — UPLOAD RELEVANT DOCUMENTS

What documents does AI need to do this project's work? Upload them. Brand guidelines, style guides, reference documents, previous outputs, client briefs.

STEP 4 — NAME CONVERSATIONS CLEARLY

Within the project, give conversations descriptive names. "Email sequence for Q4 campaign" is more useful six weeks later than "New Chat (3)."

PRODUCT NOTE (VERIFIED SEPTEMBER 2026):

ChatGPT Projects are available on free and paid plans, with file limits that vary by plan. Claude Projects are available on paid plans. Gemini offers workspaces and custom configurations, but the names, file behaviour, and availability vary by product and account. Feature details change frequently; verify current capabilities and limits in the platform's help centre before relying on a specific feature.

Volatility: HIGH

THE PROJECT FOR RECURRING WORK

Some of the most valuable AI projects are not for unique work, but for recurring tasks. If you produce the same type of document weekly or monthly — a status report, a client briefing, a team update — a project containing your template, your style reference, and your standard content elements saves significant setup time.

For recurring work, the project contains:

- A template for the document type
- Style and tone specifications
- An example of a completed version you were happy with
- Any standard information that appears in every edition (team structure, project status, recurring context)

Each week or month, you open the project, paste in the new content or data, and ask AI to produce the new version to the same standard as the example.

CHAPTER 12

Files, Documents, and Source Material

WHY YOU SHOULD GIVE AI DOCUMENTS, NOT DESCRIPTIONS

There is a consistent pattern in how people under-use AI: they describe the document instead of providing it.

"I have a market research report that says our target demographic is 25-34..."

"Based on our company values, which include innovation and customer focus..."

"Our pricing is roughly in the mid-range for the sector..."

These descriptions are always incomplete. They filter the document through the user's current focus and omit everything that doesn't seem immediately relevant — including things that might matter.

When you provide the actual document, AI can read all of it. It can notice things you didn't highlight. It can extract the specific data point you forgot to mention. It can identify tensions between different sections that your description smoothed over.

Provide documents, not descriptions of documents.

WHAT AI CAN DO WITH DOCUMENTS

The practical capabilities, as of late 2026:

Summarise — Extract key points, findings, or decisions at a specified level of detail.

Extract — Pull out specific information: all the action items from meeting notes, all the financial figures from a report, all the dates and deadlines from a contract.

Analyse — Identify patterns, inconsistencies, risks, or significant points: flag anything in this contract that differs from our standard terms; identify which sections of this report have the weakest evidence.

Compare — Compare multiple documents: these two supplier proposals against the same set of criteria; this year's report against last year's.

Answer questions — "What does this contract say about notice periods?" "What is the research paper's conclusion about the effect of X on Y?"

Rewrite or adapt — Take an existing document and adapt it: rewrite in simpler language; convert to a different format; update with new information.

WHAT TO WATCH OUT FOR WITH DOCUMENTS

AI can make mistakes when working with documents, and those mistakes can be difficult to spot:

Hallucinating content — AI may generate a plausible-sounding claim about a document that is not actually in the document. Always verify specific extracted facts against the original.

Missing important nuance — AI may summarise a complex legal or financial document in a way that loses critical nuance. "The contract allows termination with 30 days notice" may be accurate but miss that this only applies to non-performance situations.

Incorrect numbers — AI sometimes misreads or miscalculates figures in documents. Any number extracted from a document should be verified against the source.

Losing track of long documents — Very long documents may not be fully attended to. Important information in the middle or end of a very long document is more likely to be missed than

information at the beginning.

VERIFY

After any significant extraction or analysis task on a document:

- 1.** Check at least two specific extracted facts against the source document directly
- 2.** Ask AI what it considered to be the most uncertain or ambiguous parts of the document
- 3.** Ask AI to flag any information it couldn't find in the document that your prompt might have led it to expect

These three checks catch the most common document-analysis failures.

DOCUMENT TYPES AND HOW TO USE THEM

CONTRACTS AND LEGAL DOCUMENTS

Most useful for: identifying specific clauses, comparing to standard terms, extracting obligations and deadlines, generating a plain-English summary.

Always verify: any claim about rights, obligations, or liabilities against the exact contract language.

FINANCIAL REPORTS AND DATA

Most useful for: summarising findings, extracting key figures, identifying trends, generating plain-English explanation.

Always verify: specific numbers against the source.

RESEARCH PAPERS

Most useful for: summarising methodology and findings, identifying key claims, extracting conclusions.

Always verify: is this what the paper actually concludes? Has the summary missed important limitations?

MEETING NOTES AND TRANSCRIPTS

Most useful for: extracting action items, producing a summary, identifying decisions made, creating a follow-up email.

Watch out for: misattributing who said what, missing an action item from a long transcript.

INTERNAL DOCUMENTS

Most useful for: adapting existing content, maintaining consistency, finding relevant previous work.

Watch out for: applying information from outdated documents without realising it.

PRO TIP

When uploading a long document, help AI navigate it. Tell AI which sections are most relevant: "The contract is 22 pages. For this task, the most relevant sections are the Termination clause (section 8) and the Indemnities (section 14). You can skim the rest unless it's relevant." This focuses AI's attention and tends to produce more accurate extraction from the sections that matter.

CHAPTER 13

Tools — What AI Can Do Beyond Text

AI IS NOT JUST A TEXT GENERATOR

The AI tools available in late 2026 can do considerably more than generate text. Understanding which capabilities exist — and which don't — determines what you can ask for, what you need to verify externally, and where AI can genuinely replace tasks that previously required separate tools.

This chapter introduces the main tool categories. Because capabilities vary by platform, by plan level, and change frequently, all specific capabilities include a Product Note with a volatility rating. High-volatility items are not in the main body.

WEB SEARCH

Most major AI platforms can search the web during a conversation. This changes one of AI's most significant limitations: its knowledge cutoff.

Without web search, AI's knowledge has a cutoff date. If you ask about recent events, current prices, or the latest data, AI either doesn't know or (more dangerously) generates plausible-sounding but outdated information.

With web search enabled, AI can retrieve current information and incorporate it into responses. It can check whether something has changed, find recent statistics, and access information published after its training cutoff.

WHAT WEB SEARCH CHANGES:

- AI can answer questions about recent events
- AI can retrieve current data rather than estimating from outdated training data
- AI can find specific sources rather than relying on memory

WHAT WEB SEARCH DOES NOT CHANGE:

- AI still interprets and summarises what it finds, and can still get that interpretation wrong
- Not all information is publicly accessible on the web
- AI can still be wrong about which sources to trust
- You still need to verify important claims

PRODUCT NOTE (VERIFIED SEPTEMBER 2026):

OpenAI, Anthropic, and Google all document web-search or research capabilities, but access, names, limits, and default behaviour vary by product, plan, region, and rollout. Always confirm whether search was used in a given response if accuracy matters.

Volatility: MEDIUM

DEEP RESEARCH AND MULTI-STEP WEB RESEARCH

Some platforms now offer a research mode that does more than answer a single search question. It can plan a search, investigate several sources, compare what it finds, and return a cited report. The name and exact behaviour vary: a platform may call it Deep Research, Research, or use another label.

Use it when the question is broad, current, and worth the time required for a proper investigation. Do not use it as a shortcut around verification.

A RELIABLE RESEARCH BRIEF

1. State the decision or piece of work the research will support.
2. Define the date range, geography, audience, and terms that matter.
3. Ask for primary or authoritative sources where they exist.

4. Request a source-linked report that separates evidence from interpretation.
5. Review the plan or scope before the research runs, if the platform offers that option.
6. Open the important sources and verify the claims you intend to use.

WEB RESEARCH CHECKLIST

- Did I define the date, place, audience, and decision this research supports?
- Did the tool actually use web sources, or did it answer from its existing context?
- Can I open the cited source and find the claim there?
- Is the source primary or authoritative for this question?
- Did the tool distinguish evidence from interpretation?
- What remains uncertain, disputed, or outside the scope of the search?

For example:

Research the current rules for [specific question] in [jurisdiction] as of [date].

This will support [decision or document]. Prefer official government, regulator, company, or primary research sources. Give me:

1. the answer in plain English;
 2. the sources for each important claim;
 3. areas where sources disagree or the position is uncertain;
 4. a list of facts I should verify before relying on the result.
- Do not fill gaps with assumptions. Say what you could not establish.

The citations are a map back to the evidence, not a guarantee that the report is correct. Research tools can select weak sources, misread a source, or draw a conclusion the source does not support. The verification step remains yours.

PRODUCT NOTE

OpenAI, Google, and Anthropic all document research tools that can search across multiple sources and return citations, but access, limits, connected sources, and names vary by plan, region, and product. Check the current help centre before designing a workflow around one. The stable principle is: define the research task, inspect the sources, and verify the claims that matter.

Volatility: HIGH

CODE EXECUTION AND CALCULATION

Some AI platforms can write and run code within the conversation, enabling reliable calculation and data processing.

This is significant because it addresses one of AI's consistent weaknesses: arithmetic and calculation. When AI uses code execution, it doesn't estimate calculations — it runs them. The result is accurate in the same way a spreadsheet result is accurate (if the formula is right).

Practical uses:

- Accurate calculations on data you provide
- Processing and analysing datasets
- Running formulas or statistical operations
- Generating charts and graphs from data

What this changes: Calculations performed via code execution are reliable. Calculations performed without code execution — AI simply generating the answer as text — are not reliably accurate, particularly for anything beyond simple arithmetic.

The rule: If your task involves maths that matters, either use a platform with code execution enabled, or verify the calculation yourself.

PRODUCT NOTE (VERIFIED SEPTEMBER 2026):

Several major platforms support code execution, file analysis, or both, but the feature name, access, limits, and supported formats vary. Consult the current platform documentation before relying on a calculation or data workflow.

Volatility: MEDIUM

IMAGE GENERATION

Some AI tools can generate images from text descriptions. This capability is useful for:

- Rough concept visualisations for design briefs
- Placeholder images for presentations
- Illustrations for non-commercial documents
- Visual brainstorming

WHAT TO BE CAREFUL OF:

- Generated images may not be suitable for commercial use without checking the platform's terms
- Generated images of people can produce photorealistic results that may be misleading out of context
- Quality varies significantly across platforms

PRODUCT NOTE (VERIFIED SEPTEMBER 2026):

Image generation is integrated into several major AI platforms, but product names, models, access, and licensing terms vary. Verify the platform's current terms and applicable law before using generated images commercially.

Volatility: HIGH — See Appendix E for current details.

FILE CREATION AND EXPORT

AI tools increasingly support creating downloadable files: documents, spreadsheets, presentations. This is still an evolving area, with quality and reliability varying considerably.

The practical approach: use AI to generate the content and structure, then create the final formatted document in your standard tool (Word, Excel, PowerPoint, Google Workspace). Attempting to produce final, formatted output directly from AI is often slower and less reliable than the two-step approach.

WHAT AI TOOLS CANNOT RELIABLY DO (AS OF LATE 2026)

Equally important is knowing what AI tools currently cannot do well:

- **Send communications on your behalf** without explicit confirmation steps
- **Access your internal systems** (email, calendar, CRM, databases) without integration tools set up by your organisation
- **Make decisions** that require judgment, ethics, or accountability — AI can analyse and recommend; decisions remain with you
- **Replace professional advice** in law, medicine, finance, or other regulated domains — AI can inform; professionals

advise

WATCH OUT

Tool availability changes. A capability that exists today may be modified, restricted, or removed. A capability that doesn't exist today may arrive. The best practice is to verify current capabilities on the platform's own documentation before designing a workflow around a specific tool feature. This book covers capabilities as of September 2026; check Appendix E and the platform's help centre for current status.

CHAPTER 14

The Memory Problem (and How to Work Around It)

WHY AI DOESN'T RELIABLY REMEMBER YOU

Unless a platform explicitly supplies relevant saved context, project knowledge, searchable history, or connected sources, every time you start a new conversation with an AI tool, you are meeting a stranger who happens to be very knowledgeable. They do not reliably know your name, your job, your preferences, your previous conversations, or anything about your situation.

The practical effect is that you should spend time re-establishing important context unless you have confirmed which persistence feature is active and what it can access.

Some platforms have begun introducing memory features — where AI can retain facts between conversations. These features are useful but imperfect, and their availability and reliability vary. This chapter covers the memory problem and its most practical workarounds.

THE FOUR TYPES OF "MEMORY" IN CURRENT AI TOOLS

IN-CONVERSATION MEMORY (ALWAYS AVAILABLE)

Within a single conversation, AI remembers everything that has been said. You can refer to "the document we discussed" or "your earlier suggestion" and AI will understand. This is reliable and consistent.

PLATFORM MEMORY FEATURES (VARIABLE)

Some platforms offer features that carry specific facts across conversations — things you told AI about yourself, preferences you established, key information you asked it to remember. These vary in scope and reliability.

PRODUCT NOTE (VERIFIED SEPTEMBER 2026):

ChatGPT offers saved memory and project memory, with availability and settings varying by plan and workspace. Claude offers profile preferences, project instructions, project knowledge, and—where enabled—search across previous chats. These are different mechanisms, not a guarantee that an AI remembers everything. Check the current documentation and verify what context a response actually used.

Volatility: HIGH — See Appendix E.

PROJECT MEMORY (THE MOST RELIABLE APPROACH)

Working within a Project creates a persistent context that AI carries across all conversations in that project. This is currently the most robust workaround for the memory problem for regular, recurring work.

DOCUMENT MEMORY (THE UNIVERSAL WORKAROUND)

A context document you paste in at the start of any conversation. Reliable across all platforms, requires no platform feature, and completely under your control.

BUILDING A CONTEXT DOCUMENT

A context document is a brief, structured summary of what AI needs to know about you and the work it will help with. It is the universal workaround for the memory problem.

Here is a template:

CONTEXT DOCUMENT – [Your name / role]

Last updated: [Date]

About me:

[2-3 sentences: role, sector, organisation type]

Current work context:

[What I'm working on, any ongoing projects, key relationships]

My communication preferences:

[Tone, format, UK/US spelling, things I never want in output]

Key constraints:

[Anything AI must not say, recommend, or assume]

My standard audience:

[Who reads what I produce]

Things I always want AI to do:

[Default behaviours – verification flags, format defaults, etc.]

A well-written context document takes twenty minutes to produce once and can be pasted into any AI conversation in seconds. It is more reliable than any platform feature and works across all tools.

A FILLED-IN CONTEXT DOCUMENT

Consider Priya, a secondary-school teacher preparing differentiated materials. Her context document might begin:

CONTEXT DOCUMENT – Priya, Year 8 science teacher

Last updated: September 2026

I teach mixed-ability Year 8 classes in England. Materials must use plain language, support different reading levels, and align with the current specification I provide. Never invent curriculum requirements or present generated material as ready for students without teacher review. Flag scientific claims, safety issues, and assumptions.

When I ask for differentiated material, give me a core version and two accessible variations. Keep the learning objective unchanged.

The document does not make AI responsible for the lesson. It gives AI a stable starting point and tells Priya what to check. She still supplies the current specification and reviews anything pupils will see, but she no longer has to repeat the class context and safety expectations in every session.

MEMORY THAT MATTERS AND MEMORY THAT DOESN'T

Not everything needs to be remembered. The investment in persistent context should be proportional to how often you'll need it.

WORTH PERSISTING:

- Your role, audience, and communication style
- Recurring project context
- Key constraints that always apply
- Your quality standards and preferences

NOT WORTH PERSISTING:

- One-off tasks
- Information that changes frequently
- Details specific to a single conversation

The test: if you would type this information again in your next AI session, persist it. If it's specific to today's task and you won't need it tomorrow, don't bother.

CHAPTER 15

Building Your Personal AI Workspace

FROM AD HOC TO SYSTEMATIC

Part III has covered the building blocks of an AI environment: persistent instructions, projects, document management, tools, and memory. This chapter shows how to put them together into a workspace that reflects how you actually work.

A well-built AI workspace is not a technical achievement. It is a professional one: the result of thinking carefully about what you need AI for, what context it should always have, and how your recurring work is organised.

THE WORKSPACE AUDIT

Before building, take stock. Answer these questions:

WHAT DO I USE AI FOR MOST OFTEN?

List your five most common AI tasks. These are the tasks your workspace should be optimised for.

WHAT CONTEXT DO I RE-ENTER REPEATEDLY?

What information do you find yourself providing in almost every session? This is what should be in your persistent instructions.

WHAT RECURRING PROJECTS OR TASK TYPES DO I HAVE?

What work comes around every week, every month, or for every client? Each recurring type deserves its own Project.

WHAT DOCUMENTS DOES AI REGULARLY NEED?

What files do you find yourself uploading repeatedly? These belong in the relevant Project.

WHAT ARE MY QUALITY STANDARDS?

What makes output unusable for you? What do you always correct? These corrections belong in your persistent instructions as explicit rules.

A COMPLETED WORKSPACE AUDIT

Take David's consultancy as an example. His answers produce three workspaces:

- **Client work:** one Project per active client, containing the brief, approved source documents, a tone reference, previous deliverables, and client-specific constraints.
- **Recurring reports:** a Project containing the reporting template, a completed example, the standard headings, and the monthly input checklist.
- **Personal practice:** a small workspace for prompt experiments, failure notes, and reusable patterns that are not tied to a client.

His persistent instructions hold the information that applies everywhere: his role, UK spelling, plain professional language, and a default instruction to flag unsupported facts. Project instructions hold what changes by client or task. This separation prevents a client-specific constraint from leaking into unrelated work.

A SIMPLE WORKSPACE STRUCTURE FOR MOST PROFESSIONALS

```
PERSISTENT INSTRUCTIONS (applies to everything)
├ Professional context
├ Communication preferences
├ Standard audience
└ Default behaviours (verification, format, etc.)

PROJECTS
├ Project A: [Client / Ongoing work 1]
│   ├── Project instructions (context specific to this work)
│   ├── Reference documents (brand guide, style reference, etc.)
│   └ Ongoing conversations
│
├ Project B: [Client / Ongoing work 2]
│   └ [Same structure]
│
└ Project C: Recurring tasks
    ├── Weekly report template and example
    ├── Standard document formats
    └ Template prompts for recurring work
```

THE RECURRING TASK TEMPLATE

For any task you do on a regular cadence, a template prompt is worth building. A template prompt is a prompt with placeholders — variables you update with each use — built on a structure you've already tested and refined.

Format:

```
[TEMPLATE NAME]

Context (stable):
[Everything that stays the same – purpose, audience, style, constraints]

Variable inputs this week:
[PASTE CONTENT HERE]
[KEY DATES: X]
[KEY DECISIONS: X]

Output required:
[Format, length, structure – specified once, used every time]
```

This reduces a recurring AI task to: open the template, fill in the variables, run. The setup investment is front-loaded; the recurring use is fast.

PRO TIP

Keep a "templates" section within your AI workspace — either a Project dedicated to this, or a simple document you maintain. When you write a prompt that works especially well, save it. When you develop an interaction pattern that produces good results consistently, document it. Your prompt library becomes more valuable than any generic prompt template book because it is tested, specific to your work, and continuously refined.

WORKSPACE MAINTENANCE

A workspace that isn't maintained becomes a liability. Outdated context produces confidently wrong outputs. Review your workspace periodically:

- **Monthly:** Are the persistent instructions still accurate? Has anything changed about your role, audience, or constraints?
- **When a project ends:** Archive or close the project so it doesn't contaminate future work with irrelevant context.
- **When something keeps going wrong:** What's in your workspace that might be causing this? What context might AI be applying incorrectly?

The workspace is a living document of your working relationship with AI, not a one-time setup.

WATCH OUT

Workspaces created on one platform do not transfer to another. Your Claude Projects, your ChatGPT custom instructions, and your Gemini settings are all separate. If you use multiple platforms, you need to maintain context in each, or accept that some platforms know more about your preferences than others. Many users find it practical to choose a primary platform for their most important work and use secondary platforms with more ad hoc setup.

WHAT PART III HAS ESTABLISHED

You now have the framework for building an AI environment:

- Persistent instructions that carry your context forward automatically
- Projects that organise your work and maintain document access
- Document handling practices that maximise accuracy and catch common errors
- An understanding of tools and their limitations
- A memory strategy that works across platforms
- A workspace structure suited to professional use

Part IV moves from individual interactions and workspace setup to something qualitatively different: multi-step workflows that tackle complex tasks by breaking them into stages.

From Prompts to Workflows

Complex work doesn't happen in a single exchange. Part IV introduces the thinking and techniques that take AI from answering individual questions to participating in multi-stage processes — planning, drafting, reviewing, revising, and handing off.

CHAPTER 16

When One Prompt Isn't Enough

THE SINGLE-PROMPT CEILING

For simple, well-defined tasks — summarise this, draft that, answer this question — a single well-crafted prompt is usually sufficient. But professional work is rarely simple and well-defined.

A business proposal is not one task. It is: research the client, understand their problem, frame the approach, develop the argument, structure the document, draft each section, review for consistency, check for gaps, refine the language, and produce a final version. Each of these is a distinct cognitive task that benefits from distinct treatment.

Attempting to produce a business proposal in a single prompt produces a generic outline at best and a generic document at worst. Neither is what you needed.

The ceiling of single-prompt AI use is low. Most professional work requires workflows.

WHAT MAKES SOMETHING A WORKFLOW

A workflow, for our purposes, is any AI task that:

- Involves more than one distinct phase (research, then structure, then draft, then review)
- Produces output that becomes input to the next stage
- Benefits from human review and adjustment at one or more stages
- Is too complex for a single prompt to handle well

Most significant pieces of work — proposals, reports, analyses, plans, complex communications — are workflows. The shift to thinking in workflows is one of the biggest quality improvements

available to regular AI users.

WHY BREAKING TASKS INTO STAGES WORKS

When you ask AI to do too many things at once, it compromises on everything. The research is superficial because the prompt also asked for structure. The draft is generic because the prompt also asked for analysis. The analysis is shallow because the prompt also asked for recommendations.

Breaking a task into stages allows AI to do each stage well, with appropriate depth and focus. It also allows you to review and adjust at each stage before committing the next stage to an approach that might be wrong.

The best analogy is professional practice: a good architect doesn't hand you the construction drawings at the first meeting. They produce a brief, then a concept, then a scheme, then detailed drawings — each stage reviewed and approved before the next begins.

THE THREE MOST COMMON MULTI-STAGE PATTERNS

THE RESEARCH-PLAN-DRAFT PATTERN

Stage 1: Research and gather information. Stage 2: Plan the structure. Stage 3: Draft the content.

This is the most universally useful workflow. Almost any document benefits from separating the gathering of content from the writing of it.

THE DRAFT-REVIEW-REFINE PATTERN

Stage 1: Produce a complete draft. Stage 2: Review the draft against explicit criteria. Stage 3: Refine based on the review.

This is particularly valuable because Stage 2 uses a different cognitive mode than Stage 1. Writing mode and editing mode are different, and separating them — even within AI — produces

better results.

THE DECOMPOSE-EXECUTE-ASSEMBLE PATTERN

Stage 1: Break a large task into components. Stage 2: Execute each component separately. Stage 3: Assemble the components.

This is most useful for large documents or complex outputs where different sections require different content or expertise.

REAL-WORLD EXAMPLE

David needs to produce a quarterly business review for his consultancy's three largest clients.

Single-prompt approach: "Write a quarterly business review for Client X." Result: A generic template that David has to fill in entirely himself.

WORKFLOW APPROACH:

Stage 1: "I'm going to paste in my notes from the last three months of client interactions, plus their project status reports. Extract: key achievements, challenges faced, decisions made, and open items. Present as a structured list, not prose."

David reviews the extraction, adds two items AI missed, removes one that was irrelevant.

Stage 2: "Using the extracted information, create a narrative structure for the quarterly review. Three sections: What We Achieved, What We Learned, and What Comes Next. For each section, suggest what the key two or three points should be."

David adjusts the suggested structure — moves one point to a different section, adds a point AI didn't suggest.

Stage 3: "Now draft the quarterly review document using the structure we agreed. Professional tone, plain language, about 800 words. The client is the owner and will share it with their management team."

The resulting document is specific, accurate, and usable. It took less time than the single-prompt version and produced dramatically better output — not because AI was better, but because the workflow allowed human judgment to guide it at each stage.

WHY IT WORKS

Multi-stage workflows work for the same reason that professional processes work: each stage has a clear purpose, a clear deliverable, and a defined handoff point. Breaking work into stages allows you to catch errors before they compound. It allows you to apply different types of thinking to different types of sub-tasks. And it allows AI to do each sub-task well rather than doing all sub-tasks adequately.

CHAPTER 17

Breaking Tasks Into Steps

THE DECOMPOSITION SKILL

The most important skill in multi-step AI work is not writing better prompts. It is breaking tasks into the right stages.

A well-decomposed task has stages that:

- Have distinct deliverables
- Can be reviewed independently
- Build on each other in a clear sequence
- Are each achievable to a high standard in a single focused interaction

Poorly decomposed tasks break in the wrong places, create dependencies that are hard to manage, or duplicate work across stages.

This chapter teaches the decomposition patterns that work most reliably across common professional tasks.

DECOMPOSITION PATTERN 1: INFORMATION → STRUCTURE → CONTENT

The most broadly applicable pattern. Works for almost any document.

Stage 1 – Information gathering

"Extract all the relevant information from [source(s)]."

→ Output: Raw material, structured as a list

Stage 2 – Structure design

"Given this information, propose a document structure.

What are the main sections? What is the most logical order?

What is the key argument or narrative?"

→ Output: Document structure with brief section notes

[Human review: Adjust structure before committing to draft]

Stage 3 – Draft

"Draft [Document X] using the agreed structure."

→ Output: Full draft

Stage 4 – Review and refine (optional – Chapter 18)

This pattern works for: proposals, reports, articles, presentations, plans, briefing documents.

DECOMPOSITION PATTERN 2: CLARIFY → EXPLORE → RECOMMEND

Most useful for analysis and decision support.

Stage 1 – Clarify the question

"Help me define exactly what question I need to answer.

Here is the situation: [context]

Here is what I'm trying to decide: [decision]

What additional information would be most valuable? What are the key sub-questions?"

→ Output: Sharper framing of the problem

Stage 2 – Explore options and evidence

"Given this framing, what are the main options or considerations?

What does the evidence suggest about each?"

→ Output: Analysis of options

Stage 3 – Recommend

"Based on what we've established, what would you recommend and why?

What are the main risks of your recommendation?"

→ Output: Recommendation with reasoning and risk assessment

This pattern works for: decision preparation, strategy analysis, problem diagnosis, options assessment.

DECOMPOSITION PATTERN 3: COMPONENTS → DRAFT EACH → ASSEMBLE

Most useful for large documents with distinct sections.

```
Stage 1 – Define components
"This document has [N] sections: [list them].
Let's draft each one separately. Start with Section 1."

Stage 2 – Draft each component
[For each section:]
"Draft Section [N]: [title].
Key points to cover: [list].
Length: [specification].
Tone: [specification]."
```

```
Stage 3 – Assemble and review
"I'm going to paste in all sections now.
Review them as a complete document:
– Do they flow logically?
– Are there repetitions or gaps?
– Is the tone consistent throughout?
– What would you change?"
```

This pattern works for: annual reports, long proposals, training materials, multi-section documents.

MAKING HANDOFFS EXPLICIT

Between stages, always do two things:

1. REVIEW THE OUTPUT.

Does the output of Stage N accurately reflect what you wanted? Is anything missing? Is anything wrong? Correct it before it propagates to Stage N+1.

2. CARRY FORWARD WHAT MATTERS.

When moving to the next stage, either reference the previous output directly (in the same conversation, AI will remember it) or paste in what AI needs to know. Don't assume AI will apply context from Stage 1 accurately at Stage 3 if the conversation is very long.

TRY THIS

Take a piece of work you need to produce in the next week. Spend five minutes mapping it:

1. What are the distinct stages this work naturally has?
2. What is the deliverable of each stage?
3. Where is human review most important — where could an error in Stage N make Stage N+1 irretrievably wrong?

Now design the workflow before you start. The map will tell you what to ask AI for at each stage.

WATCH OUT

Stages that are too granular create overhead without benefit. If you find yourself doing ten stages for a task that could be done in three, you have over-decomposed. The right granularity is: each stage produces something that can be reviewed and corrected before the next stage begins. If a stage is so small that there's nothing meaningful to review, combine it with the adjacent stage.

CHAPTER 18

Critique, Edit, and Improve

USING AI AGAINST ITS OWN OUTPUT

One of the most underused techniques in AI work is asking AI to critique what it has produced.

This sounds circular — and it is, slightly. AI is reviewing its own work, which might seem unlikely to catch errors it didn't notice the first time. But it is more effective than it sounds, for a practical reason: critique mode is a different cognitive stance than generation mode. When you ask AI to switch from producing to evaluating, it activates different patterns and often catches problems it generated in draft mode.

THE CRITIQUE LOOP

The basic critique loop has three steps:

STEP 1 — PRODUCE A DRAFT

Ask AI to produce the draft as normal.

STEP 2 — APPLY THE CRITIQUE PROMPT

Ask AI to evaluate the draft against explicit criteria. The more specific the criteria, the more useful the critique.

STEP 3 — REVISE BASED ON CRITIQUE

Ask AI to produce a revised version addressing the issues raised.

THE CRITIQUE PROMPT TEMPLATE

Please review the [document/email/proposal/plan] you just produced.

Evaluate it against the following criteria:

- [Criterion 1: e.g., Does it clearly address the reader's main concern?]
- [Criterion 2: e.g., Is the argument in the right order?]
- [Criterion 3: e.g., Are there claims that could be questioned?]
- [Criterion 4: e.g., Is there anything missing?]
- [Criterion 5: e.g., Is the tone right for the audience?]

For each criterion: identify what is working well and what could be improved.

At the end: give a one-paragraph overall assessment.

Do not be generous. The goal is a better document, not reassurance.

The instruction "do not be generous" matters. Without it, AI tends to produce critique that finds much to praise and softens its concerns. This produces pleasant but not especially useful critique.

WHY IT WORKS

Critique mode works better than it theoretically should because AI is drawing on large bodies of material about what constitutes good work in different domains. When asked to evaluate against explicit criteria, it can apply those standards more rigorously than it applies them by default in generation mode. Think of it as asking a writer to switch hats: generation mode produces content; critique mode applies editorial judgment.

THE BEFORE → BETTER → BEST: CRITIQUE LOOP

The scenario: Sarah has used AI to draft a proposal section and wants to improve it before sending.

Stage 1 — Draft produced

[AI produces a draft proposal section about the agency's approach to the client's brief]

Stage 2 — Critique applied

Sarah sends: *"Please review the proposal section you just drafted. Evaluate it against these criteria: Does it address the client's specific situation or is it generic? Is the argument clear and logical? Are there any claims that aren't supported by what I've told you? Is the tone right for a director-level reader at a mid-size retail company? What's missing? Be direct — I need useful feedback, not reassurance."*

AI responds: *"A few issues to flag. The opening paragraph describes the agency's general approach rather than the client's specific situation — a sceptical reader will notice you haven't demonstrated understanding of their particular challenge. The 'proven results' claim on page two isn't backed by anything specific in what you've provided. The case study section is good but should come earlier — it's currently buried after the methodology, but for this type of client, demonstrated experience is probably the first thing they want to see. The tone is appropriate throughout."*

Stage 3 — Revision

Sarah: *"Revise the draft addressing these points: lead with the client's situation, move the case study up, remove the 'proven results' claim and replace it with [SPECIFIC RESULT I'LL ADD]."*

The resulting revision is substantively better. The critique took two minutes to run. The revision took thirty seconds to instruct. The improvement would have taken Sarah much longer to identify and articulate without the structured critique loop.

SPECIFIC CRITIQUE TYPES

Different types of work benefit from different critique focuses:

FOR PERSUASIVE DOCUMENTS (PROPOSALS, BIDS, PITCHES):

"Does this make the reader want to say yes? What objections might they raise? Is the value proposition clear? What is the weakest point in the argument?"

FOR EXPLANATORY DOCUMENTS (REPORTS, BRIEFINGS, GUIDES):

"Is the key finding easy to identify? Is the logic clear? Are there technical terms an educated non-specialist would need explained? Is anything assumed that shouldn't be?"

FOR COMMUNICATIONS (EMAILS, LETTERS, ANNOUNCEMENTS):

"Is the purpose of this communication immediately clear? Is the tone right for the relationship? Could anything be misinterpreted? Is there anything missing the recipient will need to act on this?"

FOR CREATIVE OR MARKETING CONTENT:

"Is the opening strong enough to keep a reader going? Is there any generic marketing language that a sceptical reader would discount? Is the call to action clear?"

CRITIQUE WITH AN ADVERSARIAL PERSONA

For high-stakes documents, a more powerful critique technique is asking AI to adopt a specific adversarial perspective:

"Review this proposal as if you are a procurement director at a large company who has seen hundreds of similar proposals and is looking for reasons to discount this one. What would make you sceptical? What would make you put this in the discard pile?"

"Review this contract clause as if you are a lawyer representing the other party. What would you challenge? What ambiguity would you exploit?"

"Review this business case as if you are a CFO who has been burned by overoptimistic projections before. What assumptions are too aggressive? What risks are underplayed?"

These adversarial critique prompts produce more rigorous feedback than neutral evaluation and are particularly valuable before anything is presented to an actual sceptical audience.

WATCH OUT

AI critique has two consistent weaknesses. First, it tends to be comprehensive in format but inconsistent in depth — it will often identify a problem without identifying the correct solution. Second, it may miss issues that require external knowledge it doesn't have: the client context it wasn't given, the regulatory nuance that requires specialist knowledge, or the political sensitivities you didn't mention. Use AI critique to catch structural, logical, and stylistic issues. Use human review to catch everything that requires knowledge or judgment AI doesn't have.

CHAPTER 19

Chains, Loops, and Handoffs

START WITH A COMPLETE WORKFLOW

Before naming the mechanics, look at what a useful workflow feels like in practice. David needs to produce a client brief from call notes, emails, and earlier project context:

1. AI extracts the client's requirements and uncertainties.
2. David reviews the extraction, corrects one mistake, and adds two missing requirements.
3. AI proposes a structure for the brief.
4. David approves the structure before AI drafts it.
5. AI drafts the brief section by section.
6. AI critiques the draft against the agreed brief.
7. AI revises the draft, and David makes the final decision about what can be sent.

The quality comes from the connections between the stages: clear outputs, deliberate human review, and a defined standard for improvement. Now we can name the three mechanics that make those connections work.

THE BUILDING BLOCKS OF COMPLEX WORKFLOWS

Chapter 17 introduced decomposition patterns. This chapter introduces the mechanics of connecting stages together — the "plumbing" of multi-step AI workflows.

Three mechanics cover most of what you will encounter:

Chains — Sequential stages where each output becomes the next input

Loops — Iterative cycles where output is reviewed and refined until it meets a standard

Handoffs — Points where the work passes from AI to a human (or vice versa) for a specific contribution

Understanding which mechanic applies to which part of a task is what turns a list of stages into a functioning workflow.

CHAINS

A chain is the simplest multi-step pattern: $A \rightarrow B \rightarrow C$. Each step's output is the input to the next.

The key discipline in a chain is maintaining coherence: each step should produce output that is complete enough and well-formed enough to be a useful input to the next step. If Step 2 depends on Step 1 and Step 1 is mediocre, Step 3 will inherit and amplify that mediocrity.

THE CHAIN DISCIPLINE:

- Review each step's output before proceeding to the next
 - If a step's output is substantially wrong, correct it before proceeding
 - Be explicit about what carries forward: "Using the outline from Step 1, now draft Section 3"
-

LOOPS

A loop is an iterative cycle: Produce $\rightarrow$ Review $\rightarrow$ Improve $\rightarrow$ (Review again) $\rightarrow$ (Improve again) $\rightarrow$ Done.

Loops are appropriate when:

- The quality bar is high and the first draft is unlikely to meet it
- The criteria for success are explicit and can be systematically applied
- Time spent iterating is less than time spent correcting downstream

The critique loop in Chapter 18 is a loop. So is any scenario where you are refining a draft toward a specific standard.

THE LOOP DISCIPLINE:

- Define what "done" looks like before you start looping — otherwise loops run indefinitely
 - Be specific about what is wrong in each review cycle — vague feedback produces marginal improvement
 - Know when to stop: a document that is 90% right and been through two loops is usually better to complete by hand than to loop a third time
-

HANDOFFS

A handoff is a point where the work passes between AI and human. Handoffs are not failures — they are deliberate design choices about where human judgment adds value that AI cannot.

The best workflows design handoffs deliberately:

AI → Human handoffs occur when:

- AI has produced material that needs human judgment to evaluate (is this claim actually correct?)
- The next step requires knowledge AI doesn't have (what does the client actually want?)
- A decision needs to be made that carries accountability (does this represent our position?)
- The work will be put in front of someone who needs to trust it

Human → AI handoffs occur when:

- A human has made a judgment that AI should now implement (here's the structure I want — now draft it)
 - New information has arrived that AI needs to incorporate (here's their response — now draft a reply)
 - A human review has identified specific changes that AI should make
-

PRO TIP

Document your best workflows. When you design a workflow that works well for a type of task, write it down as a template. The workflow structure — not just individual prompts but the sequence, the handoff points, and the review criteria — is reusable. Over time, a small library of tested workflow templates is one of the most valuable assets you can build as a professional AI user.

WATCH OUT

Over-engineered workflows add overhead without proportionate benefit. A three-paragraph email does not need a seven-stage workflow. Scale the process to the stakes and complexity of the work. A major proposal warrants multiple stages with careful review. A routine internal update does not. The workflow discipline is most valuable when the stakes are high and the output will be scrutinised — apply proportionate judgment elsewhere.

WHAT PART IV HAS ESTABLISHED

You can now design and execute multi-step AI workflows:

- Recognise when a task requires multiple stages
- Break tasks into stages with clear deliverables and handoff points
- Apply three decomposition patterns to different task types
- Use the critique loop to improve drafts systematically
- Connect stages with chains, loops, and deliberate handoffs

Part V applies these skills to twenty-five specific, tested workflows across the most common professional and personal use cases.

25 Real-World Workflows

Part V is a reference library. These twenty-five workflows are tested, structured, and ready to use — or adapt for your specific situation. Each follows the same format: Situation, Steps with exact prompts, Verification, and Customisation notes.

How to use this section: find the workflow closest to what you need. Read the Situation to confirm it's a good match. Run the Steps. Follow the Verification. Then read Customisation to adapt it to your specific circumstances.

Placeholders appear in [BRACKETS]. Replace them with your specific content before running.

Find a workflow by time and complexity

Use this as a quick starting point. The times are estimates for a prepared user with the necessary source material available; the work still needs proportionate review.

STARTING POINT	TYPICAL TIME	WORKFLOWS
Quick win — one clear task	About 5 minutes	W1, W4, W9, W19, W24
Medium build — several prompts and a review	15-30 minutes	W2, W3, W6, W8, W10, W13, W16, W20
Deep work — source material, stages, or decisions	30+ minutes	W5, W7, W14, W15, W17, W18, W21, W22, W23, W25

Start with the nearest match for your situation, then choose the simplest level that is safe for the stakes. A complex workflow is useful when the work is complex; it is unnecessary overhead for a short, low-risk task.

BUSINESS AND PROFESSIONAL

W1 — WRITE AND IMPROVE A PROFESSIONAL EMAIL

SITUATION

You need to write an email that matters — a client communication, a proposal follow-up, a difficult message to a colleague, a professional introduction. The stakes are high enough that you don't want to draft it without a second opinion.

STEP 1 — PROVIDE CONTEXT

I need to write a professional email. Here is the context:

Sender: [Your name, role, organisation]

Recipient: [Name, role, organisation, relationship to you]

Purpose of the email: [What you want to achieve with this email]

Background: [Relevant history, context, what's happened so far]

Tone required: [Formal / professional-informal / warm / firm]

Key constraints: [Anything this email must not say or do]

Desired response or outcome: [What you want the recipient to do or feel after reading it]

Length preference: [Brief / medium / thorough]

Please draft this email.

STEP 2 — CRITIQUE THE DRAFT

Review the email you just drafted. Evaluate it on:

- Is the purpose immediately clear?
- Is the tone right for the relationship?
- Could anything be misinterpreted?
- Is anything missing that the recipient needs to act on this?
- Is it the right length?

Be direct about any problems.

STEP 3 — REVISE

Revise the email addressing the issues you identified. [Add any additional changes you want: "Also, make the opening less formal" / "Remove the third paragraph" / "Add a clear call to action at the end."]

VERIFICATION

Before sending: Does it say what you intend? Read it from the recipient's perspective. Would the tone land as intended? Is there anything that could be misconstrued? If the email involves any commitment, fact, or number, verify it independently.

CUSTOMISE

- For internal communications: reduce formality, increase directness
- For client complaints: add "The tone must acknowledge their frustration without admitting liability or making commitments we haven't authorised" to Step 1
- For cold outreach: add "This person does not know me, so make the value proposition clear within the first two sentences"
- For difficult personal situations: add context about the relationship history and the sensitivity required

W2 — PREPARE FOR A DIFFICULT CONVERSATION

SITUATION

You have a conversation coming up that you are not looking forward to. Performance review. Negotiation. Complaint. Conflict. Family tension. You want to go in prepared, not reactive.

STEP 1 — DEFINE THE CONVERSATION

I have a difficult conversation to prepare for. Please help me prepare.

Situation: [What is happening and why this conversation is needed]
Who I'm talking to: [Their role, your relationship, relevant background]
What I want to achieve: [The ideal outcome of this conversation for you]
What I'm worried about: [The scenarios I'm dreading or want to avoid]
My position: [What my bottom line is, what I'm willing to do]
Their likely position: [What you expect they want or will say]

STEP 2 — MAP THEIR PERSPECTIVE

Help me understand their perspective as fully as possible.

What legitimate concerns or interests might they have?
What might they feel about this conversation before it starts?
What outcome might they be hoping for?
What would make this conversation go well from their point of view?

I want to understand their side, not just my own.

STEP 3 — PREPARE FOR LIKELY SCENARIOS

Based on what we've established, what are the three or four most likely ways this conversation could go? For each scenario, what is the key thing I need to do or say to steer it toward a good outcome?

STEP 4 — DRAFT KEY STATEMENTS

Help me draft the key things I need to say clearly:

1. My opening statement (what I'll say to start the conversation)
2. My core position, stated clearly but not aggressively
3. A response to their most likely pushback

Tone: [Firm and direct / Warm and collaborative / Measured and professional]

VERIFICATION

Is the opening statement something you could actually say without it feeling scripted or artificial? Have you considered their perspective genuinely, not just strategically? If the conversation involves any specific commitments, facts, or numbers, confirm these before you go in.

CUSTOMISE

- For performance conversations with a direct report: add HR policy context and any documentation you have
- For negotiation: add "I also want to identify possible compromise positions between our positions"
- For family situations: adjust tone instructions significantly and focus Steps 2 and 3 more heavily on empathy
- For complaints: add "I want to be assertive without being hostile"

W3 — RESEARCH A TOPIC FOR A MEETING

SITUATION

You have a meeting coming up on a topic you need to know more about. You want to be prepared enough to contribute intelligently without spending hours reading primary material.

STEP 1 — ESTABLISH THE KNOWLEDGE REQUIREMENT

I have a meeting in [timeframe] about [topic]. I need to be well-prepared.

My current knowledge level: [None / Basic / Some background]

Who else will be in the meeting: [Roles and expertise level]

The likely agenda or decisions: [What the meeting will cover]

What I need to be able to do: [Contribute / Ask intelligent questions / Understand what others are saying / Make a recommendation]

Please give me a briefing on [topic] pitched at the right level for my knowledge and this context.

STEP 2 — IDENTIFY KNOWLEDGE GAPS

Based on what you've told me, what are the three to five things I most need to understand clearly before this meeting? What are the key terms or concepts I should be familiar with? What questions should I be ready to answer?

STEP 3 — PREPARE QUESTIONS

What would be the most intelligent questions I could ask in this meeting to contribute productively and demonstrate I've prepared? Focus on questions that move the meeting forward, not just questions that demonstrate my knowledge.

VERIFICATION

AI briefings are starting points, not authoritative sources. Verify any specific facts, statistics, or claims you intend to cite in the meeting. Confirm current information for any fast-moving topics (technology, regulation, markets). If the meeting involves decisions with significant consequences, supplement AI preparation with primary sources.

CUSTOMISE

- For technical topics: ask AI explicitly to explain technical concepts in plain language before introducing the terminology

- For regulatory topics: add "Focus particularly on the current regulatory position and flag any areas of uncertainty or ongoing change"
- For commercially sensitive topics: structure the briefing around the business context, not just the technical content

W4 — SUMMARISE A LONG DOCUMENT

SITUATION

You have a document that is too long to read fully in the time available — or you need to extract specific information efficiently. The document might be a report, a policy, a contract, a research paper, or a lengthy email chain.

STEP 1 — UPLOAD AND INSTRUCT

Please read the document I'm attaching. [Attach document.]

Before summarising, tell me: how long is this document, what type of document is it, and what appear to be its main sections?

STEP 2 — TARGETED SUMMARY

Now produce a summary for the following purpose:

Audience: [Who will read or use this summary]

Purpose: [What decision, action, or understanding this summary will support]

What to include: [Key findings / Main arguments / Action items / Financial figures / Decisions made / Recommendations]

What to exclude: [Background information already known / Technical detail not needed / Appendices]

Length: [Two paragraphs / Five bullet points / One page]

Format: [Prose / Bullet points / Table / Structured sections]

STEP 3 — FLAG WHAT NEEDS ATTENTION

Now separately:

1. What are the most important facts, numbers, or claims in this document that I should verify independently?
2. Were there any sections of the document that you found ambiguous or that required assumptions?
3. Is there anything in the document that seems unusual, surprising, or that a careful reader should notice?

VERIFICATION

Check at least two specific facts extracted from the document against the source text. Verify any numbers or figures independently. If the document is a contract or legal instrument, the plain-language summary is a starting point — do not act on it without reference to the original language for anything that matters.

CUSTOMISE

- For contracts: add "Specifically flag: any obligations on us, any deadlines, any unusual clauses, and any terms that differ from standard commercial practice"
- For research papers: add "Evaluate the strength of the evidence. What is the methodology? What are the limitations the authors acknowledge?"
- For meeting notes/transcripts: add "Extract all action items with owner and deadline. Identify all decisions made."
- For financial documents: add "Flag any calculation that relies on an assumption I should verify"

W5 — DRAFT A PROJECT PROPOSAL

SITUATION

You need to write a proposal for a project — internal or external — that makes a clear case for the project, describes the approach, and asks for approval or resources.

STEP 1 — GATHER INFORMATION

I need to write a project proposal. Before drafting, help me structure the content.

Project name: [Name]

Proposing to: [Who will approve this – their role and decision-making context]

What the project will do: [Summary of the project]

The problem it solves: [What is not working now / what opportunity is being missed]

The proposed approach: [How you will do it]

Resources needed: [Time, money, people]

Expected outcomes: [What success looks like, ideally with measurable indicators]

Risks: [What might go wrong]

Timeline: [When it will happen]

Given this, what information am I missing that a strong proposal should include?

STEP 2 — STRUCTURE THE PROPOSAL

Based on what we've established, propose a document structure for this proposal. The reader is [decision-maker's role]. They will want to know: does this project make sense, is the approach credible, and is the case for investment compelling?

Suggest section headings with one sentence on what each section should achieve.

STEP 3 — DRAFT SECTION BY SECTION

Now draft Section 1: [name].
Key points: [from your structure plan]
Length: [specification]
Tone: [professional / business-like / persuasive but evidence-based]

[Repeat for each section, reviewing between sections.]

STEP 4 — CRITIQUE AND REFINE

Review the complete proposal. As a decision-maker who has seen many project proposals:

- What is the most compelling element?
- What is the weakest element?
- What questions would you still have after reading this?
- Is the case for the resources requested clearly made?

VERIFICATION

Verify any financial figures. Check that outcomes are genuinely measurable as stated. Ensure any claimed precedents or comparisons are accurate. Have someone familiar with the decision-maker's perspective review before submission.

CUSTOMISE

- For internal proposals: reduce formality, increase emphasis on operational impact and resource efficiency
- For external bids: add the client's evaluation criteria explicitly and verify the proposal addresses each one
- For new budget requests: add a cost-benefit analysis section and ensure the return on investment is explicitly stated

W6 — WRITE A JOB POSTING

SITUATION

You need to attract the right candidates to a role without writing something generic that attracts everyone and therefore no one.

STEP 1 — DEFINE THE ROLE

Help me write a job posting. Here is the role:

Role title: [Title]

Organisation: [Company name and brief description]

What this person will actually do: [Day-to-day responsibilities]

Who they will work with: [Team, reporting line, stakeholders]

What makes someone great at this role: [Skills, experience, behaviours – be specific]

What makes this role attractive: [Why would the right person want this job?]

Essential requirements: [Non-negotiable qualifications or experience]

Desirable but not essential: [Preferences]

Salary range: [If disclosable]

Location/remote/hybrid: [Working arrangement]

Any particular types of candidates you want to attract or avoid?

STEP 2 — DRAFT THE POSTING

Write the job posting.

Tone: [Professional and direct / Warm and human / Formal]

Length: [Standard job board format – concise / Detailed]

Avoid: clichés like "dynamic team player", "fast-paced environment", "passionate", "hit the ground running". Write in plain language that describes the actual role and what we're looking for.

STEP 3 — TEST THE POSTING

Read this posting as a strong candidate who is browsing job boards quickly.

- Does the role stand out?
- Is it clear what they'd be doing day-to-day?
- Is it clear what kind of person would succeed here?
- Would a great candidate self-select in – or would they skip past it?

What would make it stronger?

VERIFICATION

Have the hiring manager review for accuracy — AI may have emphasised the wrong things. Check that essential requirements are genuinely essential (unnecessarily restrictive requirements are both wasteful and a legal risk in many jurisdictions). Verify compliance with any applicable equal opportunities or advertising requirements.

CUSTOMISE

- For specialist technical roles: add a technical requirements section and be specific about tools and technologies required
- For senior leadership roles: emphasise strategic aspects and leadership behaviours over task descriptions
- For organisations with strong cultures: add context about the culture and be honest about what makes the organisation unusual

W7 — PREPARE A CLIENT REPORT

SITUATION

You need to produce a client-facing report that summarises work done, findings, or recommendations. The report must be professional, clear, and demonstrate value.

STEP 1 — GATHER AND STRUCTURE CONTENT

I'm producing a client report. Please help me structure it.

Client: [Name, type of organisation, relationship]

Purpose of the report: [What it covers and why the client needs it]

What I have to work with: [The source material I'm going to attach or describe]

What the client cares most about: [Their priorities – cost, risk, growth, efficiency, etc.]

What we want them to do or feel after reading: [Approval, action, confidence, etc.]

Format requirements: [Any specified format / Our standard format]

Length: [Approximate]

Please propose a report structure with the most important content first.

STEP 2 — DRAFT FROM SOURCE MATERIAL

[Attach or paste source material]

Draft the report using this source material. Follow the proposed structure.

Write for the client – assume they are [knowledge level] about this topic.

Professional tone, plain language, no unnecessary jargon.

STEP 3 — EXECUTIVE SUMMARY

Now write the executive summary:

- The client's situation and what we were asked to do
- The most important three to five findings or conclusions
- What we recommend they do next

Length: One page maximum. Read-first, not read-last.

STEP 4 — REVIEW

Review this report as if you are the client receiving it.

- Is the most important information easy to find?
- Is there anything that is unclear or requires background knowledge the client might not have?
- Is there anything missing that a client would reasonably expect to see?
- Is the visual structure (headings, paragraphs, any tables) easy to navigate?

VERIFICATION

Verify all statistics, financial figures, and dates against source material. Check that recommendations are logically supported by the findings. Have a colleague read it who has not been close to the work — they will catch assumptions you've made that the client won't share.

W8 — CREATE A MEETING AGENDA AND FOLLOW-UP

SITUATION

You are running a meeting that needs to produce decisions or actions, not just discussion. You want an agenda that focuses the meeting, and a follow-up that records what happened clearly.

STEP 1 — DESIGN THE AGENDA

I'm running a [type] meeting with [participants and their roles].

Purpose of this meeting: [What decisions need to be made / what needs to be resolved / what needs to be planned]

Time available: [Duration]

Background context: [What situation this meeting is responding to]

Please design an agenda that:

- Allocates time to each item proportional to its importance
- Puts decisions before discussion where possible
- Identifies who leads each item
- Ends with clear next steps

Include: a pre-read list if relevant, any materials that should be prepared in advance.

STEP 2 — POST-MEETING NOTES → FOLLOW-UP

Here are my notes from the meeting: [paste notes / attach transcript or notes document]

Please produce a meeting follow-up document containing:

1. Decisions made (what was agreed)
2. Actions (who does what by when)
3. Topics tabled for another time
4. Any open questions not resolved

Format: brief and structured. This will be sent to all participants within 24 hours of the meeting.

VERIFICATION

Circulate the actions list to attendees for confirmation before sending — memory of who committed to what diverges quickly. Verify any facts or figures noted in the meeting before recording them as agreed.

MARKETING AND CONTENT

W9 — WRITE A LINKEDIN POST

SITUATION

You want to publish something on LinkedIn that is genuinely worth reading — not corporate content or hollow inspiration, but something that says something real.

STEP 1 — DEFINE THE CONTENT

I want to write a LinkedIn post. Here is what I want to share:

My point of view: [What I think / what I've observed / what I've learned]

The insight: [Why this matters or why it's worth saying]

Any story or example: [A specific situation or experience that illustrates the point]

My audience: [Who I'm writing for — industry, role, interest]

My usual voice on LinkedIn: [Formal / conversational / technical / personal]

Engagement goal: [Spark discussion / share expertise / build visibility / announce something]

Any example of a post I've written that I liked? [Paste it if so]

STEP 2 — DRAFT

Write the LinkedIn post.

Length: [Short (100-150 words) / Medium (200-250) / Long (300-400)]

Format: [Prose / Hook + body / Numbered points / Story structure]

Do not use: hollow inspiration, vague motivational filler, "I'm excited to announce", "game-changing", or any phrase that reads like AI wrote it.

First line must make someone stop scrolling.

STEP 3 — TEST THE OPENING

Read only the first sentence. Would someone who doesn't know me stop scrolling to read the rest? If not, write three alternative opening sentences that would.

VERIFICATION

Does this sound like you? Read it aloud. Would a real person in your network read this and think it was worth their time? Is the claim you're making one you can genuinely support?

CUSTOMISE

- If announcing something: lead with the "so what" — why should the reader care before you describe what's happening
- If sharing expertise: include one specific, non-obvious insight rather than general advice
- If writing about experience: use a specific story, not a generalisation

W10 — DRAFT A BLOG POST FROM NOTES

SITUATION

You have a topic, a rough idea, or a set of notes that you want to turn into a publishable blog post. The notes are not polished prose — they are raw material.

STEP 1 — EXTRACT AND STRUCTURE FROM NOTES

I want to write a blog post. Here are my rough notes: [paste notes]

Please extract from these notes:

- The central argument or insight
- The main supporting points
- Any examples or evidence
- What is missing that would strengthen the post

Then propose a structure: what order should these ideas go in to make the strongest post?

STEP 2 — DEFINE THE POST

Using the proposed structure, I want you to draft a blog post with these specifications:

Target reader: [Who reads this – their background, what they care about]

Publication: [Where it will be published]

Tone: [my previous post example / formal-but-readable / conversational / authoritative]

Length: [word count]

Purpose: [Inform / Persuade / Entertain / Share experience]

Draft the post section by section. Start with the introduction – I'll review before you continue.

STEP 3 — OPENING AND CONCLUSION

The opening paragraph and the closing paragraph are the two most important.

Opening: Replace the current opening with three alternative versions. One should open with a question, one with a specific story, one with a direct statement of the central argument.

Closing: The conclusion should end with something the reader can do or think about – not a summary. Rewrite it with a concrete, specific takeaway.

VERIFICATION

Is the central argument clear? Has anyone made this point before in a nearly identical way — and if so, does your version add something? Are any specific claims, statistics, or examples accurate? Is the reading level and vocabulary appropriate for your audience?

W11 — REPURPOSE CONTENT ACROSS FORMATS

SITUATION

You have a piece of content — a presentation, a report, an article, a talk — that contains value you want to extract and publish in different formats.

STEP 1 — CONTENT INVENTORY

I have a [content type] that I want to repurpose. [Attach or paste the content.]

Please identify:

- The three to five most valuable insights or ideas in this content
- The most interesting or surprising point
- Any statistics or specific evidence that would stand alone
- Any stories or examples that would work as standalone content

STEP 2 — FORMAT PLANNING

I want to create content for:

[List the formats: LinkedIn post / email newsletter / short video script / Twitter/X thread / FAQ document / How-to guide / Slide deck / etc.]

For each format I've listed, suggest which part of the original content would work best and how it should be adapted.

STEP 3 — EXECUTE BY FORMAT

Let's start with [first format].

[Using the relevant content identified in Step 2, apply the appropriate workflow for that format – W9 for LinkedIn, W10 for a blog post, etc.]

VERIFICATION

Repurposed content should be adapted, not just repackaged. Verify that each piece works as a standalone without requiring the original context. Check that any statistics or claims retain their accuracy in the new format (statistics can be misleading when stripped of context).

W12 — WRITE WEBSITE COPY

SITUATION

You need to write copy for a website page — a homepage, a services page, an about page, or a landing page — that converts readers into enquiries or customers.

STEP 1 — DEFINE THE PAGE

I need to write copy for a [page type] for [business/product name].

Business description: [What the business does, for whom, and what makes it different]

Page purpose: [What this page should make the reader do or feel]

Target visitor: [Who comes to this page, what they're looking for, what their concern or question is]

Main message: [The single most important thing this page should communicate]

Tone: [Professional / Friendly / Premium / Practical / Specialist]

Any examples of websites with copy we admire: [links or descriptions]

Please draft this page.

STEP 2 — TEST THE COPY

Read this page as a sceptical first-time visitor who has never heard of [business name] and is comparing several options.

- Does the headline make it immediately clear what this is for?
- Is the value proposition obvious within the first ten seconds?
- Are there any claims that a sceptical visitor would question?
- Is there a clear and compelling reason to take the next step?
- What is the main thing that would make someone leave this page?

STEP 3 — HEADLINES

The headline is the most important element. Write ten alternative headlines for this page. Range from direct to intriguing to benefit-led to question-based. I'll choose the strongest or combine elements from several.

VERIFICATION

Does the copy accurately represent what the business actually does? Are any claims (testimonials, statistics, guarantees) genuine and verifiable? Have a target customer read it cold and ask them: "What do you think this company does? Would you contact them?"

RESEARCH AND LEARNING

W13 — RESEARCH A TOPIC YOU KNOW NOTHING ABOUT

SITUATION

You need to understand a topic you have little or no background in — a new sector, a technical subject, a policy area, a market — well enough to form a view or make a decision.

STEP 1 — ESTABLISH ORIENTATION

I need to understand [topic]. I currently know [your current level – almost nothing / the basics / a little].

Please give me:

1. A plain-English overview: what this is, why it matters, who is involved
2. The key terminology I need to know to read about this intelligently
3. The main debates, disagreements, or open questions in this area
4. What the most important things I should read or watch to go deeper

Pitch this at someone who is intelligent but not yet knowledgeable about this topic.

STEP 2 — IDENTIFY YOUR SPECIFIC NEED

My specific reason for learning about this: [Why you're researching this – decision, project, meeting, writing, general interest]

Given this context, what do I most need to understand? What can I safely deprioritise?

STEP 3 — TEST UNDERSTANDING

Now ask me five questions about what we've just covered. Start simple and get more complex. For each question I answer, tell me whether my understanding is correct and where I should think more carefully.

VERIFICATION

AI summaries of complex topics can oversimplify or — particularly for technical, scientific, or legal subjects — miss important nuance. Use this as a starting point, not a concluding source. For any topic where you will make significant decisions or public claims, follow up with primary sources. Check the AI briefing for any claims that seem surprising — they may be wrong.

CUSTOMISE

- For legal or regulatory topics: explicitly note "flag any areas of ongoing debate or uncertainty" and add "I'll need to verify the current regulatory position from official sources"
 - For scientific topics: ask for the quality of evidence behind major claims
 - For market or commercial topics: ask AI to flag the uncertainty inherent in market estimates
-

W14 — UNDERSTAND A COMPLEX DOCUMENT

SITUATION

You have a document — a contract, a policy, a technical report, a legal opinion — that you need to understand but that is written in language that does not make understanding easy.

STEP 1 — INITIAL READ

Please read this document. [Attach document.]

First pass:

- What type of document is this?
- What is its purpose?
- What are the main sections?
- Who are the parties or stakeholders involved?

Then: what does this document actually mean in plain English? Give me a summary a non-specialist could act on.

STEP 2 — SPECIFIC QUESTIONS

I want to understand the following specific aspects of this document:

1. [Your first specific question – e.g., "What are our obligations under section 4?"]
2. [Second question]
3. [Third question]

For each, tell me: what does the document say, what does it mean in plain language, and is there anything ambiguous or unclear about it?

STEP 3 — RISK AND ACTION IDENTIFICATION

From a practical standpoint:

- What are the most important obligations or deadlines in this document?
- What are the main risks this document creates or addresses?
- What actions do I or my organisation need to take as a result of this document?
- Are there any clauses or provisions that seem unusual or that I should take advice on?

VERIFICATION

For contracts, legal opinions, and regulatory documents: the plain-English summary is an aid to understanding, not a legal substitute. Any action taken on the basis of a document should reference the original text. For significant legal or financial matters, consult a qualified professional.

W15 — COMPARE OPTIONS BEFORE A DECISION

SITUATION

You need to choose between options — suppliers, approaches, products, strategies — and want a structured comparison to inform the decision.

STEP 1 — DEFINE THE DECISION

I need to compare [options] to [make a decision].

Context: [Why this decision matters and what's at stake]

Decision-maker: [Who will make the final decision and what they care about]

Options to compare: [List them]

Evaluation criteria: [What matters – list them in rough order of importance]

Constraints: [Budget, timeline, technical requirements, regulatory requirements]

Information available: [What you know about each option – attach documents if relevant]

STEP 2 — STRUCTURED COMPARISON

Compare [options] against the criteria I've specified.

Format: A table where rows are evaluation criteria and columns are options. For each cell, provide a brief assessment (not just a rating – a sentence or two of substance).

After the table: a one-paragraph narrative on each option summarising its main strengths and main weaknesses.

STEP 3 — RECOMMENDATION AND CHALLENGE

Based on the comparison, what would your recommendation be? What is the main argument for each of the other options over the recommended one?

Then: challenge the recommendation. What would have to be true for a different option to be the right choice?

VERIFICATION

Any specific claims about each option should be verified against the options' own documentation. AI comparisons are only as good as the information provided — if you have incomplete information about one option, flag this explicitly in the decision

record. For significant procurement decisions, this structured analysis is a useful starting point, not a substitute for due diligence.

W16 — LEARN A NEW SKILL FASTER

SITUATION

You want to learn something — a tool, a technique, a process, a domain — and want to use AI to accelerate the learning.

This is not limited to office or professional-services work. A teacher learning a new classroom tool, a healthcare worker learning a documentation process, or a creative professional learning unfamiliar software can use the same workflow — provided the current training, safety, or professional requirements are supplied and checked.

STEP 1 — MAP THE SKILL

I want to learn [skill/topic/tool].

My current level: [Complete beginner / Have the basics / Some intermediate experience]

My purpose: [Why I want to learn this — be specific about what I want to be able to do]

My constraints: [Time available, how I learn best, any specific challenges]

Please give me:

1. A clear learning path — what do I need to learn first, second, third?
2. The concepts that are most important to understand early (getting these wrong means struggling later)
3. The most common mistakes beginners make
4. How long should each stage realistically take?

STEP 2 — ACTIVE PRACTICE DESIGN

Design a practice plan for the next [time period] that builds this skill through doing, not just reading.

Each practice exercise should:

- Be achievable in [time per session]
- Build on the previous exercise
- Produce something I can assess myself

Include: what "good enough to move on" looks like for each stage.

STEP 3 — TEACH ME AND TEST ME

Let's work through the first concept: [topic from learning path].

Explain it clearly. Then test my understanding with a question. When I answer, tell me whether I'm right, where I've misunderstood, and what I should focus on.

Continue until I can explain this concept back to you correctly in my own words.

VERIFICATION

AI-generated learning plans may miss important nuance, especially for professional skills where regulatory, safety, or ethical dimensions matter. For professional certifications, verify that the learning path aligns with the actual certification curriculum. For technical skills, test your understanding through actual practice, not just by being able to answer AI's questions.

PERSONAL PRODUCTIVITY

W17 — PLAN A PROJECT

SITUATION

You have a project to deliver — at work or in your personal life — and you want to plan it properly before starting.

STEP 1 — DEFINE THE PROJECT

I need to plan a project. Here is what I know:

What the project produces: [The end deliverable]

Purpose: [Why this project matters / what problem it solves]

Deadline: [When it must be complete]

Resources: [What time, budget, people, or tools are available]

Constraints: [What you cannot do / what is fixed]

Dependencies: [What you're waiting for / what others are doing]

Risks: [What might go wrong]

Key stakeholders: [Who needs to be kept informed or who can block this]

STEP 2 — BUILD THE PLAN

Please develop a project plan that includes:

1. Work breakdown: what are all the tasks required to deliver this?
2. Sequence: which tasks depend on which?
3. Timeline: working back from the deadline, when does each task need to happen?
4. Critical path: which tasks, if delayed, will delay the whole project?
5. Milestones: what are the key checkpoints?

Present as a structured plan I can use to track progress.

STEP 3 — RISK AND CONTINGENCY

For this project, what are the three most likely things to go wrong? For each, what is the early warning sign, and what should I do if it happens?

VERIFICATION

Have someone with relevant experience review the plan — AI planning is good at structure but may miss domain-specific risks or dependencies. Verify timelines are realistic against your actual experience with similar work.

W18 — PROCESS A BACKLOG OF TASKS

SITUATION

You have an overwhelming list of tasks, projects, emails, or obligations and need to bring order to it.

STEP 1 — CAPTURE EVERYTHING

I'm going to paste in everything that is on my to-do list, in my inbox, or on my mind as outstanding. [Paste raw list – it doesn't need to be organised.]

First pass: read through this and tell me:

- How many distinct things are here?
- Are there any that are really the same thing?
- Are there any that seem more urgent than I might realise?

STEP 2 — TRIAGE

Now help me triage these tasks:

For each task, categorise it:

1. URGENT & IMPORTANT: Must happen this week
2. IMPORTANT but not urgent: Schedule for this month
3. URGENT but not important: Could be delegated or handled quickly
4. NEITHER: Question whether this needs to happen at all

Flag any task where you can't categorise it and tell me why.

STEP 3 — WEEK PLAN

Based on what's in Category 1 (urgent and important), help me plan this week.

My available time this week is approximately [hours/days].

Fixed commitments already in the calendar: [list them]

Which of the Category 1 tasks should happen which day? What is a realistic week that is ambitious but achievable?

VERIFICATION

Review the prioritisation against your own judgment — AI may prioritise based on what seems externally important without knowing what your actual professional context makes most critical. Check whether any deadline-driven items have been correctly identified as urgent.

W19 — DRAFT A DIFFICULT PERSONAL MESSAGE

SITUATION

You need to write something personal that is emotionally complex — an apology, a message to someone you've lost touch with, a difficult family communication, a message declining an invitation in a way that preserves the relationship.

STEP 1 — ESTABLISH THE SITUATION

I need to write a personal message and I want to get it right.

Situation: [What has happened and why this message is necessary]

Recipient: [Who this is to — relationship, history]

What I want to achieve: [The outcome I'm hoping for]

What I'm worried about: [How it might go wrong]

My honest feelings: [As much as you want to share — this helps AI understand the emotional register]

Tone required: [Warm / Apologetic / Careful / Direct / Affectionate]

Please draft this message.

STEP 2 — TONE CHECK

Read this message from the recipient's perspective.

- How might they feel receiving this?
- Does it say what I intended, or could it be read differently?
- Is there anything that might land wrong even though it wasn't intended to?
- Is the length right for the situation?

STEP 3 — ALTERNATIVES

Write two alternative versions of this message:

- One that is slightly more direct
- One that is slightly warmer

I'll choose elements from whichever feels most right.

VERIFICATION

Read the final draft aloud. Does it sound like you? Would the recipient recognise your voice in it? If the message is truly important, sleep on it before sending — AI drafts can be technically good but emotionally wrong in ways that a night's consideration reveals.

W20 — PREPARE FOR A JOB INTERVIEW

SITUATION

You have a job interview — internal or external — and want to prepare seriously rather than just hoping for the best.

STEP 1 — RESEARCH

I have a job interview for the role of [title] at [organisation].

Please help me prepare.

What I know about the organisation: [paste or link available information]

The job description: [attach or paste]

My relevant experience: [brief summary]

My main strengths for this role: [your view]

My main concerns about this role: [what you're less sure about]

Based on the job description and organisation, what should I know before this interview?

STEP 2 — ANTICIPATE QUESTIONS

Based on the role and organisation, what are the fifteen most likely questions I will be asked? Include:

- Questions about my experience and background
- Questions about why I want this role
- Competency-based questions ("tell me about a time when...")
- Questions about my knowledge of the organisation
- Difficult or challenging questions I should be ready for

For each question, note what the interviewer is likely trying to learn from it.

STEP 3 — PREPARE ANSWERS

Let's prepare my answers to the five hardest questions from the list.

For each, I'll give you my rough answer. You will:

- Identify what is strong about it
- Identify what is missing or could be clearer
- Suggest how to make it more compelling

Let's start with: [first hard question]. My rough answer: [your answer]

STEP 4 — QUESTIONS TO ASK

What are the best questions I could ask the interviewer? Suggest five to eight questions that:

- Demonstrate I've prepared and thought about the role
- Help me genuinely assess whether this is the right role for me
- Are not questions whose answers are on the website

Avoid: questions that seem presumptuous, overly salary-focused, or that suggest I haven't done basic research.

VERIFICATION

Verify all company facts before the interview — AI's information may be outdated. Check the company's website, recent news, and LinkedIn for current information. Practise answers aloud — answers that read well often sound unnatural when spoken. Time your answers to key questions.

HOME AND LIFE

W21 — PLAN A TRIP

SITUATION

You want to plan a trip — city break, holiday, business travel — and want more than a generic list of tourist attractions.

STEP 1 — DEFINE THE TRIP

I'm planning a trip and want help making it as good as possible.

Destination: [Where]

Dates/duration: [When, how long]

Who is travelling: [Adults/children/ages, special requirements]

Type of trip: [Relaxation / Culture / Adventure / City exploration / Family / Romantic / Other]

Interests: [What you genuinely enjoy]

Things you want to avoid: [Tourist traps, certain types of activity, etc.]

Budget level: [Budget / Mid-range / Premium]

Existing bookings: [What's already confirmed]

What you've already decided to do: [Any fixed elements]

STEP 2 — PERSONALISED RECOMMENDATIONS

Based on this, give me:

1. The things in [destination] that genuinely match my interests — not the standard tourist list
2. The things most people do that I might want to skip given what I've told you
3. Practical considerations specific to this trip that I should think about in advance

STEP 3 — ITINERARY

Draft a day-by-day itinerary for [number of days].

Build it around my interests and priorities, not a generic tourist circuit.

Be realistic about travel time between locations.

Flag anything that needs advance booking.

Leave space — don't over-schedule.

VERIFICATION

Verify opening hours, prices, and booking requirements from current official sources — AI information may be outdated. For international travel, check current visa, health, and entry requirements from official government sources. Book anything time-sensitive or popular in advance.

W22 — UNDERSTAND A CONTRACT OR LEGAL DOCUMENT

SITUATION

You have received a contract, terms and conditions, rental agreement, or legal document that you need to understand before signing or acting on.

STEP 1 — PLAIN-ENGLISH SUMMARY

Please read this [document type]. [Attach document.]

Give me a plain-English summary:

- What is this document?
- What does it obligate me to do?
- What does it obligate the other party to do?
- What are the most important terms?
- When does it begin and end?

STEP 2 — RED FLAG IDENTIFICATION

Now read this document more carefully for potential concerns.

Flag:

- Any unusual clauses that differ from what you'd typically expect
- Any terms that limit my rights significantly
- Any financial obligations beyond the obvious ones
- Any automatic renewal or escalation clauses
- Any clauses where my obligations are unclear or broad
- Anything I should ask the other party to clarify or change before signing

STEP 3 — QUESTIONS TO RAISE

Based on what you've found, what specific questions should I raise with [the other party / my solicitor / the landlord] before signing? List them in order of importance.

VERIFICATION

This analysis is a tool for understanding and preparation, not legal advice. For any contract involving significant financial commitment, important rights, or professional risk, consult a qualified solicitor. AI can misread or oversimplify legal language. Always read the actual contract before signing — do not rely solely on an AI summary.

W23 — GET A SECOND OPINION ON A HEALTH QUESTION

SITUATION

You have a health question — a symptom you want to understand, a diagnosis you've been given and want to learn more about, a medication you want to understand better — and want to use AI as a starting point for a better-informed conversation with your doctor.

STEP 1 — STATE THE QUESTION CLEARLY

I want to understand [health topic] better so I can have a more informed conversation with my doctor.

Context: [Brief background — age, general health context if relevant, what has already been discussed with a healthcare professional]

My specific question: [What I want to understand]

What I already know: [Any information you've already received from a healthcare professional]

Please explain [topic] in plain language. This is to help me understand, not to replace medical advice.

STEP 2 — QUESTIONS FOR YOUR DOCTOR

Based on what you've explained, what would be the most useful questions for me to ask my doctor? I want to:

- Understand my situation better
- Know what my options are
- Understand what I should be monitoring or watching for
- Know when I should seek urgent attention

VERIFICATION

AI is not a substitute for professional medical advice. For any health question with significant consequences, speak to a qualified healthcare professional. AI information about health topics may be incomplete, out of date, or not applicable to your specific situation. Never adjust medication, ignore professional advice, or make significant health decisions based solely on AI output.

WRITING AND EDITING

W24 — EDIT AND IMPROVE YOUR OWN WRITING

SITUATION

You have written something — an article, a report, a proposal, an email — and want to improve it before it goes out. You know it could be better but aren't sure exactly how.

STEP 1 — DIAGNOSTIC READ

Please read this piece of writing. [Attach or paste the document.]

Give me a diagnostic before you suggest any changes:

- What is the main argument or purpose of this piece?
- What is working well?
- What are the main weaknesses?
- Where did you lose interest or find the writing unclear?

Do not make any changes yet – just diagnose.

STEP 2 — TARGETED EDITING

Now, based on the diagnosis, help me improve this piece.

Priority: [Which weakness is most important to fix first?]

Please:

1. Rewrite the weakest paragraph – show me what better looks like
2. Suggest structural changes if needed – should anything be moved, cut, or added?
3. Flag any sentences that are unclear or that say less than they should

Do not over-edit – preserve my voice. Make it better, not different.

STEP 3 — LINE EDITING

Now do a line edit pass on this document.

Look for:

- Sentences that are too long or too complex
- Passive voice where active would be stronger
- Filler words and phrases that add length without meaning
- Any word used repeatedly that should be varied
- Anything that could be expressed more precisely

For each change: show the original and the improved version, and note briefly why the change helps.

VERIFICATION

Read the final edited version aloud. Does it still sound like you? Has AI edited out your voice along with your weaknesses? It is better to keep a few imperfections that are genuinely yours than to have a technically polished piece that reads like it was written by someone else.

W25 — TURN ROUGH NOTES INTO A POLISHED DOCUMENT

SITUATION

You have rough notes, a voice memo transcript, scattered thoughts, or a messy draft that needs to become a proper document. The content is there — it just needs structure and polish.

STEP 1 — MAKE SENSE OF THE MATERIAL

I have rough material that I need to turn into [document type]. [Attach or paste notes.]

First pass: help me understand what I have.

- What are the main ideas in this material?
- What is the implied argument or narrative?
- What is well-developed versus what is sketchy?
- What is missing that the final document will need?

STEP 2 — BUILD THE STRUCTURE

Propose a structure for [document type] based on what's in the notes.

I need the document to:

- Achieve: [purpose]
- For: [audience]
- In approximately: [word count or length]

For each proposed section, note: what content from my notes goes here, and what will need to be added.

STEP 3 — DRAFT

Now draft [document type] section by section using my notes as source material.

Preserve my ideas and voice. Add structure and polish – do not add new ideas or content I haven't indicated.

Start with [first section]. I'll review before you continue to the next.

STEP 4 — FINAL PASS

Now read the complete draft as a whole. Does it flow? Does each section lead naturally to the next? Is there anything that still sounds like rough notes rather than finished prose? Produce a final clean version.

VERIFICATION

Read the final document as if you were reading it for the first time. Is anything missing that you intended to include? Has anything been included that you didn't intend? Does the document accurately represent your thinking, or has it drifted toward something slightly different in the polishing process?

This completes the twenty-five real-world workflows. The library is designed to be used selectively — find the workflow that matches your situation, adapt the placeholders to your specifics, and use the verification steps before trusting the output.

Part VI covers the habits and practices that make these workflows reliable: how to verify output, what AI consistently gets wrong, and how to build an AI practice that serves you well over time.

The AI User's Playbook

The twenty-five workflows in Part V are only as good as the habits that surround them. A technically perfect workflow can still produce embarrassing or harmful output if the result goes unverified. A strong AI practice is built on something less glamorous than clever prompts: the discipline of knowing when to trust output, what to check, and how to catch errors before they matter.

Part VI is about that discipline. It is, in many ways, the most important part of the book.

CHAPTER 20

How to Verify AI Output

THE CONFIDENCE PROBLEM

AI writes with the same fluent confidence whether it is correct or incorrect.

This is worth pausing on. When a human expert is unsure, they usually show it: they hedge, they qualify, they say "I think" or "you'd want to check this." When AI is unsure, it typically does not. It produces the same well-structured, authoritative-sounding prose it would produce if it knew exactly what it was talking about.

This is not dishonesty. It is a feature of how AI generates text: it produces the most plausible continuation of the words before it, and confident prose is more plausible than tentative prose in most contexts. But the practical effect is that you cannot use prose confidence as a signal of factual accuracy. The two are disconnected.

This means verification is not optional for important work. It is a habit as fundamental as saving your documents.

WHAT VERIFICATION IS — AND IS NOT

Verification is not re-reading the AI output more carefully.

If the output says "the company was founded in 1987," reading it again more carefully will not tell you whether 1987 is right. It will only tell you that AI says 1987 with apparent confidence.

Verification means checking claims against sources that are independent of the AI output: the original document, the company website, a primary source, a qualified professional, your own direct knowledge of the situation.

Verification is proportional to stakes. A quick summary of a document you will read anyway needs a lighter verification pass than a public statement containing statistics that will be published under your name.

THE VERIFICATION FRAMEWORK

Before trusting and using any significant AI output, work through these eight questions:

1. IS THIS FACTUAL OR CREATIVE?

Factual content (dates, names, statistics, company information, legal positions, scientific findings) requires verification against sources. Creative content (writing style, structure, framing) can be assessed on its own merits without external sources.

Know which type of output you have. A report that reads like authoritative factual analysis but was written by AI with no source documents is creative content masquerading as factual — and requires the same verification as factual content.

2. DOES IT CONTAIN SPECIFIC NUMBERS?

Numbers are a high-risk category. AI is confidently wrong about numbers with greater frequency than almost any other category. Any specific statistic, date, financial figure, percentage, or calculation deserves independent verification before you use it.

The test: can you find this number in a source you can point to?

3. DOES IT CITE SOURCES?

If AI has cited sources, verify them. AI hallucination of citations is a well-documented and persistent failure mode. AI may produce a citation that looks real — a plausible author name, a plausible journal, a plausible year — for a paper that does not exist. Verify by finding the actual source and confirming the claim attributed to it is what the source actually says.

Do not assume that the presence of citations means the claims are verified.

4. ARE THE SOURCES REAL?

This is distinct from the previous question. Even if AI has provided real source titles or URLs, verify they are what they claim to be and that the content of the source supports the claim AI made.

5. COULD THIS INFORMATION HAVE CHANGED?

AI has a knowledge cutoff. Information about current prices, regulations, leadership, products, statistics, and events may be outdated. For any fast-moving area — technology, regulation, markets, current events — check whether the information reflects the current situation.

Even with web search enabled, AI may not surface the most recent information or may misread what it finds.

6. DID AI MAKE AN ASSUMPTION?

Complex prompts often involve AI filling gaps you didn't explicitly address. AI will do this without flagging it unless you specifically ask.

Ask: is there anything in this output that I didn't specify, and has AI assumed correctly? The missing context from Chapter 6 can produce plausible-but-wrong assumptions that make their way into the output.

A useful check: ask AI directly — "What assumptions did you make in producing this output?"

7. DID IT FOLLOW MY CONSTRAINTS?

Check that the output respects the specific constraints you gave. Does it stay within the stated word count? Does it avoid the topics you said to avoid? Is it written for the audience you specified? Is the tone what you asked for?

Constraints are sometimes partially honoured. AI may follow most of your instructions and quietly ignore one. Read explicitly for compliance.

8. WHAT EVIDENCE SUPPORTS THE CONCLUSION?

For outputs that reach conclusions, make recommendations, or offer analysis: is the reasoning sound? Does the conclusion follow from what was provided? Could the same evidence support a different conclusion?

This is the hardest verification task and the one most commonly skipped. It requires judgment, not just source-checking.

VERIFY

BEFORE YOU TRUST THE ANSWER

Ask:

- Is this factual or creative?
- Does it contain numbers?
- Does it cite sources?
- Are the sources real?
- Could the information have changed?
- Did the AI make an assumption?
- Did it follow my constraints?
- What evidence supports the conclusion?

HIGH-RISK OUTPUT CATEGORIES

Some types of AI output require particularly rigorous verification:

STATISTICS AND DATA

Verify every number against a primary source. AI confabulates plausible statistics with notable regularity. The figure "studies show that 73% of employees..." is a common pattern — verify that the study exists, that the figure is correctly stated, and that it applies to the context in which you're using it.

NAMES AND ATTRIBUTIONS

Quotes attributed to named individuals should be verified. AI can generate plausible-sounding quotes from real people that those people never said. This is a category where being wrong can cause significant professional or legal damage.

LEGAL AND REGULATORY INFORMATION

Law changes. Regulations change. Jurisdictional variations are significant. Any legal or regulatory claim should be verified against current official sources. AI may confidently state a legal position that was correct two years ago and is not correct now.

MEDICAL INFORMATION

As above, and with additional stakes. Verify medical information from qualified professionals and current clinical sources.

TECHNICAL SPECIFICATIONS

Product features, software capabilities, technical standards — all change and all deserve verification from current official documentation.

LOWER-RISK OUTPUT CATEGORIES

Verification should be proportionate. These categories typically need lighter verification:

- Structural and organisational decisions (how to order a document, what sections to include)
- Tone and style adjustments
- Creative framing and presentation

- General conceptual explanations of stable topics
- Grammar, clarity, and prose quality

Even here, read critically. Low-risk doesn't mean no-risk.

REAL-WORLD EXAMPLE

Margaret, a retired professional who has taken up writing local history articles, uses AI to help research a piece about a local Victorian factory. AI produces a confident account of the factory's founding, its owner, and the number of workers it employed at its peak.

She runs the verification checklist. The founding date differs from one she has in a local archive document. The owner's name is slightly different from the name she has found in a census record. The employment figure appears plausible but has no source she can trace.

She uses the AI output as a framework and a starting point — it helpfully organised the narrative structure and identified what she should look for. But every specific fact goes back to a source she can cite. The final article is accurate because she verified; it was well-structured because AI helped.

This is the right relationship between AI production and human verification.

BUILDING THE VERIFICATION HABIT

Verification is most useful when it is habitual rather than occasional. The goal is to make it automatic: before you share, publish, send, or act on significant AI output, you have a brief, consistent check.

The check does not need to be long. For most outputs, running through the eight questions takes two minutes. For high-stakes outputs, it takes longer. The discipline is applying it consistently, not applying it intensively.

The section that does most harm is the one you verified every part of except the one part that was wrong.

WATCH OUT

Asking AI to verify its own output has limited value. AI can tell you what it believes to be accurate, but this is circular — if AI believed the claim was wrong, it wouldn't have made it in the first place. AI self-review can catch logical inconsistencies and obvious errors, but it cannot substitute for checking claims against independent sources. Use AI critique for structure and reasoning; use external sources for facts.

CHAPTER 21

What AI Gets Wrong (Consistently)

THE FAILURE TAXONOMY

AI failures are not random. They fall into predictable categories that appear consistently across different models, different platforms, and different types of task. Knowing these categories means you can anticipate where to look, rather than discovering errors after they've caused a problem.

This chapter is a practical catalogue of known AI failure modes — not to make you sceptical of AI, but to make your scepticism accurate: applied where the risks are highest, relaxed where the risks are lower.

FAILURE CATEGORY 1 — HALLUCINATION

What it is: AI generates specific facts, names, dates, statistics, citations, or events that did not happen or do not exist — but presents them with the same confidence as accurate information.

Why it happens: AI generates plausible text, not verified text. A hallucinated statistic is generated because it is the kind of statistic that would plausibly appear in this type of document. A hallucinated citation is generated because it is the kind of citation that would plausibly appear in academic writing on this topic.

Where it appears most:

- Citations and references (highly specific hallucinations)
- Statistics and percentages (plausible but invented figures)
- Biographical facts about real but less prominent people
- Historical detail about events AI has limited training data on
- Company-specific information AI wasn't trained on

How to catch it: Verify specific claims against primary sources. Treat all citations as unverified until confirmed. Be especially sceptical of very specific numbers.

The severity: High for professional use. A hallucinated citation in a published article, a made-up statistic in a board presentation, or an incorrect legal date in a contract review — these cause real professional and sometimes legal consequences.

FAILURE CATEGORY 2 — OUTDATED INFORMATION

What it is: AI presents information that was accurate as of its training cutoff but is no longer current.

Why it happens: AI's knowledge is frozen at the point of training. The world continues to change. Models released in 2025 may have training data that lags by twelve to eighteen months before release, and then be used for years after.

Where it appears most:

- Technology product details (features, pricing, availability)
- Regulatory and legal information
- Organisational information (who holds which role, company structure)
- Market data and economic statistics
- Current events and recent developments

How to catch it: For any fast-moving topic, verify current information from official sources. Even with web search enabled, verify important current-state claims.

The severity: Medium to high, depending on context. Outdated legal information is high-severity. A slightly outdated technology feature description is lower.

FAILURE CATEGORY 3 — CONFIDENT OVER-GENERALISATION

What it is: AI produces statements that are broadly plausible but overstate how generally applicable they are — missing important exceptions, jurisdictional variations, sector-specific differences, or situational nuances.

Why it happens: AI generates the most generally accurate response to the most common interpretation of the question. It doesn't automatically surface the ways in which the answer differs for your specific situation.

Where it appears most:

- Legal questions with jurisdictional variation ("employment law says X" without specifying which country's)
- Medical questions with individual variation ("the standard treatment is X")
- Business advice without sector or context sensitivity
- Technical guidance where the "standard" approach isn't appropriate for specific circumstances

How to catch it: Ask explicitly: "In what situations might this advice not apply? What are the important exceptions?" For professional matters, verify against advice calibrated to your specific jurisdiction and circumstances.

FAILURE CATEGORY 4 — ARITHMETIC ERRORS

What it is: AI makes errors in arithmetic, calculation, and numerical reasoning — particularly on multi-step calculations, percentage calculations, and anything requiring exact arithmetic.

Why it happens: AI generates text. Numbers are text. When AI produces "the answer is 43%," it has generated the number 43 because it is plausible in context — not because it has calculated

and arrived at 43. Text generation and mathematical calculation are different processes. Without code execution, AI maths is unreliable.

Where it appears most:

- Percentage calculations in reports or analyses
- Financial modelling or estimation
- Multi-step arithmetic
- Unit conversions
- Time-based calculations (especially with complex date arithmetic)

How to catch it: Check all arithmetic independently. Use code execution when available for anything that matters. Never use an AI-generated financial figure without verifying the arithmetic.

The severity: High when numbers appear in decisions, reports, or presentations. A wrong percentage in a financial summary is a professional error regardless of where it came from.

FAILURE CATEGORY 5 — CONTEXT DRIFT

What it is: In long conversations, AI gradually loses track of constraints and context established early in the conversation, reverting to defaults it would apply without those constraints.

Why it happens: AI attends to the full conversation but with diminishing precision as length increases. Constraints stated in message one of a twenty-message conversation may be less firmly applied at message eighteen than they were at message two.

Where it appears most:

- Long editing sessions where style constraints gradually slip

- Extended research conversations where early framing becomes less influential
- Multi-stage workflows where the first stage's constraints should carry forward but don't

How to catch it: For long conversations, periodically re-state key constraints. When output seems to have drifted from what you specified, check the early conversation and explicitly re-anchor the constraint.

FAILURE CATEGORY 6 — SYCOPHANCY

What it is: AI tends to agree with and validate what the user presents, even when the user is wrong. It will often affirm an incorrect claim rather than challenge it, adjust its previous answer to accommodate the user's pushback even when its original answer was correct, and frame critique in ways that are so mild as to be useless.

Why it happens: AI is trained using human feedback, which tends to reward agreeable responses. This creates a bias toward agreement and validation.

Where it appears most:

- When you present a flawed plan and ask AI to review it
- When you push back on AI's correct assessment and it backs down
- When you ask AI to evaluate your own writing or thinking

How to catch it: Actively ask for challenge. Use adversarial critique prompts (Chapter 18). Ask "what is wrong with this?" not "what do you think of this?" When AI agrees with you completely, treat that as a flag worth checking, not a validation.

The severity: Moderate. Sycophancy is more a quality issue than a factual one, but it can lead to significant errors when you're looking for honest critique of important work.

FAILURE CATEGORY 7 — INSTRUCTION DRIFT

What it is: AI partially follows instructions rather than fully following them. It may honour the spirit of an instruction while missing a specific requirement, or follow all instructions except one.

Why it happens: Complex sets of instructions create competing pressures. AI optimises across them in ways that may not match your priority ordering. An instruction to "be brief" may be honoured by shortening paragraphs while maintaining a section structure you asked to be removed.

Where it appears most:

- Multi-constraint prompts with several specific requirements
- Editing instructions where some changes are made but others are missed
- Format specifications where the structure is approximately right but details are wrong

How to catch it: For important outputs, review explicitly against your stated constraints. Check each instruction in your prompt: was it followed? If not, restate it as the sole instruction in a follow-up.

FAILURE CATEGORY 8 — PLAUSIBLE BUT WRONG REASONING

What it is: AI produces reasoning that appears sound — logical structure, coherent argument, sensible-sounding progression — but reaches the wrong conclusion, misapplies a principle, or draws on an incorrect premise.

Why it happens: Plausible reasoning is generated in the same way as plausible text. A well-structured argument is a pattern AI has encountered many times and can reproduce convincingly. The quality of the reasoning structure does not guarantee the quality of the reasoning.

Where it appears most:

- Legal analysis
- Financial reasoning
- Causal analysis ("X happened because...")
- Risk assessment
- Strategic recommendations

How to catch it: For important analytical output, apply the verification question: "Does the conclusion actually follow from the premises?" Test the reasoning by looking for an alternative explanation or conclusion that the same evidence could support. Have a qualified person review reasoning in high-stakes domains.

PRO TIP

When you receive an output that matters, ask AI: "What are the three most important ways this analysis could be wrong? What assumptions am I relying on that might not hold?" This is not a perfect catch-all, but it surfaces failure modes AI can identify — which is often enough to direct your verification effort toward the right areas.

A PRACTICAL SUMMARY: WHERE TO BE MOST CAREFUL

OUTPUT TYPE	PRIMARY FAILURE RISK	VERIFICATION PRIORITY
Statistics and numbers	Hallucination, arithmetic error	HIGH — verify all figures
Citations and references	Hallucination	HIGH — verify all citations exist
Legal/regulatory claims	Outdated info, over-generalisation	HIGH — verify from current official sources
Medical information	Over-generalisation	HIGH — verify from clinical sources
Reasoning and analysis	Plausible-wrong reasoning, sycophancy	MEDIUM — check argument logic
Current events/market data	Outdated information	MEDIUM — verify recency
Technology specifications	Outdated information	MEDIUM — verify from official docs
Writing and structure	Instruction drift	LOW — read against your brief
Creative framing	Generally lower risk	LOW — assess on merits

CHAPTER 22

Calibrating Trust

TRUST AS A CALIBRATION PROBLEM

The wrong response to AI's failure modes is distrust of everything AI produces. The right response is calibrated trust: knowing which outputs to use confidently, which to verify before using, and which to treat as starting points only.

Calibration means your level of trust is proportionate to the actual risk profile of the output. A person who trusts everything AI produces equally — the sentence structure and the legal claim alike — will eventually rely on something they should have checked. A person who trusts nothing AI produces is throwing away most of the tool's value for no good reason.

THE TRUST CALIBRATION MATRIX

Two dimensions determine how much trust is appropriate:

Stakes — What happens if this output is wrong?

- High stakes: wrong output causes significant professional, financial, legal, or personal harm
- Medium stakes: wrong output causes embarrassment, rework, or moderate professional consequences
- Low stakes: wrong output is caught before any harm, or the consequences are easily reversible

Verifiability — How easily can you check this?

- Easy: you can verify against a source you have access to in minutes
- Moderate: verification requires more effort — finding sources, specialist input, research

- Hard: verification requires expertise or information you don't have easy access to

Use these two dimensions to calibrate your verification investment:

STAKES \ VERIFIABILITY	EASY TO VERIFY	MODERATE	HARD TO VERIFY
High stakes	Verify and document	Verify carefully, seek expert input	Do not rely on AI alone — use as one input among several
Medium stakes	Verify key claims	Spot-check critical elements	Flag uncertainty; present as AI-assisted
Low stakes	Read critically	Use with awareness	Use cautiously; review if consequences grow

BUILDING A PERSONAL TRUST CALIBRATION

Over time, you will build a sense of where AI performs reliably for your specific use cases and where it consistently needs correction.

This is worth tracking, at least informally. When you notice a consistent failure pattern — AI always overestimates timelines in project plans; AI never quite captures your organisation's formal register; AI gets financial calculations wrong — note it. That pattern becomes a standing verification rule in your practice.

When you notice consistent strengths — AI always produces excellent first-draft email structures; AI is consistently right about the logical order of a presentation; AI's summarisation of meeting notes is almost always accurate — you can extend trust in those areas and invest verification effort where it's most needed.

Calibration is personal and task-specific. It improves with experience.

REAL-WORLD EXAMPLE

David has now used AI consistently for six months across his consultancy work. He has noticed:

Reliable: Drafting client emails (he rarely changes more than a sentence). Structuring proposals. Summarising meeting notes from his own bullet-point notes (not from transcripts). Producing presentation outlines.

Needs checking: Any output that contains specific market data or statistics. Recommendations that require knowing something about a client he may not have fully briefed. Legal or regulatory claims about anything jurisdiction-specific.

Use as starting point only: Any financial projections or estimates. Technical specifications for systems he isn't deeply expert in. Anything where the quality of the output depends heavily on current information he hasn't provided.

His verification practice is proportionate: high effort where he knows errors are likely; lighter touch where AI is consistently reliable. He doesn't spend time verifying email structure and doesn't use AI financial projections without checking every number.

COMMUNICATING AI-ASSISTED WORK

In professional contexts, it is worth having a clear internal policy on how you represent AI-assisted work:

What to disclose: When AI has generated factual claims, analysis, or recommendations that appear in work presented to clients, stakeholders, or colleagues, it is good practice to note that AI tools were used in preparation and that claims have been verified. This is both honest and a natural quality signal.

What you remain responsible for: AI is a tool. You are accountable for the work. If an AI-generated figure in a report is wrong, "AI produced it" is not a professional defence. Your review process, your verification standard, and your decision to include it are yours.

The credibility calculation: In most professional contexts, "I used AI to help draft this and then carefully verified it" is now a normal and credible statement. "I relied on AI without checking" is not. The trust question is whether your process is sound, not whether AI was involved.

WATCH OUT

Trust miscalibration can go in both directions. Over-trusting AI produces errors that damage professional credibility. But over-distrusting AI — refusing to use well-verified AI output because it came from AI — is also a miscalibration. It adds unnecessary time and effort and fails to take advantage of a tool that works well in the right context. The goal is accurate calibration, not maximum scepticism.

CHAPTER 23

Building Your AI Habits

WHY HABITS MATTER MORE THAN TECHNIQUES

Every technique in this book is available to you from today. The 7 Questions, the Before → Better → Best approach, the critique loop, the verification framework — all of them can be applied immediately.

But techniques that aren't habitual are techniques that get used occasionally, in high-stakes situations, when you remember, and not used the rest of the time. The people who get the most consistent value from AI are not necessarily the most technically sophisticated. They are the people who have built consistent habits that apply the right practices automatically.

This chapter is about building those habits.

THE FOUR HABITS THAT MATTER MOST

From watching how people build effective AI practices, four habits account for most of the difference between occasional useful results and consistently good output:

HABIT 1 — PREPARE BEFORE YOU PROMPT

The single biggest quality leverage point is the thirty seconds spent applying the 7 Questions before starting a significant AI task. Most people skip this and type whatever is on their mind. Those who make it habitual find that they invest slightly more time upfront and substantially less time in revision.

Make it automatic: for any AI task that involves more than a few minutes of work, run the 7 Questions first. Even quickly. Even imperfectly. The discipline of asking "what am I trying to achieve?" before typing is worth far more than any clever prompt technique.

HABIT 2 — VERIFY BEFORE YOU USE

Establish a consistent minimum verification standard for any output you will use professionally. The specific standard depends on the stakes of your work, but the habit is non-negotiable: before you send, publish, present, or act on significant AI output, you have checked something.

For most professional use, the minimum is: read critically, check any numbers and specific facts, and confirm the output follows the constraints you set. This takes two to five minutes. Make it as automatic as saving a file.

HABIT 3 — BUILD YOUR WORKSPACE ITERATIVELY

Your AI workspace is not something you set up once and leave. It improves every time you notice a gap — every time AI gets something wrong that it would have got right with better standing instructions, every time you re-type context you've already established, every time you use a prompt that worked particularly well.

The habit: when something goes wrong that a better workspace would have prevented, fix the workspace. When you write a prompt that works especially well, save it. When you notice you're re-entering the same context repeatedly, persist it. Small iterative improvements to your workspace compound into significant efficiency gains over time.

HABIT 4 — CRITIQUE BEFORE YOU FINISH

Before finalising any significant piece of AI-assisted work, run it through a critique pass. This doesn't need to be formal — it can be as simple as asking AI "what is weak about this?" and reading the response with genuine openness.

People who skip this step send proposals with obvious structural weaknesses, publish articles with underdeveloped arguments, and submit reports with claims that don't hold up. People who make critique a habit catch these things before they matter.

BUILDING HABITS: THE PRACTICAL APPROACH

Knowing what the habits are is not the same as having them. Three approaches help:

ATTACH HABITS TO EXISTING TRIGGERS.

Link new AI habits to things you already do. Before opening ChatGPT for a task: run the 7 Questions. Before closing a document you've used AI to help produce: run the verification checklist. The existing action becomes the trigger for the new habit.

START SMALLER THAN FEELS NECESSARY.

A habit that takes thirty seconds and is done every time is more valuable than a thorough process that is done occasionally. Begin with the minimum viable version of each habit. Add complexity as it becomes automatic.

REVIEW PERIODICALLY.

Once a month, look at what you're using AI for and ask: is my practice serving me well? What am I getting that's consistently poor quality? What am I spending time verifying that I could build into my workspace? What technique from this book have I not tried yet that might help?

A WEEK IN THE LIFE OF A CALIBRATED AI USER

To make this concrete, here is what a week of integrated AI use looks like for someone with established habits:

Monday morning:

Preparing for a client review meeting. Opens the client's Project. Pastes in the latest project data and asks AI for a status summary. Reviews it against his own knowledge — catches one outdated figure. Uses the agenda workflow (W8). Both ready in thirty minutes.

Tuesday:

Needs to draft a challenging email to a supplier about a missed deadline. Runs the 7 Questions briefly. Uses W1. Reviews the draft — adjusts the tone. Sends within twenty minutes of starting.

Wednesday:

Writing a quarterly business report. Starts with the Information → Structure → Content pattern (Chapter 17). Reviews the structure before drafting. Runs a critique loop on the draft (Chapter 18). Verifies the two financial figures independently. Done in two hours rather than half a day.

Thursday:

Reviewing a contract. Uses W22. Checks the flag list AI produces against her solicitor's advice. Updates her standing workspace instructions to include a note about a contract clause type she should always flag.

Friday:

Reflecting on the week. One output needed correction — AI got a date wrong. Updates her personal verification checklist to add "verify dates in contract summaries." Small improvement. Compounds over time.

PRO TIP

The most powerful single thing you can do this week: set up your persistent instructions if you haven't already. Spend thirty minutes writing a context document that covers your professional context, communication preferences, standard audience, and key constraints. Paste it into every AI session this week and notice how much setup time you save. Then move it into your platform's persistent instructions setting so it applies automatically.

WHAT GOOD AI USE DOES NOT LOOK LIKE

It is worth being explicit about the patterns that look like productive AI use but aren't:

Prompt collecting without application. Building a library of prompts from the internet and using them once before returning to the same vague requests as before.

Volume without quality. Generating large quantities of AI output and editing it all manually because it's not quite right — faster to produce, slower overall because verification and correction eat the time savings.

Outsourcing judgment. Using AI to make decisions that require accountability. AI can inform decisions; it cannot make them on your behalf.

Ignoring failure patterns. Noticing that AI consistently gets something wrong and continuing to rely on it in that area without verification or adjustment.

Treating the first draft as final. The first AI output on a complex task is almost never the best it can be. One iteration — a critique loop, a refinement pass, a targeted revision — consistently improves it.

CHAPTER 24

Keeping Up Without Burning Out

THE PACE PROBLEM

AI capabilities are changing faster than most professional skills do. A technique that works well today may be less useful in six months because the models have improved and no longer need that scaffolding. A feature you relied on may be removed, replaced, or significantly altered. New capabilities that genuinely change what's possible appear with some regularity.

This pace is genuinely exhausting if you try to follow all of it. It is also largely unnecessary.

WHAT ACTUALLY CHANGES AND WHAT DOESN'T

The most important distinction for anyone trying to keep up: some things change rapidly, and some things don't.

THINGS THAT CHANGE RAPIDLY:

- Specific model capabilities (what a given version of ChatGPT, Claude, or Gemini can do)
- Platform features (what tools are available in which product at which price point)
- UI and interface details (where to find which setting)
- Best-in-class techniques for specific edge cases
- Pricing and availability

THINGS THAT CHANGE SLOWLY OR NOT AT ALL:

- The principles of good interaction design (context, instructions, examples, output specs)
- The 7 Questions framework — these questions are derived from human communication theory, not from AI quirks
- The verification mindset

- The workflow thinking
- The habit of preparing before prompting

The core framework in this book — the 7 Questions, the AI Interaction Stack, the verification habits, the workflow patterns — is designed to be durable. These are not tricks that work today because of a specific model behaviour. They are principles that reflect what it means to communicate a task clearly to a capable but uninformed agent. That doesn't change when the models improve.

What you need to update is the feature-level detail: which capabilities exist, what the product notes say, where volatile information lives. That is why this book puts product-specific detail in Appendix E rather than in the main text.

A SUSTAINABLE INFORMATION DIET

You do not need to read every AI newsletter. You do not need to watch every YouTube tutorial about the latest model release. A sustainable information diet for a professional AI user has three elements:

1. A QUARTERLY CHECK-IN ON MAJOR DEVELOPMENTS.

Once a quarter, spend an hour reading about what has changed. Look for: new capabilities that might affect your use cases, changes to platforms you rely on, anything that suggests a technique you use needs updating. This is sufficient for most professional users.

2. PLATFORM-SPECIFIC NOTIFICATIONS FOR THINGS YOU RELY ON.

If you use a specific AI tool for important work, follow that tool's official announcements. When it announces significant changes, read them. This catches the changes that affect your specific workflow.

3. PRACTICAL EXPERIMENTATION OVER PASSIVE CONSUMPTION.

The most efficient way to learn about AI improvements is to use them, not to read about them. When a new capability is announced that seems relevant to your work, try it on an actual task. Fifteen minutes of practical use teaches you more than an hour of reading reviews.

What to largely ignore: hype cycles, "the new model is 1000% better" content, and AI influencer output that is primarily optimised for engagement rather than practical value.

WHEN TO UPDATE YOUR PRACTICE

You need to update your practice when:

- A tool you rely on has changed significantly (feature removed, new capability worth using)
- You encounter a consistent failure mode that a new technique might address
- Your work situation has changed and your AI use should reflect that
- Something in your workspace is producing consistently poor output

You do not need to update your practice because:

- A new model was released
 - Someone on the internet claims a new technique is revolutionary
 - Your colleague uses AI differently than you
-

THE TRAJECTORY OF A SKILLED AI USER

The early period of AI use is typically characterised by inconsistency: some interactions work well, most don't, and the reasons feel opaque. This book is intended to take you out of that phase.

The middle period — where good habits are established and the core framework is internalised — is where most of the professional value lives. Consistent results, predictable quality, reliable verification. This is not exciting. It is useful.

The advanced period involves increasing sophistication in workflow design, more precise calibration of trust, and the ability to identify the right tool and technique for novel situations without referring back to a guide. It comes with accumulated practice, not with reading more AI content.

The trajectory is this: from opaque and inconsistent to systematic and reliable, then from reliable to sophisticated. Part I and Part II get you to systematic. Part III and Part IV get you to reliable. The habits in Part VI get you to sophisticated over time.

WHY IT WORKS

The people who develop the best AI practices are not the ones who read the most about AI. They are the ones who use it most consistently, reflect on what works, and gradually refine their approach. The 7 Questions and the verification habits give you a framework to do that reflection productively — so that each iteration of your practice is an improvement on the last.

A FINAL WORD

Learning to use AI well is a practical craft: techniques can be taught, but judgment develops through practice and reflection.

You now have the framework, the techniques, and the workflows. The rest is practice.

Apply the 7 Questions. Build your workspace. Design workflows that match your work. Verify before you trust. Iterate until it's right. Reflect on what worked. Update what didn't.

And, in time, it becomes less like following a framework and more like simply knowing how to work.

APPENDICES

Appendix A

The 7 Questions — Quick Reference Card

Print this. Photograph it. Keep it visible until the questions become automatic.

BEFORE EVERY SIGNIFICANT AI INTERACTION

- 1. What am I trying to achieve?** — State the real goal, not just the surface task.
- 2. What does AI need to know?** — Add the background, audience, constraints, and relationships that matter.
- 3. What information can I provide?** — Attach or paste documents, data, examples, and previous work.
- 4. What exactly should AI do?** — Use a clear action verb and define the scope.
- 5. What should the result look like?** — Specify format, length, tone, structure, and audience.
- 6. How will I check it?** — Identify the facts, numbers, assumptions, and reasoning that need review.
- 7. What should happen next?** — Decide whether this is the final result or the input to another stage.

THE INTERACTION STACK, COMPRESSED

Intent · Context · Instructions · Examples · Source material ·
Workspace · Tools · Workflow · Output · Verification · Iteration

BEFORE YOU TRUST THE OUTPUT

Factual or creative? Contains numbers? Citations present?
Sources real? Could the information have changed? Did AI make
an assumption? Were the constraints followed? Does the
conclusion follow from the evidence?

The full explanations appear in Chapters 2, 3, and 20. This page
is designed as a quick reference, not as a replacement for those
chapters.

Appendix B

The Before → Better → Best Worksheet

*Use this when you want to systematically improve a prompt.
Work through the questions before redrafting.*

YOUR CURRENT PROMPT

Write out your current prompt exactly as you usually send it:

[Your current prompt here]

STAGE 1 — DIAGNOSE THE GAPS

Answer these questions about your current prompt:

JOB CLARITY

What specific action have you asked AI to take?

If the answer is vague ("help me," "work on," "assist with") — rewrite with a clear verb and scope.

CONTEXT PROVIDED

What background have you given AI that it couldn't know on its own?

What is the most important thing AI doesn't know that would change its answer?

INFORMATION PROVIDED

What documents, data, or examples have you attached or pasted?

Is there source material that would make the output more accurate or specific?

OUTPUT SPECIFIED

Have you specified: format, length, tone, audience?

If any of these are unspecified, what is AI likely to default to — and is that what you want?

VERIFICATION PLANNED

What specifically might be wrong in the output?

How will you check it?

STAGE 2 — BUILD THE BETTER VERSION

Using your diagnosis, write the improved prompt:

ROLE (if relevant):

You are a [role] with expertise in [domain].

CONTEXT:

[Background information – situation, history, constraints, relationships]

TASK:

[Specific action verb] [specific deliverable] for [purpose].

SOURCE MATERIAL:

[Attached: / Using: / Based on:]

OUTPUT REQUIREMENTS:

Format: [prose / bullet points / table / structured document]

Length: [word count / page count / number of points]

Tone: [formal / professional-informal / conversational]

Audience: [who this is for and what they know]

CONSTRAINTS:

Do not: [explicit exclusions]

Must include: [required elements]

VERIFICATION REQUEST:

Please flag any claims you are uncertain about and any assumptions you have made.

STAGE 3 — THE BEST VERSION

After receiving output from the Better version:

WHAT WORKED WELL?

WHAT STILL NEEDS IMPROVEMENT?

WHAT FOLLOW-UP INSTRUCTION WOULD MAKE IT RIGHT?

WHAT SHOULD I ADD TO MY STANDING WORKSPACE TO AVOID THIS GAP NEXT TIME?

EXAMPLE: BEFORE → BETTER → BEST

STAGE	PROMPT	WHAT CHANGED
Before	"Help me with a presentation for our board."	Vague task, no context, no output specs
Better	"Create a presentation outline for our Q3 board update covering financial performance, key decisions, and outlook."	Specific task, content defined
Best	"You are preparing a board presentation for [Company]. Audience: eight non-executive directors, one of whom is new. Format: twelve slides maximum. Content: Q3 financial performance (headline numbers I'll provide), two strategic decisions made this quarter, and a twelve-month outlook. Tone: direct and professional. Lead with the most important finding, not with agenda. Do not include operational detail — board level only. Flag any section where you've made assumptions about emphasis that I should review."	Role, specific deliverable, audience, slide count, content list, tone, structural requirement, exclusion, verification request

Appendix C

Verification Checklist

A single-page reference for use before relying on significant AI output.

THE EIGHT-QUESTION VERIFICATION CHECKLIST

Work through these before sending, publishing, presenting, or acting on any significant AI output.

☐ **1. Factual or creative?**

Identify which type of output this is. Factual claims require source verification. Creative output (structure, framing, style) can be assessed on its own merits.

☐ **2. Numbers present?**

Identify every specific number, percentage, statistic, or calculation in the output.

For each one: can you find it in a source you can point to?

If no: do not use it until you can.

☐ **3. Citations present?**

If AI has cited sources, references, or research:

For each one: does the source exist? Does it actually say what AI claims?

Search for the source independently. Do not assume.

☐ **4. Sources real?**

Beyond checking that sources exist: does the source's content actually support the claim AI has attributed to it? A real paper may be misquoted or misapplied.

☐ **5. Could this have changed?**

For any claim about current facts — who holds a role, what a product costs, what a regulation says, what an organisation does — is this current? AI knowledge has a cutoff. For fast-moving topics, verify from official current sources.

☐ **6. Assumptions made?**

Ask AI directly: "What assumptions did you make in producing this output?" Review the assumptions. Are they correct? Are there important exceptions in your situation?

☐ **7. Constraints followed?**

Check each instruction you gave against the output. Were all followed? Is anything missing that you required? Is anything present that you excluded?

☐ **8. Reasoning sound?**

For analytical outputs: does the conclusion actually follow from the evidence? Could the same evidence support a different conclusion? What would have to be true for the conclusion to be wrong?

RISK-CALIBRATED VERIFICATION GUIDE

OUTPUT TYPE	PRIORITY CHECKS
Report with statistics	2 (numbers), 5 (currency), 8 (reasoning)
Academic or research content	2, 3 (citations), 4 (sources), 8
Legal or regulatory content	5 (currency), 6 (assumptions), 8
Medical information	5, 6, 8 — plus qualified professional review
Client-facing communications	6 (assumptions), 7 (constraints), 8
Internal communications	7 (constraints), light check on others
Creative/editorial content	Read critically; 7 (constraints)

WHEN YOU FIND AN ERROR

- 1.** Correct the specific error

2. Ask whether the error is isolated or a signal of a broader problem with this output
 3. Consider whether your workflow should have caught this earlier
 4. Update your standing verification rules if this is a pattern
-

Appendix D

AI Use Cases by Profession

A practical starting point for common professional contexts. Use this to identify which workflows and techniques are most relevant to your work, and to think about what your workspace should contain.

This is a starting point, not an exhaustive list. Your specific situation will always have uses and constraints not captured here.

Marketing and Communications

MOST VALUABLE USE CASES:

- Drafting and editing client communications, campaigns, and copy (W1, W9, W10, W11, W12)
- Research and briefing preparation for campaigns (W3, W13)
- Repurposing content across channels (W11)
- Preparing for client presentations (W3, W7)
- Producing first drafts from briefs (W25)

KEY WORKSPACE ELEMENTS:

- Brand guidelines document (uploaded to each client project)

- Tone-of-voice reference (example copy that captures the right register)
- Target audience profiles
- Standard exclusions per client (regulatory constraints, competitor references, etc.)

WHERE TO BE MOST CAREFUL:

Statistics and market claims in copy — these are high-visibility and often checked by clients or regulators. Verify every specific figure before it goes into client-facing material.

Small Business Owners and Entrepreneurs

MOST VALUABLE USE CASES:

- Client proposals and pitches (W5)
- Client and supplier communications (W1, W2)
- Summarising information for decisions (W4, W15)
- Business planning and project management (W17)
- Marketing content (W9, W10, W12)

KEY WORKSPACE ELEMENTS:

- Company background and positioning statement
- Target customer description
- Standard terms and constraints
- Tone guide (how the business communicates)
- Any regulatory constraints on the business

WHERE TO BE MOST CAREFUL:

Legal and financial claims in business documents, and any content that will be seen by clients before internal review.

HR and People Professionals

MOST VALUABLE USE CASES:

- Job descriptions and postings (W6)
- Employee communications
- Policy and process documentation
- Meeting notes and follow-ups (W8)
- Difficult conversation preparation (W2)

KEY WORKSPACE ELEMENTS:

- Organisation name and structure
- HR policy documents (relevant sections uploaded for reference)
- Employment law jurisdiction
- Standard communication tone
- Constraints (legal language to avoid in employee communications, etc.)

WHERE TO BE MOST CAREFUL:

Employment law and HR policy claims are jurisdiction-specific and change regularly. Never rely on AI for specific legal positions without verification from qualified legal sources and your current policy documents.

Legal Professionals

MOST VALUABLE USE CASES:

- Document summarisation and review (W4, W14)
- Research preparation and briefing (W3, W13)
- Drafting non-legal client communications (W1)
- Meeting preparation (W3, W8)
- Summarising case materials for internal briefing

KEY WORKSPACE ELEMENTS:

- Jurisdiction and area of practice
- Firm style and communication standards
- Standard disclaimers and qualification language
- What AI output is and is not used for (internal note: AI is a drafting aid; all legal content is reviewed by qualified solicitors)

WHERE TO BE MOST CAREFUL:

AI should not be used to generate legal advice or opinions without qualified review. Legal knowledge changes; AI may state outdated positions confidently. AI summaries of contracts and legal documents should always be checked against the original for anything that matters.

Finance and Accounting Professionals

MOST VALUABLE USE CASES:

- Summarising reports and financial documents (W4)
- Client communication drafts (W1, W7)
- Research preparation (W3, W13)
- Project and planning support (W17)
- Preparing for meetings (W3, W8)

KEY WORKSPACE ELEMENTS:

- Regulatory jurisdiction and relevant standards
- Client confidentiality constraints
- Appropriate disclaimer language for client-facing communications
- Technical terminology explanations for non-specialist communications

WHERE TO BE MOST CAREFUL:

All financial figures in AI output must be verified against source data. AI arithmetic is unreliable without code execution. Any regulatory or tax claim must be verified against current HMRC/relevant authority guidance. Never include AI-generated financial projections in client work without independent verification.

Teachers and Education Professionals

MOST VALUABLE USE CASES:

- Lesson planning and resource development (W16, W25)
- Differentiated materials for different ability levels
- Parent communications (W1)
- Research and preparation for topics (W13)
- Assessment question development

KEY WORKSPACE ELEMENTS:

- Year group / age range and subject
- Curriculum framework (e.g. National Curriculum in England)
- School's communication style and tone
- Any relevant pastoral or safeguarding context that affects communications

WHERE TO BE MOST CAREFUL:

Curriculum-specific content should be verified against current specifications. Content for student consumption should be reviewed by a teacher before use. Subject matter accuracy should be verified, especially in science, history, and other factually complex areas.

Researchers and Academics

MOST VALUABLE USE CASES:

- Literature orientation for new topics (W13)
- Summarising papers and documents (W4, W14)
- Structuring arguments and papers (W17, W25)
- Editing and improving prose (W24)
- Preparing for seminars and presentations (W3)

KEY WORKSPACE ELEMENTS:

- Field and sub-discipline
- Preferred citation style
- Institutional guidelines on AI use (these vary significantly; check your institution's current policy)
- Level of technical language appropriate for target publication

WHERE TO BE MOST CAREFUL:

All citations must be verified independently — AI hallucination of academic references is well-documented. AI summaries of research papers may miss important methodological nuance. AI should not be used to generate empirical claims or research findings. Follow your institution's and target journal's guidelines on AI disclosure.

General Professionals (Broad Business Contexts)

MOST VALUABLE STARTING WORKFLOWS:

- W1 (Professional email) — universally applicable
- W3 (Meeting preparation) — very high return
- W4 (Summarising documents) — immediate time saving
- W8 (Meeting agenda and follow-up) — immediately practical

- W15 (Building a workspace) — the highest long-term leverage point

MINIMUM VIABLE WORKSPACE:

- Your role and organisation type (two sentences)
 - Your standard communication tone
 - UK/US English preference
 - Things AI should never include in your output (jargon, specific competitors, etc.)
 - Default verification instruction ("please flag any claims I should verify independently")
-

Appendix E

Product Notes — Platform-Specific Information

This appendix contains product-specific information that is subject to change. All entries are verified as of September 2026. Volatility ratings indicate how likely information is to have changed since verification.

Before relying on any of this information for a workflow or decision, check the platform's current documentation. AI product features, pricing, and availability change frequently and sometimes without announcement.

HOW TO CHECK CURRENT INFORMATION

For each major platform:

- **ChatGPT:** help.openai.com
- **Claude:** support.anthropic.com and docs.anthropic.com
- **Gemini:** support.google.com/gemini

- **Microsoft Copilot:** support.microsoft.com

For pricing specifically, always check the platform's pricing page directly. Prices listed by third parties (including this book) become outdated.

Useful official starting points:

- ChatGPT Projects:
<https://help.openai.com/en/articles/10169521-projects-in-chatgpt>
- ChatGPT Custom Instructions:
<https://help.openai.com/en/articles/8096356-chat-preferences-for-chatgpt>
- ChatGPT GPTs: <https://help.openai.com/en/articles/8798878-building-and-publishing-a-gpt>
- Claude Projects and personalisation:
<https://support.anthropic.com/en/articles/9519177-how-can-i-create-and-manage-projects>
- Gemini Deep Research:
<https://support.google.com/gemini/answer/15719111>

E1 — ChatGPT (OpenAI)

Plans available (September 2026): Free, Plus, Pro, Team, Enterprise

PROJECTS FEATURE:

Available on free and paid plans. Projects group conversations, project instructions, and reference files; memory settings and file limits vary by plan and workspace. Check current documentation for the limits that apply to your account.

Volatility: MEDIUM

MEMORY FEATURE:

Saved memory and project memory can carry selected context across conversations, but availability, settings, and behaviour vary by plan, workspace, and region. Treat memory as a convenience, not as complete or infallible recall.

Volatility: MEDIUM

WEB SEARCH:

Available across plans (with limitations on Free). Search quality varies; results should be treated as a starting point for verification, not a verified source.

Volatility: MEDIUM

CODE EXECUTION (ADVANCED DATA ANALYSIS):

Some ChatGPT plans and tools support code execution, file analysis, and chart generation. Availability and limits vary. Most reliable arithmetic is obtained when the calculation is actually executed and the inputs and formula are checked.

Volatility: LOW

CUSTOM INSTRUCTIONS (PERSISTENT INSTRUCTIONS):

Available across ChatGPT plans, with limits and interface details varying by plan and platform. They apply broadly, while project instructions are scoped to a particular project and may take precedence within it.

Volatility: LOW

GPTS:

You can use existing GPTs where your account and workspace permit it. Creating or publishing a new GPT is subject to current account, plan, workspace, and eligibility rules; personal ChatGPT accounts may not be able to create or publish new GPTs. Check OpenAI's current GPT documentation before promising this capability to every reader.

Volatility: HIGH

IMAGE GENERATION (CHATGPT IMAGES):

Integrated into ChatGPT, with capabilities and access changing over time. Commercial use is subject to OpenAI's current terms and applicable law — verify both before using generated images in commercial projects.

Volatility: HIGH — terms and capabilities change

FILE UPLOAD:

PDF, Word, Excel, images and other formats. Size and count limits apply — check current documentation.

Volatility: MEDIUM

E2 — Claude (Anthropic)

Plans available (September 2026): Free, Pro, Team, Enterprise

PROJECTS FEATURE:

Available on Pro and above. Similar to ChatGPT Projects — groups conversations, maintains uploaded documents, supports project-level instructions. One of the more mature implementations of this feature.

Volatility: MEDIUM

SYSTEM PROMPTS / PERSISTENT INSTRUCTIONS:

Available via Projects and via platform settings. Robust and reliable as of verification date.

Volatility: LOW

WEB SEARCH:

Available on Pro and above, via tool integration. As with all AI web search: verify important findings from primary sources.

Volatility: MEDIUM

FILE UPLOAD:

PDFs, Word documents, images, and other formats supported. Claude has historically handled long documents well — but verify large document behaviour from current documentation.

Volatility: MEDIUM

MEMORY:

Claude provides several forms of persistence and personalisation, including profile preferences, project instructions, project knowledge, and—where enabled—search across previous chats. These mechanisms differ from one another and do not guarantee complete recall across conversations. Check current availability and settings in Claude's documentation.

Volatility: HIGH

EXTENDED CONTEXT WINDOW:

Claude models have historically supported long context windows, enabling analysis of very long documents. Specific limits depend on the current model version.

Volatility: MEDIUM

E3 — Gemini (Google)

Plans available (September 2026): Gemini (free), Gemini Advanced (paid subscription), Gemini for Workspace (business integration)

GOOGLE WORKSPACE INTEGRATION:

Gemini integrates with Google Docs, Gmail, Drive, Sheets, and Slides for users with Workspace accounts. For professionals already using Google Workspace, this integration is significant for productivity. Depth and reliability of integration varies by product.

Volatility: HIGH — integration features are actively expanding

WEB SEARCH:

Gemini has native Google Search integration, which is a meaningful advantage for tasks requiring current information. Search quality is generally strong; verify important claims as normal.

Volatility: LOW (integration is core to the product)

GEMS (CUSTOM AI CONFIGURATIONS):

Gemini offers custom configurations called Gems. Creation, file support, and availability vary by account, plan, region, and product version; check the current Gemini help centre before relying on a specific capability.

Volatility: MEDIUM

FILE UPLOAD:

Images, PDFs, and other formats supported. Google Workspace document integration provides additional access to Drive files.

Volatility: MEDIUM

E4 — Microsoft Copilot

Plans available (September 2026): Copilot (free, integrated into Windows and Edge), Copilot Pro (paid), Microsoft 365 Copilot (business, requires M365 subscription)

PRIMARY USE CASE:

Copilot's primary value for professional users is its integration with Microsoft 365 applications — Word, Excel, Outlook, Teams, PowerPoint. For organisations running M365, this integration can be significant.

MICROSOFT 365 COPILOT:

If your organisation has deployed M365 Copilot, AI assistance is available within the applications you already use. Capabilities include: summarising emails and meetings, drafting documents

within Word, analysing data in Excel, generating presentations in PowerPoint.

Volatility: HIGH — M365 Copilot is rapidly evolving

DATA SECURITY:

For organisations using M365 Copilot under an enterprise agreement, data handling follows Microsoft's enterprise data protection terms. This is significant for regulated industries.

Verify current terms with your IT and legal teams.

Volatility: MEDIUM

E5 — Cross-Platform Notes

CHOOSING A PRIMARY PLATFORM:

For most professional users, the choice of primary AI platform should be based on: integration with your existing tools (if you use Google Workspace, Gemini's integration is worth considering; if you use M365, Copilot may be relevant); the specific capabilities you need most (long document analysis, code execution, image generation); and the plan that provides the right features for your use.

USING MULTIPLE PLATFORMS:

Many users find it practical to use more than one platform. A common pattern: one platform as the primary for ongoing projects (with workspace set up), a second for specific tasks it handles better (e.g., using web-search-heavy research tasks with a strong search-integrated tool).

PRIVACY AND DATA:

All major AI platforms have privacy policies governing how they use your inputs. Key questions to check in current documentation:

- Is your data used to train models?
- What is the opt-out process?

- For enterprise accounts: what are the data retention and usage terms?

For professional use involving client data, sensitive business information, or personally identifiable information, review the platform's current privacy and data terms before processing that information.

Volatility: HIGH — policies change; verify before processing sensitive data

E6 — Image Generation (All Platforms)

CURRENT STATE (SEPTEMBER 2026):

Image generation is integrated into ChatGPT through its current image-generation experience and is available through several other platforms. Product names, models, access, and quality change frequently.

COMMERCIAL USE:

Before using AI-generated images in commercial contexts (client work, marketing, publications), verify the current terms of service of the specific platform. Terms for commercial use vary and have changed multiple times. Verify before use.

COPYRIGHT CONSIDERATIONS:

The legal status of AI-generated images for commercial use is an evolving area in multiple jurisdictions. Current legal guidance in the UK and elsewhere is subject to ongoing development. Check current legal guidance before relying on AI-generated images in commercial contexts.

Volatility: VERY HIGH — both platform terms and legal frameworks are actively changing

Appendix E was verified in September 2026. Product features, pricing, availability, and terms change frequently. Always verify from the platform's own current documentation before relying on specific features or capabilities for professional use.

A Note on Sources

This book does not use footnotes in the main text, because most of the principles it describes are based on documented patterns of AI behaviour rather than a single research paper or study. The following notes point to the most useful reference points for readers who want to go further.

ON AI FAILURE MODES AND HALLUCINATION:

The phenomenon of AI hallucination is well-documented in academic literature. A practical starting point is work published by researchers at major AI labs themselves — Anthropic, OpenAI, and Google DeepMind all publish research findings on model behaviour that is accessible to non-specialists on their respective research websites.

ON PROMPT ENGINEERING:

The field has produced several peer-reviewed studies on what makes prompts effective. The most reliable summary of current understanding remains the documentation produced by each AI lab — particularly Anthropic's documentation on working with Claude and OpenAI's guidance on ChatGPT. These are updated more frequently than academic publications.

ON VERIFICATION PRACTICES:

Standards for AI output verification in professional contexts are still emerging. In the UK, professional bodies in law, accounting, and medicine have begun to publish guidance — check your relevant professional body's current guidance.

ON PLATFORM FEATURES:

All platform-specific information should be verified from the platform's own documentation. The most current information is always on the platform's help centre, not in any third-party publication.

ABOUT THE AUTHOR



Eleanor Mercer is a professional development trainer and consultant who has worked with professionals across law, financial services, and management consulting for fifteen years. Her original practice focused on communication, writing, and presentation — the skills that separate people who can do excellent work from people who can also explain it clearly.

When AI tools entered professional workplaces in 2022, she began testing them, helping clients use them, and watching closely what happened when capable professionals sat down with capable technology and still got mediocre results.

The pattern was consistent: the gap was almost never technical. It was what people gave AI before they started — the context, the instructions, the job description, the source material, and the verification process.

The 7 Questions framework and the AI Interaction Stack in this book were developed through that observation, tested in professional training practice, and refined over two years of real use before appearing in print.

The Art of AI is the first book in a planned series on professional AI skills for non-technical users. Eleanor writes, trains, and publishes through The Art of AI.

About This Book

The Art of AI, Book 1, is the first in a planned series covering practical AI skills for non-technical professionals. Future books in the series will cover AI for specific professional contexts, advanced workflow design, and team-level AI integration.

AI tools were used in the research, drafting, and editing of this book. All factual claims were verified against primary sources at the time of writing. All examples were tested. All editorial decisions — structure, voice, inclusion, emphasis — are the author's own.

Published 2026. First edition.
