



THE ART OF AI — BOOK 1

The Art of AI

How to Get Better Results From Every AI
Conversation

A Practical Guide to Better Prompts, Better Context, and Better
Workflows

MITESH MAHARAJ

FIRST EDITION · 2026

*For every person who typed something into
ChatGPT, got a vague wall of text, and quietly
wondered whether everyone else was getting
something they weren't.*

COPYRIGHT

Copyright © 2026 The Art of AI

All rights reserved. No part of this publication may be reproduced, distributed, or transmitted in any form or by any means, including photocopying, recording, or other electronic or mechanical methods, without the prior written permission of the publisher, except in the case of brief quotations embodied in critical reviews and certain other non-commercial uses permitted by copyright law.

Published by: The Art of AI

Series: The Art of AI, Book 1

First edition: 2026

AI Disclosure: AI tools were used in the research, drafting, and editing of this book. All factual claims were verified against primary sources at the time of writing. All examples were tested. Editorial judgment, structure, and final decisions are the author's own.

Amazon KDP AI disclosure completed in the publishing dashboard per current requirements. KDP's AI content policy should be rechecked before any future edition.

CONTENTS

| INTRODUCTION

| You Are Not Bad at AI

| PART I

| A Different Way to Think About AI

| Chapter 1

| Why AI Gives You Generic Answers (and How to Change That)

| Chapter 2

| The AI Interaction Stack — What Actually Determines Output Quality

| Chapter 3

| The 7 Questions — Your Portable AI Framework

| Chapter 4

| How to Talk to an AI (It's Not Google)

| PART II

| Build Better Interactions

| Chapter 5

| Give AI a Job, Not Just a Question

| Chapter 6

| Context Is Everything

| Chapter 7

| Instructions That Actually Work

| Chapter 8

| Examples Change Everything

| Chapter 9

| Control the Output

| PART III

| Build an AI Environment

| Chapter 10

| Stop Starting From Zero Every Time

Chapter 11

Projects and Persistent Instructions

Chapter 12

Files, Documents, and Source Material

Chapter 13

Tools — What AI Can Do Beyond Text

Chapter 14

The Memory Problem (and How to Work Around It)

Chapter 15

Building Your Personal AI Workspace

PART IV

From Prompts to Workflows

Chapter 16

When One Prompt Isn't Enough

Chapter 17

Breaking Tasks Into Steps

Chapter 18

Critique, Edit, and Improve

Chapter 19

Chains, Loops, and Handoffs

PART V

25 Real-World Workflows

BUSINESS AND PROFESSIONAL

W1 — Write and Improve a Professional Email

W2 — Prepare for a Difficult Conversation

W3 — Research a Topic for a Meeting

W4 — Summarise a Long Document

W5 — Draft a Project Proposal

W6 — Write a Job Posting

W7 — Prepare a Client Report

W8 — Create a Meeting Agenda and Follow-Up

MARKETING AND CONTENT

W9 — Write a LinkedIn Post

W10 — Draft a Blog Post from Notes

W11 — Repurpose Content Across Formats

W12 — Write Website Copy

RESEARCH AND LEARNING

W13 — Research a Topic You Know Nothing About

W14 — Understand a Complex Document

W15 — Compare Options Before a Decision

W16 — Learn a New Skill Faster

PERSONAL PRODUCTIVITY

W17 — Plan a Project

W18 — Process a Backlog of Tasks

W19 — Draft a Difficult Personal Message

W20 — Prepare for a Job Interview

HOME AND LIFE

W21 — Plan a Trip

W22 — Understand a Contract or Legal Document

W23 — Get a Second Opinion on a Health Question

WRITING AND EDITING

W24 — Edit and Improve Your Own Writing

W25 — Turn Rough Notes into a Polished Document

PART VI

The AI User's Playbook

Chapter 20

How to Verify AI Output

Chapter 21

What AI Gets Wrong (Consistently)

Chapter 22

Calibrating Trust

Chapter 23

Building Your AI Habits

Chapter 24

Keeping Up Without Burning Out

APPENDICES

Appendix A

The 7 Questions Quick Reference Card

Appendix B

The Before → Better → Best Worksheet

Appendix C

Verification Checklist

Appendix D

AI Use Cases by Profession

Appendix E

Product Notes — Platform-Specific Information

E1 — Persistent Instructions and Custom Instructions

E2 — Projects

E3 — File and Document Upload

E4 — Web Search

E5 — Voice and Audio

E6 — Memory

E7 — Model differences (brief reference)

A Note on What Comes Next

What you will actually learn

By the end of Part I, you will understand why AI produces generic output and exactly what changes when it produces something good. That understanding alone is worth more than any list of prompts, because it means you can diagnose what went wrong in any interaction and fix it — rather than just trying different words until something works.

By the end of Part II, you will know how to build an AI interaction deliberately: giving AI the role, the context, the instructions, the examples, and the output format that get you the result you need.

By the end of Part III, you will know how to build an AI environment — a persistent workspace where you stop repeating yourself, where your files and preferences carry forward, and where AI actually knows something about your situation.

By the end of Part IV, you will be able to design multi-step workflows: decomposing a complex task into stages, running them in sequence, and using AI's output from one step as the input to the next.

Part V gives you twenty-five tested workflows to use immediately or adapt. Part VI gives you the habits and verification practices that distinguish people who use AI reliably from people who use it hopefully.

THE FRAMEWORK

Running through all of it is a single seven-question framework you can carry in your head. Before any significant AI interaction, these seven questions identify what you need to provide. They work on every AI platform — ChatGPT, Claude, Gemini, Copilot, and every tool built on these models. They will still work when the current models are replaced by the next generation.

The framework is called The 7 Questions. You will encounter it in Chapter 3, and you will use it for the rest of the book.

A WORD ABOUT TECHNICAL KNOWLEDGE

None of this requires technical knowledge. You do not need to understand how large language models work. You do not need to know what a token is, or an API, or an embedding. When technical concepts appear in this book, they appear because they directly affect how you use AI, and they are always explained in ordinary language before the technical term is introduced.

What you need is the willingness to think about AI interactions differently — not as typing questions into a search engine, but as briefing a capable colleague who starts every conversation knowing nothing about you, your work, or what you need.

Once you understand that, everything else follows.

Let's begin.

The four chapters in this part build the mental model that makes everything else in the book work. They are short by design. If you understand what this part is teaching, you will already be getting noticeably better results before you reach Part II.

The problem everyone has but nobody explains

Ask ten people what frustrates them about AI and you will hear the same answers:

"It's too generic."

"It doesn't understand what I actually want."

"It's confident but wrong."

"The answer doesn't fit my situation."

"It gives me five paragraphs when I needed two sentences."

These are not problems with the AI. They are problems with the interaction.

This chapter explains why generic inputs produce generic outputs — and what specifically changes when they don't.

THE SLOT MACHINE MODEL (AND WHY IT'S WRONG)

Most people approach AI the way they might approach a vending machine: insert request, press button, hope for the right thing to fall out. When it doesn't, they try slightly different words. Sometimes this works. Usually it's a matter of luck.

This approach treats AI as a slot machine. It also produces slot-machine results: random, unpredictable, and mostly disappointing.

The people who get consistently good results from AI are not luckier. They are not using secret prompts. They understand something that slot-machine users don't: AI output is shaped by input, and almost every dimension of that input is under your control.

WHY AI IS GENERIC BY DEFAULT

When you send AI a vague request — "Write me a summary of this report" — it has to make assumptions about everything it doesn't know:

- Who is going to read this summary?
- How long should it be?
- What are the most important points to include?

- What should be left out?
- What tone is appropriate?
- What does the reader already know?
- What is the purpose of the summary?

AI will fill in all of these gaps. It will fill them in with the most statistically average answers it can construct — the kind of summary that would be acceptable to the largest number of readers in the largest number of situations. This is why AI output so often feels as though it was written for everyone and therefore fits no one particularly well.

The good news: every assumption AI makes is one you could have specified yourself. The summary would look completely different if you told AI: *"This is for our CEO, who hasn't read the report and needs to decide whether to approve a £200,000 budget. She wants two paragraphs, direct language, and the headline number up front."*

Same document. Same AI. Dramatically different output. The only difference is what you provided.

WHY IT WORKS

AI tools generate text by predicting what words should follow what came before, based on patterns learned from enormous quantities of written material. This means the quality and specificity of what they produce is directly shaped by the quality and specificity of what they receive. Providing more relevant information doesn't confuse AI — it constrains the prediction space to something closer to what you actually need.

THE THREE MOST COMMON INPUT FAILURES

After looking at hundreds of AI interactions that produced poor results, three input problems appear again and again:

1. THE VAGUE JOB DESCRIPTION

"Help me with my presentation" tells AI you want a presentation. It doesn't tell it the topic, the audience, the desired length, your current draft, what help means, or what "better" would look like. AI has no choice but to respond generically to a generic request.

2. THE MISSING CONTEXT

"Reply to this email" gives AI a text to work with. It doesn't tell AI who sent it, what the relationship is, what a good reply should accomplish, whether you need to decline or accept, how formal the tone should be, or what your constraints are. AI will guess. Sometimes it guesses right. Usually the first reply it produces needs significant editing.

3. THE WRONG ASSUMED AUDIENCE

"Explain compound interest" gives AI a topic. But who is asking? A secondary school student? A financial professional? Someone deciding whether to take out a mortgage? AI will produce an explanation pitched at a generic reader — which means it may be too basic, too technical, or simply beside the point for your actual need.

REAL-WORLD EXAMPLE

Sarah is a marketing manager who needs to send a follow-up email to a potential agency partner after an introductory call. She types:

"Write a follow-up email after a business meeting."

AI produces a template with placeholders like [Name] and [Company] and generic phrases like "It was great connecting with you." It's technically a follow-up email. It has nothing to do with Sarah's actual situation, the tone of the meeting, what was discussed, or what the next step should be.

She then provides: who the meeting was with, what was discussed, what her specific goal for the follow-up is, and how she wants to come across. The email AI produces next is one she could send without changes.

Same AI. Same model. The variable was what she gave it.

WHAT ACTUALLY CHANGES GOOD AI OUTPUT

Five inputs move the needle most consistently. They appear in greater detail throughout this book, but in brief:

1. **A clear job:** What specifically do you want AI to do? (Not "help with" — do.)
2. **Relevant context:** What does AI need to know about your situation, audience, and constraints?
3. **Source material:** What documents, notes, or data should AI work from?
4. **Output specifications:** What should the result look like — length, format, tone, structure?
5. **Verification instructions:** How should AI flag uncertainty or potential errors?

This is not a complete list. The next chapter introduces the full model. But if you apply only these five, you will notice an immediate and significant improvement.

TRY THIS

Take a prompt you have used before that produced a disappointing result. Write it down.

Now ask yourself: What did I assume AI would know that it couldn't have known? What context did I leave out? What does "good" actually look like for this task?

Rewrite the prompt with those gaps filled in and run both versions. Notice what changed.

WHAT THIS CHAPTER ESTABLISHED

Generic inputs produce generic outputs. This is a feature, not a bug — it means every improvement you make to your input produces a corresponding improvement in output. You are not at the mercy of the AI. You are in control of the most important variable in every interaction.

The next chapter introduces a model for understanding exactly which variables are in play — and which ones matter most for any given task.

WATCH OUT

More detail is not always better. Adding irrelevant information doesn't improve output — it can dilute it. The skill is not writing longer prompts; it's identifying what AI actually needs to know to do this specific task well. You'll develop this judgment through the exercises in Part II.

Introducing a map

When a piece of work goes wrong, it helps to know where in the process the problem occurred. A cake that doesn't rise failed at a different point than a cake that doesn't taste right. The diagnosis matters.

AI interactions fail at predictable places. This chapter introduces a model — the AI Interaction Stack — that maps those places. It gives you a vocabulary for diagnosing what went wrong and a checklist for getting it right before you start.

THE AI INTERACTION STACK

The AI Interaction Stack has eleven layers. Each layer represents something that shapes what AI produces. When an interaction produces poor results, at least one of these layers is inadequate or missing.

You do not need all eleven layers for every task. A quick factual question might only need layers 1–4. A complex document workflow might need all eleven. The point is knowing which layers exist so you can deliberately choose which ones to activate.

Here they are, in order:

LAYER 1 — INTENT

What are you actually trying to achieve? Not the surface task ("write a summary") but the underlying goal ("help my CEO decide quickly"). These are often different, and the difference matters.

A summary intended to help someone decide looks different from a summary intended to brief someone who already decided. Same task, different intent, different output.

Before you start: Can you state your actual goal in one sentence?

LAYER 2 — CONTEXT

What does AI need to know about your situation? Who is involved? What is the background? What constraints apply? What has already been tried?

Context is the layer most people consistently underinvest in. It is also the layer where providing more almost always pays off.

Before you start: What does AI not know that would change its answer?

LAYER 3 — INSTRUCTIONS

What exactly should AI do? This is distinct from intent. Intent is the goal; instructions are the specific action you want AI to take to achieve it.

"Summarise this document" is an instruction. "Summarise this document in three bullet points, prioritising financial implications, written for a non-specialist reader" is a better one.

Before you start: How specific are your instructions? Could a capable human follow them without asking a question?

LAYER 4 — EXAMPLES

What does good look like? Sometimes the most efficient way to specify a tone, structure, or style is to show AI an example of it — a previous email you wrote, a format you liked, a sample of the style you want.

Examples do more than description. Telling AI to be "warm but professional" is imprecise. Showing it a warm-but-professional email is not.

Before you start: Is there an example that would make the target output clearer?

LAYER 5 — SOURCE MATERIAL

What documents, data, notes, or materials should AI work from? AI can read files — contracts, reports, research papers, meeting notes, spreadsheets — and its output will be grounded in what you provide rather than in generalisation.

When you ask AI to help with a specific document or dataset, provide it. Don't describe it — provide it.

Before you start: What source materials could improve the specificity and accuracy of the output?

LAYER 6 — WORKSPACE

Have you set up a persistent context — a project, a set of standing instructions, a shared background — so you don't have to repeat essential information every time?

For tasks you do repeatedly, a workspace eliminates repetition and builds consistency. Part III is dedicated to this layer.

Before you start: Is this a recurring task where a workspace would save significant setup time?

LAYER 7 — TOOLS

What capabilities does AI have access to in this context? Can it search the web? Run calculations? Read uploaded files? Generate images? Access calendar or email?

Knowing which tools are available — and which aren't — shapes what you can reasonably ask for and where you need to verify externally.

Before you start: What tools are available? Does this task benefit from any of them?

LAYER 8 — WORKFLOW

Is this a single-step task or a multi-step process? Complex work almost always benefits from decomposition: breaking the task into stages, running each stage separately, reviewing and adjusting between stages.

Single-prompt attempts at complex tasks produce mediocre results that are hard to improve. Staged workflows produce better outputs at each stage and make errors easier to catch.

Before you start: Is this a single-step task, or should it be broken into stages?

LAYER 9 — OUTPUT

What should the result look like? Format, length, structure, tone, audience, medium — all of these shape the output and all are specifiable.

Without output specifications, AI defaults to whatever format its training suggests is most common for this type of request. That default is often wrong for your specific situation.

Before you start: Have you specified format, length, tone, and structure?

LAYER 10 — VERIFICATION

How will you check the output? What specifically might be wrong? AI makes errors — factual, logical, and interpretive — with calm confidence. Your verification approach should match the stakes of the task and the risk profile of the content.

Telling AI to flag its own uncertainties is a useful first step. It doesn't replace your review, but it helps.

Before you start: How will you verify this output before you use or share it?

LAYER 11 — ITERATION

How will you use this output as the starting point for the next stage? AI interactions are rarely single exchanges — they are conversations where each response informs the next prompt. Good output at layer 11 becomes input at layer 2 of the next pass.

Before you start: Is this a one-shot task, or will you need to iterate? What will the next step be?

WHY IT WORKS

The stack is not a checklist to complete in full before every interaction. It is a diagnostic framework. When output is disappointing, the stack tells you where to look. When you're designing a complex task, the stack tells you what to prepare. Most interactions need four or five layers well-activated; a few need all eleven.

THE STACK IN PRACTICE

Here is the same task at two different levels of stack activation:

LOW STACK (LAYERS 1, 3 ONLY):

"Summarise this meeting."

Result: A generic summary with equal weight given to all topics discussed. Possibly bulleted, possibly in paragraphs. Pitched at a generic reader. Takes five minutes to edit into something useful.

HIGH STACK (LAYERS 1-5, 9-10):

"I need a meeting summary for three people who weren't present: my manager, the client's account manager, and one of our developers. Each person needs to see different things. My manager wants the key decisions and next steps only. The client account manager wants what the client committed to. The developer wants the technical requirements we discussed. The attached transcript is the source. Write three separate summaries — one for each person — no longer than half a page each. Flag anything that seemed unresolved or ambiguous."

Result: Three targeted summaries, immediately usable, with potential problems flagged.

The difference is not length of prompt. It is activation of the right layers for the task.

TRY THIS

Think of an AI task you need to do this week. Write down what layers of the stack you would normally activate — most people use layers 1, 3, and maybe 9 by default. Now consider which additional layers are relevant to this task. What would you add if you invested two more minutes in the setup?

WHICH LAYERS MATTER MOST?

The research on what makes AI interactions succeed points consistently to three layers as the highest-leverage investments for most people in most tasks:

Context (Layer 2) is the most consistently underinvested layer and the one with the highest return. Providing relevant context costs little and changes output significantly.

Instructions (Layer 3) are the most frequently underdeveloped layer. Most instructions are too vague. The skill of writing specific, actionable instructions is developed in Chapter 7.

Verification (Layer 10) is the most commonly skipped layer. Part VI covers it in depth. Skipping it is also the source of most AI-related professional embarrassment.

WATCH OUT

The stack is a mental model, not a procedure. It does not mean you should spend twenty minutes preparing every AI prompt. A quick factual question needs layers 1 and 3, briefly addressed, and that's correct. The stack is useful when you are planning a complex task or diagnosing a poor result — not as mandatory overhead for every interaction.

A framework you can carry in your head

The AI Interaction Stack gives you a complete map. But maps aren't always what you need. Sometimes you need a quick checklist — something you can run through in sixty seconds before an important interaction.

The 7 Questions is that checklist.

It is derived from the stack but simplified into seven questions that cover the most important dimensions of any significant AI interaction. It is short enough to memorise. It is structured enough to catch the most common input failures. And it is abstract enough to work on any AI platform, with any model, now or in the future.

THE 7 QUESTIONS

Before any significant AI interaction:

1. WHAT AM I TRYING TO ACHIEVE?

State the underlying goal, not just the surface task. "I want my CEO to approve the budget" is more useful than "I want a summary." The goal shapes every other decision.

2. WHAT DOES AI NEED TO KNOW?

What background information, constraints, relationships, or context is essential to getting a good output? What would a capable human assistant need to know before starting this task?

3. WHAT INFORMATION CAN I PROVIDE?

What documents, data, notes, examples, or materials exist that would make AI's output more specific and more accurate? What can you attach or paste in?

4. WHAT EXACTLY SHOULD AI DO?

What is the specific action? Not "help with" — do. Write, summarise, extract, compare, analyse, rewrite, draft, plan. Be precise about the verb and the scope.

5. WHAT SHOULD THE RESULT LOOK LIKE?

Format, length, tone, structure, audience, medium. If you do not specify, AI will choose. Specify.

6. HOW WILL I CHECK IT?

What might be wrong? What specifically needs verification before you use or share this output? Are there numbers, dates, names, claims, or recommendations that require checking?

7. WHAT SHOULD HAPPEN NEXT?

Is this one step in a larger process? What will you do with this output? Will it become input to another AI prompt, a document you share, or a decision you make?

WHY IT WORKS

Each question addresses a different failure mode. Question 1 catches the wrong goal. Question 2 catches the missing context. Question 3 catches unused resources. Question 4 catches the vague instruction. Question 5 catches the wrong format. Question 6 catches the unverified output. Question 7 catches the single-step tunnel vision.

You can have a perfect answer to six questions and fail on the seventh. The framework is complete by design.

USING THE 7 QUESTIONS: A WORKED EXAMPLE

Sarah needs to create a briefing document on a potential new supplier for her director.

Let's run through the questions:

Q1 — WHAT AM I TRYING TO ACHIEVE?

My director needs to decide whether to invite this supplier to a formal tender. I need to give her the information required to make that decision, not a complete company profile.

Q2 — WHAT DOES AI NEED TO KNOW?

The decision context (tender shortlist), the supplier's name (TechFlow Ltd), what the tender is for (CRM software implementation), our current system (Salesforce), our rough budget (£80,000), and what factors matter most to us (integration capability, support quality, UK presence).

Q3 — WHAT INFORMATION CAN I PROVIDE?

I have the supplier's website, a brochure they sent us, and notes from a call I took with their sales contact last week.

Q4 — WHAT EXACTLY SHOULD AI DO?

Create a supplier briefing document using the materials I provide. Organised under the headings: Company Overview, Relevant Experience, Key Capabilities, Potential Risks, Recommended Questions to Ask.

Q5 — WHAT SHOULD THE RESULT LOOK LIKE?

One page, maximum. Professional but plain language. My director doesn't have time for anything longer. Bullet points under each heading.

Q6 — HOW WILL I CHECK IT?

I'll verify any specific facts against the brochure and my call notes. I'll flag for my director any claims from the supplier that couldn't be independently verified.

Q7 — WHAT SHOULD HAPPEN NEXT?

This document will go to my director for her review. If she decides to proceed, I'll use AI to help draft the invitation to tender based on the same materials.

Now Sarah turns this into a prompt:

"I need a one-page supplier briefing document for my director, who is deciding whether to include TechFlow Ltd in a CRM software tender shortlist. The tender is for an implementation project with a rough budget of £80,000. Our current system is Salesforce."

*

The document should be organised under these five headings: Company Overview, Relevant Experience, Key Capabilities, Potential Risks, Recommended Questions to Ask.

*

Please work from the materials I am attaching: [brochure], [call notes]. Plain language, bullet points, maximum one page. Flag any claims in the supplier's materials that you cannot independently verify from what I've provided."

This is a complete interaction. It took about three minutes to prepare using the 7 Questions. It would have taken considerably longer to produce a good result any other way.

REAL-WORLD EXAMPLE

James is a recent graduate preparing for a job interview. He runs through the 7 Questions:

Q1 — I want to walk into this interview well-prepared, able to give specific answers about why I want this role.

Q2 — The company is Meridian Digital, a mid-size marketing agency. The role is Junior Account Manager. My background is a marketing degree and six months of internship experience. The interview is in two days.

Q3 — I have the job description, the company's website, their LinkedIn page, and a Glassdoor page with recent employee reviews.

Q4 — I want AI to help me prepare for likely interview questions, draft answers to the hardest ones, and identify things about the company I should know.

Q5 — I want a structured prep document: company background summary, ten likely questions with draft answers, and a section on questions I should ask them.

Q6 — I'll verify any company facts against their current website before the interview.

Q7 — I'll use the document to do a practice run the night before.

The resulting prompt takes forty seconds to write. The output he gets is interview prep he can actually use, not generic "tell me about a time you overcame a challenge" boilerplate.

THE 7 QUESTIONS AS A PORTABLE CARD

The 7 Questions are reproduced as a reference card in Appendix A. You can photograph it, print it, or simply remember it. The goal is that it becomes automatic — something you run through mentally before any AI interaction that matters.

You do not need to answer every question in full every time. For a simple, low-stakes task, three or four seconds on each question is sufficient. For a complex, high-stakes task — a document going to a senior stakeholder, a recommendation with significant consequences, a piece of work that will be published — the full seven questions deserve a full minute each.

PRO TIP

Once you have a well-developed answer to the 7 Questions for a recurring task — a weekly report, a client email format, a type of document you produce regularly — save it as a template. The 7 Questions become a workspace setup (Chapter 11). Run once; reuse indefinitely.

WATCH OUT

The 7 Questions are a preparation tool, not a rigid protocol. If you find yourself spending ten minutes answering the questions before a thirty-second task, you are misapplying the framework. Use judgment about scale. The questions exist to catch the gaps that consistently cause problems — not to add overhead to interactions that don't need it.

The search engine habit

Most people's first instinct when using AI is to type a search query. Short. Keyword-heavy. Punchy.

"best email subject lines"

"how to negotiate salary"

"presentation tips"

This works for Google because Google is retrieving existing documents that match your keywords. The shorter and cleaner the query, the more precisely it matches what's out there.

AI tools are not search engines. They are not retrieving existing content. They are generating new content based on what you've given them. The search-engine reflex is the single most common reason people get worse results than they could.

WHAT AI ACTUALLY DOES

Without requiring a technical explanation: an AI language model generates text by working out what words are most likely to follow what came before, based on patterns learned from enormous quantities of text. It is not looking up answers. It is constructing them.

This has three practical implications:

1. MORE RELEVANT INPUT, MORE RELEVANT OUTPUT.

A longer, more specific prompt does not confuse AI. It constrains the range of possible responses toward what you actually need. Specificity is an advantage.

2. AI HAS NO PERSISTENT MEMORY BETWEEN CONVERSATIONS.

Every time you start a new conversation, AI begins with no knowledge of who you are, what you've discussed before, or what your preferences are. This is why context must be provided explicitly — and why workspaces (Part III) matter so much for regular users.

3. AI GENERATES PLAUSIBLE TEXT, NOT NECESSARILY ACCURATE TEXT.

AI's training teaches it what good writing looks like in a given context. This means it can generate fluent, confident text on subjects it is wrong about. The confidence in the prose is not correlated with the accuracy of the content. This is the most important habit to build: treating AI output as a draft that requires verification, not a source that can be trusted.

WHY IT WORKS

Understanding the generation model — rather than thinking of AI as an oracle or a search engine — changes how you interact with it. It explains why providing examples works (they constrain the generated text toward the style you want). It explains why AI can be confidently wrong (plausibility and accuracy are different things). It explains why asking AI to explain its reasoning helps (it forces more systematic text generation). These are not quirks to work around; they are features of the model to use deliberately.

THE CONVERSATION, NOT THE QUERY

AI tools work best as conversations, not as single-query exchanges. The first response is rarely the last one.

This is not because AI needs multiple chances to get things right (though sometimes it does). It is because a conversation allows you to:

- Add context you forgot in the first prompt
- Adjust the output format based on what you see
- Request a deeper dive on one section
- Push back on something that seems wrong
- Ask AI to regenerate with a different constraint

Think of it as briefing a capable but uninformed colleague. The first briefing gets you a first draft. The subsequent conversation refines it. This is normal, efficient, and often faster than trying to write the perfect prompt in one go.

THE ANATOMY OF AN EFFECTIVE PROMPT

Not every prompt needs all of these elements. But a complete prompt for a significant task includes:

ROLE (optional but often useful):

Act as a [role] with expertise in [domain].

CONTEXT:

Background information AI needs to know.

Who is involved, what the situation is, what constraints apply.

TASK:

Exactly what you want AI to do.

Use a clear action verb. Be specific.

SOURCE MATERIAL (when relevant):

[Attached document / pasted text / data]

OUTPUT REQUIREMENTS:

Format, length, tone, audience, structure.

VERIFICATION REQUEST (for factual or high-stakes output):

Please flag any claims you are uncertain about,
any assumptions you've made, and any information

I should verify independently.

This is not a template to apply rigidly. It is a map of what a complete prompt includes. For quick tasks, you might cover all of this in two sentences. For complex tasks, each element might be a paragraph.

TRY THIS

Take a standard query you might type into a search engine and rewrite it as an AI prompt using the anatomy above. You don't need all six elements — pick the three most relevant.

For example:

Search query: "how to negotiate salary"

AI prompt: "I am about to have a salary negotiation conversation with my manager. I have been in this role for two years, my performance reviews have both been 'exceeds expectations', and I know from a conversation with a colleague that my salary is about 12% below the market rate for this role in London. I am asking for a 15% increase. Help me prepare for this conversation: what arguments are strongest, what objections might I face, and how should I handle the moment when they counter with a lower number? Tone: calm and factual, not aggressive."

Notice the difference in what AI has to work with.

ASKING AI TO THINK, NOT JUST WRITE

One of the most underused techniques is asking AI to reason through a problem before producing an answer.

When AI generates text quickly, it sometimes takes the most obvious route. Asking AI to think first — to break down the problem, identify considerations, or outline before drafting — tends to produce more thoughtful, more accurate results.

Phrases that invoke this:

- *"Before answering, think through the main considerations."*
- *"First give me an outline, then draft the full document."*
- *"Walk me through your reasoning."*
- *"What are the main risks or failure points here?"*

This works because it changes the pattern of text generation — the model is working through intermediate steps rather than jumping to a conclusion. You can think of it as asking someone to show their work rather than just hand you an answer.

PRO TIP

When you receive an output that is close to what you want but not quite right, do not start again from scratch. Respond to the output with a specific refinement instruction: "That's good but the tone is too formal — can you make it sound more conversational?" or "The summary is accurate but too long — cut it to three bullet points and keep only the most important point in each." Iteration on an existing draft is almost always faster than rewriting.

WATCH OUT

AI confidently produces wrong answers. The fluency of AI prose — the fact that it reads well, sounds authoritative, and doesn't hedge — can make it easy to accept without checking. Any output that contains specific facts, numbers, names, dates, quotes, citations, or recommendations deserves verification before you rely on it. Chapter 20 covers verification in detail. For now, adopt the rule: polished prose is not evidence of accuracy.

WHAT PART I HAS ESTABLISHED

You now have the foundation:

- Generic inputs produce generic outputs because AI fills gaps with averages
- The AI Interaction Stack identifies the eleven layers that determine output quality
- The 7 Questions give you a portable framework for activating the right layers

- AI is a generation tool, not a retrieval tool — and conversation is the natural mode of interaction

Part II takes each element of a strong interaction and develops it into a skill.

The five chapters in Part II develop the core skills that apply to every AI interaction you will ever have. Each chapter takes one element of the AI Interaction Stack and teaches you to use it deliberately. Together they account for the majority of the quality difference between people who get good results from AI and people who don't.

The question versus the job

There is a meaningful difference between asking AI a question and giving AI a job.

A question invites a response. A job assigns a task with a clear deliverable, a scope, and implicit quality criteria.

Question: "What do I need to know about redundancy letters?"

Job: "You are an HR communications specialist. Draft a redundancy letter for a member of our sales team who has been with the company for three years. The

reason is restructuring, not performance. Tone: respectful and clear. The letter must include: confirmation of the decision, notice period (one month), details of the settlement package (to be inserted), available support resources, and the process for appeal. UK employment law applies."

The question gets you information about redundancy letters — a textbook answer pitched at an anonymous reader. The job gets you a document you can actually edit and use.

WHY IT WORKS

Role assignment and job framing do not change what AI "knows." They change the frame through which AI selects and structures what it produces. A prompt that says "you are an experienced financial analyst" causes AI to draw more heavily from the patterns associated with financial analysis — more technical vocabulary, more rigorous structure, more appropriate caution. A prompt that specifies a deliverable (a letter, a report, an agenda) constrains the format. Together, role and job framing channel AI's generation toward something more precisely useful.

THE BEFORE → BETTER → BEST: GIVING AI A JOB

BEFORE (the vague question)

"What should I include in a client proposal?"

What AI produces: A generic list of proposal sections that any business book might contain. Executive summary, scope, timeline, pricing, terms. Correct. Applicable to no one's specific situation.

Why it's weak: AI has no information about the client, the industry, the type of work, the relationship, or what distinguishes this proposal from a hundred others.

BETTER (the specific task)

"Write an outline for a proposal for a digital marketing retainer, targeting a mid-size UK retail company. The agency specialises in SEO and paid social. The prospect has expressed interest but is also talking to two other agencies. The proposal needs to address: our methodology, relevant case studies, pricing structure (monthly retainer, tiered options), and what makes us different."

What AI produces: A structured outline specific to this type of proposal with relevant sections and brief notes on what each should contain.

What improved: AI now has the industry, the type of work, the competitive context, and the key elements to address.

BEST (role + job + context + output specifications)

"You are a senior account director at a boutique UK digital marketing agency. You are writing a proposal for a monthly SEO and paid social retainer for Thornton's, a mid-size UK garden centre chain with 12 locations. The prospect is also considering two larger agencies. The decision maker is the Marketing Director, who has told us she values transparency and long-term partnerships over flashy pitches."

Write a proposal outline that: leads with their specific situation rather than our credentials; demonstrates understanding of the garden centre sector and its seasonal patterns; presents three pricing tiers (Starter / Growth / Full Service); and ends with a section on how we measure success and report to clients.

Structure: section headings with one-sentence notes on content. Not the full proposal — the outline I will then develop. Max two pages when written up."

What AI produces: A highly specific outline that could only work for this client, this agency, and this competitive context. The structure respects the decision maker's stated values. The pricing presentation is appropriately tiered. The language is right for the relationship.

WHAT CHANGED?

WHAT WAS ADDED	WHY IT MATTERED
Role: senior account director	Contextualised the voice and level of expertise
Specific client name and context	Anchored content in a real situation
Decision maker's stated values	Shaped the strategic angle of the proposal
Sector knowledge requirement	Prompted sector-relevant thinking
Clear deliverable	"Outline I will develop" prevented over-generation
Output constraints	Kept length appropriate

CHOOSING THE RIGHT ROLE

Role assignment is optional but often valuable. The rule of thumb: assign a role when the task has a domain, a voice, or a level of expertise that you want AI to reflect.

Useful role frames:

- Professional identity: *"You are an experienced employment solicitor"*
- Expertise level: *"You are a secondary school science teacher"*
- Perspective: *"You are a sceptical reader encountering this argument for the first time"*
- Function: *"You are a copy editor reviewing this draft"*

Role frames that don't help much: overly elaborate backstories, fictional identities with no relevance to the task, vague superlatives ("you are the world's greatest expert"). Keep roles grounded in what kind of expertise or perspective would actually be useful for this task.

ASSIGNING A SPECIFIC, ACTIONABLE JOB

The instruction "help me with X" is not a job. It is an open invitation that AI will answer with whatever seems most plausibly helpful — which is usually the most generic interpretation of the task.

A job has:

- A clear action verb: write, summarise, extract, compare, identify, structure, rewrite, analyse
- A clear deliverable: a document, a list, a table, a set of questions, a draft
- A clear scope: this document, these notes, this specific situation

Compare:

- *"Help me with my CV"* → AI will ask what kind of help or produce something generic
- *"Rewrite the experience section of my CV to emphasise project management and stakeholder communication, using active verbs, targeting senior operations manager roles in the financial services sector"* → AI has a job

TRY THIS

Write down a task you regularly do that takes more than an hour. Now describe it as a job for AI using this structure:

[Role, if relevant]. [Action verb] [specific deliverable] using [source material]. The output should [format/length/tone specifications]. The intended reader is [audience]. Do not [explicit exclusion if needed].

Notice how much this forces you to clarify what you actually want before AI starts.

WATCH OUT

Assigning an expert role does not make AI's output more accurate. It may make it more fluently confident — which, as noted in Chapter 4, can be more convincing and equally wrong. Role framing improves format, voice, and structure. It does not improve factual accuracy. Verify subject-matter claims regardless of the role you assigned.

The context gap

Here is the uncomfortable truth about most AI interactions: AI starts every conversation knowing nothing about you, your work, your organisation, your history with this client, your professional constraints, or the specific situation you are dealing with.

You know all of these things. They seem obvious. So you don't mention them.

This is the context gap — the difference between what AI needs to know to give you a genuinely useful answer, and what you actually provide. It is the single most valuable gap to close.

WHAT COUNTS AS CONTEXT

Context is not just background information. It is any fact about your situation that would change how a knowledgeable person would approach the task. It includes:

Situational context — What is happening? What is the history? Why does this matter now?

Audience context — Who is going to read, use, or act on this output? What do they know? What do they need? What are their constraints?

Relationship context — What is the relationship between parties? Long-standing client? New contact? Internal colleague? Adversarial negotiation?

Constraint context — What must the output not include? What word count, format, or style requirements apply? What legal, regulatory, or policy constraints apply?

Resource context — What documents, data, or previous work is relevant? What has already been tried?

Goal context — What does success look like? What decision will this output inform? What action will it prompt?

REAL-WORLD EXAMPLE

David runs a small consulting firm. He is writing a proposal for a client renewal and asks AI to help draft the pricing section.

Without context: AI produces a generic template for presenting consulting fees — day rates, retainer structures, and payment terms that could apply to any consulting firm anywhere.

With context: David tells AI that this client has been with the firm for four years, has three times mentioned "feeling nickel-and-dimed" by itemised billing, is a family business where the owner and CFO are the same person, and that David wants to move to a monthly retainer to create predictable revenue on both sides. He also notes the client's previous annual spend was around £42,000.

The output AI produces with this context is a specifically targeted pricing section that addresses the client's stated frustration, proposes a retainer at a price point consistent with their history, and frames the retainer in terms of what the client gets (predictability, priority access, no surprises) rather than what David gets. It is a usable draft. The generic version was not.

The difference was entirely in the context David provided.

THE CONTEXT TEST

A useful diagnostic: before sending a prompt, ask yourself:

"If I had just hired a very good human assistant who was new to this job, what would they need to know before I asked them to do this?"

That question surfaces most of the context you are probably omitting. A good assistant would ask about the audience, the history, the constraints, the goal, and the format before starting work. Providing that context in your prompt removes the need for them to ask — and removes the iterations required when AI guesses wrong.

THE BEFORE → BETTER → BEST: CONTEXT

BEFORE

"Write an email to our supplier about a late delivery."

Output: A generic, somewhat formal complaint email about a late delivery with placeholder information. Polite but weightless.

BETTER

"Write an email to our supplier, BuildRight Ltd, about a delivery that is now 12 days late. The order was for 500 units of packaging material needed for a Christmas promotional campaign. The delay means we cannot start production on time and we risk missing the campaign window."

Output: A specific email that references the order type, the operational impact, and the time pressure. Usable with minor edits.

BEST

"Write an email to BuildRight Ltd, our packaging supplier of five years. Our contact is James Thornton in their sales team. The delivery of 500 units of promotional packaging (order number BRTX-2241) is now 12 days late. This delivery was critical for our Christmas campaign — production was due to start on Monday and we have an external print deadline of 18th November.

We do not want to damage the relationship, which has been excellent until this point. The tone should be: firm and clear about the business impact, but not accusatory. We want two things: a confirmed delivery date in writing and an explanation of what happened.

Do not threaten to cancel the order or use legal language at this stage."

Output: An email that reflects the history of the relationship, is specific about the business impact and timeline, makes two clear requests, and observes the explicit constraint about not escalating. Might be sent with no editing.

WHAT CHANGED?

CONTEXT ADDED	WHY IT MATTERED
Five-year relationship	Shaped the non-accusatory tone
Named contact	Made the email personal
Order number	Made the reference specific and professional
Production deadline	Explained why timing matters so much
Print deadline date	Created concrete urgency without exaggeration
Two specific requests	Focused the email on clear outcomes
Explicit constraint (no legal language)	Prevented escalation AI might otherwise include

PROVIDING CONSTRAINT CONTEXT

Constraints are particularly easy to omit because they feel obvious to you. "Obviously I can't mention the legal proceedings." "Obviously I shouldn't mention the budget number." "Obviously the tone needs to be formal."

These are obvious to you. They are invisible to AI. State every constraint that would make the output unusable if violated.

Useful constraint phrasing:

- *"Do not include [X]."*
- *"This cannot mention [X] for legal reasons."*
- *"The tone must not suggest [X] — this would damage the relationship."*
- *"UK spelling and law applies throughout."*
- *"This will be reviewed by [X] so should not contain [Y]."*

PRO TIP

For recurring tasks where the context is always the same, you don't need to retype it every time. Chapter 11 covers persistent instructions — a way of establishing your standard context once so it applies automatically to every conversation. The context you would type for every client email, every report, every team communication becomes a standing brief that AI carries forward.

WATCH OUT

Providing irrelevant context can dilute the output as much as providing too little. If you paste a twenty-page document when AI only needs three paragraphs, you risk getting a response that is unfocused across all twenty pages. Context should be relevant to this task. If in doubt, introduce the document and tell AI which section matters: *"I'm attaching our full company handbook, but for this task only the redundancy policy section (pages 14-17) is relevant."*

Why most instructions are too vague

Instructions are where most users put most of their effort — and still get disappointing results. The reason is usually not that the instructions are wrong. It is that they are underspecified.

"Make it more professional" is an instruction. But "more professional" in what dimension? More formal vocabulary? Shorter sentences? No contractions? Different paragraph structure? Fewer qualifications? More precise language?

AI will choose. And it will choose the version of "more professional" that appears most frequently in its training data for this type of text — which may not be what you meant at all.

Effective instructions are specific enough that a capable human could follow them without asking a clarifying question.

THE SPECIFICITY TEST

Before sending an instruction, ask: *"Could two different people interpret this differently?"*

If yes, it is underspecified.

"Make it shorter" → How much shorter? Half? A third? A specific word count?

"Make it more engaging" → More engaging how? More conversational? Shorter paragraphs? An opening question? Personal examples?

"Make it suitable for our audience" → Who is the audience and what do they need?

"Improve the structure" → More headings? Different order? Remove sections? Add an introduction?

Each of these pairs shows an instruction that fails the specificity test, and the questions that reveal what was actually meant. Answer those questions in the instruction.

POSITIVE AND NEGATIVE INSTRUCTIONS

Instructions can specify what to do (positive) or what not to do (negative). Both are important. Most people use only positive instructions and then wonder why the output contains things they didn't want.

Negative instructions are particularly useful for:

- Tone and register: *"Do not use exclamation marks. Do not use filler phrases like 'I hope this finds you well.'"*
- Content constraints: *"Do not mention the merger — this is not public."*
- Style: *"Do not use bullet points — this should be flowing prose."*
- Claims: *"Do not make specific claims about delivery times — these vary by region."*

A well-instructed prompt often includes both: *"Write this in a direct, confident tone. Do not use passive voice. Do not hedge claims with phrases like 'it could be argued' or 'some might say.'"*

INSTRUCTION TYPES AND WHEN TO USE THEM

Different tasks benefit from different instruction types:

Format instructions — What structure should the output take?

"Use a numbered list." / "Present this as a table with columns for X, Y, Z." / "Use three short paragraphs, no headings." / "One page maximum, A4, with a heading and four bullet points."

Tone instructions — How should it sound?

"Direct and professional." / "Warm but concise." / "Technical but accessible — assume the reader has a science degree but is not a specialist in this field." / "In the voice of the existing document I'm attaching."

Priority instructions — What matters most?

"Lead with the cost saving, not the process improvement." / "The most important point must be in the first sentence." / "The client's concern about timeline is more important than the features list — address it first."

Exclusion instructions — What to leave out?

"Do not include implementation details — this is for a non-technical audience." /

"Exclude anything older than 2024 — the reader will already know the history." /

"Do not include a summary — the reader will read the whole thing."

Role instructions — From whose perspective?

"Write this as if the reader has never heard of our company." / "The reader is a

sceptic — address their likely objections before they ask."

THE BEFORE → BETTER → BEST: INSTRUCTIONS

BEFORE

"Rewrite this paragraph to make it better."

[Paragraph: "Our team has extensive experience in the delivery of complex technology projects. We have worked with a wide range of clients across multiple sectors and have a proven track record of delivering projects on time and on budget."]

Output: AI produces a slightly more polished version of the same paragraph. Same ideas, slightly varied vocabulary. The reader would not notice a meaningful difference.

BETTER

"Rewrite this paragraph to remove clichés and make it more specific and credible. Replace phrases like 'extensive experience' and 'proven track record' with concrete language."

Output: AI produces a version that avoids the worst clichés and becomes more specific in tone, though it invents specifics it doesn't have.

BEST

"Rewrite this paragraph for inclusion in a proposal to a housing association considering a digital transformation project. The audience is a director who has seen many similar proposals and is sceptical of vague claims.

Instructions:

- Remove all clichés: 'extensive experience', 'wide range', 'proven track record', 'multiple sectors' must all go.*
- Replace vague claims with specific patterns of evidence. I will provide specifics in a moment — for now, leave [BRACKETS] as placeholders for: number of years in the sector, number of completed projects, one specific client sector, and one specific result we achieved.*
- First sentence should establish relevance to housing sector specifically, not generic technology experience.*
- Maximum two sentences. Punchy, not promotional."*

Output: A tightly constructed two-sentence paragraph that is specific in structure, leaves appropriate placeholders, and is pitched correctly at the sceptical reader.

WHAT CHANGED?

INSTRUCTION ELEMENT	WHY IT MATTERED
Audience specified	Shaped the pitch level and scepticism antidote
Banned phrases listed explicitly	Prevented the most common failure
Placeholder pattern	Kept AI honest about missing specifics
Structural requirement (first sentence)	Prevented the generic opener
Length constraint (two sentences)	Forced concision
Tone instruction (not promotional)	Prevented the default enthusiastic pitch

TRY THIS

Take the last AI output you were dissatisfied with. Identify one thing that was wrong with it — too long, wrong tone, too generic, incorrect emphasis, wrong format.

Write a follow-up instruction that specifically addresses that one thing, using the specificity test: could two people interpret this instruction differently?

If yes, refine it until the answer is no.

WATCH OUT

Instructions that conflict produce confused output. If you say "be brief and comprehensive," "be formal but conversational," or "be specific but applicable to all situations," AI will attempt to honour all instructions simultaneously and usually succeed at none of them. When instructions conflict, decide which takes priority and say so: *"Be brief — prioritise this over comprehensiveness if necessary."*

Showing versus telling

There are two ways to specify what you want from AI. You can describe it, or you can demonstrate it.

Describing works for many things. But when tone, format, or style is what you care about, demonstrating is almost always more effective.

This is because language about language is inherently imprecise. "Warm but professional" means something different to you than it does to AI. "Clear and engaging" is interpreted through thousands of examples in AI's training data, most of which are not exactly what you meant. The gap between your description and AI's interpretation is where most style failures live.

An example collapses that gap. Instead of asking AI to match your description of a style, you are showing it the style and asking it to match that.

TYPES OF EXAMPLES

Examples can take several forms:

tone examples

A piece of your own writing, or a piece of writing in the voice you want. "Write in the same tone as this email I wrote last year" or "match the style of this example I'm pasting in."

format examples

A previous document in the format you need. "Use the same structure as this report section." "Produce output that looks like this."

quality standard examples

A high-quality output you want to match or improve on. "This is an example of the level of specificity I want."

negative examples

Something you don't want. "Do not write like this" is sometimes as effective as "write like this."

WHY IT WORKS

Providing examples is what AI researchers call "few-shot prompting" — giving AI a small number of examples from which to infer the desired pattern. The technical reason it works is that examples constrain the generation space toward outputs with similar characteristics: similar vocabulary choices, similar sentence length, similar structural decisions. Your description of "professional" is abstract; an example of it is concrete.

THE BEFORE → BETTER → BEST: EXAMPLES

The scenario: Margaret wants a weekly AI-generated summary of news relevant to her industry (sustainable packaging) to send to her team. She wants it to sound as if she wrote it, not AI.

BEFORE

"Write a summary of this week's sustainability news for my team."

Output: A formal, comprehensive news summary. Heavy paragraphs. AI-standard prose. Margaret would have to rewrite it entirely before it sounds like her.

BETTER

"Write a summary of this week's sustainability packaging news for my team. Tone: conversational, direct, not overly formal."

Output: Less formal, but still not quite right. "Conversational" is interpreted differently by different people, and AI doesn't know what Margaret sounds like.

BEST

"I write a weekly news summary for my team of 8. Here is an example of one from three months ago that captured the tone exactly right:

[EXAMPLE SUMMARY]

"Three things from this week:

- 1. The EU is tightening its position on recycled content claims — this matters for our German clients more than anywhere else. Worth watching.*
- 2. Paperboard prices stabilised in October after six months of increases. Our procurement team will already know this, but the sector analysis I've linked to is good background.*

3. *Maersk announced a new bio-methanol shipping route. Too early to know if this is real progress or PR. I'll share more when I know more.'*

Write this week's summary in the same format and voice. Source material I'm attaching covers: [pasted or uploaded articles]. Keep it to three items unless something genuinely warrants a fourth."

Output: A summary that reads like it could have been written by Margaret — same three-item structure, same direct tone, same characteristic qualification ("too early to know"). She edits two words and sends it.

WHAT CHANGED?

WHAT WAS ADDED	WHY IT MATTERED
Example of actual past output	Gave AI a concrete target instead of abstract description
Format from example	Three numbered items, short, no preamble
Voice characteristics from example	Opinionated, direct, honest about uncertainty
Source material provided	Grounded in real content rather than general knowledge
Constraint on length	Prevented overgeneration

USING YOUR OWN WRITING AS AN EXAMPLE

One of the most practical applications of example-giving is using your own previous writing to teach AI your voice. This is particularly useful for:

- Regular communications: weekly updates, team emails, client newsletters
- Documents with a house style: reports, proposals in a format you've established
- Personal correspondence: emails where you want AI to draft and you want it to sound like you

The method: provide two or three examples of your writing in the relevant genre, and ask AI to match the style. Longer examples give AI more to work from. If you have a document that perfectly captures your voice or style, use that as the primary

reference.

PRO TIP

If you save a "style reference" document in your AI workspace (Chapter 11), you don't need to paste examples into every prompt. You reference the document once in your standing instructions and it's available for every subsequent interaction.

WHEN NOT TO USE EXAMPLES

Examples constrain AI toward the example. This is usually what you want — but not always.

If you want AI to be creative, novel, or to produce something better than your existing examples, providing examples can limit the output. In those cases, describe the qualities you want and explicitly say: *"Feel free to go beyond this starting point."*

Similarly: if the examples you have are not actually that good, AI may faithfully replicate their weaknesses. Use examples you would genuinely like to match.

WATCH OUT

Examples are most effective for style and format. They are less effective for accuracy. Providing an example of a data table you want does not teach AI to fill it with correct data. You are specifying format, not content. Verify content as normal regardless of how good the format looks.

The output problem

You have provided excellent context, a clear job description, specific instructions, and a strong example. AI produces a response.

It is too long.

It is the wrong format.

The most important point is buried in paragraph four.

It uses a slightly different structure than you specified.

This happens because output defaults are baked into AI's training. Without explicit specifications, AI produces the type, length, and format of output most commonly associated with your type of request. That default is often close but rarely exactly right.

Controlling output is the final step in building an effective interaction — specifying not just what to produce, but what the result should look like when it arrives.

THE FIVE OUTPUT DIMENSIONS

Five dimensions of output are almost always worth specifying for significant tasks:

1. LENGTH

The most commonly underspecified dimension. How long? Not "short" — how many words, sentences, bullet points, or paragraphs? "A one-page briefing," "a three-bullet summary," "a maximum of 150 words," "five sentences" — these are unambiguous. "Brief" is not.

2. FORMAT

How should the information be structured? Prose paragraphs, numbered list, bullet points, table, structured document with headings, slide deck content, script, dialogue? Different formats serve different purposes. The format should match how the reader will use the output.

3. STRUCTURE

If the output has sections or parts, what are they and in what order? Specifying structure prevents AI from inventing a structure that may not serve your purpose. "The document should have four sections: Background, Key Findings, Recommendation, Next Steps" is better than allowing AI to choose.

4. TONE

How should it sound? The vocabulary level, the degree of formality, the use of first person or third person, the amount of hedging or qualification, the presence or absence of technical language — all are specifiable. Use examples when description alone is insufficient (Chapter 8).

5. AUDIENCE

Who will read this? What do they know? What do they need? The same information presented to a technical expert and a first-time reader should be substantially different. Specifying audience tells AI how to calibrate complexity, assumed knowledge, and depth of explanation.

WHY IT WORKS

AI has no inherent preference for one format over another. It defaults to the format most common in its training data for a given type of request. Explicit output specifications don't constrain AI in ways that affect quality — they redirect its generation toward the format you need. Specifying output is not about limiting AI; it is about using what AI has already generated to its best effect for your purpose.

THE BEFORE → BETTER → BEST: OUTPUT CONTROL

The scenario: James needs a summary of a forty-page government report on graduate employment for a blog post he is writing.

BEFORE

"Summarise this report on graduate employment."

Output: Six paragraphs, dense, comprehensive, presenting findings in the order they appear in the report. Difficult to read. Not formatted for a blog audience. James would have to restructure it entirely.

BETTER

"Summarise the key findings of this report for a general audience interested in graduate careers. Keep it short and clear."

Output: Shorter, more readable, but still not structured for a blog post. "Short and clear" was under-specified.

BEST

"Summarise the key findings of this graduate employment report for a blog post aimed at final-year students. They are interested in practical takeaways about their job prospects, not policy analysis."

Format: Three sections with short headings:

- 1. What the data says (two to three bullet points with the most significant statistics)*
- 2. What this means for you (two to three practical implications in plain language)*
- 3. What to do about it (one or two concrete suggestions)*

Maximum 300 words total. Active voice. No jargon. Read level: a confident 20-year-old who reads the Guardian.

Do not include methodology, policy recommendations, or anything that doesn't directly affect a job-seeking student."

Output: A structured, audience-appropriate blog summary. The sections are clear, the language is right for the audience, and the length is correct. James reads it once, adds a paragraph of his own at the end, and posts it.

WHAT CHANGED?

OUTPUT SPECIFICATION	WHY IT MATTERED
Three named sections	Imposed a usable structure
Bullet count per section	Prevented overgeneration
Maximum 300 words	Made length unambiguous
"Active voice"	Improved readability
Reading level specified	Calibrated vocabulary and complexity
Explicit exclusion list	Prevented irrelevant policy content

SPECIFYING FORMAT: A PRACTICAL REFERENCE

When you need:

THIS	SAY
A concise summary	"Three to five bullet points, one sentence each"
A structured document	List the sections explicitly with brief descriptions
A table	"Present this as a table with columns for [X], [Y], [Z]"
Prose, not bullet points	"Write in flowing paragraphs — no bullet points"
An email	"Subject line + body, professional tone, under 150 words"
An outline	"Numbered outline with second-level sub-points"
A list	"Numbered list, [n] items, no more"
A one-pager	"One A4 page of content: heading, three sections"

CONTROLLING OUTPUT MID-CONVERSATION

Output control is not only an upfront specification. It is also a live tool in conversation.

When the output is close but not quite right, you can redirect it with a specific instruction:

- *"Good structure but too long — cut it by half."*
- *"Remove the executive summary section — I don't need it."*
- *"This is in bullet points — rewrite as prose paragraphs."*
- *"The third point is the most important — lead with it."*
- *"Change the tone — this reads too formally for the audience."*

Each of these is a specific, actionable correction. AI will act on them precisely. The instruction "make it better" is not a specific correction. Be specific about the dimension you want changed.

TRY THIS

For your next significant AI task, write out your output specification separately before you write the main prompt. Answer these five questions:

1. How long?
2. What format?
3. What structure?
4. What tone?
5. Who is the audience?

Then build those answers into the prompt. Notice whether the first output is closer to what you need before you have to ask for revisions.

WATCH OUT

Excessive output specification can overconstrain the response. If you specify every sentence, paragraph, and word choice, you are writing the document yourself and asking AI to type it. Specify the frame — length, format, structure, tone, audience — and leave AI to fill it in appropriately. The goal is a usable first draft, not a perfect document that required more work to specify than it would have taken to write.

WHAT PART II HAS ESTABLISHED

You can now build an effective AI interaction from the ground up:

- Give AI a clear role and a specific job, not a vague question
- Provide the context that closes the gap between what AI assumes and what you actually need
- Write instructions specific enough that a capable person could follow them without asking clarifying questions
- Show AI examples when style, format, or tone matter
- Specify the output in terms of length, format, structure, tone, and audience

Part III moves from individual interactions to the environment in which they happen — and the significant gains available from building that environment deliberately.

Most people use AI as if it were a calculator: pick it up, use it once, put it down. Every new calculation starts from zero. Part III shows you how to change that — building a persistent

environment where AI carries your context forward, knows your situation, and becomes genuinely more useful the more you invest in setting it up.

The repetition problem

If you use AI more than occasionally, you have encountered this situation: you are about to start a new conversation and you find yourself typing, again, the things AI should already know.

"I work in marketing at a mid-size UK financial services company. I write for a non-specialist audience. I prefer a formal-but-accessible tone. UK spelling throughout. Never use exclamation marks..."

This is not how a useful working relationship operates. A good colleague doesn't need to be reminded of your name and job every time you speak to them.

The reason this happens is that AI has no persistent memory across conversations by default. Every new conversation begins with a blank slate. The context you built in Tuesday's session is gone on Wednesday. You are starting from zero, every time.

This is one of AI's most significant practical limitations — and one that is largely solvable.

THREE STRATEGIES FOR NOT STARTING FROM ZERO

STRATEGY 1 — SYSTEM INSTRUCTIONS (PERSISTENT INSTRUCTIONS)

Most AI platforms now offer a way to set persistent instructions that apply to every conversation. Sometimes called "custom instructions," "system prompts," or "memory settings," these are a set of facts and preferences you provide once and

that AI carries forward automatically.

This is the most powerful strategy for regular users. Chapter 11 covers it in depth.

STRATEGY 2 — PROJECTS

ChatGPT, Claude, and several other platforms offer a "Projects" feature that groups conversations, persists context, and maintains uploaded documents within a dedicated workspace. For recurring work — client accounts, ongoing projects, regular document types — projects eliminate most setup repetition.

STRATEGY 3 — CONTEXT DOCUMENTS

When platform features are unavailable or insufficient, a context document you paste in at the start of relevant conversations serves the same purpose. A single well-crafted paragraph covering who you are, what you're doing, and what your preferences are can be copied and pasted in seconds.

WHY IT WORKS

Persistent context doesn't improve AI's underlying capability. It changes what AI has to work with. An AI that knows your role, your audience, your style preferences, and your constraints produces better first drafts and needs fewer iterations. The investment is front-loaded — you write the context once — and the return compounds with every subsequent conversation.

WHAT BELONGING TO THE SAME WORKSPACE GIVES YOU

When you work within a project or persistent environment:

- Your uploaded documents remain available without re-uploading
- Your style and tone preferences apply without re-specifying
- Your role and context are understood without re-explaining
- Outputs from previous conversations can inform new ones
- Multiple AI conversations about the same work form a coherent body of output

This is a qualitative shift in how AI integrates into professional work. It moves AI from an occasional tool you use for one-off tasks to something closer to a working partner for ongoing projects.

REAL-WORLD EXAMPLE

Sarah manages a marketing account for three different clients. She sets up a separate Project for each:

Project: Hartley Financial — contains the brand guidelines document, a style reference email, a note on the client's preferences ("informal tone in client-facing copy; formal in regulatory documents"), the previous quarter's campaign materials, and a standing instruction that the client's compliance team must review any copy before publication.

Project: Meridian Retail — contains the target audience profile, previous copy that worked well, a note on what the Marketing Director has rejected in the past, and the seasonal campaign calendar.

When Sarah opens the Hartley Financial project to write new email copy, AI already knows the brand, the style, the constraints, and the history. The first draft is usable. Without the project, she would spend the first five minutes of every session re-briefing context she's already established.

WATCH OUT

Persistent context can become stale. If your role changes, a client's preferences shift, or a project evolves, the standing context needs to be updated. Review your projects and persistent instructions every few months. Outdated context is better than no context, but not by much — AI will confidently apply the wrong assumptions.

Setting up persistent instructions

Persistent instructions are the highest-leverage single investment most regular AI users can make. Writing them well takes thirty minutes. The return on those thirty minutes compounds with every subsequent AI conversation.

WHAT TO INCLUDE IN PERSISTENT INSTRUCTIONS

The most valuable things to specify in persistent instructions:

1. YOUR PROFESSIONAL CONTEXT

Who you are, what you do, what organisation or sector you work in. This calibrates AI's assumptions about what background knowledge is relevant and what terminology is appropriate.

Example: "I am a senior solicitor at a UK mid-market law firm. My clients are typically owner-managed businesses. I communicate with clients, not other lawyers."

2. YOUR COMMUNICATION PREFERENCES

Tone, vocabulary level, things you never want to see in output. This is where you codify the things you would otherwise correct in every first draft.

Example: "Write in plain English. No jargon unless I request it. No excessive hedging. No bullet lists unless I specify them. UK English throughout. Professional but accessible."

3. YOUR STANDARD AUDIENCE

Who reads what you produce? This prevents AI from pitching responses at the wrong level.

Example: "Assume I am writing for a non-specialist reader unless I specify otherwise. Never assume technical knowledge of legal process or financial instruments."

4. KEY CONSTRAINTS THAT ALWAYS APPLY

Anything that, if AI ignored it, would make the output unusable. Legal constraints, brand constraints, regulatory constraints, relationship sensitivities.

Example: "I work in financial services. Do not make specific investment recommendations. Always include appropriate caveats where content could be interpreted as financial advice."

5. THINGS YOU WANT AI TO ALWAYS DO

Behaviours you want as default, not just when you remember to ask.

Example: "At the end of any factual output, flag any claims I should verify independently. If you are uncertain about any fact, say so explicitly rather than stating it with false confidence."

A TEMPLATE FOR PERSISTENT INSTRUCTIONS

Here is a starting template. Adapt it to your situation:

```
About me:
[Your role, sector, organisation type, years of experience if relevant]

What I typically need AI for:
[Your most common use cases – writing, analysis, research, editing,
planning]

My communication preferences:
[Tone, vocabulary level, formatting preferences, things you never want]

My standard audience:
[Who reads your output – clients, colleagues, the public, specialists,
generalists]

Key constraints that always apply:
[Legal, regulatory, brand, relationship, or accuracy constraints]

What I always want AI to do:
[Default behaviours you want in every interaction – verification flags,
format defaults, etc.]
```

TRY THIS

Set up your persistent instructions today. Use the template above. Start with the sections that feel most obvious — your professional context and communication preferences — and add the others as you notice what you keep re-specifying in individual conversations.

After two weeks, review the instructions. What are you still correcting in every conversation? That's a gap in your persistent instructions. Add it.

SETTING UP A PROJECT

Projects are more powerful than persistent instructions because they can also hold documents, ongoing outputs, and accumulated context. Here is how to approach setting one up:

STEP 1 — DEFINE THE PROJECT'S PURPOSE

What work is this project for? One client? One recurring task type? One ongoing work project? The answer determines what belongs in it.

STEP 2 — WRITE THE PROJECT INSTRUCTIONS

These are like persistent instructions but scoped to this project. Include the context specific to this work that doesn't belong in your general instructions.

STEP 3 — UPLOAD RELEVANT DOCUMENTS

What documents does AI need to do this project's work? Upload them. Brand guidelines, style guides, reference documents, previous outputs, client briefs.

STEP 4 — NAME CONVERSATIONS CLEARLY

Within the project, give conversations descriptive names. "Email sequence for Q4 campaign" is more useful six weeks later than "New Chat (3)."

PRODUCT NOTE (VERIFIED SEPTEMBER 2026):

Projects are available in ChatGPT (all paid plans), Claude (Pro and above), and Gemini Advanced. Feature details, including what types of documents can be uploaded and how many, vary by platform and are subject to change. Verify current capabilities and limits on the platform's help centre before relying on specific features.

Volatility: HIGH

THE PROJECT FOR RECURRING WORK

Some of the most valuable AI projects are not for unique work, but for recurring tasks. If you produce the same type of document weekly or monthly — a status report, a client briefing, a team update — a project containing your template, your style reference, and your standard content elements saves significant setup time.

For recurring work, the project contains:

- A template for the document type
- Style and tone specifications
- An example of a completed version you were happy with
- Any standard information that appears in every edition (team structure, project status, recurring context)

Each week or month, you open the project, paste in the new content or data, and ask AI to produce the new version to the same standard as the example.

Why you should give AI documents, not descriptions

There is a consistent pattern in how people under-use AI: they describe the document instead of providing it.

"I have a market research report that says our target demographic is 25-34..."

"Based on our company values, which include innovation and customer focus..."

"Our pricing is roughly in the mid-range for the sector..."

These descriptions are always incomplete. They filter the document through the user's current focus and omit everything that doesn't seem immediately relevant — including things that might matter.

When you provide the actual document, AI can read all of it. It can notice things you didn't highlight. It can extract the specific data point you forgot to mention. It can identify tensions between different sections that your description smoothed over.

Provide documents, not descriptions of documents.

WHAT AI CAN DO WITH DOCUMENTS

The practical capabilities, as of late 2026:

Summarise — Extract key points, findings, or decisions at a specified level of detail.

Extract — Pull out specific information: all the action items from meeting notes, all the financial figures from a report, all the dates and deadlines from a contract.

Analyse — Identify patterns, inconsistencies, risks, or significant points: flag anything in this contract that differs from our standard terms; identify which sections of this report have the weakest evidence.

Compare — Compare multiple documents: these two supplier proposals against the same set of criteria; this year's report against last year's.

Answer questions — "What does this contract say about notice periods?" "What is the research paper's conclusion about the effect of X on Y?"

Rewrite or adapt — Take an existing document and adapt it: rewrite in simpler language; convert to a different format; update with new information.

WHAT TO WATCH OUT FOR WITH DOCUMENTS

AI can make mistakes when working with documents, and those mistakes can be difficult to spot:

Hallucinating content — AI may generate a plausible-sounding claim about a document that is not actually in the document. Always verify specific extracted facts against the original.

Missing important nuance — AI may summarise a complex legal or financial document in a way that loses critical nuance. "The contract allows termination with 30 days notice" may be accurate but miss that this only applies to non-performance situations.

Incorrect numbers — AI sometimes misreads or miscalculates figures in documents. Any number extracted from a document should be verified against the source.

Losing track of long documents — Very long documents may not be fully attended to. Important information in the middle or end of a very long document is more likely to be missed than information at the beginning.

VERIFY

After any significant extraction or analysis task on a document:

- 1.** Check at least two specific extracted facts against the source document directly
- 2.** Ask AI what it considered to be the most uncertain or ambiguous parts of the document
- 3.** Ask AI to flag any information it couldn't find in the document that your prompt might have led it to expect

These three checks catch the most common document-analysis failures.

DOCUMENT TYPES AND HOW TO USE THEM

CONTRACTS AND LEGAL DOCUMENTS

Most useful for: identifying specific clauses, comparing to standard terms, extracting obligations and deadlines, generating a plain-English summary.
Always verify: any claim about rights, obligations, or liabilities against the exact contract language.

FINANCIAL REPORTS AND DATA

Most useful for: summarising findings, extracting key figures, identifying trends, generating plain-English explanation.
Always verify: specific numbers against the source.

RESEARCH PAPERS

Most useful for: summarising methodology and findings, identifying key claims, extracting conclusions.
Always verify: is this what the paper actually concludes? Has the summary missed important limitations?

MEETING NOTES AND TRANSCRIPTS

Most useful for: extracting action items, producing a summary, identifying decisions made, creating a follow-up email.

Watch out for: misattributing who said what, missing an action item from a long transcript.

INTERNAL DOCUMENTS

Most useful for: adapting existing content, maintaining consistency, finding relevant previous work.

Watch out for: applying information from outdated documents without realising it.

PRO TIP

When uploading a long document, help AI navigate it. Tell AI which sections are most relevant: "The contract is 22 pages. For this task, the most relevant sections are the Termination clause (section 8) and the Indemnities (section 14). You can skim the rest unless it's relevant." This focuses AI's attention and tends to produce more accurate extraction from the sections that matter.

AI is not just a text generator

The AI tools available in late 2026 can do considerably more than generate text. Understanding which capabilities exist — and which don't — determines what you can ask for, what you need to verify externally, and where AI can genuinely replace tasks that previously required separate tools.

This chapter introduces the main tool categories. Because capabilities vary by platform, by plan level, and change frequently, all specific capabilities include a Product Note with a volatility rating. High-volatility items are not in the main body.

WEB SEARCH

Most major AI platforms can search the web during a conversation. This changes one of AI's most significant limitations: its knowledge cutoff.

Without web search, AI's knowledge has a cutoff date. If you ask about recent events, current prices, or the latest data, AI either doesn't know or (more dangerously) generates plausible-sounding but outdated information.

With web search enabled, AI can retrieve current information and incorporate it into responses. It can check whether something has changed, find recent statistics, and access information published after its training cutoff.

WHAT WEB SEARCH CHANGES:

- AI can answer questions about recent events
- AI can retrieve current data rather than estimating from outdated training data
- AI can find specific sources rather than relying on memory

WHAT WEB SEARCH DOES NOT CHANGE:

- AI still interprets and summarises what it finds, and can still get that interpretation wrong
- Not all information is publicly accessible on the web
- AI can still be wrong about which sources to trust
- You still need to verify important claims

PRODUCT NOTE (VERIFIED SEPTEMBER 2026):

Web search is available in ChatGPT (all plans), Claude (Sonnet and above on Pro plans), and Gemini (all plans). Default on/off settings and search quality vary. Always confirm whether search was used in a given response if accuracy matters.

Volatility: MEDIUM

CODE EXECUTION AND CALCULATION

Some AI platforms can write and run code within the conversation, enabling reliable calculation and data processing.

This is significant because it addresses one of AI's consistent weaknesses: arithmetic and calculation. When AI uses code execution, it doesn't estimate calculations — it runs them. The result is accurate in the same way a spreadsheet result is accurate (if the formula is right).

Practical uses:

- Accurate calculations on data you provide
- Processing and analysing datasets
- Running formulas or statistical operations
- Generating charts and graphs from data

What this changes: Calculations performed via code execution are reliable. Calculations performed without code execution — AI simply generating the answer as text — are not reliably accurate, particularly for anything beyond simple arithmetic.

The rule: If your task involves maths that matters, either use a platform with code execution enabled, or verify the calculation yourself.

PRODUCT NOTE (VERIFIED SEPTEMBER 2026):

Code execution ("Advanced Data Analysis" in ChatGPT, native in Claude) is available on paid plans of major platforms. CSV and spreadsheet upload for analysis is broadly supported. Consult current platform documentation for file size limits and supported formats.

Volatility: MEDIUM

IMAGE GENERATION

Some AI tools can generate images from text descriptions. This capability is useful for:

- Rough concept visualisations for design briefs
- Placeholder images for presentations
- Illustrations for non-commercial documents
- Visual brainstorming

WHAT TO BE CAREFUL OF:

- Generated images may not be suitable for commercial use without checking the platform's terms
- Generated images of people can produce photorealistic results that may be misleading out of context
- Quality varies significantly across platforms

PRODUCT NOTE (VERIFIED SEPTEMBER 2026):

Image generation is integrated into ChatGPT (DALL·E) and available as a separate step in several other platforms. Licensing terms for commercial use vary and should be verified on the platform's own documentation before use.

Volatility: HIGH — See Appendix E for current details.

FILE CREATION AND EXPORT

AI tools increasingly support creating downloadable files: documents, spreadsheets, presentations. This is still an evolving area, with quality and reliability varying considerably.

The practical approach: use AI to generate the content and structure, then create the final formatted document in your standard tool (Word, Excel, PowerPoint, Google Workspace). Attempting to produce final, formatted output directly from AI is often slower and less reliable than the two-step approach.

WHAT AI TOOLS CANNOT RELIABLY DO (AS OF LATE 2026)

Equally important is knowing what AI tools currently cannot do well:

- **Send communications on your behalf** without explicit confirmation steps
- **Access your internal systems** (email, calendar, CRM, databases) without integration tools set up by your organisation
- **Make decisions** that require judgment, ethics, or accountability — AI can analyse and recommend; decisions remain with you
- **Replace professional advice** in law, medicine, finance, or other regulated domains — AI can inform; professionals advise

WATCH OUT

Tool availability changes. A capability that exists today may be modified, restricted, or removed. A capability that doesn't exist today may arrive. The best practice is to verify current capabilities on the platform's own documentation before designing a workflow around a specific tool feature. This book covers capabilities as of September 2026; check Appendix E and the platform's help centre for current status.

Why AI doesn't remember you

Every time you start a new conversation with an AI tool, you are meeting a stranger who happens to be very knowledgeable. They have never met you before. They do not know your name, your job, your preferences, your previous conversations, or anything about your situation.

This is technically described as AI having no persistent memory across conversations. The practical effect is that you spend time at the start of every conversation re-establishing context you established before.

Some platforms have begun introducing memory features — where AI can retain facts between conversations. These features are useful but imperfect, and their availability and reliability vary. This chapter covers the memory problem and its most practical workarounds.

THE FOUR TYPES OF "MEMORY" IN CURRENT AI TOOLS

IN-CONVERSATION MEMORY (ALWAYS AVAILABLE)

Within a single conversation, AI remembers everything that has been said. You can refer to "the document we discussed" or "your earlier suggestion" and AI will understand. This is reliable and consistent.

PLATFORM MEMORY FEATURES (VARIABLE)

Some platforms offer features that carry specific facts across conversations — things you told AI about yourself, preferences you established, key information you asked it to remember. These vary in scope and reliability.

PRODUCT NOTE (VERIFIED SEPTEMBER 2026):

ChatGPT has a "Memory" feature (available on paid plans) that can retain user-specified facts across conversations. Claude can be configured with persistent context through Projects and system prompts. These features evolve rapidly — check current documentation.

Volatility: HIGH — See Appendix E.

PROJECT MEMORY (THE MOST RELIABLE APPROACH)

Working within a Project creates a persistent context that AI carries across all conversations in that project. This is currently the most robust workaround for the memory problem for regular, recurring work.

DOCUMENT MEMORY (THE UNIVERSAL WORKAROUND)

A context document you paste in at the start of any conversation. Reliable across all platforms, requires no platform feature, and completely under your control.

BUILDING A CONTEXT DOCUMENT

A context document is a brief, structured summary of what AI needs to know about you and the work it will help with. It is the universal workaround for the memory problem.

Here is a template:

```
CONTEXT DOCUMENT – [Your name / role]
Last updated: [Date]

About me:
[2-3 sentences: role, sector, organisation type]

Current work context:
[What I'm working on, any ongoing projects, key relationships]

My communication preferences:
[Tone, format, UK/US spelling, things I never want in output]

Key constraints:
[Anything AI must not say, recommend, or assume]

My standard audience:
[Who reads what I produce]

Things I always want AI to do:
[Default behaviours – verification flags, format defaults, etc.]
```

A well-written context document takes twenty minutes to produce once and can be pasted into any AI conversation in seconds. It is more reliable than any platform feature and works across all tools.

MEMORY THAT MATTERS AND MEMORY THAT DOESN'T

Not everything needs to be remembered. The investment in persistent context should be proportional to how often you'll need it.

WORTH PERSISTING:

- Your role, audience, and communication style
- Recurring project context
- Key constraints that always apply
- Your quality standards and preferences

NOT WORTH PERSISTING:

- One-off tasks

- Information that changes frequently
- Details specific to a single conversation

The test: if you would type this information again in your next AI session, persist it. If it's specific to today's task and you won't need it tomorrow, don't bother.

From ad hoc to systematic

Part III has covered the building blocks of an AI environment: persistent instructions, projects, document management, tools, and memory. This chapter shows how to put them together into a workspace that reflects how you actually work.

A well-built AI workspace is not a technical achievement. It is a professional one: the result of thinking carefully about what you need AI for, what context it should always have, and how your recurring work is organised.

THE WORKSPACE AUDIT

Before building, take stock. Answer these questions:

WHAT DO I USE AI FOR MOST OFTEN?

List your five most common AI tasks. These are the tasks your workspace should be optimised for.

WHAT CONTEXT DO I RE-ENTER REPEATEDLY?

What information do you find yourself providing in almost every session? This is what should be in your persistent instructions.

WHAT RECURRING PROJECTS OR TASK TYPES DO I HAVE?

What work comes around every week, every month, or for every client? Each recurring type deserves its own Project.

WHAT DOCUMENTS DOES AI REGULARLY NEED?

What files do you find yourself uploading repeatedly? These belong in the relevant Project.

WHAT ARE MY QUALITY STANDARDS?

What makes output unusable for you? What do you always correct? These corrections belong in your persistent instructions as explicit rules.

A SIMPLE WORKSPACE STRUCTURE FOR MOST PROFESSIONALS

PERSISTENT INSTRUCTIONS (applies to everything)

- ├ Professional context
- ├ Communication preferences
- ├ Standard audience
- └ Default behaviours (verification, format, etc.)

PROJECTS

- ├ Project A: [Client / Ongoing work 1]
 - ├ Project instructions (context specific to this work)
 - ├ Reference documents (brand guide, style reference, etc.)
 - └ Ongoing conversations
- ├ Project B: [Client / Ongoing work 2]
 - └ [Same structure]
- └ Project C: Recurring tasks
 - ├ Weekly report template and example
 - ├ Standard document formats
 - └ Template prompts for recurring work

THE RECURRING TASK TEMPLATE

For any task you do on a regular cadence, a template prompt is worth building. A template prompt is a prompt with placeholders — variables you update with each use — built on a structure you've already tested and refined.

Format:

[TEMPLATE NAME]

Context (stable):

[Everything that stays the same – purpose, audience, style, constraints]

Variable inputs this week:

[PASTE CONTENT HERE]

[KEY DATES: X]

[KEY DECISIONS: X]

Output required:

[Format, length, structure – specified once, used every time]

This reduces a recurring AI task to: open the template, fill in the variables, run. The setup investment is front-loaded; the recurring use is fast.

PRO TIP

Keep a "templates" section within your AI workspace — either a Project dedicated to this, or a simple document you maintain. When you write a prompt that works especially well, save it. When you develop an interaction pattern that produces good results consistently, document it. Your prompt library becomes more valuable than any generic prompt template book because it is tested, specific to your work, and continuously refined.

WORKSPACE MAINTENANCE

A workspace that isn't maintained becomes a liability. Outdated context produces confidently wrong outputs. Review your workspace periodically:

- **Monthly:** Are the persistent instructions still accurate? Has anything changed about your role, audience, or constraints?
- **When a project ends:** Archive or close the project so it doesn't contaminate future work with irrelevant context.
- **When something keeps going wrong:** What's in your workspace that might be causing this? What context might AI be applying incorrectly?

The workspace is a living document of your working relationship with AI, not a one-time setup.

WATCH OUT

Workspaces created on one platform do not transfer to another. Your Claude Projects, your ChatGPT custom instructions, and your Gemini settings are all separate. If you use multiple platforms, you need to maintain context in each, or accept that some platforms know more about your preferences than others. Many users find it practical to choose a primary platform for their most important work and use secondary platforms with more ad hoc setup.

WHAT PART III HAS ESTABLISHED

You now have the framework for building an AI environment:

- Persistent instructions that carry your context forward automatically
- Projects that organise your work and maintain document access
- Document handling practices that maximise accuracy and catch common errors
- An understanding of tools and their limitations
- A memory strategy that works across platforms
- A workspace structure suited to professional use

Part IV moves from individual interactions and workspace setup to something qualitatively different: multi-step workflows that tackle complex tasks by breaking them into stages.

Complex work doesn't happen in a single exchange. Part IV introduces the

thinking and techniques that take AI from answering individual questions to participating in multi-stage processes — planning, drafting, reviewing, revising, and handing off.

The single-prompt ceiling

For simple, well-defined tasks — summarise this, draft that, answer this question — a single well-crafted prompt is usually sufficient. But professional work is rarely simple and well-defined.

A business proposal is not one task. It is: research the client, understand their problem, frame the approach, develop the argument, structure the document, draft each section, review for consistency, check for gaps, refine the language, and produce a final version. Each of these is a distinct cognitive task that benefits from distinct treatment.

Attempting to produce a business proposal in a single prompt produces a generic outline at best and a generic document at worst. Neither is what you needed.

The ceiling of single-prompt AI use is low. Most professional work requires workflows.

WHAT MAKES SOMETHING A WORKFLOW

A workflow, for our purposes, is any AI task that:

- Involves more than one distinct phase (research, then structure, then draft, then review)
- Produces output that becomes input to the next stage

- Benefits from human review and adjustment at one or more stages
- Is too complex for a single prompt to handle well

Most significant pieces of work — proposals, reports, analyses, plans, complex communications — are workflows. The shift to thinking in workflows is one of the biggest quality improvements available to regular AI users.

WHY BREAKING TASKS INTO STAGES WORKS

When you ask AI to do too many things at once, it compromises on everything. The research is superficial because the prompt also asked for structure. The draft is generic because the prompt also asked for analysis. The analysis is shallow because the prompt also asked for recommendations.

Breaking a task into stages allows AI to do each stage well, with appropriate depth and focus. It also allows you to review and adjust at each stage before committing the next stage to an approach that might be wrong.

The best analogy is professional practice: a good architect doesn't hand you the construction drawings at the first meeting. They produce a brief, then a concept, then a scheme, then detailed drawings — each stage reviewed and approved before the next begins.

THE THREE MOST COMMON MULTI-STAGE PATTERNS

THE RESEARCH-PLAN-DRAFT PATTERN

Stage 1: Research and gather information. Stage 2: Plan the structure. Stage 3: Draft the content.

This is the most universally useful workflow. Almost any document benefits from separating the gathering of content from the writing of it.

THE DRAFT-REVIEW-REFINE PATTERN

Stage 1: Produce a complete draft. Stage 2: Review the draft against explicit criteria. Stage 3: Refine based on the review.

This is particularly valuable because Stage 2 uses a different cognitive mode than Stage 1. Writing mode and editing mode are different, and separating them — even within AI — produces better results.

THE DECOMPOSE-EXECUTE-ASSEMBLE PATTERN

Stage 1: Break a large task into components. Stage 2: Execute each component separately. Stage 3: Assemble the components.

This is most useful for large documents or complex outputs where different sections require different content or expertise.

REAL-WORLD EXAMPLE

David needs to produce a quarterly business review for his consultancy's three largest clients.

Single-prompt approach: "Write a quarterly business review for Client X."

Result: A generic template that David has to fill in entirely himself.

WORKFLOW APPROACH:

Stage 1: "I'm going to paste in my notes from the last three months of client interactions, plus their project status reports. Extract: key achievements, challenges faced, decisions made, and open items. Present as a structured list, not prose."

David reviews the extraction, adds two items AI missed, removes one that was irrelevant.

Stage 2: "Using the extracted information, create a narrative structure for the quarterly review. Three sections: What We Achieved, What We Learned, and What Comes Next. For each section, suggest what the key two or three points should be."

David adjusts the suggested structure — moves one point to a different section, adds a point AI didn't suggest.

Stage 3: "Now draft the quarterly review document using the structure we agreed. Professional tone, plain language, about 800 words. The client is the owner and will share it with their management team."

The resulting document is specific, accurate, and usable. It took less time than the single-prompt version and produced dramatically better output — not because AI was better, but because the workflow allowed human judgment to guide it at each stage.

WHY IT WORKS

Multi-stage workflows work for the same reason that professional processes work: each stage has a clear purpose, a clear deliverable, and a defined handoff point. Breaking work into stages allows you to catch errors before they compound. It allows you to apply different types of thinking to different types of sub-tasks. And it allows AI to do each sub-task well rather than doing all sub-tasks adequately.

The decomposition skill

The most important skill in multi-step AI work is not writing better prompts. It is breaking tasks into the right stages.

A well-decomposed task has stages that:

- Have distinct deliverables
- Can be reviewed independently
- Build on each other in a clear sequence
- Are each achievable to a high standard in a single focused interaction

Poorly decomposed tasks break in the wrong places, create dependencies that are hard to manage, or duplicate work across stages.

This chapter teaches the decomposition patterns that work most reliably across common professional tasks.

DECOMPOSITION PATTERN 1: INFORMATION → STRUCTURE → CONTENT

The most broadly applicable pattern. Works for almost any document.

Stage 1 – Information gathering

"Extract all the relevant information from [source(s)]."

→ Output: Raw material, structured as a list

Stage 2 – Structure design

"Given this information, propose a document structure."

What are the main sections? What is the most logical order?

What is the key argument or narrative?"

→ Output: Document structure with brief section notes

[Human review: Adjust structure before committing to draft]

Stage 3 – Draft

"Draft [Document X] using the agreed structure."

→ Output: Full draft

Stage 4 – Review and refine (optional – Chapter 18)

This pattern works for: proposals, reports, articles, presentations, plans, briefing documents.

DECOMPOSITION PATTERN 2: CLARIFY → EXPLORE → RECOMMEND

Most useful for analysis and decision support.

Stage 1 – Clarify the question

"Help me define exactly what question I need to answer."

Here is the situation: [context]

Here is what I'm trying to decide: [decision]

What additional information would be most valuable? What are the key sub-questions?"

→ Output: Sharper framing of the problem

Stage 2 – Explore options and evidence

"Given this framing, what are the main options or considerations?"

What does the evidence suggest about each?"

→ Output: Analysis of options

Stage 3 – Recommend

"Based on what we've established, what would you recommend and why?"

What are the main risks of your recommendation?"

→ Output: Recommendation with reasoning and risk assessment

This pattern works for: decision preparation, strategy analysis, problem diagnosis, options assessment.

DECOMPOSITION PATTERN 3: COMPONENTS → DRAFT EACH → ASSEMBLE

Most useful for large documents with distinct sections.

```
Stage 1 – Define components
"This document has [N] sections: [list them].
Let's draft each one separately. Start with Section 1."

Stage 2 – Draft each component
[For each section:]
"Draft Section [N]: [title].
Key points to cover: [list].
Length: [specification].
Tone: [specification]."
```

```
Stage 3 – Assemble and review
"I'm going to paste in all sections now.
Review them as a complete document:
– Do they flow logically?
– Are there repetitions or gaps?
– Is the tone consistent throughout?
– What would you change?"
```

This pattern works for: annual reports, long proposals, training materials, multi-section documents.

MAKING HANDOFFS EXPLICIT

Between stages, always do two things:

1. REVIEW THE OUTPUT.

Does the output of Stage N accurately reflect what you wanted? Is anything missing? Is anything wrong? Correct it before it propagates to Stage N+1.

2. CARRY FORWARD WHAT MATTERS.

When moving to the next stage, either reference the previous output directly (in the same conversation, AI will remember it) or paste in what AI needs to know. Don't assume AI will apply context from Stage 1 accurately at Stage 3 if the conversation is very long.

TRY THIS

Take a piece of work you need to produce in the next week. Spend five minutes mapping it:

1. What are the distinct stages this work naturally has?
2. What is the deliverable of each stage?
3. Where is human review most important — where could an error in Stage N make Stage N+1 irretrievably wrong?

Now design the workflow before you start. The map will tell you what to ask AI for at each stage.

WATCH OUT

Stages that are too granular create overhead without benefit. If you find yourself doing ten stages for a task that could be done in three, you have over-decomposed. The right granularity is: each stage produces something that can be reviewed and corrected before the next stage begins. If a stage is so small that there's nothing meaningful to review, combine it with the adjacent stage.

Using AI against its own output

One of the most underused techniques in AI work is asking AI to critique what it has produced.

This sounds circular — and it is, slightly. AI is reviewing its own work, which might seem unlikely to catch errors it didn't notice the first time. But it is more effective than it sounds, for a practical reason: critique mode is a different cognitive stance than generation mode. When you ask AI to switch from producing to evaluating, it activates different patterns and often catches problems it generated in draft mode.

THE CRITIQUE LOOP

The basic critique loop has three steps:

STEP 1 — PRODUCE A DRAFT

Ask AI to produce the draft as normal.

STEP 2 — APPLY THE CRITIQUE PROMPT

Ask AI to evaluate the draft against explicit criteria. The more specific the criteria, the more useful the critique.

STEP 3 — REVISE BASED ON CRITIQUE

Ask AI to produce a revised version addressing the issues raised.

THE CRITIQUE PROMPT TEMPLATE

Please review the [document/email/proposal/plan] you just produced.

Evaluate it against the following criteria:

- [Criterion 1: e.g., Does it clearly address the reader's main concern?]
- [Criterion 2: e.g., Is the argument in the right order?]
- [Criterion 3: e.g., Are there claims that could be questioned?]
- [Criterion 4: e.g., Is there anything missing?]
- [Criterion 5: e.g., Is the tone right for the audience?]

For each criterion: identify what is working well and what could be improved.

At the end: give a one-paragraph overall assessment.

Do not be generous. The goal is a better document, not reassurance.

The instruction "do not be generous" matters. Without it, AI tends to produce critique that finds much to praise and softens its concerns. This produces pleasant but not especially useful critique.

WHY IT WORKS

Critique mode works better than it theoretically should because AI is drawing on large bodies of material about what constitutes good work in different domains. When asked to evaluate against explicit criteria, it can apply those standards more rigorously than it applies them by default in generation mode. Think of it as asking a writer to switch hats: generation mode produces content; critique mode applies editorial judgment.

THE BEFORE → BETTER → BEST: CRITIQUE LOOP

The scenario: Sarah has used AI to draft a proposal section and wants to improve it before sending.

Stage 1 — Draft produced

[AI produces a draft proposal section about the agency's approach to the client's brief]

Stage 2 — Critique applied

Sarah sends: *"Please review the proposal section you just drafted. Evaluate it against these criteria: Does it address the client's specific situation or is it generic? Is the argument clear and logical? Are there any claims that aren't supported by what I've told you? Is the tone right for a director-level reader at a mid-size retail company? What's missing? Be direct — I need useful feedback, not reassurance."*

AI responds: *"A few issues to flag. The opening paragraph describes the agency's general approach rather than the client's specific situation — a sceptical reader will notice you haven't demonstrated understanding of their particular challenge. The 'proven results' claim on page two isn't backed by anything specific in what you've provided. The case study section is good but should come earlier — it's currently buried after the methodology, but for this type of client, demonstrated experience is probably the first thing they want to see. The tone is appropriate throughout."*

Stage 3 — Revision

Sarah: *"Revise the draft addressing these points: lead with the client's situation, move the case study up, remove the 'proven results' claim and replace it with [SPECIFIC RESULT I'LL ADD]."*

The resulting revision is substantively better. The critique took two minutes to run. The revision took thirty seconds to instruct. The improvement would have taken Sarah much longer to identify and articulate without the structured critique loop.

SPECIFIC CRITIQUE TYPES

Different types of work benefit from different critique focuses:

FOR PERSUASIVE DOCUMENTS (PROPOSALS, BIDS, PITCHES):

"Does this make the reader want to say yes? What objections might they raise? Is the value proposition clear? What is the weakest point in the argument?"

FOR EXPLANATORY DOCUMENTS (REPORTS, BRIEFINGS, GUIDES):

"Is the key finding easy to identify? Is the logic clear? Are there technical terms an educated non-specialist would need explained? Is anything assumed that shouldn't be?"

FOR COMMUNICATIONS (EMAILS, LETTERS, ANNOUNCEMENTS):

"Is the purpose of this communication immediately clear? Is the tone right for the relationship? Could anything be misinterpreted? Is there anything missing the recipient will need to act on this?"

FOR CREATIVE OR MARKETING CONTENT:

"Is the opening strong enough to keep a reader going? Is there any generic marketing language that a sceptical reader would discount? Is the call to action clear?"

CRITIQUE WITH AN ADVERSARIAL PERSONA

For high-stakes documents, a more powerful critique technique is asking AI to adopt a specific adversarial perspective:

"Review this proposal as if you are a procurement director at a large company who has seen hundreds of similar proposals and is looking for reasons to discount this one. What would make you sceptical? What would make you put this in the discard pile?"

"Review this contract clause as if you are a lawyer representing the other party. What would you challenge? What ambiguity would you exploit?"

"Review this business case as if you are a CFO who has been burned by overoptimistic projections before. What assumptions are too aggressive? What risks are underplayed?"

These adversarial critique prompts produce more rigorous feedback than neutral evaluation and are particularly valuable before anything is presented to an actual sceptical audience.

WATCH OUT

AI critique has two consistent weaknesses. First, it tends to be comprehensive in format but inconsistent in depth — it will often identify a problem without identifying the correct solution. Second, it may miss issues that require external knowledge it doesn't have: the client context it wasn't given, the regulatory nuance that requires specialist knowledge, or the political sensitivities you didn't mention. Use AI critique to catch structural, logical, and stylistic issues. Use human review to catch everything that requires knowledge or judgment AI doesn't have.

The building blocks of complex workflows

Chapter 17 introduced decomposition patterns. This chapter introduces the mechanics of connecting stages together — the "plumbing" of multi-step AI workflows.

Three mechanics cover most of what you will encounter:

Chains — Sequential stages where each output becomes the next input

Loops — Iterative cycles where output is reviewed and refined until it meets a standard

Handoffs — Points where the work passes from AI to a human (or vice versa) for a specific contribution

Understanding which mechanic applies to which part of a task is what turns a list of stages into a functioning workflow.

CHAINS

A chain is the simplest multi-step pattern: $A \rightarrow B \rightarrow C$. Each step's output is the input to the next.

The key discipline in a chain is maintaining coherence: each step should produce output that is complete enough and well-formed enough to be a useful input to the next step. If Step 2 depends on Step 1 and Step 1 is mediocre, Step 3 will inherit and amplify that mediocrity.

THE CHAIN DISCIPLINE:

- Review each step's output before proceeding to the next
 - If a step's output is substantially wrong, correct it before proceeding
 - Be explicit about what carries forward: "Using the outline from Step 1, now draft Section 3"
-

LOOPS

A loop is an iterative cycle: Produce $\rightarrow$ Review $\rightarrow$ Improve $\rightarrow$ (Review again) $\rightarrow$ (Improve again) $\rightarrow$ Done.

Loops are appropriate when:

- The quality bar is high and the first draft is unlikely to meet it
- The criteria for success are explicit and can be systematically applied
- Time spent iterating is less than time spent correcting downstream

The critique loop in Chapter 18 is a loop. So is any scenario where you are refining a draft toward a specific standard.

THE LOOP DISCIPLINE:

- Define what "done" looks like before you start looping — otherwise loops run indefinitely
 - Be specific about what is wrong in each review cycle — vague feedback produces marginal improvement
 - Know when to stop: a document that is 90% right and been through two loops is usually better to complete by hand than to loop a third time
-

HANDOFFS

A handoff is a point where the work passes between AI and human. Handoffs are not failures — they are deliberate design choices about where human judgment adds value that AI cannot.

The best workflows design handoffs deliberately:

AI → Human handoffs occur when:

- AI has produced material that needs human judgment to evaluate (is this claim actually correct?)
- The next step requires knowledge AI doesn't have (what does the client actually want?)
- A decision needs to be made that carries accountability (does this represent our position?)
- The work will be put in front of someone who needs to trust it

Human → AI handoffs occur when:

- A human has made a judgment that AI should now implement (here's the structure I want — now draft it)
- New information has arrived that AI needs to incorporate (here's their response — now draft a reply)
- A human review has identified specific changes that AI should make

A COMPLETE WORKFLOW EXAMPLE: WRITING A CLIENT BRIEF

To illustrate chains, loops, and handoffs working together:

GOAL: Produce a client brief for an upcoming project

STAGE 1 [Chain – AI]

Input: Call notes, client emails, previous project context

Task: Extract all relevant client requirements, stated and implied

Output: Structured requirements list

STAGE 2 [Handoff – Human reviews]

Task: Review requirements list

- Add two requirements AI missed
- Flag one requirement AI stated incorrectly
- Approve as the basis for the brief

STAGE 3 [Chain – AI]

Input: Approved requirements list

Task: Propose brief structure and section headings with brief notes

Output: Proposed brief structure

STAGE 4 [Handoff – Human reviews]

Task: Review structure

- Adjust one section order
- Approve

STAGE 5 [Chain – AI]

Input: Approved structure

Task: Draft the brief, section by section

Output: Full draft brief

STAGE 6 [Loop – AI critiques]

Task: Review draft against brief criteria

Output: Critique identifying two structural issues and one unclear section

STAGE 7 [Loop – AI revises]

Input: Draft + critique

Task: Revise draft addressing identified issues

Output: Revised draft

STAGE 8 [Handoff – Human final review]

Task: Review revised draft

- Final edits (two sentences)
- Approve for client distribution

This workflow, with appropriate tools, might take ninety minutes instead of half a day. The work is better — more complete, better structured, with fewer omissions — because each stage is focused and each handoff is deliberate.

PRO TIP

Document your best workflows. When you design a workflow that works well for a type of task, write it down as a template. The workflow structure — not just individual prompts but the sequence, the handoff points, and the review criteria — is reusable. Over time, a small library of tested workflow templates is one of the most valuable assets you can build as a professional AI user.

WATCH OUT

Over-engineered workflows add overhead without proportionate benefit. A three-paragraph email does not need a seven-stage workflow. Scale the process to the stakes and complexity of the work. A major proposal warrants multiple stages with careful review. A routine internal update does not. The workflow discipline is most valuable when the stakes are high and the output will be scrutinised — apply proportionate judgment elsewhere.

WHAT PART IV HAS ESTABLISHED

You can now design and execute multi-step AI workflows:

- Recognise when a task requires multiple stages
- Break tasks into stages with clear deliverables and handoff points
- Apply three decomposition patterns to different task types
- Use the critique loop to improve drafts systematically
- Connect stages with chains, loops, and deliberate handoffs

Part V applies these skills to twenty-five specific, tested workflows across the most common professional and personal use cases.

Part V is a reference library. These twenty-five workflows are tested, structured, and ready to use — or adapt for your specific situation. Each follows the same format: Situation, Steps with exact prompts, Verification, and Customisation notes.

How to use this section: find the workflow closest to what you need. Read the Situation to confirm it's a good match. Run the Steps. Follow the Verification. Then read Customisation to adapt it to your specific circumstances.

Placeholders appear in [BRACKETS]. Replace them with your specific content before running.

BUSINESS AND PROFESSIONAL

Situation

You need to write an email that matters — a client communication, a proposal follow-up, a difficult message to a colleague, a professional introduction. The stakes are high enough that you don't want to draft it without a second opinion.

STEP 1 — PROVIDE CONTEXT

I need to write a professional email. Here is the context:

Sender: [Your name, role, organisation]

Recipient: [Name, role, organisation, relationship to you]

Purpose of the email: [What you want to achieve with this email]

Background: [Relevant history, context, what's happened so far]

Tone required: [Formal / professional-informal / warm / firm]

Key constraints: [Anything this email must not say or do]

Desired response or outcome: [What you want the recipient to do or feel after reading it]

Length preference: [Brief / medium / thorough]

Please draft this email.

STEP 2 — CRITIQUE THE DRAFT

Review the email you just drafted. Evaluate it on:

- Is the purpose immediately clear?
- Is the tone right for the relationship?
- Could anything be misinterpreted?
- Is anything missing that the recipient needs to act on this?
- Is it the right length?

Be direct about any problems.

STEP 3 — REVISE

Revise the email addressing the issues you identified. [Add any additional changes you want: "Also, make the opening less formal" / "Remove the third paragraph" / "Add a clear call to action at the end."]

VERIFICATION

Before sending: Does it say what you intend? Read it from the recipient's perspective. Would the tone land as intended? Is there anything that could be misconstrued? If the email involves any commitment, fact, or number, verify it independently.

CUSTOMISE

- For internal communications: reduce formality, increase directness
- For client complaints: add "The tone must acknowledge their frustration without admitting liability or making commitments we haven't authorised" to Step 1
- For cold outreach: add "This person does not know me, so make the value proposition clear within the first two sentences"
- For difficult personal situations: add context about the relationship history and the sensitivity required

Situation

You have a conversation coming up that you are not looking forward to. Performance review. Negotiation. Complaint. Conflict. Family tension. You want to go in prepared, not reactive.

STEP 1 — DEFINE THE CONVERSATION

I have a difficult conversation to prepare for. Please help me prepare.

Situation: [What is happening and why this conversation is needed]
who I'm talking to: [Their role, your relationship, relevant background]
what I want to achieve: [The ideal outcome of this conversation for you]
what I'm worried about: [The scenarios I'm dreading or want to avoid]
My position: [What my bottom line is, what I'm willing to do]
Their likely position: [What you expect they want or will say]

STEP 2 — MAP THEIR PERSPECTIVE

Help me understand their perspective as fully as possible.

what legitimate concerns or interests might they have?
what might they feel about this conversation before it starts?
what outcome might they be hoping for?
what would make this conversation go well from their point of view?

I want to understand their side, not just my own.

STEP 3 — PREPARE FOR LIKELY SCENARIOS

Based on what we've established, what are the three or four most likely ways this conversation could go? For each scenario, what is the key thing I need to do or say to steer it toward a good outcome?

STEP 4 — DRAFT KEY STATEMENTS

Help me draft the key things I need to say clearly:

1. My opening statement (what I'll say to start the conversation)
2. My core position, stated clearly but not aggressively
3. A response to their most likely pushback

Tone: [Firm and direct / Warm and collaborative / Measured and professional]

VERIFICATION

Is the opening statement something you could actually say without it feeling scripted or artificial? Have you considered their perspective genuinely, not just strategically? If the conversation involves any specific commitments, facts, or numbers, confirm these before you go in.

CUSTOMISE

- For performance conversations with a direct report: add HR policy context and any documentation you have
- For negotiation: add "I also want to identify possible compromise positions between our positions"
- For family situations: adjust tone instructions significantly and focus Steps 2 and 3 more heavily on empathy
- For complaints: add "I want to be assertive without being hostile"

Situation

You have a meeting coming up on a topic you need to know more about. You want to be prepared enough to contribute intelligently without spending hours reading primary material.

STEP 1 — ESTABLISH THE KNOWLEDGE REQUIREMENT

I have a meeting in [timeframe] about [topic]. I need to be well-prepared.

My current knowledge level: [None / Basic / Some background]

Who else will be in the meeting: [Roles and expertise level]

The likely agenda or decisions: [What the meeting will cover]

What I need to be able to do: [Contribute / Ask intelligent questions / Understand what others are saying / Make a recommendation]

Please give me a briefing on [topic] pitched at the right level for my knowledge and this context.

STEP 2 — IDENTIFY KNOWLEDGE GAPS

Based on what you've told me, what are the three to five things I most need to understand clearly before this meeting? What are the key terms or concepts I should be familiar with? What questions should I be ready to answer?

STEP 3 — PREPARE QUESTIONS

What would be the most intelligent questions I could ask in this meeting to contribute productively and demonstrate I've prepared? Focus on questions that move the meeting forward, not just questions that demonstrate my knowledge.

VERIFICATION

AI briefings are starting points, not authoritative sources. Verify any specific facts, statistics, or claims you intend to cite in the meeting. Confirm current information for any fast-moving topics (technology, regulation, markets). If the meeting involves decisions with significant consequences, supplement AI preparation with primary sources.

CUSTOMISE

- For technical topics: ask AI explicitly to explain technical concepts in plain language before introducing the terminology
- For regulatory topics: add "Focus particularly on the current regulatory position and flag any areas of uncertainty or ongoing change"

- For commercially sensitive topics: structure the briefing around the business context, not just the technical content

Situation

You have a document that is too long to read fully in the time available — or you need to extract specific information efficiently. The document might be a report, a policy, a contract, a research paper, or a lengthy email chain.

STEP 1 — UPLOAD AND INSTRUCT

Please read the document I'm attaching. [Attach document.]

Before summarising, tell me: how long is this document, what type of document is it, and what appear to be its main sections?

STEP 2 — TARGETED SUMMARY

Now produce a summary for the following purpose:

Audience: [Who will read or use this summary]

Purpose: [What decision, action, or understanding this summary will support]

What to include: [Key findings / Main arguments / Action items / Financial figures / Decisions made / Recommendations]

What to exclude: [Background information already known / Technical detail not needed / Appendices]

Length: [Two paragraphs / Five bullet points / One page]

Format: [Prose / Bullet points / Table / Structured sections]

STEP 3 — FLAG WHAT NEEDS ATTENTION

Now separately:

1. What are the most important facts, numbers, or claims in this document that I should verify independently?
2. Were there any sections of the document that you found ambiguous or that required assumptions?
3. Is there anything in the document that seems unusual, surprising, or that a careful reader should notice?

VERIFICATION

Check at least two specific facts extracted from the document against the source text. Verify any numbers or figures independently. If the document is a contract or legal instrument, the plain-language summary is a starting point — do not act on it without reference to the original language for anything that matters.

CUSTOMISE

- For contracts: add "Specifically flag: any obligations on us, any deadlines, any unusual clauses, and any terms that differ from standard commercial practice"
- For research papers: add "Evaluate the strength of the evidence. What is the methodology? What are the limitations the authors acknowledge?"
- For meeting notes/transcripts: add "Extract all action items with owner and deadline. Identify all decisions made."
- For financial documents: add "Flag any calculation that relies on an assumption I should verify"

Situation

You need to write a proposal for a project — internal or external — that makes a clear case for the project, describes the approach, and asks for approval or resources.

STEP 1 — GATHER INFORMATION

I need to write a project proposal. Before drafting, help me structure the content.

Project name: [Name]

Proposing to: [Who will approve this — their role and decision-making context]

What the project will do: [Summary of the project]

The problem it solves: [What is not working now / what opportunity is being missed]

The proposed approach: [How you will do it]

Resources needed: [Time, money, people]

Expected outcomes: [What success looks like, ideally with measurable indicators]

Risks: [What might go wrong]

Timeline: [When it will happen]

Given this, what information am I missing that a strong proposal should include?

STEP 2 — STRUCTURE THE PROPOSAL

Based on what we've established, propose a document structure for this proposal. The reader is [decision-maker's role]. They will want to know: does this project make sense, is the approach credible, and is the case for investment compelling?

Suggest section headings with one sentence on what each section should achieve.

STEP 3 — DRAFT SECTION BY SECTION

Now draft Section 1: [name].

Key points: [from your structure plan]

Length: [specification]

Tone: [professional / business-like / persuasive but evidence-based]

[Repeat for each section, reviewing between sections.]

STEP 4 — CRITIQUE AND REFINE

Review the complete proposal. As a decision-maker who has seen many project proposals:

- What is the most compelling element?
- What is the weakest element?
- What questions would you still have after reading this?
- Is the case for the resources requested clearly made?

VERIFICATION

Verify any financial figures. Check that outcomes are genuinely measurable as stated. Ensure any claimed precedents or comparisons are accurate. Have someone familiar with the decision-maker's perspective review before submission.

CUSTOMISE

- For internal proposals: reduce formality, increase emphasis on operational impact and resource efficiency
- For external bids: add the client's evaluation criteria explicitly and verify the proposal addresses each one
- For new budget requests: add a cost-benefit analysis section and ensure the return on investment is explicitly stated

Situation

You need to attract the right candidates to a role without writing something generic that attracts everyone and therefore no one.

STEP 1 — DEFINE THE ROLE

Help me write a job posting. Here is the role:

Role title: [Title]

Organisation: [Company name and brief description]

What this person will actually do: [Day-to-day responsibilities]

Who they will work with: [Team, reporting line, stakeholders]

What makes someone great at this role: [Skills, experience, behaviours – be specific]

What makes this role attractive: [Why would the right person want this job?]

Essential requirements: [Non-negotiable qualifications or experience]

Desirable but not essential: [Preferences]

Salary range: [If disclosable]

Location/remote/hybrid: [Working arrangement]

Any particular types of candidates you want to attract or avoid?

STEP 2 — DRAFT THE POSTING

Write the job posting.

Tone: [Professional and direct / Warm and human / Formal]

Length: [Standard job board format – concise / Detailed]

Avoid: clichés like "dynamic team player", "fast-paced environment", "passionate", "hit the ground running". Write in plain language that describes the actual role and what we're looking for.

STEP 3 — TEST THE POSTING

Read this posting as a strong candidate who is browsing job boards quickly.

- Does the role stand out?
- Is it clear what they'd be doing day-to-day?
- Is it clear what kind of person would succeed here?
- Would a great candidate self-select in – or would they skip past it?

What would make it stronger?

VERIFICATION

Have the hiring manager review for accuracy — AI may have emphasised the wrong things. Check that essential requirements are genuinely essential (unnecessarily restrictive requirements are both wasteful and a legal risk in many jurisdictions). Verify compliance with any applicable equal opportunities or advertising requirements.

CUSTOMISE

- For specialist technical roles: add a technical requirements section and be specific about tools and technologies required
- For senior leadership roles: emphasise strategic aspects and leadership behaviours over task descriptions
- For organisations with strong cultures: add context about the culture and be honest about what makes the organisation unusual

Situation

You need to produce a client-facing report that summarises work done, findings, or recommendations. The report must be professional, clear, and demonstrate value.

STEP 1 — GATHER AND STRUCTURE CONTENT

I'm producing a client report. Please help me structure it.

Client: [Name, type of organisation, relationship]

Purpose of the report: [What it covers and why the client needs it]

What I have to work with: [The source material I'm going to attach or describe]

What the client cares most about: [Their priorities – cost, risk, growth, efficiency, etc.]

What we want them to do or feel after reading: [Approval, action, confidence, etc.]

Format requirements: [Any specified format / Our standard format]

Length: [Approximate]

Please propose a report structure with the most important content first.

STEP 2 — DRAFT FROM SOURCE MATERIAL

[Attach or paste source material]

Draft the report using this source material. Follow the proposed structure.

Write for the client — assume they are [knowledge level] about this topic. Professional tone, plain language, no unnecessary jargon.

STEP 3 — EXECUTIVE SUMMARY

Now write the executive summary:

- The client's situation and what we were asked to do
- The most important three to five findings or conclusions
- What we recommend they do next

Length: One page maximum. Read-first, not read-last.

STEP 4 — REVIEW

Review this report as if you are the client receiving it.

- Is the most important information easy to find?
- Is there anything that is unclear or requires background knowledge the client might not have?
- Is there anything missing that a client would reasonably expect to see?
- Is the visual structure (headings, paragraphs, any tables) easy to navigate?

VERIFICATION

Verify all statistics, financial figures, and dates against source material. Check that recommendations are logically supported by the findings. Have a colleague read it who has not been close to the work — they will catch assumptions you've made that the client won't share.

Situation

You are running a meeting that needs to produce decisions or actions, not just discussion. You want an agenda that focuses the meeting, and a follow-up that records what happened clearly.

STEP 1 — DESIGN THE AGENDA

I'm running a [type] meeting with [participants and their roles].

Purpose of this meeting: [What decisions need to be made / what needs to be resolved / what needs to be planned]

Time available: [Duration]

Background context: [What situation this meeting is responding to]

Please design an agenda that:

- Allocates time to each item proportional to its importance
- Puts decisions before discussion where possible
- Identifies who leads each item
- Ends with clear next steps

Include: a pre-read list if relevant, any materials that should be prepared in advance.

STEP 2 — POST-MEETING NOTES → FOLLOW-UP

Here are my notes from the meeting: [paste notes / attach transcript or notes document]

Please produce a meeting follow-up document containing:

1. Decisions made (what was agreed)
2. Actions (who does what by when)
3. Topics tabled for another time
4. Any open questions not resolved

Format: brief and structured. This will be sent to all participants within 24 hours of the meeting.

VERIFICATION

Circulate the actions list to attendees for confirmation before sending — memory of who committed to what diverges quickly. Verify any facts or figures noted in the meeting before recording them as agreed.

MARKETING AND CONTENT

Situation

You want to publish something on LinkedIn that is genuinely worth reading — not corporate content or hollow inspiration, but something that says something real.

STEP 1 — DEFINE THE CONTENT

I want to write a LinkedIn post. Here is what I want to share:

My point of view: [What I think / what I've observed / what I've learned]

The insight: [Why this matters or why it's worth saying]

Any story or example: [A specific situation or experience that illustrates the point]

My audience: [Who I'm writing for — industry, role, interest]

My usual voice on LinkedIn: [Formal / conversational / technical / personal]

Engagement goal: [Spark discussion / share expertise / build visibility / announce something]

Any example of a post I've written that I liked? [Paste it if so]

STEP 2 — DRAFT

Write the LinkedIn post.

Length: [Short (100-150 words) / Medium (200-250) / Long (300-400)]

Format: [Prose / Hook + body / Numbered points / Story structure]

Do not use: hollow inspiration, vague motivational filler, "I'm excited to announce", "game-changing", or any phrase that reads like AI wrote it.

First line must make someone stop scrolling.

STEP 3 — TEST THE OPENING

Read only the first sentence. Would someone who doesn't know me stop scrolling to read the rest? If not, write three alternative opening sentences that would.

VERIFICATION

Does this sound like you? Read it aloud. Would a real person in your network read this and think it was worth their time? Is the claim you're making one you can genuinely support?

CUSTOMISE

- If announcing something: lead with the "so what" — why should the reader care before you describe what's happening
 - If sharing expertise: include one specific, non-obvious insight rather than general advice
 - If writing about experience: use a specific story, not a generalisation
-

Situation

You have a topic, a rough idea, or a set of notes that you want to turn into a publishable blog post. The notes are not polished prose — they are raw material.

STEP 1 — EXTRACT AND STRUCTURE FROM NOTES

I want to write a blog post. Here are my rough notes: [paste notes]

Please extract from these notes:

- The central argument or insight
- The main supporting points
- Any examples or evidence
- What is missing that would strengthen the post

Then propose a structure: what order should these ideas go in to make the strongest post?

STEP 2 — DEFINE THE POST

Using the proposed structure, I want you to draft a blog post with these specifications:

Target reader: [Who reads this – their background, what they care about]

Publication: [Where it will be published]

Tone: [my previous post example / formal-but-readable / conversational / authoritative]

Length: [word count]

Purpose: [Inform / Persuade / Entertain / Share experience]

Draft the post section by section. Start with the introduction – I'll review before you continue.

STEP 3 — OPENING AND CONCLUSION

The opening paragraph and the closing paragraph are the two most important.

Opening: Replace the current opening with three alternative versions. One should open with a question, one with a specific story, one with a direct statement of the central argument.

Closing: The conclusion should end with something the reader can do or think about — not a summary. Rewrite it with a concrete, specific takeaway.

VERIFICATION

Is the central argument clear? Has anyone made this point before in a nearly identical way — and if so, does your version add something? Are any specific claims, statistics, or examples accurate? Is the reading level and vocabulary appropriate for your audience?

Situation

You have a piece of content — a presentation, a report, an article, a talk — that contains value you want to extract and publish in different formats.

STEP 1 — CONTENT INVENTORY

I have a [content type] that I want to repurpose. [Attach or paste the content.]

Please identify:

- The three to five most valuable insights or ideas in this content
- The most interesting or surprising point
- Any statistics or specific evidence that would stand alone
- Any stories or examples that would work as standalone content

STEP 2 — FORMAT PLANNING

I want to create content for:

[List the formats: LinkedIn post / email newsletter / short video script / Twitter/X thread / FAQ document / How-to guide / Slide deck / etc.]

For each format I've listed, suggest which part of the original content would work best and how it should be adapted.

STEP 3 — EXECUTE BY FORMAT

Let's start with [first format].

[Using the relevant content identified in Step 2, apply the appropriate workflow for that format – W9 for LinkedIn, W10 for a blog post, etc.]

VERIFICATION

Repurposed content should be adapted, not just repackaged. Verify that each piece works as a standalone without requiring the original context. Check that any statistics or claims retain their accuracy in the new format (statistics can be misleading when stripped of context).

Situation

You need to write copy for a website page — a homepage, a services page, an about page, or a landing page — that converts readers into enquiries or customers.

STEP 1 — DEFINE THE PAGE

I need to write copy for a [page type] for [business/product name].

Business description: [What the business does, for whom, and what makes it different]

Page purpose: [What this page should make the reader do or feel]

Target visitor: [Who comes to this page, what they're looking for, what their concern or question is]

Main message: [The single most important thing this page should communicate]

Tone: [Professional / Friendly / Premium / Practical / Specialist]

Any examples of websites with copy we admire: [links or descriptions]

Please draft this page.

STEP 2 — TEST THE COPY

Read this page as a sceptical first-time visitor who has never heard of [business name] and is comparing several options.

- Does the headline make it immediately clear what this is for?
- Is the value proposition obvious within the first ten seconds?
- Are there any claims that a sceptical visitor would question?
- Is there a clear and compelling reason to take the next step?
- What is the main thing that would make someone leave this page?

STEP 3 — HEADLINES

The headline is the most important element. Write ten alternative headlines for this page. Range from direct to intriguing to benefit-led to question-based. I'll choose the strongest or combine elements from several.

VERIFICATION

Does the copy accurately represent what the business actually does? Are any claims (testimonials, statistics, guarantees) genuine and verifiable? Have a target customer read it cold and ask them: "What do you think this company does? Would you contact them?"

RESEARCH AND LEARNING

Situation

You need to understand a topic you have little or no background in — a new sector, a technical subject, a policy area, a market — well enough to form a view or make a decision.

STEP 1 — ESTABLISH ORIENTATION

I need to understand [topic]. I currently know [your current level – almost nothing / the basics / a little].

Please give me:

1. A plain-English overview: what this is, why it matters, who is involved
2. The key terminology I need to know to read about this intelligently
3. The main debates, disagreements, or open questions in this area
4. What the most important things I should read or watch to go deeper

Pitch this at someone who is intelligent but not yet knowledgeable about this topic.

STEP 2 — IDENTIFY YOUR SPECIFIC NEED

My specific reason for learning about this: [Why you're researching this – decision, project, meeting, writing, general interest]

Given this context, what do I most need to understand? What can I safely deprioritise?

STEP 3 — TEST UNDERSTANDING

Now ask me five questions about what we've just covered. Start simple and get more complex. For each question I answer, tell me whether my understanding is correct and where I should think more carefully.

VERIFICATION

AI summaries of complex topics can oversimplify or — particularly for technical, scientific, or legal subjects — miss important nuance. Use this as a starting point, not a concluding source. For any topic where you will make significant decisions or public claims, follow up with primary sources. Check the AI briefing for any claims that seem surprising — they may be wrong.

CUSTOMISE

- For legal or regulatory topics: explicitly note "flag any areas of ongoing debate or uncertainty" and add "I'll need to verify the current regulatory position from official sources"
- For scientific topics: ask for the quality of evidence behind major claims
- For market or commercial topics: ask AI to flag the uncertainty inherent in market estimates

Situation

You have a document — a contract, a policy, a technical report, a legal opinion — that you need to understand but that is written in language that does not make understanding easy.

STEP 1 — INITIAL READ

Please read this document. [Attach document.]

First pass:

- What type of document is this?
- What is its purpose?
- What are the main sections?
- Who are the parties or stakeholders involved?

Then: what does this document actually mean in plain English? Give me a summary a non-specialist could act on.

STEP 2 — SPECIFIC QUESTIONS

I want to understand the following specific aspects of this document:

1. [Your first specific question – e.g., "What are our obligations under section 4?"]
2. [Second question]
3. [Third question]

For each, tell me: what does the document say, what does it mean in plain language, and is there anything ambiguous or unclear about it?

STEP 3 — RISK AND ACTION IDENTIFICATION

From a practical standpoint:

- What are the most important obligations or deadlines in this document?
- What are the main risks this document creates or addresses?
- What actions do I or my organisation need to take as a result of this document?
- Are there any clauses or provisions that seem unusual or that I should take advice on?

VERIFICATION

For contracts, legal opinions, and regulatory documents: the plain-English summary is an aid to understanding, not a legal substitute. Any action taken on the basis of a document should reference the original text. For significant legal or financial matters, consult a qualified professional.

Situation

You need to choose between options — suppliers, approaches, products, strategies — and want a structured comparison to inform the decision.

STEP 1 — DEFINE THE DECISION

I need to compare [options] to [make a decision].

Context: [Why this decision matters and what's at stake]

Decision-maker: [Who will make the final decision and what they care about]

Options to compare: [List them]

Evaluation criteria: [What matters – list them in rough order of importance]

Constraints: [Budget, timeline, technical requirements, regulatory requirements]

Information available: [What you know about each option – attach documents if relevant]

STEP 2 — STRUCTURED COMPARISON

Compare [options] against the criteria I've specified.

Format: A table where rows are evaluation criteria and columns are options. For each cell, provide a brief assessment (not just a rating – a sentence or two of substance).

After the table: a one-paragraph narrative on each option summarising its main strengths and main weaknesses.

STEP 3 — RECOMMENDATION AND CHALLENGE

Based on the comparison, what would your recommendation be? What is the main argument for each of the other options over the recommended one?

Then: challenge the recommendation. What would have to be true for a different option to be the right choice?

VERIFICATION

Any specific claims about each option should be verified against the options' own documentation. AI comparisons are only as good as the information provided — if you have incomplete information about one option, flag this explicitly in the decision record. For significant procurement decisions, this structured analysis is a useful starting point, not a substitute for due diligence.

Situation

You want to learn something — a tool, a technique, a process, a domain — and want to use AI to accelerate the learning.

STEP 1 — MAP THE SKILL

I want to learn [skill/topic/tool].

My current level: [Complete beginner / Have the basics / Some intermediate experience]

My purpose: [Why I want to learn this – be specific about what I want to be able to do]

My constraints: [Time available, how I learn best, any specific challenges]

Please give me:

1. A clear learning path – what do I need to learn first, second, third?
2. The concepts that are most important to understand early (getting these wrong means struggling later)
3. The most common mistakes beginners make
4. How long should each stage realistically take?

STEP 2 — ACTIVE PRACTICE DESIGN

Design a practice plan for the next [time period] that builds this skill through doing, not just reading.

Each practice exercise should:

- Be achievable in [time per session]
- Build on the previous exercise
- Produce something I can assess myself

Include: what "good enough to move on" looks like for each stage.

STEP 3 — TEACH ME AND TEST ME

Let's work through the first concept: [topic from learning path].

Explain it clearly. Then test my understanding with a question. When I answer, tell me whether I'm right, where I've misunderstood, and what I should focus on.

Continue until I can explain this concept back to you correctly in my own words.

VERIFICATION

AI-generated learning plans may miss important nuance, especially for professional skills where regulatory, safety, or ethical dimensions matter. For professional certifications, verify that the learning path aligns with the actual certification curriculum. For technical skills, test your understanding through actual practice, not just by being able to answer AI's questions.

PERSONAL PRODUCTIVITY

Situation

You have a project to deliver — at work or in your personal life — and you want to plan it properly before starting.

STEP 1 — DEFINE THE PROJECT

I need to plan a project. Here is what I know:

What the project produces: [The end deliverable]

Purpose: [Why this project matters / what problem it solves]

Deadline: [When it must be complete]

Resources: [What time, budget, people, or tools are available]

Constraints: [What you cannot do / what is fixed]

Dependencies: [What you're waiting for / what others are doing]

Risks: [What might go wrong]

Key stakeholders: [Who needs to be kept informed or who can block this]

STEP 2 — BUILD THE PLAN

Please develop a project plan that includes:

1. Work breakdown: what are all the tasks required to deliver this?
2. Sequence: which tasks depend on which?
3. Timeline: working back from the deadline, when does each task need to happen?
4. Critical path: which tasks, if delayed, will delay the whole project?
5. Milestones: what are the key checkpoints?

Present as a structured plan I can use to track progress.

STEP 3 — RISK AND CONTINGENCY

For this project, what are the three most likely things to go wrong? For each, what is the early warning sign, and what should I do if it happens?

VERIFICATION

Have someone with relevant experience review the plan — AI planning is good at structure but may miss domain-specific risks or dependencies. Verify timelines are realistic against your actual experience with similar work.

Situation

You have an overwhelming list of tasks, projects, emails, or obligations and need to bring order to it.

STEP 1 — CAPTURE EVERYTHING

I'm going to paste in everything that is on my to-do list, in my inbox, or on my mind as outstanding. [Paste raw list – it doesn't need to be organised.]

First pass: read through this and tell me:

- How many distinct things are here?
- Are there any that are really the same thing?
- Are there any that seem more urgent than I might realise?

STEP 2 — TRIAGE

Now help me triage these tasks:

For each task, categorise it:

1. URGENT & IMPORTANT: Must happen this week
2. IMPORTANT but not urgent: Schedule for this month
3. URGENT but not important: Could be delegated or handled quickly
4. NEITHER: Question whether this needs to happen at all

Flag any task where you can't categorise it and tell me why.

STEP 3 — WEEK PLAN

Based on what's in Category 1 (urgent and important), help me plan this week.

My available time this week is approximately [hours/days].

Fixed commitments already in the calendar: [list them]

Which of the Category 1 tasks should happen which day? What is a realistic week that is ambitious but achievable?

VERIFICATION

Review the prioritisation against your own judgment — AI may prioritise based on what seems externally important without knowing what your actual professional context makes most critical. Check whether any deadline-driven items have been correctly identified as urgent.

Situation

You need to write something personal that is emotionally complex — an apology, a message to someone you've lost touch with, a difficult family communication, a message declining an invitation in a way that preserves the relationship.

STEP 1 — ESTABLISH THE SITUATION

I need to write a personal message and I want to get it right.

Situation: [What has happened and why this message is necessary]

Recipient: [Who this is to – relationship, history]

What I want to achieve: [The outcome I'm hoping for]

What I'm worried about: [How it might go wrong]

My honest feelings: [As much as you want to share – this helps AI understand the emotional register]

Tone required: [Warm / Apologetic / Careful / Direct / Affectionate]

Please draft this message.

STEP 2 — TONE CHECK

Read this message from the recipient's perspective.

- How might they feel receiving this?
- Does it say what I intended, or could it be read differently?
- Is there anything that might land wrong even though it wasn't intended to?
- Is the length right for the situation?

STEP 3 — ALTERNATIVES

Write two alternative versions of this message:

- One that is slightly more direct
- One that is slightly warmer

I'll choose elements from whichever feels most right.

VERIFICATION

Read the final draft aloud. Does it sound like you? Would the recipient recognise your voice in it? If the message is truly important, sleep on it before sending — AI drafts can be technically good but emotionally wrong in ways that a night's consideration reveals.

Situation

You have a job interview — internal or external — and want to prepare seriously rather than just hoping for the best.

STEP 1 — RESEARCH

I have a job interview for the role of [title] at [organisation].

Please help me prepare.

What I know about the organisation: [paste or link available information]

The job description: [attach or paste]

My relevant experience: [brief summary]

My main strengths for this role: [your view]

My main concerns about this role: [what you're less sure about]

Based on the job description and organisation, what should I know before this interview?

STEP 2 — ANTICIPATE QUESTIONS

Based on the role and organisation, what are the fifteen most likely questions I will be asked? Include:

- Questions about my experience and background
- Questions about why I want this role
- Competency-based questions ("tell me about a time when...")
- Questions about my knowledge of the organisation
- Difficult or challenging questions I should be ready for

For each question, note what the interviewer is likely trying to learn from it.

STEP 3 — PREPARE ANSWERS

Let's prepare my answers to the five hardest questions from the list.

For each, I'll give you my rough answer. You will:

- Identify what is strong about it
- Identify what is missing or could be clearer
- Suggest how to make it more compelling

Let's start with: [first hard question]. My rough answer: [your answer]

STEP 4 — QUESTIONS TO ASK

What are the best questions I could ask the interviewer? Suggest five to eight questions that:

- Demonstrate I've prepared and thought about the role
- Help me genuinely assess whether this is the right role for me
- Are not questions whose answers are on the website

Avoid: questions that seem presumptuous, overly salary-focused, or that suggest I haven't done basic research.

VERIFICATION

Verify all company facts before the interview — AI's information may be outdated. Check the company's website, recent news, and LinkedIn for current information. Practise answers aloud — answers that read well often sound unnatural when spoken. Time your answers to key questions.

HOME AND LIFE

Situation

You want to plan a trip — city break, holiday, business travel — and want more than a generic list of tourist attractions.

STEP 1 — DEFINE THE TRIP

I'm planning a trip and want help making it as good as possible.

Destination: [Where]

Dates/duration: [When, how long]

Who is travelling: [Adults/children/ages, special requirements]

Type of trip: [Relaxation / Culture / Adventure / City exploration / Family / Romantic / Other]

Interests: [What you genuinely enjoy]

Things you want to avoid: [Tourist traps, certain types of activity, etc.]

Budget level: [Budget / Mid-range / Premium]

Existing bookings: [What's already confirmed]

What you've already decided to do: [Any fixed elements]

STEP 2 — PERSONALISED RECOMMENDATIONS

Based on this, give me:

1. The things in [destination] that genuinely match my interests – not the standard tourist list
2. The things most people do that I might want to skip given what I've told you
3. Practical considerations specific to this trip that I should think about in advance

STEP 3 — ITINERARY

Draft a day-by-day itinerary for [number of days].

Build it around my interests and priorities, not a generic tourist circuit.

Be realistic about travel time between locations.

Flag anything that needs advance booking.

Leave space — don't over-schedule.

VERIFICATION

Verify opening hours, prices, and booking requirements from current official sources — AI information may be outdated. For international travel, check current visa, health, and entry requirements from official government sources. Book anything time-sensitive or popular in advance.

Situation

You have received a contract, terms and conditions, rental agreement, or legal document that you need to understand before signing or acting on.

STEP 1 — PLAIN-ENGLISH SUMMARY

Please read this [document type]. [Attach document.]

Give me a plain-English summary:

- What is this document?
- What does it obligate me to do?
- What does it obligate the other party to do?
- What are the most important terms?
- When does it begin and end?

STEP 2 — RED FLAG IDENTIFICATION

Now read this document more carefully for potential concerns.

Flag:

- Any unusual clauses that differ from what you'd typically expect
- Any terms that limit my rights significantly
- Any financial obligations beyond the obvious ones
- Any automatic renewal or escalation clauses
- Any clauses where my obligations are unclear or broad
- Anything I should ask the other party to clarify or change before signing

STEP 3 — QUESTIONS TO RAISE

Based on what you've found, what specific questions should I raise with [the other party / my solicitor / the landlord] before signing? List them in order of importance.

VERIFICATION

This analysis is a tool for understanding and preparation, not legal advice. For any contract involving significant financial commitment, important rights, or professional risk, consult a qualified solicitor. AI can misread or oversimplify legal language. Always read the actual contract before signing — do not rely solely on an AI summary.

Situation

You have a health question — a symptom you want to understand, a diagnosis you've been given and want to learn more about, a medication you want to understand better — and want to use AI as a starting point for a better-informed conversation with your doctor.

STEP 1 — STATE THE QUESTION CLEARLY

I want to understand [health topic] better so I can have a more informed conversation with my doctor.

Context: [Brief background — age, general health context if relevant, what has already been discussed with a healthcare professional]

My specific question: [What I want to understand]

What I already know: [Any information you've already received from a healthcare professional]

Please explain [topic] in plain language. This is to help me understand, not to replace medical advice.

STEP 2 — QUESTIONS FOR YOUR DOCTOR

Based on what you've explained, what would be the most useful questions for me to ask my doctor? I want to:

- Understand my situation better
- Know what my options are
- Understand what I should be monitoring or watching for
- Know when I should seek urgent attention

VERIFICATION

AI is not a substitute for professional medical advice. For any health question with significant consequences, speak to a qualified healthcare professional. AI information about health topics may be incomplete, out of date, or not applicable to your specific situation. Never adjust medication, ignore professional advice, or make significant health decisions based solely on AI output.

WRITING AND EDITING

Situation

You have written something — an article, a report, a proposal, an email — and want to improve it before it goes out. You know it could be better but aren't sure exactly how.

STEP 1 — DIAGNOSTIC READ

Please read this piece of writing. [Attach or paste the document.]

Give me a diagnostic before you suggest any changes:

- What is the main argument or purpose of this piece?
- What is working well?
- What are the main weaknesses?
- Where did you lose interest or find the writing unclear?

Do not make any changes yet – just diagnose.

STEP 2 — TARGETED EDITING

Now, based on the diagnosis, help me improve this piece.

Priority: [Which weakness is most important to fix first?]

Please:

1. Rewrite the weakest paragraph – show me what better looks like
2. Suggest structural changes if needed – should anything be moved, cut, or added?
3. Flag any sentences that are unclear or that say less than they should

Do not over-edit – preserve my voice. Make it better, not different.

STEP 3 — LINE EDITING

Now do a line edit pass on this document.

Look for:

- Sentences that are too long or too complex
- Passive voice where active would be stronger
- Filler words and phrases that add length without meaning
- Any word used repeatedly that should be varied
- Anything that could be expressed more precisely

For each change: show the original and the improved version, and note briefly why the change helps.

VERIFICATION

Read the final edited version aloud. Does it still sound like you? Has AI edited out your voice along with your weaknesses? It is better to keep a few imperfections that are genuinely yours than to have a technically polished piece that reads like it was written by someone else.

Situation

You have rough notes, a voice memo transcript, scattered thoughts, or a messy draft that needs to become a proper document. The content is there — it just needs structure and polish.

STEP 1 — MAKE SENSE OF THE MATERIAL

I have rough material that I need to turn into [document type]. [Attach or paste notes.]

First pass: help me understand what I have.

- What are the main ideas in this material?
- What is the implied argument or narrative?
- What is well-developed versus what is sketchy?
- What is missing that the final document will need?

STEP 2 — BUILD THE STRUCTURE

Propose a structure for [document type] based on what's in the notes.

I need the document to:

- Achieve: [purpose]
- For: [audience]
- In approximately: [word count or length]

For each proposed section, note: what content from my notes goes here, and what will need to be added.

STEP 3 — DRAFT

Now draft [document type] section by section using my notes as source material.

Preserve my ideas and voice. Add structure and polish – do not add new ideas or content I haven't indicated.

Start with [first section]. I'll review before you continue to the next.

STEP 4 — FINAL PASS

Now read the complete draft as a whole. Does it flow? Does each section lead naturally to the next? Is there anything that still sounds like rough notes rather than finished prose? Produce a final clean version.

VERIFICATION

Read the final document as if you were reading it for the first time. Is anything missing that you intended to include? Has anything been included that you didn't intend? Does the document accurately represent your thinking, or has it drifted toward something slightly different in the polishing process?

This completes the twenty-five real-world workflows. The library is designed to be used selectively — find the workflow that matches your situation, adapt the placeholders to your specifics, and use the verification steps before

trusting the output.

Part VI covers the habits and practices that make these workflows reliable: how to verify output, what AI consistently gets wrong, and how to build an AI practice that serves you well over time.

Part V gave you twenty-five workflows to use. Part VI gives you the habits that make them reliable.

This is the part most books skip. They teach you how to get output. They don't teach you what to do with it — how to check it, what to watch for, how much to trust it, and how to build the kind of AI practice that compounds in value over time rather than staying hit-or-miss.

That's what the next five chapters cover.

The problem with confident prose

AI writes well. This is, in a specific sense, a problem.

When a source is hesitant, poorly worded, or obviously uncertain, you naturally approach it with scepticism. When a source is fluent, structured, and confident, you're inclined to trust it. This is a reasonable heuristic for human writing. It is a poor heuristic for AI.

AI produces confident prose regardless of whether the underlying content is correct. It does not get uncertain when it reaches the edge of its knowledge. It does not hedge more on the questions it genuinely doesn't know the answer to. It hedges when you ask it to hedge. Without that instruction, it states.

This is not a bug that will be fixed in the next model release. It is structural. Understanding it changes how you use AI — from reading AI output as you would read a newspaper article, to reading it the way you'd read a first draft from a capable but occasionally unreliable colleague.

That change in posture — from consumer to editor — is the foundation of everything in this chapter.

WHAT VERIFICATION IS NOT

Verification does not mean checking every sentence of every AI output before you use it.

That would be more work than doing the task yourself, which defeats the purpose.

Verification is proportionate. It is targeted. It is a set of quick, learnable habits that catch the errors that matter — the ones that would embarrass you, cost you money, damage a relationship, or lead to a bad decision — before they do.

The goal is not perfect scepticism. The goal is calibrated trust.

THE TWO CATEGORIES OF AI ERROR

AI errors fall into two broad categories, and they require different responses.

CATEGORY 1 — ERRORS OF FACT

AI states something that is not true. A date is wrong. A statistic is invented. A product no longer exists. A law changed. A person's title is incorrect. A quote is misattributed.

These errors can look indistinguishable from correct information. They are checked by going to primary sources.

CATEGORY 2 — ERRORS OF REASONING OR FIT

AI's information may be correct, but its analysis, recommendation, or application to your specific situation is wrong. It has misunderstood the task. It has applied a general principle incorrectly. It has drawn a conclusion that doesn't follow from the evidence.

These errors are checked by your own judgment. Does this actually make sense? Does this conclusion follow from what the AI has laid out? Does this recommendation fit my actual situation?

Both categories require different verification techniques.

THE BEFORE YOU TRUST CHECKLIST

For any output that matters — that you will act on, share, or use in a significant decision — run through these eight questions:

VERIFY — BEFORE YOU TRUST THE ANSWER

1. Is this factual or creative?

If the task was creative (write a story, draft an email, generate ideas), most of this checklist doesn't apply. If the task required facts, it does.

1. Does it contain numbers?

Statistics, percentages, counts, dates, prices, and measurements are the most common location of confident errors. Treat any number that matters as unverified until you check it.

1. Does it cite sources?

AI may produce citations that look real and are not. A convincing title, a plausible author, a credible journal — all fabricated. Do not assume a citation is real because it looks professional.

1. Are the sources real?

If sources are cited: search for them. The combination of author, title, and year should return the actual source. If you can't find it, it may not exist.

1. Could the information have changed?

Regulations, product features, prices, personnel, company structures, research findings — these change. AI's knowledge has a cutoff date and may lag even before that. For fast-moving topics, verify from a current source.

1. Did AI make an assumption?

AI fills gaps with reasonable-sounding assumptions. If you didn't specify something, check what AI assumed. Ask: "What assumptions did you make to answer this?"

1. Did it follow my constraints?

Read the output against your original instructions. Length, format, tone, what to include, what to exclude. AI occasionally drops constraints, particularly in longer outputs or when the constraints conflict.

1. What evidence supports the conclusion?

If AI has reached a conclusion or recommendation, ask: where does that follow from? Is the reasoning shown? Would a different set of assumptions lead to a different conclusion?

VERIFICATION BY TASK TYPE

Not every task carries the same risk profile. Here is a practical guide:

TASK TYPE	MAIN RISK	VERIFICATION APPROACH
Factual research	Invented facts, outdated information	Cross-reference key claims with primary sources
Summarisation	Omissions, misemphasis	Check summary against original document
Document analysis (contract, legal)	Misreading, oversimplification	Read original before acting; professional advice for high stakes
Medical information	Incomplete, not personalised, outdated	Treat as background only; discuss with healthcare professional
Financial information	Not personalised, not current, not regulated	Supplement with regulated advice for significant decisions
Data analysis	Wrong calculations, misinterpretation of the task	Check key figures; test with simple examples
Writing (emails, reports, proposals)	Tone wrong, commitment made you didn't intend, assumption not correct	Read carefully before sending; check facts in the text
Creative tasks	Lower verification burden, but check for unintended content	Read for anything that lands differently than intended

THE ONE-VERIFICATION-STEP DISCIPLINE

For most everyday AI use, the full checklist is unnecessary. But one verification habit that protects you in almost every scenario: **read the output as if you wrote it and are responsible for it.**

Not as someone reviewing AI. As someone who owns what it says.

This shift in perspective catches a surprising proportion of errors. The email AI drafted that makes a commitment you didn't authorise. The summary that emphasises the wrong point. The report section that draws a conclusion you didn't ask for. The sentence that could be read in a way you didn't intend.

Reading your own name on something changes how carefully you read it.

REAL-WORLD EXAMPLE

Marcus is a financial analyst using AI to prepare a briefing on an emerging market for a client presentation. AI produces a thorough, well-structured overview with economic figures, market size estimates, and growth projections — all cited to what appear to be reputable sources.

Marcus runs the Before You Trust checklist. The prose sounds authoritative. The structure is good. But the checklist surfaces three issues:

- Two statistics come from sources he can't locate when he searches. He flags these.
- One growth projection is clearly based on a report from 2023. He checks whether a more recent estimate exists. It does, and the figures have changed significantly.
- The conclusion recommends market entry, but re-reading the criteria Marcus specified reveals that AI overlooked the regulatory constraint he mentioned in Step 1. The conclusion doesn't fit the brief.

Marcus catches all three issues before the presentation. The AI briefing is 80% of the way there — useful, structured, and a significant time saving. But the 20% that needed human judgment could have caused real damage if it had gone unchecked.

WATCH OUT

The more impressive the output looks, the easier it is to skip verification. This is backwards. A beautifully formatted, confident, comprehensive AI response is exactly the situation that calls for checking — because the polish can mask errors that would have been obvious in a rougher draft.

WHEN TO ASK AI TO VERIFY ITSELF

AI cannot reliably verify its own factual claims — it has no external access to confirm whether something is true. But it can be a useful diagnostic tool:

Review what you just told me.

What are you most confident about in this response?

What are you least confident about?

Are there any areas where you might be working from incomplete or outdated information?

What would I need to check before acting on this?

This does not produce perfect self-criticism. But it often surfaces genuine uncertainties that were not flagged in the original response — and it tells you where to focus your independent verification.

PRO TIP

Build a personal "high-stakes list" — the specific types of output you will always verify before using. For most professionals, this includes: statistics and figures you will quote, legal and regulatory statements, financial projections, and any claim that will go in front of a client, employer, or decision-maker. Knowing your own list in advance means you verify consistently, not only when something happens to feel uncertain.

The honest chapter

Every book that teaches you how to use AI better should also tell you, plainly, what AI is not good at. This is that chapter.

These are not edge cases or obscure failure modes. They are patterns — things that happen often enough that a thoughtful AI user learns to anticipate them. Knowing them in advance doesn't mean avoiding AI for these tasks. It means approaching them differently: with more careful setup, more targeted verification, and a clear understanding of where the human judgment needs to be applied.

FAILURE 1 — IT INVENTS FACTS WITH FULL CONFIDENCE

This is the most important failure mode to understand, and the one that has caused the most damage to AI's reputation in serious use.

AI is trained to produce plausible, coherent text. When asked for a fact it doesn't have, it does not always say "I don't know." It often produces a fact-shaped sentence — a number, a name, a date, a citation — that fits the context and sounds correct.

These invented facts are sometimes called "hallucinations" — a term that is evocative but slightly misleading, because it implies something random and dream-like. In practice, AI's invented facts tend to be plausible in the specific way that makes them hardest to catch: they're in the right ballpark, they fit the pattern, they have the right shape.

A real statistic about a real topic from the right approximate era. A paper that could plausibly exist, by an author who does exist, in a journal that does exist. A court case that didn't happen, but probably should have.

When to be especially alert: specific numbers, citations and references, historical details, quotes attributed to real people, claims about specific laws or regulations, product specifications, recent events.

FAILURE 2 — ITS KNOWLEDGE HAS A CUTOFF (AND LAGS BEFORE IT)

All AI models have a training cutoff — a date beyond which they have no information. But even before that date, knowledge can be uneven: events from the months immediately before the cutoff may be underrepresented simply because there hadn't been time for them to generate extensive written material.

The practical implication: for fast-moving topics, AI may be out of date even when you think its information should be current.

When to be especially alert: anything that changes frequently — regulations, product features, pricing, personnel, company structures, market conditions, medical guidelines, software versions.

The workaround: use AI for context and framework, then verify time-sensitive specifics from current sources.

FAILURE 3 — IT TELLS YOU WHAT YOU WANT TO HEAR

AI is trained partly on human feedback, and humans tend to prefer responses that agree with them, confirm their views, and are positive about their ideas. This creates a consistent bias toward validation.

If you present a plan and ask whether it's good, AI will typically find things to praise — even if the plan has significant problems. If you present an argument and ask whether it's convincing, AI will often agree that it is. If you say "I'm thinking about doing X — what do you think?" the response will usually be more supportive than critical.

This is a problem when you need genuine critique. It is especially dangerous when the decision matters.

How to counteract it: ask explicitly for criticism, not balance. "What is wrong with this plan?" is better than "What do you think of this plan?" "What would someone argue against this?" is better than "Is this a good argument?" Structuring the prompt to make critique the explicit task breaks the validation bias more reliably than asking for "honest feedback."

TRY THIS

Test this yourself. Take something you've decided to do — a business decision, a plan, a piece of writing you're proud of. Ask AI: "What do you think of this?" Note the response. Then ask: "What are the strongest arguments against this?" Compare. The second response will almost always be more useful.

FAILURE 4 — IT LOSES TRACK IN LONG CONVERSATIONS

AI's effective attention has limits. In long conversations — particularly those covering multiple topics, with many back-and-forth exchanges, or involving complex instructions given early on — AI can lose track of earlier context.

This manifests as: forgetting a constraint you specified at the start, contradicting something it established earlier, gradually drifting away from the specified format, or simply not carrying an important piece of context into a later response.

When to be especially alert: long conversations, complex multi-part tasks, any situation where you specified important constraints or context early and haven't referenced them since.

The workaround: for long or complex work, periodically restate key constraints. For multi-step workflows, include a brief context reminder at the start of each major step. If something important disappears, simply re-establish it.

FAILURE 5 — IT APPLIES GENERAL KNOWLEDGE INCORRECTLY TO SPECIFIC SITUATIONS

AI knows a lot about how things work in general. It knows less about your specific situation, your specific industry's norms, your specific organisation's constraints, and the specific context you're operating in.

When AI applies general knowledge to a specific situation, the gap between the general and the specific is where errors live.

The legal principle that applies generally may not apply in your jurisdiction. The marketing approach that works in most industries may not work in yours. The management technique that AI recommends may be appropriate for most workplace cultures but not yours.

When to be especially alert: any time AI is making a recommendation or applying a framework to your specific situation, particularly in regulated fields, specialist industries, or contexts with unusual constraints.

The workaround: provide the specific context explicitly. Tell AI what makes your situation different from the general case. Then ask it to apply its knowledge to those specifics — and assess whether the result accounts for them.

FAILURE 6 — IT CAN'T TELL YOU WHAT IT DOESN'T KNOW

Ask AI a question it can't answer well, and it will usually attempt to answer it anyway. The absence of knowledge is rarely clearly flagged.

You can prompt AI to be more explicit about uncertainty — "Please tell me where you're uncertain about this" or "What are the limits of your knowledge here?" — and this helps. But it doesn't fully solve the problem. AI's metacognition (its ability to know what it doesn't know) is imperfect.

The implication: you cannot rely on AI to raise its own hand when it's out of its depth. You have to build verification habits that don't depend on AI alerting you to problems.

FAILURE 7 — MATHS AND CALCULATION

This has improved significantly with newer models and tool access, but remains worth flagging: AI language models are not calculators, and numerical reasoning — particularly multi-step arithmetic, complex calculations, and data interpretation — can produce errors that are presented with complete confidence.

For simple calculations, AI is generally reliable. For complex ones, use a calculator or spreadsheet and have AI work through the reasoning rather than just produce a number.

When to be especially alert: any output where a calculation determines a recommendation, projection, or conclusion. Check the arithmetic independently.

WHY IT WORKS

Understanding these failure patterns doesn't make you cynical about AI — it makes you calibrated. A tool used with clear awareness of its failure modes is far more valuable than the same tool used without that awareness. A surgeon who understands exactly what can go wrong in a procedure is more reliable than one who hasn't thought about it, not less. The same logic applies here.

A SUMMARY TABLE

WHAT AI GETS WRONG	WHY IT HAPPENS	HOW TO PROTECT YOURSELF
Inventing facts	Trained to produce plausible text	Verify specific claims from primary sources
Outdated information	Knowledge cutoff and lag	Check time-sensitive facts from current sources
Telling you what you want to hear	Human feedback bias	Ask explicitly for criticism, not balance
Losing track in long conversations	Attention limits	Restate key constraints; use structured workflows
Misapplying general knowledge	Can't know your specific situation	Provide specific context; verify recommendations fit your case
Not flagging uncertainty	Metacognition limits	Don't rely on AI to signal its own limits
Numerical errors	Language models aren't calculators	Verify significant calculations independently

Trust is not binary

The question people most often ask about AI reliability is the wrong question. They ask: "Can I trust it?" — expecting a yes or no.

The useful question is: "How much trust is appropriate for this specific type of output, from this specific model, for this specific task?" And the answer varies considerably.

AI is not uniformly reliable or uniformly unreliable. It is excellent in some areas, unreliable in others, and variable in the middle. Calibrated trust — knowing roughly where the edges of reliability are — is one of the most practically valuable skills you can develop as an AI user.

This chapter gives you a framework for developing it.

THE THREE TRUST ZONES

Think of AI output as falling into three broad zones:

ZONE 1 — HIGH RELIABILITY (USE WITH NORMAL CARE)

Tasks where AI performs consistently well and errors, when they occur, are usually obvious:

- Drafting and editing prose where you provide the substance
- Brainstorming and generating options
- Explaining concepts in plain language
- Summarising content you have provided (with caveats for long documents)
- Formatting and restructuring text
- Generating questions, frameworks, and checklists
- Creative tasks with clear parameters
- Translation (major languages, non-specialist content)

For Zone 1 tasks: read the output, apply your own judgment, send it. Occasional light verification is sufficient.

ZONE 2 — MODERATE RELIABILITY (VERIFY BEFORE ACTING)

Tasks where AI is often useful but errors are consequential and not always obvious:

- Research on topics where facts matter and you will use them
- Summarising documents you haven't read yourself
- Analysis and recommendations where the reasoning matters
- Professional writing that will carry your name or organisation's name
- Any output that includes specific numbers, names, dates, or citations

For Zone 2 tasks: run the Before You Trust checklist. Verify key facts. Read carefully before sending or acting.

ZONE 3 — HIGH RISK (USE AS INPUT, NOT OUTPUT)

Tasks where AI can assist but human expertise is essential before acting:

- Legal analysis — use AI to understand the terrain, consult a professional before acting
- Medical information — use AI to prepare questions, speak to a healthcare professional for decisions
- Financial advice — use AI for general orientation, consult a regulated professional for significant decisions
- Regulatory compliance — use AI to identify the relevant framework, verify current requirements from official sources
- Any situation where being wrong has significant professional, financial, legal, or health consequences

For Zone 3 tasks: treat AI output as useful background and a starting point for the right questions. Do not act on it directly for high-stakes decisions.

WATCH OUT

Zone 3 doesn't mean "don't use AI here." It means: use AI to be better informed, to ask better questions, and to understand the landscape — then take the decision or action with appropriate professional support. AI can make your conversation with your solicitor, doctor, or financial adviser significantly more productive. It is not a substitute for those conversations.

HOW TO BUILD YOUR PERSONAL TRUST CALIBRATION

Calibration comes from experience — specifically, from noticing when AI is right and when it isn't. This requires a deliberate habit:

1. **When AI gets something wrong, note it.** Not obsessively — just a mental note, or a brief log. What type of task was it? What type of error? This pattern recognition compounds over time.
 1. **When AI gets something unexpectedly right, note that too.** Calibration isn't just about catching failures. Understanding where AI reliably exceeds your expectations is equally valuable.
 1. **Test AI in areas where you already know the answer.** For any domain where you have expertise, run some test prompts. Where does it perform well? Where does it make errors? This is much more informative than discovering failures in unfamiliar territory where you can't recognise the error.
 1. **Ask AI for its confidence.** Not because the self-assessment is perfectly accurate, but because the pattern of what AI claims confidence in versus where it hedges gives you useful signal. A response full of hedges is worth reading differently from one delivered with total certainty.
-

THE MODEL MATTERS (BUT LESS THAN YOU THINK)

Different AI models have different strengths, weaknesses, and reliability profiles. Claude tends toward more careful hedging. GPT models have historically been more prone to confident invention but have improved significantly. Gemini has particular strengths in real-time information access.

These differences are real and worth being aware of. But they're smaller than the effect of how you prompt. A well-structured interaction on a mid-tier model will outperform a poorly structured interaction on the most capable model available. The quality of the interaction setup — clear task, relevant context, explicit constraints, structured verification — matters more than the model choice for most everyday tasks.

The exception: for genuinely difficult analytical tasks, complex reasoning, or work that requires sustained accuracy across a long document, the model choice starts to matter more. Choose accordingly.

PRO TIP

Run the same prompt on two different models for any output you'll rely on heavily. Not because you'll always get different answers — often you won't — but because where the models disagree is exactly where you should look most carefully. Disagreement between models is a flag, not a tiebreaker. It means the question deserves more investigation.

THE TRUST-BUILDING LOOP

Trust, properly calibrated, is not a fixed setting. It develops through a loop:

1. Use AI for a task
2. Verify the output
3. Note where it was right and where it wasn't
4. Adjust your trust calibration for that task type
5. Design your prompts and verification accordingly

This loop, repeated across different task types over time, produces something genuinely useful: a personal model of how AI performs for your specific work. That model — the accumulated experience of knowing where to push hard and where to verify carefully — is one of the most valuable things you can develop as a practitioner.

The habits gap

There is a gap between people who use AI occasionally and people who use it well. It is not, primarily, a knowledge gap. Most people who use AI regularly have read the same articles, tried the same features, seen the same demonstrations.

The gap is a habits gap. People who get consistently good results from AI have built small, consistent practices that compound over time. People who get inconsistent results haven't.

This chapter is about the habits that bridge that gap.

THE FOUR HABITS THAT MATTER MOST

HABIT 1 — FRONT-LOAD THE CONTEXT

The most common reason AI produces a mediocre first response is insufficient context. The most common fix — "make it more specific," "add my role," "be less generic" — is a correction to a problem that didn't need to happen.

The habit: before you start typing a prompt, spend thirty seconds asking yourself: what does AI need to know to do this well? Then put that in the prompt.

This is the 7 Questions framework in its most everyday form. You don't need to run all seven questions before every interaction. You do need the instinct — pausing before you prompt to ask: have I given AI what it needs?

This habit alone, consistently applied, produces a measurable improvement in first-draft quality. It reduces the edit-and-retry cycle. It compounds: the time saved on not re-prompting is time you keep.

HABIT 2 — SEPARATE CREATION FROM CRITIQUE

One of the most consistent ways to produce better AI-assisted output is to separate the creation step from the critique step.

In a single-pass approach, you ask AI to produce something and evaluate what comes back. The problem with this is that evaluation is harder when you're still in generation mode. You're reading to understand, not to critique.

In a two-pass approach, AI creates. Then — in a fresh perspective — AI critiques. Or you do. This separation almost always produces better output.

The habit: for anything that matters, build in a critique step. After getting a first draft, send it back with a specific critique brief: "What's weak here? What's missing? Where does the argument not follow?" The critique step takes two minutes and routinely finds something worth fixing.

HABIT 3 — KEEP A PROMPT LIBRARY

Every time you find a prompt structure that works well — for a type of task you do regularly, a format that reliably produces the right output, a critique framework that catches what matters — write it down.

Not the specific content. The structure. The prompt template that you can reuse, adapting the specifics to each new situation.

This is not complicated. A shared note, a simple document, a section of a personal reference file. Ten prompt templates that cover your most common AI tasks are worth more than a hundred prompts you'll never use again.

The habit: when you produce an interaction that works particularly well, before you close the tab, copy the prompt structure to your library. Over three months, this library becomes one of the most practically useful things you own as an AI user.

HABIT 4 — REVIEW, DON'T JUST GENERATE

The instinct of most AI users is: prompt, read, use. The practice of consistently effective AI users is: prompt, read, verify, edit, use.

The word that bridges the gap is *editor*. Effective AI users are editors of AI output, not just consumers of it.

This doesn't mean correcting every sentence. It means reading the output as if you wrote it: looking for the thing that's slightly off, the conclusion that doesn't follow, the tone that's not quite right, the sentence that makes a commitment you didn't authorise.

The habit: before you send any AI-assisted work out into the world, read it once as its author, not as its reader. What would you change if this were your first draft? Usually it's one or two things. Make those changes.

BUILDING THE HABITS: A PRACTICAL APPROACH

Habits don't form through intention. They form through repetition in specific contexts. Here is a practical approach to building each of the four:

Habit 1 (Front-load context): Put a sticky note next to your screen for one week: "What does AI need to know?" Read it before you start every prompt. After a week, the question will have become automatic.

Habit 2 (Separate creation and critique): For the next ten pieces of writing AI helps you with, run a critique step before you finalise. Note whether it finds anything. After ten times, you'll know whether it's worth the two minutes (it is).

Habit 3 (Prompt library): Create a document called "AI Prompts" in the same place you keep other working documents. For the next month, add a prompt template every time you find one that works well. After a month, decide how useful it is.

Habit 4 (Review as editor): Before your next five AI-assisted emails, read them once as the recipient before sending. Notice how often you catch something small. That small catch is the habit paying off.

REAL-WORLD EXAMPLE

David runs a small consulting firm. He started using AI for client proposals and found the output reasonable but generic.

He builds four habits over eight weeks:

- He creates a project for each active client, loaded with brief context documents: the client's priorities, the project scope, his firm's standard approach.
- He builds a proposal template prompt: a structured brief that covers context, audience, key messages, and constraints. It takes ninety seconds to fill in.
- After every AI proposal draft, he runs a two-question critique: "What's weak?" and "What's missing?"
- He reads every draft once from the client's perspective before it goes out.

Three months later: his proposals take half the time to produce and he has received more unsolicited positive feedback from clients about their quality than in the previous two years.

The AI hasn't changed. The habits have.

WHY IT WORKS

Each of these habits addresses a different failure mode. Front-loading context prevents mediocre first drafts. Separating creation from critique catches the problems that good first drafts obscure. A prompt library prevents starting from zero on repeat tasks. Reviewing as editor catches the specific kind of drift — confident, plausible, subtly wrong — that AI output is prone to. Together, they close most of the gap between occasional and reliable results.

WHAT THESE HABITS ARE NOT

They are not rigid procedures. They are defaults — starting positions that you adapt to the task.

A quick email to a colleague doesn't need a full context front-load. A routine internal summary doesn't need a formal critique step. The habits are most valuable when the stakes are high, the task is complex, or the output will carry your name somewhere it matters.

Part of developing as an AI user is knowing when the full habit stack applies and when a lighter touch is appropriate. That judgment develops with practice.

The pace problem

AI is developing faster than any of us can track. New models, new features, new platforms, new capabilities — often multiple announcements in a single week. For people who use AI professionally, this creates a particular kind of anxiety: the fear of being perpetually behind.

This chapter is about that anxiety — and why it is, in most practical respects, unnecessary.

WHAT ACTUALLY CHANGES (AND WHAT DOESN'T)

Here is something that the breathless pace of AI announcements obscures: the things that make AI interactions work well — clear task definition, relevant context, specific instructions, calibrated verification — are not features. They are practices. They don't become obsolete when a new model releases.

The 7 Questions framework will work on models that don't exist yet. Context-building habits will transfer across platforms. Verification instincts don't expire. The difference between a vague prompt and a well-structured one will remain relevant regardless of how capable models become, because the improvement from structured prompting comes from the quality of the information you provide, not from the model's limitations.

What does change:

- Specific features and where to find them in interfaces
- Which tasks models can now do that they couldn't before
- The quality ceiling — the best output available for a given task
- Platform-specific tooling (projects, memory, integrations)

What doesn't change:

- The underlying principles of what makes a good interaction
- The verification habits that protect you from errors
- The workflow structures that organise complex tasks
- The judgment about what AI is and isn't appropriate for

If you have internalised this book's framework, you are not starting from zero every time a new model drops. You are updating the features while keeping the foundations.

A SUSTAINABLE INFORMATION DIET

You don't need to read every AI newsletter, watch every YouTube video about the latest model, or test every new feature the week it launches. Here is a sustainable approach:

TIER 1 — QUARTERLY ATTENTION (30-60 MINUTES)

Once per quarter, spend an hour reviewing what has changed across your primary AI platforms. Focus on: new features relevant to how you actually use AI, any changes to the reliability or behaviour of models you rely on, anything that might change how you structure your regular workflows.

This is enough to stay genuinely current without AI becoming a part-time hobby.

TIER 2 — OPPORTUNISTIC LEARNING (AS IT ARISES)

When you read something specific that sounds relevant to a task you actually do — a new feature, a better approach to a problem you've encountered — test it on that specific task. Don't add it to your mental backlog. Test it, assess it, adopt it or discard it.

TIER 3 — PASSIVE EXPOSURE (AS BACKGROUND)

A single good source — one newsletter, one podcast, one informed person in your network — is sufficient background. You don't need to keep up with the field. You need to keep up with the practical changes to tools you use.

PRO TIP

The most important question for any AI news item is: does this change anything I do today? Most announcements, even genuinely significant ones, don't change anything about your practice immediately. Filter ruthlessly. When the answer is "yes, this changes how I should do X" — pay attention. When the answer is "interesting, but no practical implication" — move on.

TESTING NEW CAPABILITIES INTELLIGENTLY

When a new model or significant feature launches, here is a practical testing approach:

1. Identify two or three tasks from your real work that you regularly use AI for.
2. Run those tasks on the new model or with the new feature exactly as you would normally.
3. Compare the output to what you've been getting. Is it better? In what specific ways?
4. If better: update your practice. If not: keep what you have.

This is faster and more useful than testing on artificial benchmark tasks or reading evaluation reports. You learn what the change means for your specific work, not for some averaged user.

THE EXPERIENCE CURVE

There is a compounding that happens with AI practice that is easy to miss while you're in it.

The first few months of deliberate use — building prompt habits, setting up workspaces, running workflows, verifying output — feel effortful. You're building processes where there were none.

By six months, most of those processes are instinctive. You don't think about how to set up context; you just do it. You don't consciously decide to run a critique step; it happens automatically. The cognitive overhead drops, and the benefit stays.

By a year, your AI practice has a shape: a set of trusted workflows, a prompt library that reflects your actual work, a personal calibration of what AI does well for you and where to be careful. This is not something a new model release can disrupt, because it is built on practice, not features.

The people who benefit most from AI over time are not the ones who adopted it earliest or follow its development most closely. They're the ones who built consistent habits and let those habits compound.

WHEN TO REASSESS

There are moments that do warrant a more deliberate reassessment of your AI practice. Not because you're behind, but because the landscape has genuinely shifted in ways that affect your work:

- A task you previously found AI unreliable for now has significantly better model support
- A new category of tool (multi-modal, agentic, integrated) becomes stable and practical
- Your work changes significantly and the workflows you've built no longer fit
- A major platform change affects the tools you rely on

At these moments, run the equivalent of a quarterly review, but more deliberately: what has changed, what does it mean for how I work, what should I change?

This is not anxiety. It is the normal maintenance of a professional practice.

WHAT THE AI USER'S PLAYBOOK ACTUALLY IS

The five chapters of Part VI reduce to a few core commitments:

Verify proportionately. Know what AI gets wrong. Calibrate your trust to the task. Build four habits and let them compound. Keep up without chasing.

None of these is complicated. All of them take practice. The payoff is a relationship with AI that is reliable, efficient, and genuinely useful — not a slot machine you occasionally win at, but a tool you have learned to use well.

This completes Part VI. The Appendices that follow are reference materials: the 7 Questions quick reference card, the Before → Better → Best worksheet, the verification checklist, the use-cases guide, and the platform product notes.

APPENDICES

Appendix A

The 7 Questions Quick Reference Card

Cut out or save this. Use it before any AI interaction that matters.

THE 7 QUESTIONS

Ask these before every significant AI interaction.

Q1 — WHAT AM I TRYING TO ACHIEVE?

State the outcome, not the task. Not "write a summary" but "summarise this report so that my manager can decide whether to read the full document."

Q2 — WHAT DOES AI NEED TO KNOW?

Background. Context. Your role. The audience. The purpose. What would help AI understand this situation the way you do?

Q3 — WHAT INFORMATION CAN I PROVIDE?

Documents, data, previous work, examples, reference material. What do you have that AI should use? What should AI not rely on from its own training?

Q4 — WHAT EXACTLY SHOULD AI DO?

Be specific about the task. Not "help me" but "draft a three-paragraph response that acknowledges the concern, explains our position, and ends with a clear next step."

Q5 — WHAT SHOULD THE RESULT LOOK LIKE?

Format. Length. Tone. Structure. Audience. What does good look like for this specific output?

Q6 — HOW WILL I CHECK IT?

Before you prompt: know how you'll verify the output. Numbers? Sources? Fit with constraints? What will you check and how?

Q7 — WHAT SHOULD HAPPEN NEXT?

Will this output feed into something else? Does AI need to know about the next step to make this one useful? Is this a one-off or part of a workflow?

QUICK REFERENCE — THE AI INTERACTION STACK

LAYER	QUESTION IT ANSWERS
Intent	What am I trying to achieve?
Context	What does AI need to know?
Instructions	What exactly should AI do?
Examples	What does good look like?
Source material	What information can I provide?
Workspace	Where is this work living?
Tools	Does AI need any capabilities?
Workflow	What happens before and after?
Output	What should the result look like?
Verification	How will I check it?
Iteration	What happens next?

THE BEFORE YOU TRUST CHECKLIST

Before acting on AI output that matters:

- [] Is this factual or creative?
- [] Does it contain numbers? Verified?
- [] Are citations real? Checked?
- [] Could the information have changed?
- [] Did AI make assumptions I should check?
- [] Did it follow my constraints?
- [] Does the conclusion follow from the evidence?
- [] Have I read it as if I wrote it?

Appendix B

The Before → Better → Best Worksheet

Use this template to apply the Before → Better → Best method to any task.

YOUR TASK

What do you want AI to produce?

STEP 1 — BEFORE (THE NATURAL REQUEST)

Write the prompt you would have sent before reading this book:

[Your instinctive first prompt]

What is missing from this prompt?

- Role/context: _____
- Audience: _____
- Format: _____
- Constraints: _____
- Examples: _____
- Verification: _____

STEP 2 — BETTER (ADD THE MOST IMPORTANT MISSING ELEMENT)

Identify the single most important thing missing from the BEFORE prompt. Add it:

[Improved prompt with the key missing element added]

What changed and why does it matter?

STEP 3 — BEST (APPLY THE FULL 7 QUESTIONS)

Work through each question:

QUESTION	YOUR ANSWER
Q1: What am I trying to achieve?	
Q2: What does AI need to know?	
Q3: What information can I provide?	
Q4: What exactly should AI do?	
Q5: What should the result look like?	
Q6: How will I check it?	
Q7: What should happen next?	

Now write the BEST prompt incorporating all of the above:

[Your best prompt]

STEP 4 — COMPARE

Run all three prompts. Record:

BEFORE	BETTER	BEST
Quality (1-5)		
Usefulness (1-5)		
Needed editing?		
Verified?		
Would use?		

What did you learn?

REUSABILITY

Could this BEST prompt be turned into a reusable template?

Template name: _____

Placeholders to add (things that change each time):

Where you'll save it:

Appendix C

Verification Checklist

Use this checklist for any AI output you will act on, share, or use in a significant decision.

PART 1 — CONTENT VERIFICATION

FACTUAL CLAIMS

- ☐ Key facts identified (numbers, dates, names, statistics, citations)
- ☐ High-priority facts checked against primary sources
- ☐ Sources cited by AI verified as real and accessible
- ☐ Time-sensitive information checked from a current source
- ☐ Any claim I would quote or share publicly verified independently

REASONING AND CONCLUSIONS

- ☐ AI's reasoning is shown (not just the conclusion)
 - ☐ The conclusion follows from the evidence provided
 - ☐ I've identified what assumptions AI made
 - ☐ The recommendation or conclusion fits my specific situation (not just the general case)
 - ☐ I've considered alternative conclusions and whether AI's is the best-supported
-

PART 2 — INSTRUCTION ADHERENCE

- ☐ Output is the length specified
 - ☐ Output uses the format specified
 - ☐ Tone matches what was requested
 - ☐ All constraints from the prompt are observed
 - ☐ Nothing is included that was explicitly excluded
 - ☐ Nothing important is omitted that was explicitly requested
-

PART 3 — TASK-SPECIFIC CHECKS

Complete the rows relevant to your task.

TASK TYPE	SPECIFIC CHECK
Professional email	Read from recipient's perspective. Any unintended commitments? Tone correct?
Research summary	Key claims checked. Summary matches source material accurately.
Document analysis	Read original document for anything AI may have misread or simplified.
Data or calculation	Key figures checked. Calculation steps verified.
Legal or regulatory content	Not relied upon without professional advice for high-stakes decisions.
Medical information	Not acted upon without healthcare professional consultation.
Creative writing	Read for anything that lands differently than intended. Voice preserved.
Proposal or recommendation	Reasoning reviewed. Fit with specific situation confirmed.

PART 4 — FINAL REVIEW

- ☐ Read the output once as its author (not its reader)
- ☐ Anything you would not have written this way? Address it.
- ☐ Anything missing that you intended? Add it.
- ☐ Does it represent your position accurately?
- ☐ Are you comfortable with your name on it?

HIGH-STAKES SIGN-OFF (FOR OUTPUT IN REGULATED, LEGAL, MEDICAL, OR FINANCIAL AREAS)

- ☐ Treated as background and context, not authoritative guidance
- ☐ Professional advice obtained where required
- ☐ Any specific claims verified from official current sources
- ☐ Not relied upon as a substitute for professional judgment

Appendix D

AI Use Cases by Profession

This guide maps common professional roles to the AI workflows most likely to be useful. It is a starting point for identifying where AI can integrate into your work — not a comprehensive list.

HOW TO USE THIS APPENDIX

Find your profession or the closest equivalent. Review the use cases listed. Cross-reference with the Part V workflows that match each use case. For each use case marked with a , apply the high-stakes verification checklist from Appendix C before acting on output.

MARKETING AND COMMUNICATIONS

USE CASE	RELEVANT WORKFLOWS	NOTES
Email and copy drafting	W1, W24, W25	Review for tone, brand fit, any claims that need verification
Campaign brief writing	W25	Provide existing brand guidelines as source material
Research and competitor analysis	W3, W5, W13	Verify time-sensitive market data from current sources
Social and content production	W9, W10, W11	Review for accuracy before publication
Presentation preparation	W8, W3	Check facts; review from audience perspective
Meeting preparation	W3	Supplement with primary research for critical decisions

MANAGEMENT AND LEADERSHIP

USE CASE	RELEVANT WORKFLOWS	NOTES
Difficult conversation preparation	W2	Your judgment on tone and specifics is essential
Report and proposal writing	W7, W25	Review conclusions carefully before sharing
Decision research	W13, W14	Verify key facts; apply your contextual knowledge
Team communication drafting	W1, W24	Review for tone; you know the relationship context
Meeting preparation	W3	Check current data for any fast-moving areas
Performance frameworks	W6, W7	Adapt general frameworks to your organisation's context

FINANCE AND ACCOUNTING

USE CASE	RELEVANT WORKFLOWS	NOTES
Document summarisation	W4, W5	⚠ Verify figures against source documents
Research for client briefing	W3, W13	⚠ Verify from regulated, current sources
Report drafting	W7, W25	⚠ All figures and claims verified independently
Regulatory framework overview ⚠	W13	Starting point only; verify from official current guidance
Client communication	W1, W24	Review for any inadvertent commitments or advice
Presentation and briefing preparation	W8, W3	⚠ All figures and sources verified

LEGAL AND COMPLIANCE

USE CASE	RELEVANT WORKFLOWS	NOTES
Document review and summarisation △	W4, W22	Starting point; read original; professional review for high stakes
Contract red-flag identification △	W22	Use to prepare for professional review, not replace it
Research briefing △	W13, W14	Verify from current primary sources
Client communication drafting	W1, W24	Review carefully; legal professional review where required
Regulatory overview △	W13	Starting point only; verify from current official guidance

Note: AI is not a substitute for legal professional advice. Use cases marked △ should be treated as research and preparation tools, not as authoritative legal guidance.

EDUCATION AND ACADEMIA

USE CASE	RELEVANT WORKFLOWS	NOTES
Research and literature overview	W13, W14	△ Verify citations independently; check sources exist
Writing and editing	W24, W25	Check for voice preservation; verify any claims
Lesson and course planning	W6, W7, W25	Adapt AI output to your specific context and students
Summarisation and note-making	W4, W5	Cross-reference against original materials
Learning and understanding	W15, W16	Excellent use case; verify technical claims in specialist areas
Feedback drafting	W24	Review carefully; you know the individual student's context

SMALL BUSINESS OWNERS

USE CASE	RELEVANT WORKFLOWS	NOTES
Proposal and bid writing	W7, W25	Review for your specific situation; verify pricing figures
Client communication	W1, W24	Review for tone; add personal knowledge of client relationship
Research	W3, W13	Verify key market and competitor data from current sources
Marketing content	W9, W10, W11	Review for brand fit and accuracy
Financial planning overview [△]	W14	Starting point; consult a financial professional for decisions
Contract review [△]	W22	Preparation tool; professional review for significant contracts

HEALTHCARE PROFESSIONALS

USE CASE	RELEVANT WORKFLOWS	NOTES
Summarising clinical documents	W4, W5	[△] Verify against source; professional judgment applies
Patient communication drafting	W1, W24	Review carefully; clinical accuracy is your responsibility
Research and literature overview	W13, W14	[△] Verify from current clinical sources
Administrative writing	W7, W25	Standard professional review applies
Training material development	W6, W7	Review for clinical accuracy

Note: AI should not be used to generate clinical guidance, diagnoses, or treatment plans. These use cases cover administrative and professional communication tasks.

HOME AND PERSONAL USE

USE CASE	RELEVANT WORKFLOWS	NOTES
Travel planning	W21	△ Verify booking details, visa requirements, and entry requirements from official sources
Understanding contracts and documents △	W22	Preparation tool; professional advice for significant decisions
Health questions △	W23	Background information only; consult a healthcare professional
Learning something new	W15, W16	Excellent use case; verify technical claims in specialist areas
Personal writing	W24, W25	Good use case; check for voice preservation
Household decisions	W17, W18, W19, W20	Use as a thinking partner; apply your own judgment

Appendix E

Product Notes — Platform-Specific Information

This appendix contains platform-specific feature information that was current at the time of writing. AI products change frequently. All product notes should be treated as a starting point — verify current availability and behaviour from the platform's official documentation before relying on any specific feature.

IMPORTANT

The information in this appendix was verified in September 2026. Feature availability, naming, pricing, and behaviour may have changed. Before relying on any feature described here:

- 1. Check the platform's current help documentation
- 2. Look for any announcements about feature changes
- 3. Test the feature yourself before building it into your workflow

HOW TO FIND CURRENT DOCUMENTATION

PLATFORM	DOCUMENTATION URL
ChatGPT / OpenAI	help.openai.com
Claude / Anthropic	support.anthropic.com
Gemini / Google	support.google.com/gemini
Microsoft Copilot	support.microsoft.com/copilot

For any specific feature mentioned in this appendix, search the relevant help documentation using the feature name.

E1 — Persistent Instructions and Custom Instructions

WHAT THIS FEATURE DOES

Allows you to set instructions that apply to every conversation without having to retype them. The AI reads these instructions at the start of each session.

WHERE TO FIND IT (SEPTEMBER 2026)

PLATFORM	FEATURE NAME	LOCATION
ChatGPT	Custom Instructions	Settings → Personalisation → Custom Instructions
Claude	Personal preferences / system prompt	Settings → Profile
Gemini	Personalisation settings	Profile menu → Personalisation
Microsoft Copilot	Varies by version	Settings or account preferences

Volatility: MEDIUM — Feature names and locations change with interface updates. The concept is stable; the location may have moved.

WHAT TO INCLUDE IN PERSISTENT INSTRUCTIONS

- Your name and role
- Your typical audience (who you write for)
- Your preferred tone and style
- Any consistent constraints (UK/US spelling, never use X format, always include Y)
- The context of your work that applies broadly

WHAT NOT TO INCLUDE

- Project-specific information (use Projects for this)
- Instructions that only apply to some conversations
- Very long documents or detailed reference material

E2 — Projects

WHAT THIS FEATURE DOES

Groups conversations under a named project, with persistent uploaded documents and context that is shared across all conversations in that project. Designed for ongoing work where you return to the same context repeatedly.

WHERE TO FIND IT (SEPTEMBER 2026)

PLATFORM	FEATURE NAME	NOTES
ChatGPT	Projects	Available in the left sidebar
Claude	Projects	Available in the main navigation
Gemini	Gems	Named AI configurations with system instructions
Microsoft Copilot	Varies	Check current feature availability

Volatility: MEDIUM — Projects features were in active development as of September 2026 across all major platforms. Feature behaviour, capacity limits, and capabilities were subject to change.

BEST FOR

- Recurring client accounts
- Ongoing research projects
- Regular document types with consistent context
- Any work where you return to the same AI context more than a few times

E3 — File and Document Upload

WHAT THIS FEATURE DOES

Upload documents, images, spreadsheets, and other files for AI to read and reference within the conversation.

SEPTEMBER 2026 STATUS

PLATFORM	SUPPORTED FILE TYPES (APPROXIMATE)	NOTES
ChatGPT	PDF, Word, Excel, images, CSV, code files	Size limits apply; check current documentation
Claude	PDF, Word, images, text files	Context window limits determine effective size
Gemini	PDF, images, some office formats	Check current support list
Microsoft Copilot	Microsoft Office formats; varies by version	Tightest integration with Word/Excel/PowerPoint

Volatility: HIGH — Supported file types, size limits, and how documents are processed change frequently. Check current documentation for limits.

KEY BEHAVIOUR TO UNDERSTAND

AI reads the document you upload and uses its contents to answer your questions or complete the task. It does not store the document permanently — it is available for the current conversation (or project, if uploaded there). For large documents, AI may not be able to process the full content at once; for very long documents, directing AI to specific sections is more reliable than assuming it has read every page.

E4 — Web Search

WHAT THIS FEATURE DOES

Allows AI to search the web for current information during a conversation, accessing content beyond its training data cutoff.

SEPTEMBER 2026 STATUS

PLATFORM	STATUS	NOTES
ChatGPT	Available (with search tool enabled)	Can be toggled on in conversation
Claude	Available via web search tool	Check current feature availability
Gemini	Integrated by default	Built into responses where relevant
Microsoft Copilot	Integrated	Bing search integrated throughout
Perplexity	Core feature	Built on web search as primary capability

Volatility: MEDIUM — Web search availability and integration depth have been changing rapidly. Current status may differ.

WHEN IT HELPS

- Current events and recent news
- Current prices, product availability, or specifications
- Recent regulatory or policy changes
- Any information that changes frequently

LIMITATIONS TO UNDERSTAND

Even with web search, AI may not access paywalled content, may retrieve outdated pages if not directed to current sources, and should still be verified for high-stakes factual claims.

E5 — Voice and Audio

WHAT THIS FEATURE DOES

Allows conversation with AI by voice rather than text, with AI responding by voice or text.

SEPTEMBER 2026 STATUS

PLATFORM	FEATURE NAME	NOTES
ChatGPT	Voice mode / Advanced Voice	Available in mobile app and some desktop versions
Claude	Voice (mobile)	Available in Claude mobile apps
Gemini	Voice input	Available in the Gemini app
Apple Intelligence	Siri with AI integration	Built into Apple devices
Google Assistant	AI integration	Available on Android

Volatility: HIGH — Voice features were in rapid development as of September 2026. Capabilities, languages supported, and availability change frequently.

BEST USED FOR

- Hands-free situations (commuting, exercising, cooking)
- Thinking out loud and capturing ideas
- When typing is inconvenient
- Accessibility needs

LIMITATIONS

Voice mode typically has less precise instruction-following than text. For complex or structured prompts, text input remains more reliable.

E6 — Memory

WHAT THIS FEATURE DOES

Some platforms allow AI to remember facts across conversations — so that something you tell AI in one session is available in a future session without you restating it.

SEPTEMBER 2026 STATUS

PLATFORM	MEMORY FEATURE	NOTES
ChatGPT	Memory	Can be reviewed, edited, and deleted in settings
Claude	Memory (in Claude.ai)	User-controlled; viewable in settings
Gemini	Personalisation	Some cross-session personalisation

Volatility: HIGH — Memory features were actively developing as of September 2026. Behaviour, what is remembered, and privacy controls may have changed.

PRIVACY CONSIDERATION

Memory stores information from your conversations for future use. Review what is stored periodically (most platforms provide this in settings). Delete anything you don't want retained.

BEST USED FOR

- Persistent professional context that doesn't change often
- Preferences that apply across all your AI use
- Background about yourself you always want AI to know

NOT A SUBSTITUTE FOR

Project-specific context, document uploads, or detailed instructions. Memory stores facts and preferences; Projects stores ongoing work context and documents.

E7 — Model differences (brief reference)

A very brief guide to how the major models differ in behaviour as of September 2026. Model behaviour changes with updates — this is a starting orientation, not a permanent characterisation.

MODEL FAMILY	CHARACTERISTIC BEHAVIOUR	BEST FOR
GPT-4o / GPT-5 class (OpenAI)	Strong all-round capability; good instruction following; creative tasks	General professional use; coding; complex reasoning
Claude (Anthropic)	Careful, thorough responses; tends to flag uncertainty; good at document tasks	Document analysis; careful writing; nuanced tasks
Gemini (Google)	Good real-time information access; strong with Google Workspace	Research with current information; Workspace integration
Microsoft Copilot	Integrated into Microsoft 365; strong with Office document workflows	Office document tasks; enterprise workflows

Volatility: VERY HIGH — Model capabilities change significantly with each major release. This table reflects September 2026 characteristics only. Current comparisons should be assessed from recent evaluations and your own testing.

This completes the Appendices.

A Note on What Comes Next

This is the first book in the Art of AI series. Book 1 has covered the foundations: the mental model, the interaction skills, the environment, the workflows, and the habits.

Books in development for the series include:

BOOK 2 — THE ART OF AI AT WORK

Deep-dive workflows for professional use cases: managing teams, client communication, research, analysis, and professional writing at a higher level of sophistication.

BOOK 3 — THE ART OF AI FOR BUSINESS OWNERS

AI systems for small and growing businesses: marketing, proposals, operations, customer communication, and financial planning support.

Visit [the series page] for updates on publication dates and to join the reader list.

The Art of AI, Book 1

First edition, 2026

All factual claims verified at time of writing. Platform-specific information in Appendix E should be verified from current official sources before reliance. AI tools were used in the research, drafting, and editing of this manuscript. All editorial judgment, structure, and final decisions are the author's own.