

CHAPTER 9

Control the Output



THE OUTPUT PROBLEM You have provided excellent context, a clear job description, specific instructions, and a strong example. AI produces a response.

It is too long. It is the wrong format. The most important point is buried in paragraph four. It uses a slightly different structure than you specified.

This happens because output defaults are baked into AI's training. Without explicit specifications, AI produces the type, length, and format of output most commonly associated with your type of request. That default is often close but rarely exactly right.

Controlling output is the final step in building an effective interaction, specifying not just what to produce, but what the result should look like when it arrives.

THE FIVE OUTPUT DIMENSIONS Five dimensions of output are almost always worth specifying for significant tasks:

1. LENGTH The most commonly underspecified dimension. How long? Not "short", how many words, sentences, bullet points, or paragraphs? "A one-page briefing," "a three-bullet summary," "a maximum of 150 words," "five sentences", these are unambiguous. "Brief" is not.

- FORMAT

- How should the information be structured? Prose paragraphs,
- numbered list, bullet points, table, structured document with
- headings, slide deck content, script, dialogue? Different formats
- serve different purposes. The format should match how the
- reader will use the output.

- STRUCTURE

- If the output has sections or parts, what are they and in what order? Specifying structure prevents AI from inventing a structure that may not serve your purpose. "The document should have four sections: Background, Key Findings, Recommendation, Next Steps" is better than allowing AI to choose.

- TONE

- How should it sound? The vocabulary level, the degree of formality, the use of first person or third person, the amount of hedging or qualification, the presence or absence of technical language, all are specifiable. Use examples when description alone is insufficient (Chapter 8).

- AUDIENCE

- Who will read this? What do they know? What do they need? The same information presented to a technical expert and a first-time reader should be substantially different. Specifying audience tells AI how to calibrate complexity, assumed knowledge, and depth of explanation.



PART II - BUILD BETTER INTERACTIONS

WHY IT WORKS AI HAS NO INHERENT PREFERENCE FOR ONE FORMAT OVER ANOTHER. IT DEFAULTS TO THE FORMAT MOST

THE BEFORE -> BETTER -> BEST: OUTPUT CONTROL The scenario: James needs a summary of a forty-page government report on graduate employment for a blog post he is writing.

BEFORE

"Summarise this report on graduate employment."

Output: Six paragraphs, dense, comprehensive, presenting findings in the order they appear in the report. Difficult to read. Not formatted for a blog audience. James would have to restructure it entirely.

BETTER

"Summarise the key findings of this report for a general audience interested in graduate careers. Keep it short and clear."

Output: Shorter, more readable, but still not structured for a blog post. "Short and clear" was under-specified.

BEST

Chapter 9 - Control the Output

"Summarise the key findings of this graduate employment report for a blog post aimed at final-year students. They are interested in practical takeaways about their job prospects, not policy analysis.

Format: Three sections with short headings:

- What the data says (two to three bullet points with the most significant statistics)
- What this means for you (two to three practical implications in plain language)
- What to do about it (one or two concrete suggestions)

Maximum 300 words total. Active voice. No jargon. Read level: a confident 20-year-old who reads the Guardian.

Do not include methodology, policy recommendations, or anything that doesn't directly affect a job-seeking student."

Output: A structured, audience-appropriate blog summary. The sections are clear, the language is right for the audience, and the length is correct. James reads it once, adds a paragraph of his own at the end, and posts it.



WHAT CHANGED?



PART II - BUILD BETTER INTERACTIONS



OUTPUT SPECIFICATION WHY IT MATTERED

Three named sections Imposed a usable structure

Bullet count per section Prevented overgeneration

Maximum 300 words Made length unambiguous

"Active voice" Improved readability

Reading level specified Calibrated vocabulary and complexity

Explicit exclusion list Prevented irrelevant policy content

SPECIFYING FORMAT: A PRACTICAL REFERENCE When you need:

THIS SAY

A concise summary "Three to five bullet points, one sentence each"

A structured List the sections explicitly with brief descriptions document

A table "Present this as a table with columns for [X], [Y], [Z]"

Prose, not bullet "Write in flowing paragraphs, no bullet points" points

An email "Subject line + body, professional tone, under 150 words"

An outline "Numbered outline with second-level sub-points"

A list "Numbered list, [n] items, no more"

A one-pager "One A4 page of content: heading, three sections"

CONTROLLING OUTPUT MID-CONVERSATION Output control is not only an upfront specification. It is also a live tool in conversation.

When the output is close but not quite right, you can redirect it with a specific instruction:

"Good structure but too long, cut it by half."

Chapter 9 - Control the Output

"Remove the executive summary section, I don't need it." "This is in bullet points, rewrite as prose paragraphs." "The third point is the most important, lead with it." "Change the tone, this reads too formally for the audience."

Each of these is a specific, actionable correction. AI will act on them precisely. The instruction "make it better" is not a specific correction. Be specific about the dimension you want changed.

TRY THIS FOR YOUR NEXT SIGNIFICANT AI TASK, WRITE OUT YOUR OUTPUT SPECIFICATION SEPARATELY BEFORE YOU WRITE

1. How long? 2. What format? 3. What structure? 4. What tone? 5. Who is the audience?

Then build those answers into the prompt. Notice whether the first output is closer to what you need before you have to ask for revisions.



PART II - BUILD BETTER INTERACTIONS

WATCH OUT EXCESSIVE OUTPUT SPECIFICATION CAN OVERCONSTRAIN THE RESPONSE. IF YOU SPECIFY EVERY SENTENCE,

WHAT PART II HAS ESTABLISHED You can now build an effective AI interaction from the ground up:

Give AI a clear role and a specific job, not a vague question Provide the context that closes the gap between what AI assumes and what you actually need Write instructions specific enough that a capable person could follow them without asking clarifying questions Show AI examples when style, format, or tone matter Specify the output in terms of length, format, structure, tone, and audience

Part III moves from individual interactions to the environment in which they happen, and the significant gains available from building that environment deliberately.

Chapter 9 - Control the Output

PART III Build an AI Environment

Most people use AI as if it were a calculator: pick it up, use it once, put it down. Every new calculation starts from zero. Part III shows you how to change that, building a persistent environment where AI carries your context forward, knows your situation, and becomes genuinely more useful the more you invest in setting it up.



IN THIS PART

CHAPTER 10 CHAPTER 11 Stop Starting From Zero Every Time Projects and Persistent Instructions

CHAPTER 12 CHAPTER 13 Files, Documents, and Source Tools, What AI Can Do Beyond Text Material

CHAPTER 14 CHAPTER 15 The Memory Problem (and How to Building Your Personal AI Workspace Work Around It)