


How to Verify AI Output



THE CONFIDENCE PROBLEM AI writes with the same fluent confidence whether it is correct or incorrect.

This is worth pausing on. When a human expert is unsure, they usually show it: they hedge, they qualify, they say "I think" or "you'd want to check this." When AI is unsure, it typically does not. It produces the same well-structured, authoritative-sounding prose it would produce if it knew exactly what it was talking about.

This is not dishonesty. It is a feature of how AI generates text: it produces the most plausible continuation of the words before it, and confident prose is more plausible than tentative prose in most contexts. But the practical effect is that you cannot use prose confidence as a signal of factual accuracy. The two are disconnected.

This means verification is not optional for important work. It is a habit as fundamental as saving your documents.

WHAT VERIFICATION IS, AND IS NOT Verification is not re-reading the AI output more carefully.

If the output says "the company was founded in 1987," reading it again more carefully will not tell you whether 1987 is right. It will only tell you that AI says 1987 with apparent confidence.

Verification means checking claims against sources that are independent of the AI output: the original document, the company website, a primary source, a qualified professional, your own direct knowledge of the situation.

Verification is proportional to stakes. A quick summary of a document you will read anyway needs a lighter verification pass than a public statement containing statistics that will be published under your

name.

THE VERIFICATION FRAMEWORK Before trusting and using any significant AI output, work through these eight questions:



VERIFICATION IS PROPORTIONAL TO STAKES

Lower stakes Important work Quick read-through and obvious- Check sources, numbers, error check assumptions, and constraints

High consequence Hard to verify Independent sources, expert Do not rely on fluent prose alone; review, and explicit obtain evidence or qualified accountability review



CONSEQUENCES IF WRONG ->

- IS THIS FACTUAL OR CREATIVE?

- Factual content (dates, names, statistics, company information, legal positions, scientific findings) requires verification against sources. Creative content (writing style, structure, framing) can be assessed on its own merits without external sources.

Know which type of output you have. A report that reads like authoritative factual analysis but was written by AI with no source documents is creative content masquerading as factual, and requires the same verification as factual content.

- DOES IT CONTAIN SPECIFIC NUMBERS?

- Numbers are a high-risk category. AI is confidently wrong about numbers with greater frequency than almost any other category.
- Any specific statistic, date, financial figure, percentage, or calculation deserves independent verification before you use it.



PART VI - THE AI USER'S PLAYBOOK

The test: can you find this number in a source you can point to?

- DOES IT CITE SOURCES?

- If AI has cited sources, verify them. AI hallucination of citations is a well-documented and persistent failure mode. AI may produce a citation that looks real, a plausible author name, a plausible journal, a plausible year, for a paper that does not exist.
- Verify by finding the actual source and confirming the claim attributed to it is what the source actually says.

Do not assume that the presence of citations means the claims are verified.

- ARE THE SOURCES REAL?

- This is distinct from the previous question. Even if AI has
- provided real source titles or URLs, verify they are what they
- claim to be and that the content of the source supports the claim
- AI made.

- COULD THIS INFORMATION HAVE CHANGED?

- AI has a knowledge cutoff. Information about current prices,
- regulations, leadership, products, statistics, and events may be
- outdated. For any fast-moving area, technology, regulation,
- markets, current events, check whether the information reflects
- the current situation.

Even with web search enabled, AI may not surface the most recent information or may misread what it finds.

- DID AI MAKE AN ASSUMPTION?

- Complex prompts often involve AI filling gaps you didn't
- explicitly address. AI will do this without flagging it unless you
- specifically ask.

Ask: is there anything in this output that I didn't specify, and has AI assumed correctly? The missing context from Chapter 6 can produce plausible-but-wrong assumptions that make their way into the output.

Chapter 20 - How to Verify AI Output

A useful check: ask AI directly, "What assumptions did you make in producing this output?"

- DID IT FOLLOW MY CONSTRAINTS?

- Check that the output respects the specific constraints you gave.
- Does it stay within the stated word count? Does it avoid the
- topics you said to avoid? Is it written for the audience you
- specified? Is the tone what you asked for?

Constraints are sometimes partially honoured. AI may follow most of your instructions and quietly ignore one. Read explicitly for compliance.

- WHAT EVIDENCE SUPPORTS THE CONCLUSION?

- For outputs that reach conclusions, make recommendations, or
- offer analysis: is the reasoning sound? Does the conclusion follow
- from what was provided? Could the same evidence support a
- different conclusion?

This is the hardest verification task and the one most commonly skipped. It requires judgment, not just source-checking.

VERIFY

BEFORE YOU TRUST THE ANSWER Ask: , Is this factual or creative? , Does it contain numbers? , Does it cite sources? , Are the sources real? , Could the information have changed? , Did the AI make an assumption? , Did it follow my constraints? , What evidence supports the conclusion?



PART VI - THE AI USER'S PLAYBOOK

HIGH-RISK OUTPUT CATEGORIES Some types of AI output require particularly rigorous verification:

STATISTICS AND DATA Verify every number against a primary source. AI confabulates plausible statistics with notable regularity. The figure "studies show that 73% of employees..." is a common pattern, verify that the study exists, that the figure is correctly stated, and that it applies to the context in which you're using it.

NAMES AND ATTRIBUTIONS Quotes attributed to named individuals should be verified. AI can generate plausible-sounding quotes from real people that those people never said. This is a category where being wrong can cause significant professional or legal damage.

LEGAL AND REGULATORY INFORMATION Law changes. Regulations change. Jurisdictional variations are significant. Any legal or regulatory claim should be verified against current official sources. AI may confidently state a legal position that was correct two years ago and is not correct now.

MEDICAL INFORMATION As above, and with additional stakes. Verify medical information from qualified professionals and current clinical sources.

TECHNICAL SPECIFICATIONS Product features, software capabilities, technical standards, all change and all deserve verification from current official documentation.

LOWER-RISK OUTPUT CATEGORIES Verification should be proportionate. These categories typically need lighter verification:

Structural and organisational decisions (how to order a document, what sections to include)

Chapter 20 - How to Verify AI Output

Tone and style adjustments Creative framing and presentation General conceptual explanations of stable topics Grammar, clarity, and prose quality

Even here, read critically. Low-risk doesn't mean no-risk.

REAL-WORLD EXAMPLE Margaret, a retired professional who has taken up writing local history articles, uses AI to help research a piece about a local Victorian factory. AI produces a confident account of the factory's founding, its owner, and the number of workers it employed at its peak.

She runs the verification checklist. The founding date differs from one she has in a local archive document. The owner's name is slightly different from the name she has found in a census record. The employment figure appears plausible but has no source she can trace.

She uses the AI output as a framework and a starting point, it helpfully organised the narrative structure and identified what she should look for. But every specific fact goes back to a source she can cite. The final article is accurate because she verified; it was well-structured because AI helped.

This is the right relationship between AI production and human verification.

BUILDING THE VERIFICATION HABIT Verification is most useful when it is habitual rather than occasional. The goal is to make it automatic: before you share, publish, send, or act on significant AI output, you have a brief, consistent check.



PART VI - THE AI USER'S PLAYBOOK

The check does not need to be long. For most outputs, running through the eight questions takes two minutes. For high-stakes outputs, it takes longer. The discipline is applying it consistently, not applying it intensively.

The section that does most harm is the one you verified every part of except the one part that was wrong.

WATCH OUT ASKING AI TO VERIFY ITS OWN OUTPUT HAS LIMITED VALUE. AI CAN TELL YOU WHAT IT BELIEVES TO BE ACCURATE

Chapter 20 - How to Verify AI Output