


CHAPTER 24

Keeping Up Without Burning Out



THE PACE PROBLEM AI capabilities are changing faster than most professional skills do. A technique that works well today may be less useful in six months because the models have improved and no longer need that scaffolding. A feature you relied on may be removed, replaced, or significantly altered. New capabilities that genuinely change what's possible appear with some regularity.

This pace is genuinely exhausting if you try to follow all of it. It is also largely unnecessary.

WHAT ACTUALLY CHANGES AND WHAT DOESN'T The most important distinction for anyone trying to keep up: some things change rapidly, and some things don't.

THINGS THAT CHANGE RAPIDLY: Specific model capabilities (what a given version of ChatGPT, Claude, or Gemini can do) Platform features (what tools are available in which product at which price point) UI and interface details (where to find which setting) Best-in-class techniques for specific edge cases Pricing and availability

THINGS THAT CHANGE SLOWLY OR NOT AT ALL: The principles of good interaction design (context, instructions, examples, output specs) The 7 Questions framework, these questions are derived from human communication theory, not from AI quirks The verification mindset The workflow thinking The habit of preparing before prompting

The core framework in this book, the 7 Questions, the AI Interaction Stack, the verification habits, the workflow patterns, is designed to be durable. These are not tricks that work today because of a specific model behaviour. They are principles that reflect what it means to communicate a task clearly to a capable but uninformed agent. That doesn't change when the models improve.

What you need to update is the feature-level detail: which capabilities exist, what the product notes say, where volatile information lives. That is why this book puts product-specific detail in Appendix E rather than in the main text.

A SUSTAINABLE INFORMATION DIET You do not need to read every AI newsletter. You do not need to watch every YouTube tutorial about the latest model release. A sustainable information diet for a professional AI user has three elements:

- **A QUARTERLY CHECK-IN ON MAJOR DEVELOPMENTS.**
 - Once a quarter, spend an hour reading about what has changed.
 - Look for: new capabilities that might affect your use cases,
 - changes to platforms you rely on, anything that suggests a
 - technique you use needs updating. This is sufficient for most
 - professional users.
- **PLATFORM-SPECIFIC NOTIFICATIONS FOR THINGS YOU RELY ON.**
 - If you use a specific AI tool for important work, follow that tool's
 - official announcements. When it announces significant changes,
 - read them. This catches the changes that affect your specific
 - workflow.
- **PRACTICAL EXPERIMENTATION OVER PASSIVE CONSUMPTION.**
 - The most efficient way to learn about AI improvements is to use
 - them, not to read about them. When a new capability is
 - announced that seems relevant to your work, try it on an actual
 - task. Fifteen minutes of practical use teaches you more than an
 - hour of reading reviews.



PART VI - THE AI USER'S PLAYBOOK

What to largely ignore: hype cycles, "the new model is 1000% better" content, and AI influencer output that is primarily optimised for engagement rather than practical value.

WHEN TO UPDATE YOUR PRACTICE You need to update your practice when:

A tool you rely on has changed significantly (feature removed, new capability worth using) You encounter a consistent failure mode that a new technique might address Your work situation has changed and your AI use should reflect that Something in your workspace is producing consistently poor output

You do not need to update your practice because:

A new model was released Someone on the internet claims a new technique is revolutionary Your colleague uses AI differently than you

THE TRAJECTORY OF A SKILLED AI USER The early period of AI use is typically characterised by inconsistency: some interactions work well, most don't, and the reasons feel opaque. This book is intended to take you out of that phase.

The middle period, where good habits are established and the core framework is internalised, is where most of the professional value lives. Consistent results, predictable quality, reliable verification. This is not exciting. It is useful.

The advanced period involves increasing sophistication in workflow design, more precise calibration of trust, and the ability to identify the right tool and technique for novel situations

Chapter 24 - Keeping Up Without Burning Out

without referring back to a guide. It comes with accumulated practice, not with reading more AI content.

The trajectory is this: from opaque and inconsistent to systematic and reliable, then from reliable to sophisticated. Part I and Part II get you to systematic. Part III and Part IV get you to reliable. The habits in Part VI get you to sophisticated over time.

WHY IT WORKS THE PEOPLE WHO DEVELOP THE BEST AI PRACTICES ARE NOT THE ONES WHO READ THE MOST ABOUT AI. T

A FINAL WORD Learning to use AI well is a practical craft: techniques can be taught, but judgment develops through practice and reflection.

You now have the framework, the techniques, and the workflows. The rest is practice.

Apply the 7 Questions. Build your workspace. Design workflows that match your work. Verify before you trust. Iterate until it's right. Reflect on what worked. Update what didn't.

And, in time, it becomes less like following a framework and more like simply knowing how to work.



PART VI - THE AI USER'S PLAYBOOK



APPENDICES APPENDIX A

The 7 Questions, Quick Reference Card Print this. Photograph it. Keep it visible until the questions become automatic.

BEFORE EVERY SIGNIFICANT AI INTERACTION 1. What am I trying to achieve?, State the real goal, not just the surface task. 2. What does AI need to know?, Add the background, audience, constraints, and relationships that matter. 3. What information can I provide?, Attach or paste documents, data, examples, and previous work. 4. What exactly should AI do?, Use a clear action verb and define the scope. 5. What should the result look like?, Specify format, length, tone, structure, and audience. 6. How will I check it?, Identify the facts, numbers, assumptions, and reasoning that need review. 7. What should happen next?, Decide whether this is the final result or the input to another stage.

THE INTERACTION STACK, COMPRESSED Intent - Context - Instructions - Examples
- Source material - Workspace - Tools - Workflow - Output - Verification - Iteration

BEFORE YOU TRUST THE OUTPUT Factual or creative? Contains numbers? Citations present? Sources real? Could the information have changed? Did AI make an assumption? Were the constraints followed? Does the conclusion follow from the evidence?

The full explanations appear in Chapters 2, 3, and 20. This page is designed as a quick reference, not as a replacement for those chapters.



APPENDIX B

The Before -> Better -> Best Worksheet Use this when you want to systematically improve a prompt. Work through the questions before redrafting.

YOUR CURRENT PROMPT Write out your current prompt exactly as you usually send it:

[Your current prompt here]

STAGE 1, DIAGNOSE THE GAPS Answer these questions about your current prompt:

JOB CLARITY What specific action have you asked AI to take? If the answer is vague ("help me," "work on," "assist with"), rewrite with a clear verb and scope.

CONTEXT PROVIDED What background have you given AI that it couldn't know on its own? What is the most important thing AI doesn't know that would change its answer?

INFORMATION PROVIDED What documents, data, or examples have you attached or pasted? Is there source material that would make the output more accurate or specific?

OUTPUT SPECIFIED Have you specified: format, length, tone, audience? If any of these are unspecified, what is AI likely to default to, and is that what you want?

VERIFICATION PLANNED What specifically might be wrong in the output? How will you check it?

STAGE 2, BUILD THE BETTER VERSION Using your diagnosis, write the improved prompt:

ROLE (if relevant): You are a [role] with expertise in [domain].

CONTEXT: [Background information, situation, history, constraints, relationships]

TASK: [Specific action verb] [specific deliverable] for [purpose].

SOURCE MATERIAL: [Attached: / Using: / Based on:]

OUTPUT REQUIREMENTS: Format: [prose / bullet points / table / structured document] Length: [word count / page count / number of points] Tone: [formal / professional-informal / conversational] Audience: [who this is for and what they know]

CONSTRAINTS: Do not: [explicit exclusions] Must include: [required elements]

VERIFICATION REQUEST: Please flag any claims you are uncertain about and any assumptions you have made.

STAGE 3, THE BEST VERSION After receiving output from the Better version:

Appendix B - The Before -> Better -> Best Worksheet



WHAT WORKED WELL ?



WHAT STILL NEEDS IMPROVEMENT ?



WHAT FOLLOW - UP INSTRUCTION WOULD MAKE IT RIGHT ?



WHAT SHOULD I ADD TO MY STANDING WORKSPACE TO AVOID THIS GAP NEXT TIME ?



EXAMPLE : BEFORE -> BETTER -> BEST



STAGE PROMPT WHAT CHANGED

Before "Help me with a presentation for our board." Vague task, no context, no output specs

Better "Create a presentation outline for our Q3 Specific task, board update covering financial performance, content defined key decisions, and outlook."

Best "You are preparing a board presentation for Role, specific [Company]. Audience: eight non-executive deliverable, directors, one of whom is new. Format: twelve audience, slide slides maximum. Content: Q3 financial count, content list, performance (headline numbers I'll provide), tone, structural two strategic decisions made this quarter, and requirement, a twelve-month outlook. Tone: direct and exclusion, professional. Lead with the most important verification request finding, not with agenda. Do not include operational detail, board level only. Flag any section where you've made assumptions about emphasis that I should review."



APPENDIX C

Verification Checklist A single-page reference for use before relying on significant AI output.

THE EIGHT-QUESTION VERIFICATION CHECKLIST Work through these before sending, publishing, presenting, or acting on any significant AI output.

- Factual or creative?

Identify which type of output this is. Factual claims require source verification. Creative output (structure, framing, style) can be assessed on its own merits.

- Numbers present?

Identify every specific number, percentage, statistic, or calculation in the output. For each one: can you find it in a source you can point to? If no: do not use it until you can.

- Citations present?

If AI has cited sources, references, or research: For each one: does the source exist? Does it actually say what AI claims? Search for the source independently. Do not assume.

- Sources real?

Beyond checking that sources exist: does the source's content actually support the claim AI has attributed to it? A real paper may be misquoted or misapplied.

- Could this have changed?

For any claim about current facts, who holds a role, what a product costs, what a regulation says, what an organisation does, is this current? AI knowledge has a cutoff. For fast-moving topics, verify from official current sources.

- Assumptions made?

Ask AI directly: "What assumptions did you make in producing this output?" Review the assumptions. Are they correct? Are there important exceptions in your situation?

- Constraints followed?

Check each instruction you gave against the output. Were all followed? Is anything missing that you required? Is anything present that you excluded?

Appendix C - Verification Checklist

- Reasoning sound?

For analytical outputs: does the conclusion actually follow from the evidence? Could the same evidence support a different conclusion? What would have to be true for the conclusion to be wrong?



RISK-CALIBRATED VERIFICATION GUIDE



OUTPUT TYPE PRIORITY CHECKS

Report with statistics 2 (numbers), 5 (currency), 8 (reasoning)

Academic or research content 2, 3 (citations), 4 (sources), 8

Legal or regulatory content 5 (currency), 6 (assumptions), 8

Medical information 5, 6, 8, plus qualified professional review

Client-facing communications 6 (assumptions), 7 (constraints), 8

Internal communications 7 (constraints), light check on others

Creative/editorial content Read critically; 7 (constraints)

WHEN YOU FIND AN ERROR 1. Correct the specific error 2. Ask whether the error is isolated or a signal of a broader problem with this output 3. Consider whether your workflow should have caught this earlier 4. Update your standing verification rules if this is a pattern



APPENDIX D

AI Use Cases by Profession A practical starting point for common professional contexts. Use this to identify which workflows and techniques are most relevant to your work, and to think about what your workspace should contain.

This is a starting point, not an exhaustive list. Your specific situation will always have uses and constraints not captured here.

Marketing and Communications

MOST VALUABLE USE CASES: Drafting and editing client communications, campaigns, and copy (W1, W9, W10, W11, W12) Research and briefing preparation for campaigns (W3, W13) Repurposing content across channels (W11) Preparing for client presentations (W3, W7) Producing first drafts from briefs (W25)

KEY WORKSPACE ELEMENTS: Brand guidelines document (uploaded to each client project) Tone-of-voice reference (example copy that captures the right register) Target audience profiles Standard exclusions per client (regulatory constraints, competitor references, etc.)

WHERE TO BE MOST CAREFUL: Statistics and market claims in copy, these are high-visibility and often checked by clients or regulators. Verify every specific figure before it goes into client-facing material.

Small Business Owners and Entrepreneurs

MOST VALUABLE USE CASES: Client proposals and pitches (W5) Client and supplier communications (W1, W2) Summarising information for decisions (W4, W15) Business planning and project management (W17)

Appendix D - AI Use Cases by Profession

Marketing content (W9, W10, W12)

KEY WORKSPACE ELEMENTS: Company background and positioning statement Target customer description Standard terms and constraints Tone guide (how the business communicates) Any regulatory constraints on the business

WHERE TO BE MOST CAREFUL: Legal and financial claims in business documents, and any content that will be seen by clients before internal review.

HR and People Professionals

MOST VALUABLE USE CASES: Job descriptions and postings (W6) Employee communications Policy and process documentation Meeting notes and follow-ups (W8) Difficult conversation preparation (W2)

KEY WORKSPACE ELEMENTS: Organisation name and structure HR policy documents (relevant sections uploaded for reference) Employment law jurisdiction Standard communication tone Constraints (legal language to avoid in employee communications, etc.)

APPENDICES

WHERE TO BE MOST CAREFUL: Employment law and HR policy claims are jurisdiction-specific and change regularly. Never rely on AI for specific legal positions without verification from qualified legal sources and your current policy documents.

Legal Professionals

MOST VALUABLE USE CASES: Document summarisation and review (W4, W14) Research preparation and briefing (W3, W13) Drafting non-legal client communications (W1) Meeting preparation (W3, W8) Summarising case materials for internal briefing

KEY WORKSPACE ELEMENTS: Jurisdiction and area of practice Firm style and communication standards Standard disclaimers and qualification language What AI output is and is not used for (internal note: AI is a drafting aid; all legal content is reviewed by qualified solicitors)

WHERE TO BE MOST CAREFUL: AI should not be used to generate legal advice or opinions without qualified review. Legal knowledge changes; AI may state outdated positions confidently. AI summaries of contracts and legal documents should always be checked against the original for anything that matters.

Finance and Accounting Professionals

MOST VALUABLE USE CASES: Summarising reports and financial documents (W4) Client communication drafts (W1, W7)

Appendix D - AI Use Cases by Profession

Research preparation (W3, W13) Project and planning support (W17) Preparing for meetings (W3, W8)

KEY WORKSPACE ELEMENTS: Regulatory jurisdiction and relevant standards Client confidentiality constraints Appropriate disclaimer language for client-facing communications Technical terminology explanations for non-specialist communications

WHERE TO BE MOST CAREFUL: All financial figures in AI output must be verified against source data. AI arithmetic is unreliable without code execution. Any regulatory or tax claim must be verified against current HMRC/relevant authority guidance. Never include AI-generated financial projections in client work without independent verification.

Teachers and Education Professionals

MOST VALUABLE USE CASES: Lesson planning and resource development (W16, W25) Differentiated materials for different ability levels Parent communications (W1) Research and preparation for topics (W13) Assessment question development

KEY WORKSPACE ELEMENTS: Year group / age range and subject Curriculum framework (e.g. National Curriculum in England) School's communication style and tone

APPENDICES

Any relevant pastoral or safeguarding context that affects communications

WHERE TO BE MOST CAREFUL: Curriculum-specific content should be verified against current specifications. Content for student consumption should be reviewed by a teacher before use. Subject matter accuracy should be verified, especially in science, history, and other factually complex areas.

Researchers and Academics

MOST VALUABLE USE CASES: Literature orientation for new topics (W13) Summarising papers and documents (W4, W14) Structuring arguments and papers (W17, W25) Editing and improving prose (W24) Preparing for seminars and presentations (W3)

KEY WORKSPACE ELEMENTS: Field and sub-discipline Preferred citation style Institutional guidelines on AI use (these vary significantly; check your institution's current policy) Level of technical language appropriate for target publication

WHERE TO BE MOST CAREFUL: All citations must be verified independently, AI hallucination of academic references is well-documented. AI summaries of research papers may miss important methodological nuance. AI should not be used to generate empirical claims or research findings. Follow your institution's and target journal's guidelines on AI disclosure.

Appendix D - AI Use Cases by Profession

General Professionals (Broad Business Contexts)

M O S T V A L U A B L E S T A R T I N G W O R K F L O W S : W1 (Professional email), universally applicable W3 (Meeting preparation), very high return W4 (Summarising documents), immediate time saving W8 (Meeting agenda and follow-up), immediately practical W15 (Building a workspace), the highest long-term leverage point

M I N I M U M V I A B L E W O R K S P A C E : Your role and organisation type (two sentences) Your standard communication tone UK/US English preference Things AI should never include in your output (jargon, specific competitors, etc.) Default verification instruction ("please flag any claims I should verify independently")



APPENDIX E

Product Notes, Platform-Specific Information This appendix contains product-specific information that is subject to change. All entries are verified as of September 2026. Volatility ratings indicate how likely information is to have changed since verification.

Before relying on any of this information for a workflow or decision, check the platform's current documentation. AI product features, pricing, and availability change frequently and sometimes without announcement.

HOW TO CHECK CURRENT INFORMATION For each major platform:

ChatGPT: help.openai.com Claude: support.anthropic.com and docs.anthropic.com Gemini: support.google.com/gemini Microsoft Copilot: support.microsoft.com

For pricing specifically, always check the platform's pricing page directly. Prices listed by third parties (including this book) become outdated.

Useful official starting points:

ChatGPT Projects: <https://help.openai.com/en/articles/10169521-projects-in-chatgpt> ChatGPT Custom Instructions: <https://help.openai.com/en/articles/8096356-chat-preferences-for-chatgpt> ChatGPT GPTs: <https://help.openai.com/en/articles/8798878-building-and-publishing-a-gpt> Claude Projects and personalisation: <https://support.anthropic.com/en/articles/9519177-how-can-i-create-and-manage-projects> Gemini <https://support.google.com/gemini/answer/15719111> Deep Research:

E1, ChatGPT (OpenAI) Plans available (September 2026): Free, Plus, Pro, Team, Enterprise

PROJECTS FEATURE: Available on free and paid plans. Projects group conversations, project instructions, and reference files; memory settings and file limits vary by plan and workspace. Check current documentation for the limits that apply to your account. Volatility: MEDIUM

Appendix E - Product Notes, Platform-Specific Information

MEMORY FEATURE: Saved memory and project memory can carry selected context across conversations, but availability, settings, and behaviour vary by plan, workspace, and region. Treat memory as a convenience, not as complete or infallible recall. Volatility: MEDIUM

WEB SEARCH: Available across plans (with limitations on Free). Search quality varies; results should be treated as a starting point for verification, not a verified source. Volatility: MEDIUM

CODE EXECUTION (ADVANCED DATA ANALYSIS): Some ChatGPT plans and tools support code execution, file analysis, and chart generation. Availability and limits vary. Most reliable arithmetic is obtained when the calculation is actually executed and the inputs and formula are checked. Volatility: LOW

CUSTOM INSTRUCTIONS (PERSISTENT INSTRUCTIONS): Available across ChatGPT plans, with limits and interface details varying by plan and platform. They apply broadly, while project instructions are scoped to a particular project and may take precedence within it. Volatility: LOW

GPTS: You can use existing GPTs where your account and workspace permit it. Creating or publishing a new GPT is subject to current account, plan, workspace, and eligibility rules; personal ChatGPT accounts may not be able to create or publish new GPTs. Check OpenAI's current GPT documentation before promising this capability to every reader. Volatility: HIGH

APPENDICES

IMAGE GENERATION (CHATGPT IMAGES): Integrated into ChatGPT, with capabilities and access changing over time. Commercial use is subject to OpenAI's current terms and applicable law, verify both before using generated images in commercial projects. Volatility: HIGH, terms and capabilities change

FILE UPLOAD: PDF, Word, Excel, images and other formats. Size and count limits apply, check current documentation. Volatility: MEDIUM

E2, Claude (Anthropic) Plans available (September 2026): Free, Pro, Team, Enterprise

PROJECTS FEATURE: Available on Pro and above. Similar to ChatGPT Projects, groups conversations, maintains uploaded documents, supports project- level instructions. One of the more mature implementations of this feature. Volatility: MEDIUM

SYSTEM PROMPTS / PERSISTENT INSTRUCTIONS: Available via Projects and via platform settings. Robust and reliable as of verification date. Volatility: LOW

WEB SEARCH: Available on Pro and above, via tool integration. As with all AI web search: verify important findings from primary sources. Volatility: MEDIUM

FILE UPLOAD: PDFs, Word documents, images, and other formats supported. Claude has historically handled long documents well, but verify large document behaviour from current documentation.

Appendix E - Product Notes, Platform-Specific Information

Volatility: MEDIUM

MEMORY: Claude provides several forms of persistence and personalisation, including profile preferences, project instructions, project knowledge, and where enabled search across previous chats. These mechanisms differ from one another and do not guarantee complete recall across conversations. Check current availability and settings in Claude's documentation. Volatility: HIGH

EXTENDED CONTEXT WINDOW: Claude models have historically supported long context windows, enabling analysis of very long documents. Specific limits depend on the current model version. Volatility: MEDIUM

E3, Gemini (Google) Plans available (September 2026): Gemini (free), Gemini Advanced (paid subscription), Gemini for Workspace (business integration)

GOOGLE WORKSPACE INTEGRATION: Gemini integrates with Google Docs, Gmail, Drive, Sheets, and Slides for users with Workspace accounts. For professionals already using Google Workspace, this integration is significant for productivity. Depth and reliability of integration varies by product. Volatility: HIGH, integration features are actively expanding

WEB SEARCH: Gemini has native Google Search integration, which is a meaningful advantage for tasks requiring current information. Search quality is generally strong; verify important claims as normal. Volatility: LOW (integration is core to the product)

APPENDICES

GEMS (CUSTOM AI CONFIGURATIONS): Gemini offers custom configurations called Gems. Creation, file support, and availability vary by account, plan, region, and product version; check the current Gemini help centre before relying on a specific capability. Volatility: MEDIUM

FILE UPLOAD: Images, PDFs, and other formats supported. Google Workspace document integration provides additional access to Drive files. Volatility: MEDIUM

E4, Microsoft Copilot Plans available (September 2026): Copilot (free, integrated into Windows and Edge), Copilot Pro (paid), Microsoft 365 Copilot (business, requires M365 subscription)

PRIMARY USE CASE: Copilot's primary value for professional users is its integration with Microsoft 365 applications, Word, Excel, Outlook, Teams, PowerPoint. For organisations running M365, this integration can be significant.

MICROSOFT 365 COPILOT: If your organisation has deployed M365 Copilot, AI assistance is available within the applications you already use. Capabilities include: summarising emails and meetings, drafting documents within Word, analysing data in Excel, generating presentations in PowerPoint. Volatility: HIGH, M365 Copilot is rapidly evolving

DATA SECURITY: For organisations using M365 Copilot under an enterprise agreement, data handling follows Microsoft's enterprise data protection terms. This is significant for regulated industries. Verify current terms with your IT and legal teams. Volatility: MEDIUM

Appendix E - Product Notes, Platform-Specific Information

E5, Cross-Platform Notes

CHOOSING A PRIMARY PLATFORM: For most professional users, the choice of primary AI platform should be based on: integration with your existing tools (if you use Google Workspace, Gemini's integration is worth considering; if you use M365, Copilot may be relevant); the specific capabilities you need most (long document analysis, code execution, image generation); and the plan that provides the right features for your use.

USING MULTIPLE PLATFORMS: Many users find it practical to use more than one platform. A common pattern: one platform as the primary for ongoing projects (with workspace set up), a second for specific tasks it handles better (e.g., using web-search-heavy research tasks with a strong search-integrated tool).

PRIVACY AND DATA: All major AI platforms have privacy policies governing how they use your inputs. Key questions to check in current documentation:

Is your data used to train models? What is the opt-out process? For enterprise accounts: what are the data retention and usage terms?

For professional use involving client data, sensitive business information, or personally identifiable information, review the platform's current privacy and data terms before processing that information.

Volatility: HIGH, policies change; verify before processing sensitive data

APPENDICES

E6, Image Generation (All Platforms)

CURRENT STATE (SEPTEMBER 2026): Image generation is integrated into ChatGPT through its current image-generation experience and is available through several other platforms. Product names, models, access, and quality change frequently.

COMMERCIAL USE: Before using AI-generated images in commercial contexts (client work, marketing, publications), verify the current terms of service of the specific platform. Terms for commercial use vary and have changed multiple times. Verify before use.

COPYRIGHT CONSIDERATIONS: The legal status of AI-generated images for commercial use is an evolving area in multiple jurisdictions. Current legal guidance in the UK and elsewhere is subject to ongoing development. Check current legal guidance before relying on AI-generated images in commercial contexts.

Volatility: VERY HIGH, both platform terms and legal frameworks are actively changing

Appendix E was verified in September 2026. Product features, pricing, availability, and terms change frequently. Always verify from the platform's own current documentation before relying on specific features or capabilities for professional use.

Appendix E - Product Notes, Platform-Specific Information

INDEX

7 Questions 8, 10, 23, 25, 27, 28, 34, 174, 176, 177, 179, 180, 182, 183 AI Interaction Stack 10, 16, 17, 23, 34, 35, 180 chains 93, 108, 109, 112 code execution 77, 80, 81, 165, 192, 196, 200 context 7, 8, 13, 14, 15, 17, 19, 22, 24, 25, 30, 31, 35, 37, 38, 41, 42, 43, 44, 45, 46, 57, 63, 64, 65, 66, 67, 68, 69, 71, 72, 79, 82, 84, 85, 86, 87, 88, 90, 91, 92, 100, 102, 107, 108, 115, 116, 117, 118, 119, 122, 125, 128, 133, 135, 137, 143, 149, 155, 157, 159, 162, 164, 165, 172, 173, 175, 177, 179, 183, 184, 185, 186, 188, 193, 194, 196, 198, 201 context documents 66, 85, 86, 177 context drift 162, 165 constraints 7, 13, 14, 17, 24, 31, 38, 41, 43, 45, 48, 67, 68, 70, 86, 87, 89, 90, 91, 115, 137, 139, 141, 156, 158, 165, 166, 167, 175, 177, 183, 185, 187, 188, 189, 190, 192 Deep Research 78, 195

examples 8, 17, 18, 24, 30, 35, 47, 52, 53, 55, 56, 58, 63, 131, 132, 133, 179, 183, 184 few-shot prompting 53 hallucination 157, 162, 163, 169, 193 handoffs 93, 98, 101, 108, 109, 110, 111, 112 instructions 7, 8, 14, 17, 18, 19, 22, 35, 45, 47, 48, 49, 50, 51, 55, 57, 63, 64, 65, 66, 68, 69, 70, 71, 85, 88, 89, 90, 91, 92, 118, 158, 167, 175, 177, 179, 183, 195, 196, 197, 198 iteration 17, 20, 33, 43, 67, 108, 109, 175, 178, 182, 183 memory 30, 64, 65, 78, 84, 85, 87, 88, 92, 128, 195, 196, 198 output format 8, 31 persistent instructions 45, 64, 65, 68, 69, 70, 71, 88, 89, 90, 91, 92, 177, 196, 197 privacy 200 Projects 49, 50, 64, 65, 66, 67, 68, 69, 71, 72, 86, 88, 90, 92, 142, 195, 197, 200 role 8, 27, 31, 32, 36, 37, 38, 39, 40, 49, 63, 67, 68, 70, 86, 87, 89, 91, 115, 116, 118, 122, 123, 124, 125, 128, 129, 144, 145, 163, 185, 186, 187, 194 source material 14, 17, 18, 31, 40, 54, 55, 64, 73, 77, 113, 126, 127, 152, 183, 184, 185 sycophancy 162, 166, 169

INDEX

verification 8, 14, 17, 19, 20, 22, 24, 26, 27, 30, 31, 33, 40, 56, 70, 72, 74, 75, 77, 78, 79, 80, 81, 82, 83, 85, 86, 90, 113, 115, 117, 119, 120, 121, 123, 125, 127, 128, 130, 132, 133, 134, 135, 137, 138, 140, 141, 143, 144, 145, 147, 148, 149, 151, 153, 154, 155, 156, 157, 158, 159, 160, 161, 163, 164, 165, 168, 169, 170, 171, 172, 173, 174, 175, 176, 177, 178, 179, 180, 181, 182, 183, 185, 186, 187, 188, 189, 191, 192, 194, 196, 197, 198, 199, 200, 201 web search 77, 78, 157, 164, 196, 197, 198 workflows 8, 16, 17, 19, 80, 81, 83, 92, 93, 94, 95, 97, 98, 102, 108, 109, 110, 111, 112, 113, 114, 133, 139, 153, 154, 166, 176, 179, 180, 181, 182, 183, 188, 194

A Note on Sources

This book does not use footnotes in the main text, because most of the principles it describes are based on documented patterns of AI behaviour rather than a single research paper or study. The following notes point to the most useful reference points for readers who want to go further.

ON AI FAILURE MODES AND HALLUCINATION : The phenomenon of AI hallucination is well-documented in academic literature. A practical starting point is work published by researchers at major AI labs themselves, Anthropic, OpenAI, and Google DeepMind all publish research findings on model behaviour that is accessible to non-specialists on their respective research websites.

ON PROMPT ENGINEERING : The field has produced several peer-reviewed studies on what makes prompts effective. The most reliable summary of current understanding remains the documentation produced by each AI lab, particularly Anthropic's documentation on working with Claude and OpenAI's guidance on ChatGPT. These are updated more frequently than academic publications.

ON VERIFICATION PRACTICES : Standards for AI output verification in professional contexts are still emerging. In the UK, professional bodies in law, accounting, and medicine have begun to publish guidance, check your relevant professional body's current guidance.

ON PLATFORM FEATURES : All platform-specific information should be verified from the platform's own documentation. The most current information is always on the platform's help centre, not in any third-party publication.



ABOUT THE AUTHOR

Eleanor Mercer is a professional development trainer and consultant who has worked with professionals across law, financial services, and management consulting for fifteen years. Her original practice focused on communication, writing, and presentation, the skills that separate people who can do excellent work from people who can also explain it clearly.

When AI tools entered professional workplaces in 2022, she began testing them, helping clients use them, and watching closely what happened when capable professionals sat down with capable

technology and still got mediocre results.

The pattern was consistent: the gap was almost never technical. It was what people gave AI before they started, the context, the instructions, the job description, the source material, and the verification process.

The 7 Questions framework and the AI Interaction Stack in this book were developed through that observation, tested in professional training practice, and refined over two years of real use before appearing in print.

The Art of AI is the first book in a planned series on professional AI skills for non-technical users. Eleanor writes, trains, and publishes through The Art of AI.



THEARTOFAI.COM

About This Book

The Art of AI, Book 1, is the first in a planned series covering practical AI skills for non-technical professionals. Future books in the series will cover AI for specific professional contexts, advanced workflow design, and team-level AI integration.

AI tools were used in the research, drafting, and editing of this book. All factual claims were verified against primary sources at the time of writing. All examples were tested. All editorial decisions, structure, voice, inclusion, emphasis, are the author's own.

Published 2026. First edition.