

CHAPTER 1

Why AI Gives You Generic Answers (and How to Change That)



THE PROBLEM EVERYONE HAS BUT NOBODY EXPLAINS

Ask ten people what frustrates them most about AI and you will hear the same answers. The complaints are remarkably consistent — and they all point to the same root cause.

<i>"It's too generic."</i>	<i>"It doesn't understand what I actually want."</i>
<i>"It's confident but wrong."</i>	<i>"The answer doesn't fit my situation."</i>
<i>"It gives me five paragraphs when I needed two sentences."</i>	



*These are not problems with the
AI. They are problems with the interaction.*



THE SLOT MACHINE MODEL (AND WHY IT'S WRONG)

Most people approach AI the way they might approach a vending machine: insert request, press button, hope for the right thing to fall out. When it doesn't, they try slightly different words. This approach treats AI as a slot machine — and produces slot-machine results: random, unpredictable, and mostly disappointing.

The people who consistently get excellent results are not luckier. They are not using secret prompts. They understand something slot-machine users don't: **AI output is shaped by input, and almost every dimension of that input is under your control.**




Figure 1.1 — Two fundamentally different approaches to AI



WHY AI IS GENERIC BY DEFAULT

When you send AI a vague request — *"Write me a summary of this report"* — it has to make assumptions about everything it doesn't know. The AI fills in those gaps with the most statistically average answers it can construct: the kind of response acceptable to the largest number of readers in the largest number of situations. This is why AI output so often feels written for everyone and fits no one particularly well.

Gaps AI must fill when you're vague:

■ Who will read this?	■ Length required?	■ Most important points?
■ What to leave out?	■ Appropriate tone?	■ Reader's prior knowledge?

WHY IT WORKS

AI generates text by predicting what words should follow what came before, based on patterns learned from enormous quantities of written material. Providing more relevant information doesn't confuse AI — it **constrains the prediction space** to something closer to what you actually need.



THE THREE MOST COMMON INPUT FAILURES



Figure 1.2 — The three failure patterns that produce poor AI output

1. The Vague Job Description

"Help me with my presentation" tells AI you want a presentation. It doesn't tell it the topic, the audience, the desired length, your current draft, what help means, or what "better" would look like. AI has no choice but to respond generically to a generic request.

2. The Missing Context

"Reply to this email" gives AI a text to work with. It doesn't tell AI who sent it, what the relationship is, what a good reply should accomplish, how formal the tone should be, or what your constraints are. AI will guess. Sometimes it guesses right. Usually the first reply needs significant editing.

3. The Wrong Assumed Audience

"Explain compound interest" gives AI a topic. But who is asking? A secondary school student? A financial professional? Someone deciding whether to take out a mortgage? AI will produce an explanation pitched at a generic reader — which may be too basic, too technical, or simply beside the point.



REAL-WORLD EXAMPLE: SARAH'S EMAIL

Sarah is a marketing manager who needs to send a follow-up email to a potential agency partner after an introductory call.




Figure 1.3 — Same AI, same model. The only variable was what Sarah provided.



WHAT ACTUALLY CHANGES GOOD AI OUTPUT

Five inputs move the needle most consistently. Mastering them alone will produce an immediate and significant improvement in every AI interaction you have.




Figure 1.4 — The five core inputs that transform AI output quality

1	A Clear Job	What specifically do you want AI to DO? (Not 'help with' — do.)
2	Relevant Context	What does AI need to know about your situation, audience & constraints?
3	Source Material	What documents, notes, or data should AI work from?
4	Output Specifications	What should the result look like — length, format, tone, structure?
5	Verification Instructions	How should AI flag uncertainty or potential errors?

— TRY THIS

Take a prompt you have used before that produced a disappointing result.

Write it down. Then ask yourself:

- What did I assume AI would know that it couldn't have known?
- What context did I leave out?
- What does 'good' actually look like for this task?

Rewrite the prompt with those gaps filled in. Run both versions.

Notice what changed.



WHAT THIS CHAPTER ESTABLISHED

Generic inputs produce generic outputs. This is a feature, not a bug — it means every improvement you make to your input produces a corresponding improvement in output. You are not at the mercy of the AI. **You are in control of the most important variable in every interaction.**

The next chapter introduces a model for understanding exactly which variables are in play — and which ones matter most for any given task.

■ WATCH OUT

More detail is not always better. Adding irrelevant information doesn't improve output — it can dilute it. The skill is not writing longer prompts; it's identifying what AI actually needs to know to do this specific task well.

You'll develop this judgment through the exercises in Part II.



Same AI. Same model. The variable was what you gave it.